



Computer Graphics (Graphische Datenverarbeitung)

- Textures & Filtering -

WS 2021/2022



Corona

- Regular random lookup of the 3G certificates
- Contact tracing: We need to know who is in the class room
 - New ILIAS group for every lecture slot
 - Register via ILIAS or this QR code (only if you are present in this room)





Overview

- Last time
 - BRDFs
- Today
 - Textures to modify surface properties
 - Texture Parameterization
 - Procedural Shading
 - Texturing Filtering

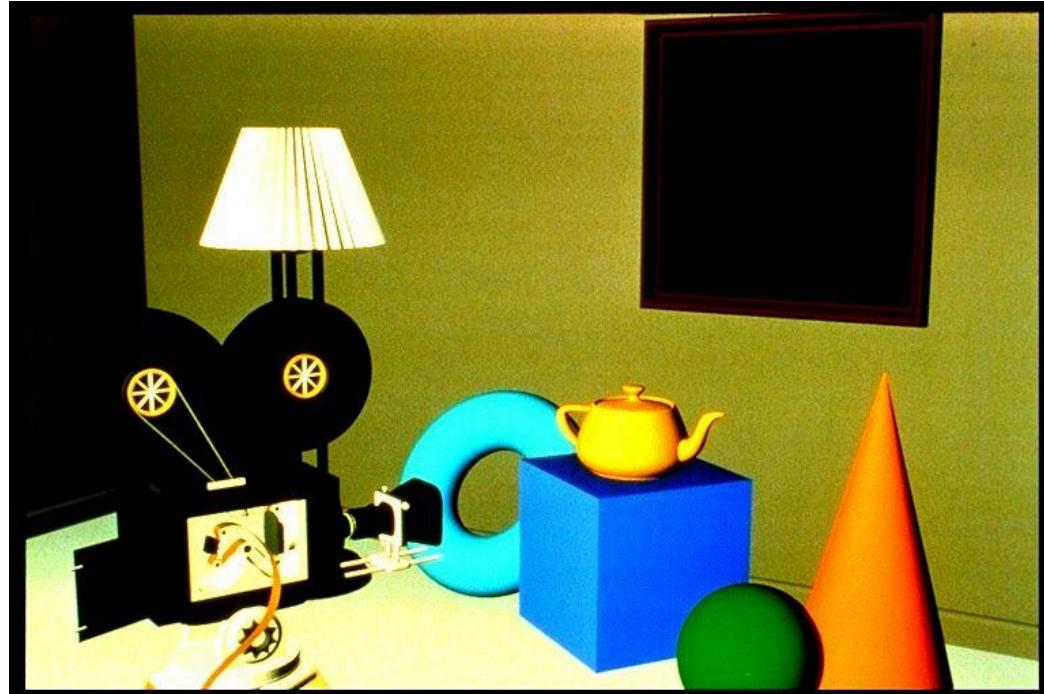
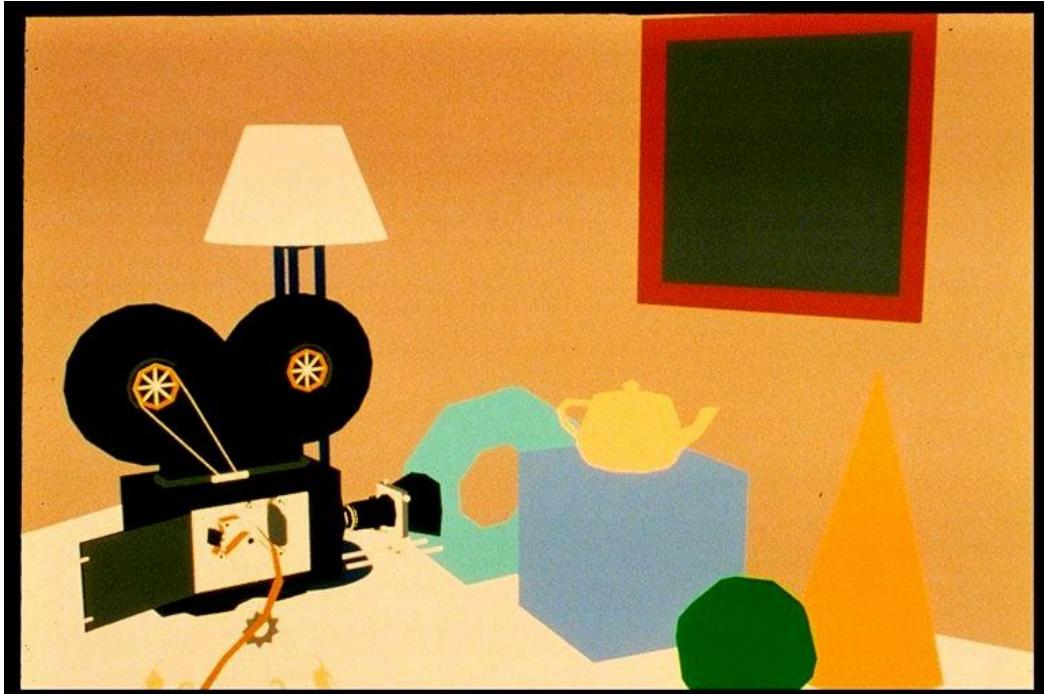


Texturing

How can we make surfaces non-uniform



Simple Illumination



- No illumination
- Constant colors
- Parallel light
- Diffuse reflection



Standard Illumination

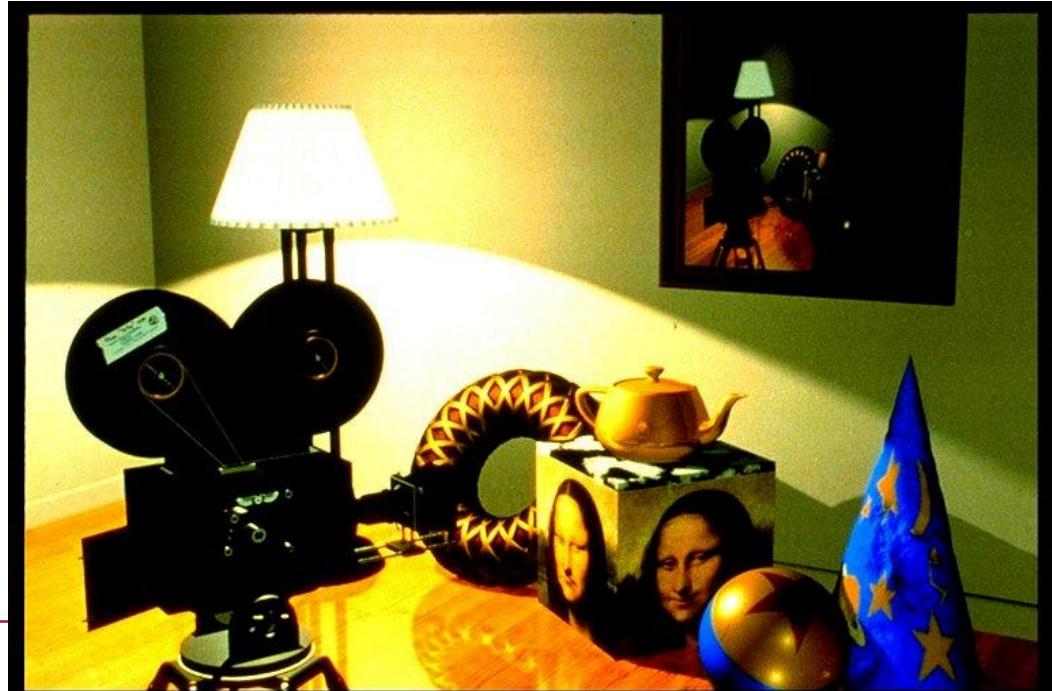


- Parallel light
- Specular reflection
- Multiple local light sources
- Different BRDFs

Object properties constant over surface

Texturing

- Locally varying object characteristics
- 2D Image Textures
- Shadows
- Bump-Mapping
- Reflection textures



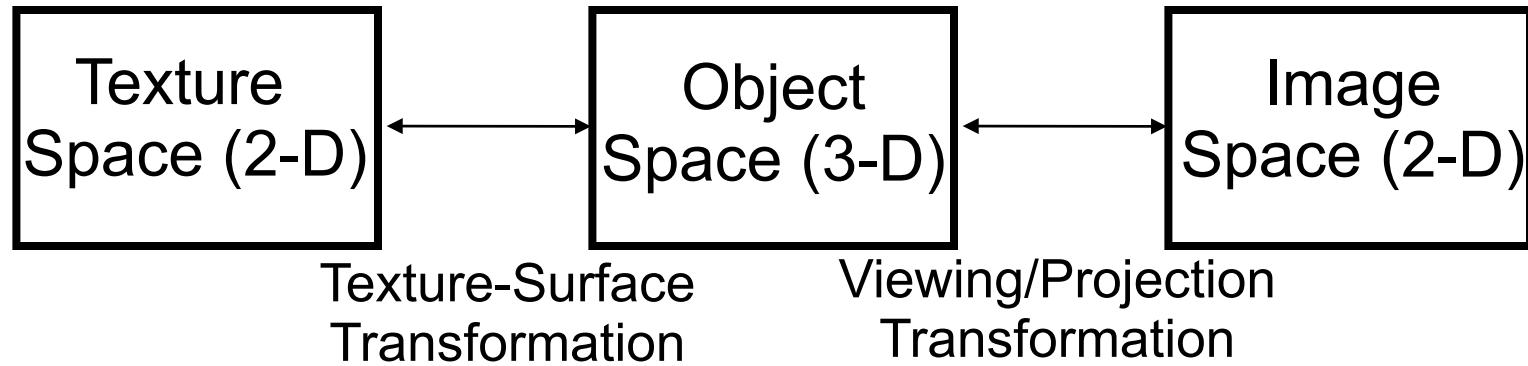


Texture-modulated Quantities

- Modulation of object surface properties
- Reflectance
 - Color (RGB), diffuse reflection coefficient k_d
 - Specular reflection coefficient k_s
- Opacity (α)
- Normal vector
 - $N(P) = N(P + t N)$ or $N = N + dN$
 - virtually moving point vs. changing normal
 - Bump mapping or Normal mapping
- Geometry
 - $P = P + dP$
 - Displacement mapping
- Distant illumination
 - Environment mapping, Reflection mapping



Texture Mapping Transformations

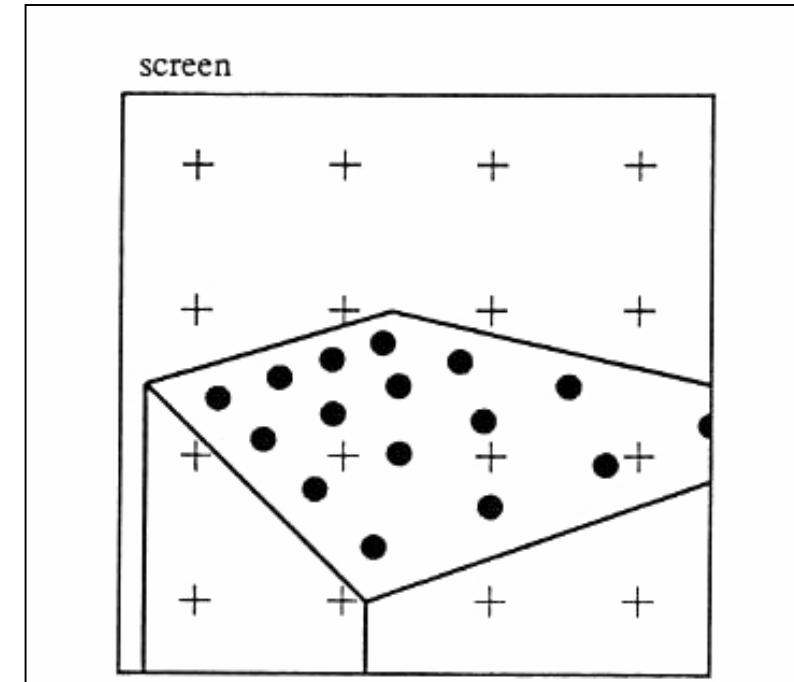
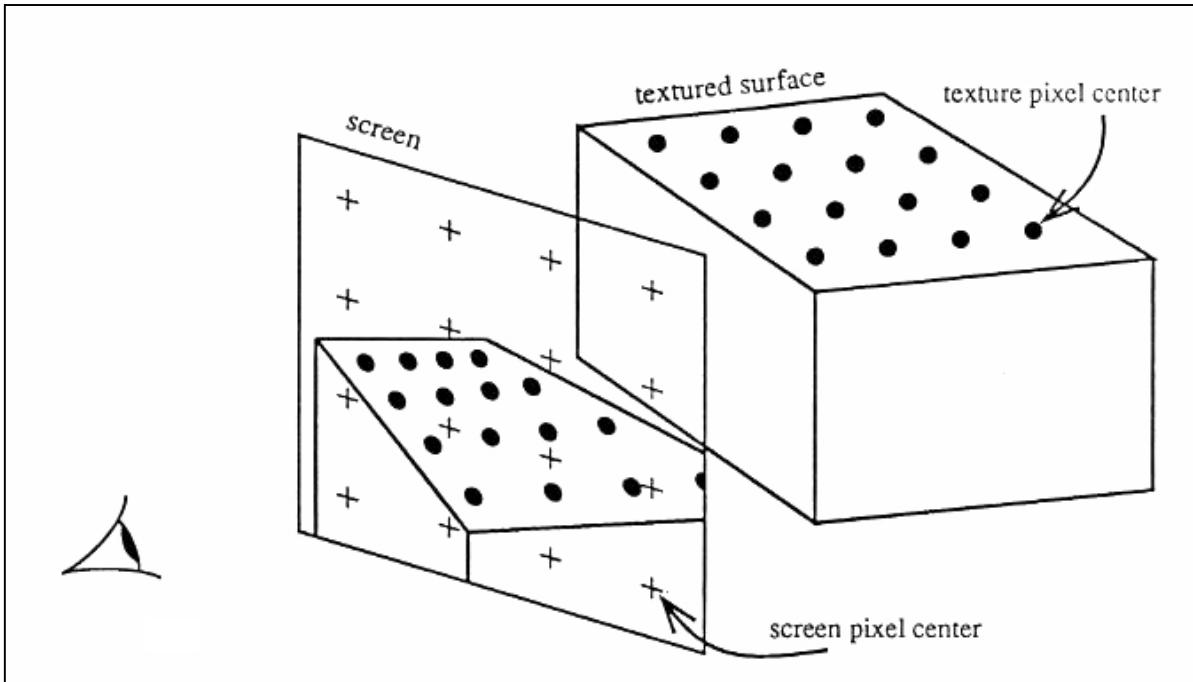


The texture is mapped onto a surface in 3-D object space, which is then mapped to the screen by the viewing projection. These two mappings are composed to find the overall 2-D texture space to 2-D image space mapping, and the intermediate 3-D space is often forgotten. This simplification suggests texture mapping's close ties with image warping and geometric distortion.

Texture space (u, v)
Object space (x_o, y_o, z_o)
Screen space (x, y)



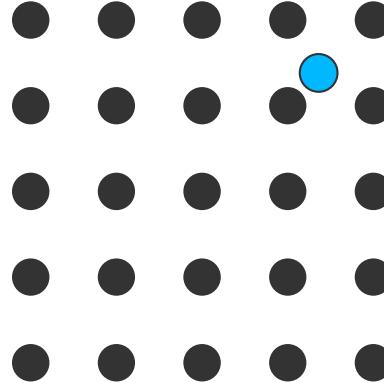
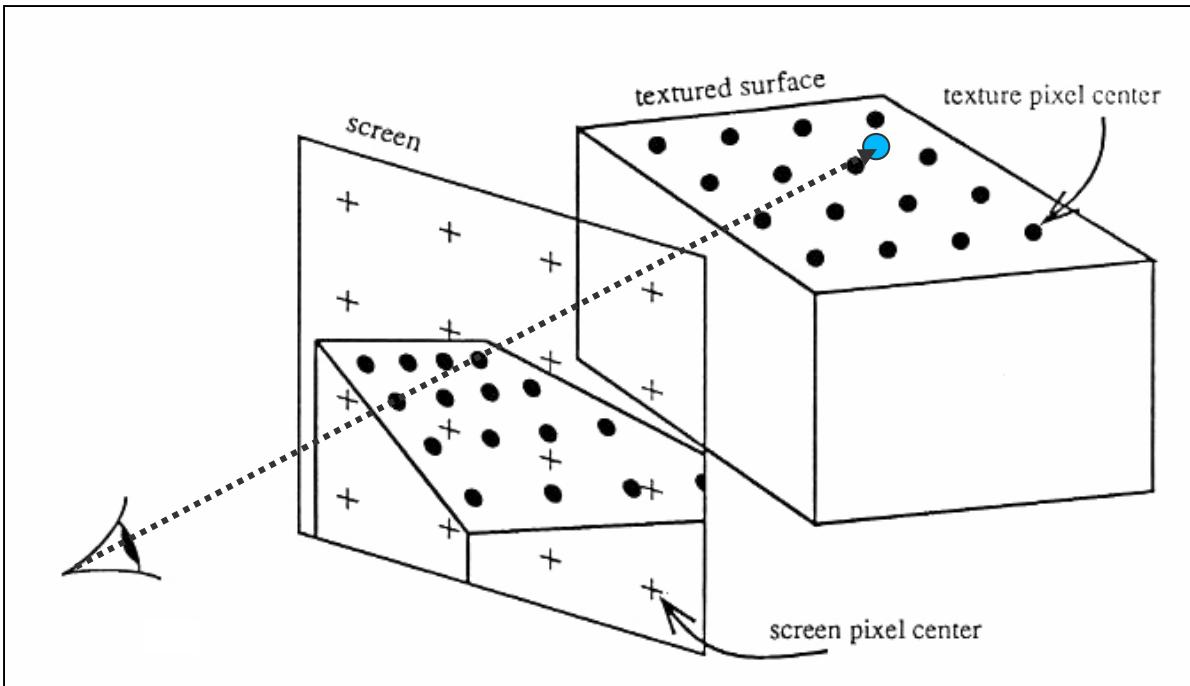
2D Texturing



- 2D texture mapped onto object
- Object projected onto 2D screen
- 2D → 2D: warping operation
- Uniform sampling ?
- Hole-filling/blending ?



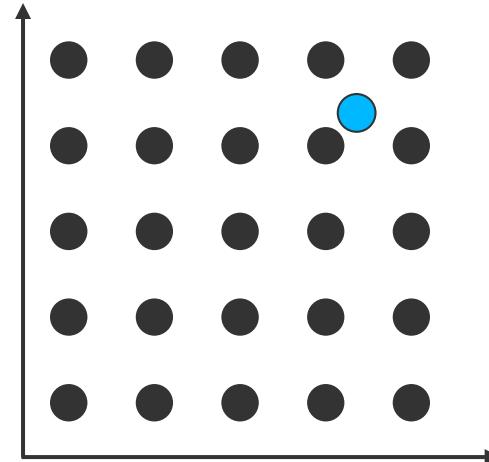
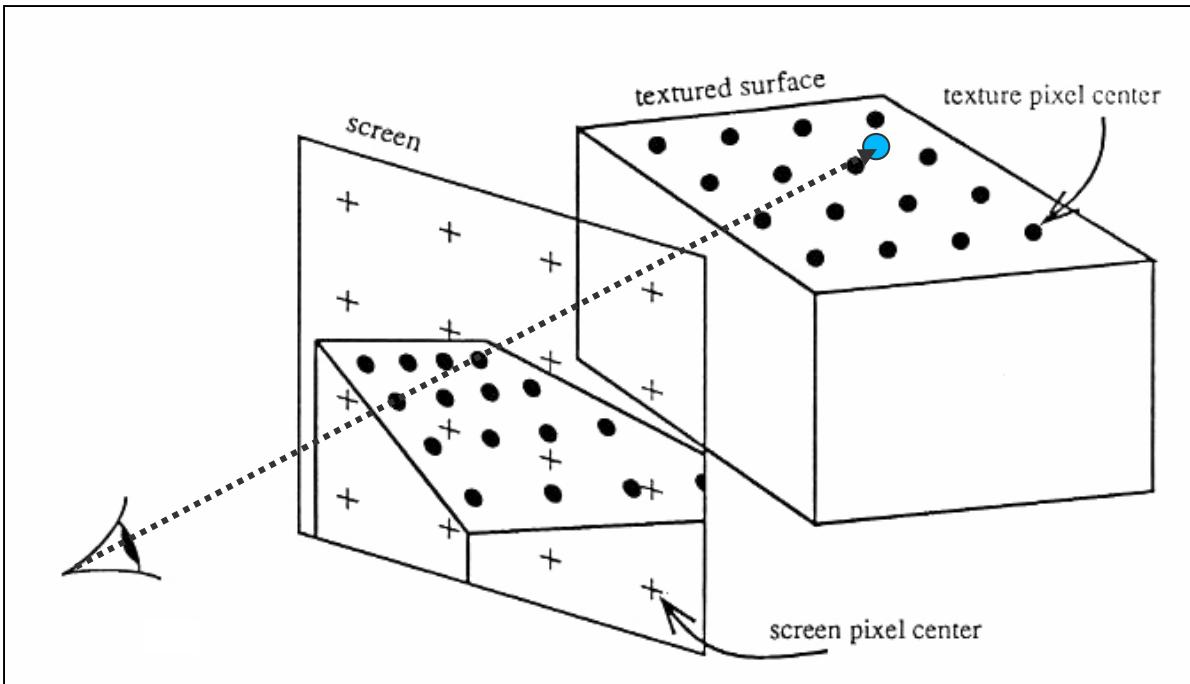
Texture Mapping in a Ray Tracer



- approximation:
 - ray hits surface
 - surface location corresponds to coordinate inside a texture



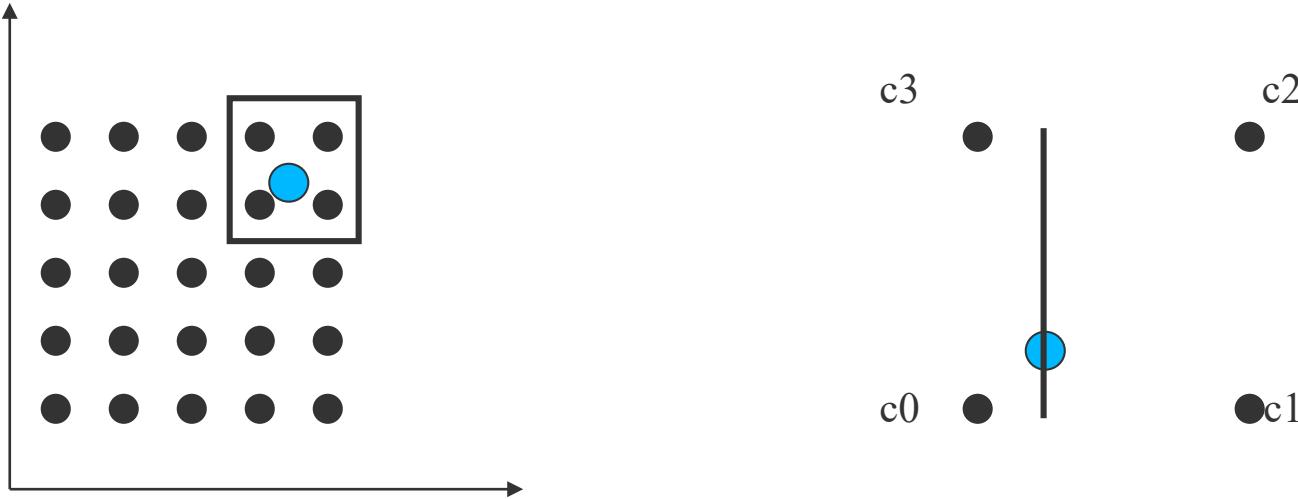
Texture Mapping in a Ray Tracer



- approximation:
 - ray hits surface
 - surface location corresponds to coordinate inside a texture



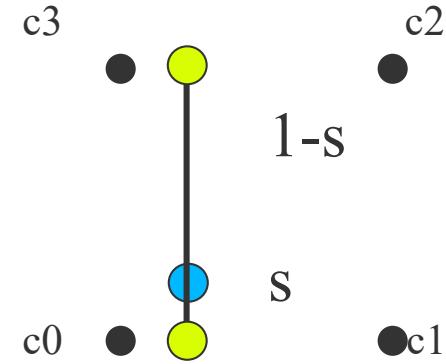
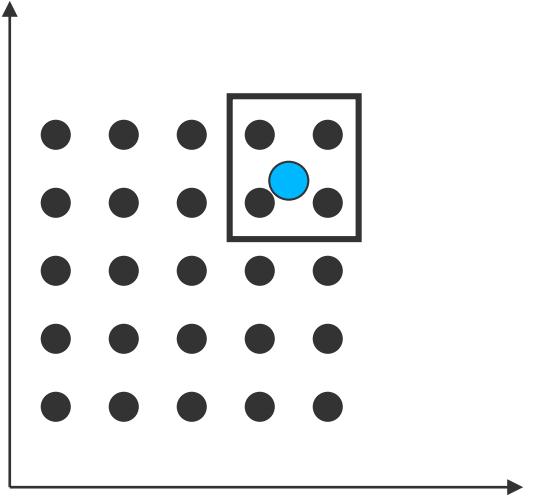
Interpolation 1D



- How to interpolate the color of the pixel?



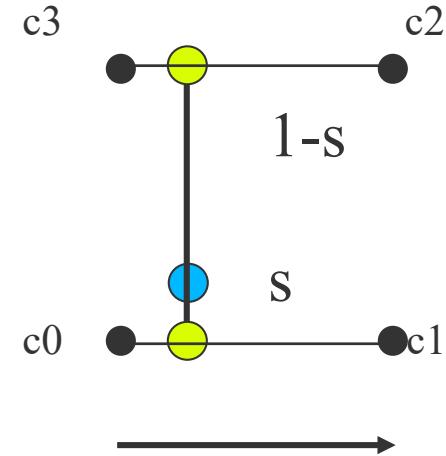
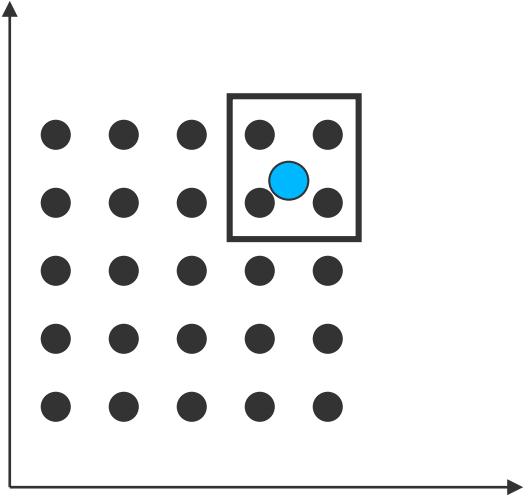
Interpolation 1D



- How to interpolate the color of the pixel?

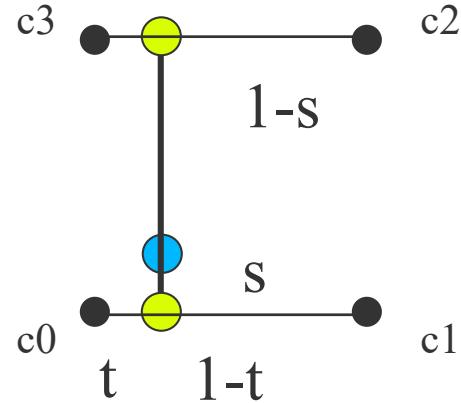
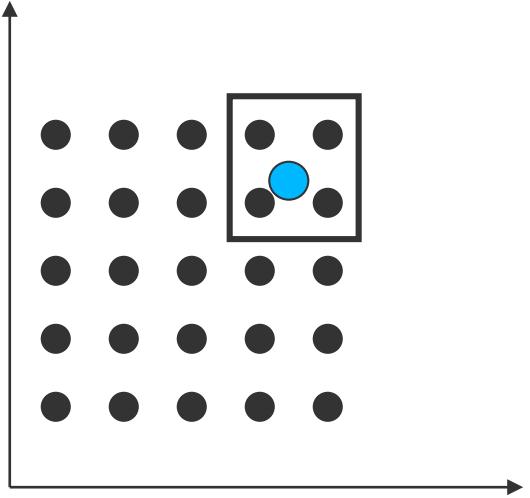


Interpolation 2D



- How to interpolate the color of the pixel?

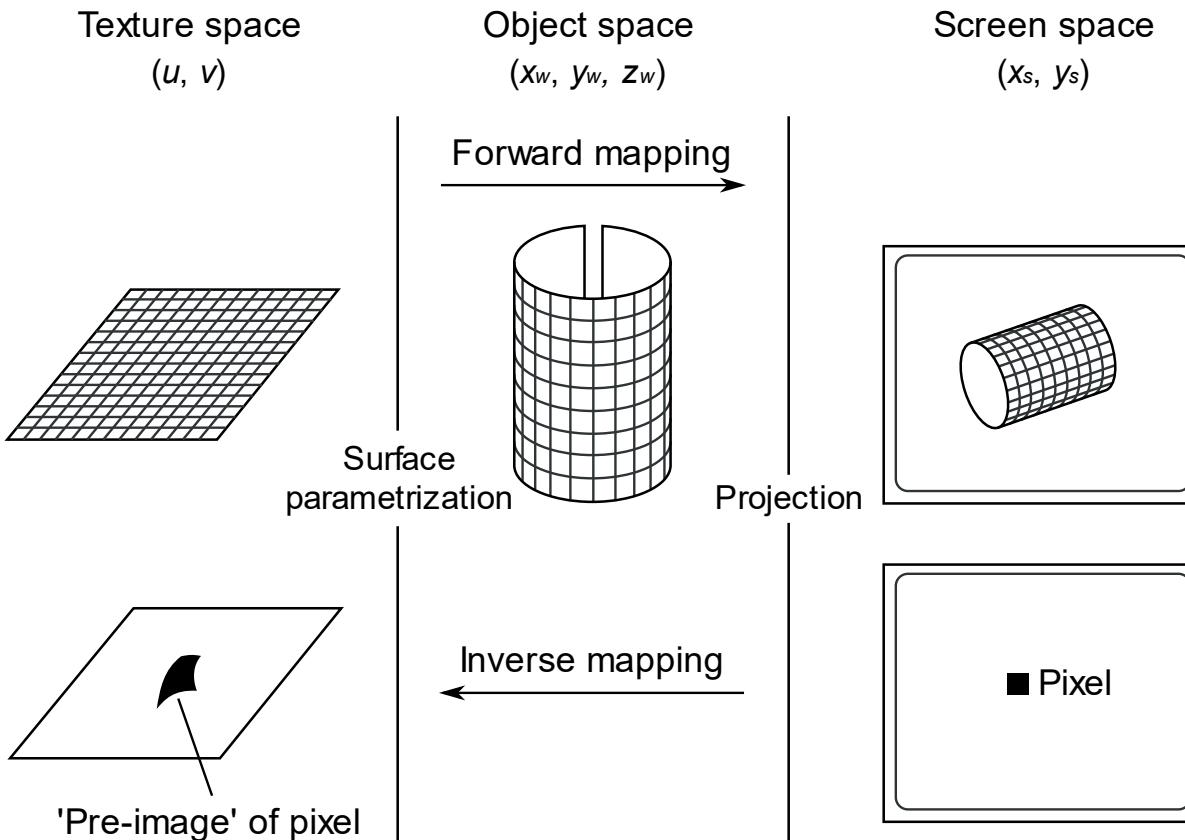
Interpolation 2D



- How to interpolate the color of the pixel?
- Lookup of four surrounding points
- $2 \times 1D:$ $i_0 = (1-t)c_0 + tc_1$
 $i_1 = (1-t)c_3 + tc_2$
- $1 \times 2D:$ $c = (1-s)i_0 + s i_1$



2D Texture Mapping



- Forward mapping
 - Object surface parameterization
 - Projective transformation
- Inverse mapping
 - Find corresponding pre-image/footprint of each pixel in texture
 - Integrate over pre-image

Forward Mapping

- Maps each texel to its position in the image
- Uniform sampling of texture space does not guarantee uniform sampling in screen space
- Possibly used if
 - The texture-to-screen mapping is difficult to invert
 - The texture image does not fit into memory

Texture scanning:

for v

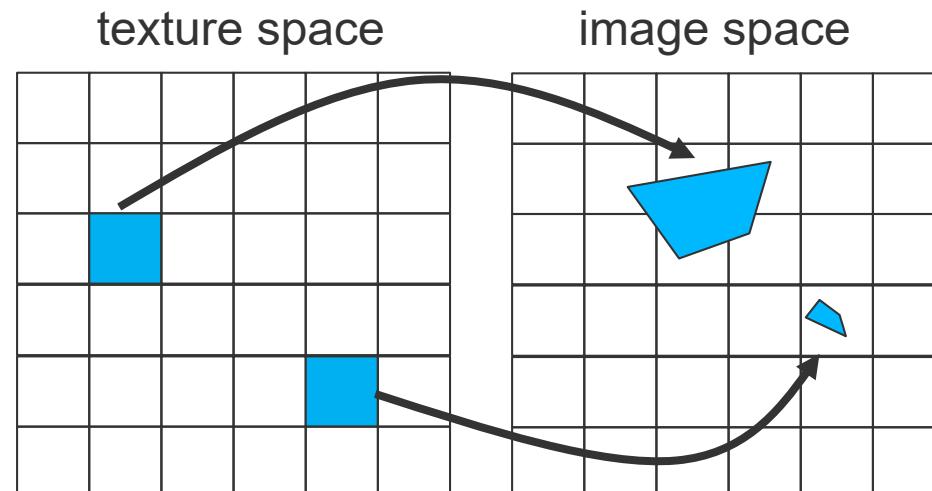
 for u

 compute $x(u,v)$ and $y(u,v)$

 copy $\text{TEX}[u,v]$ to $\text{SCR}[x,y]$

(or in general

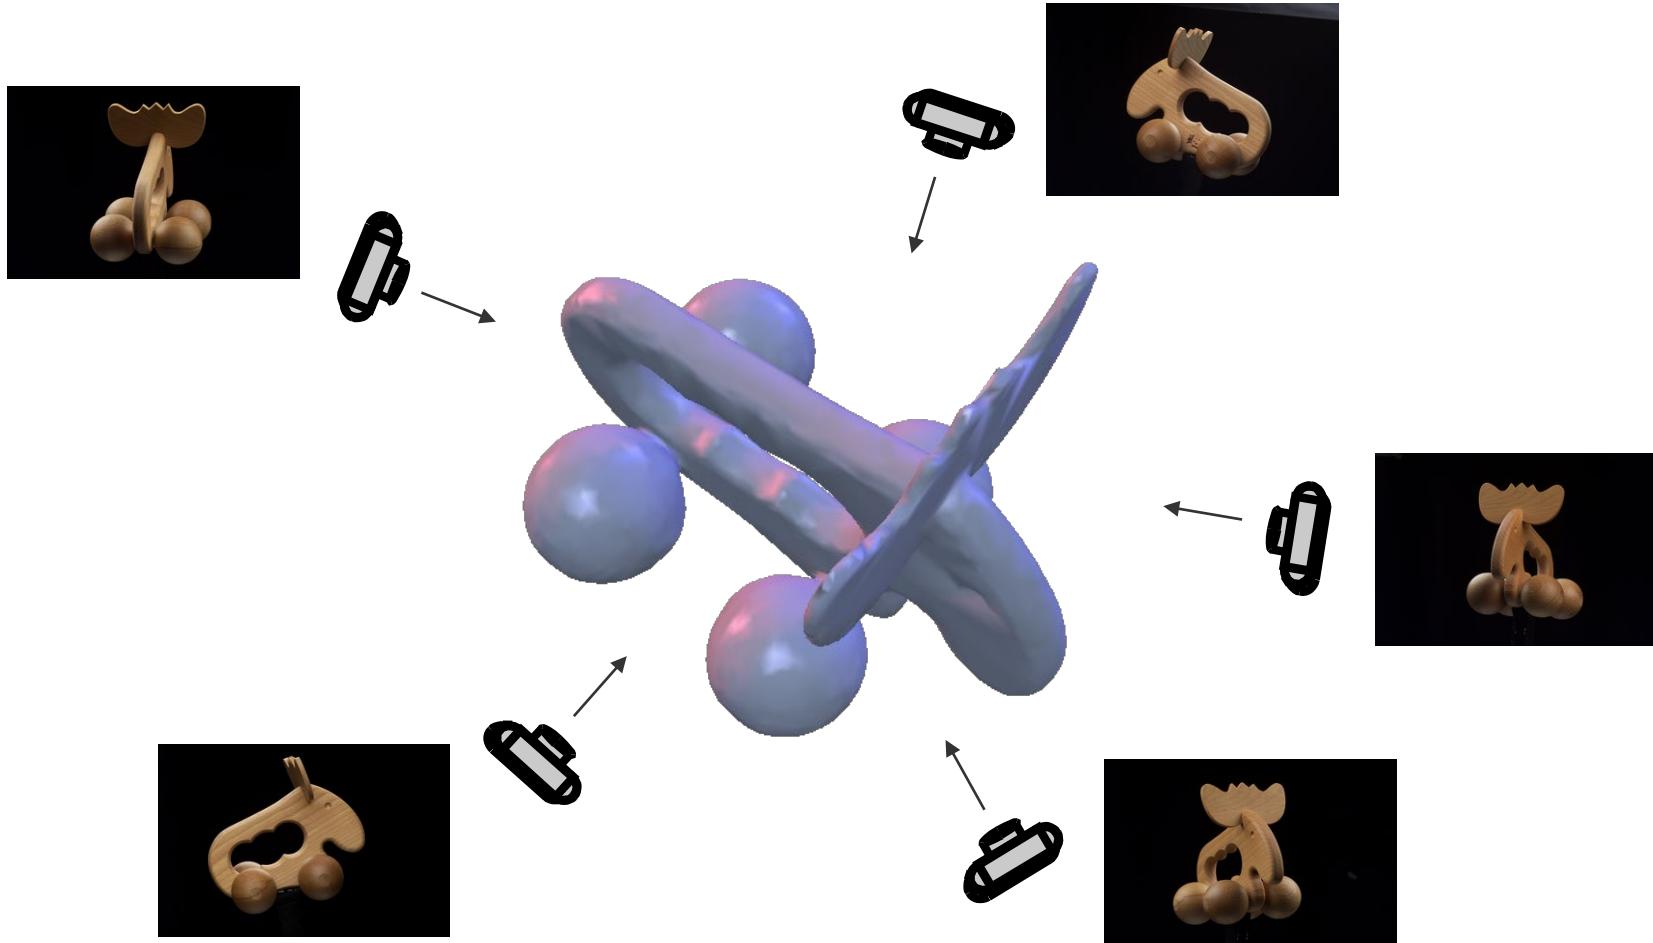
 rasterize image of $\text{TEX}[u,v]$)





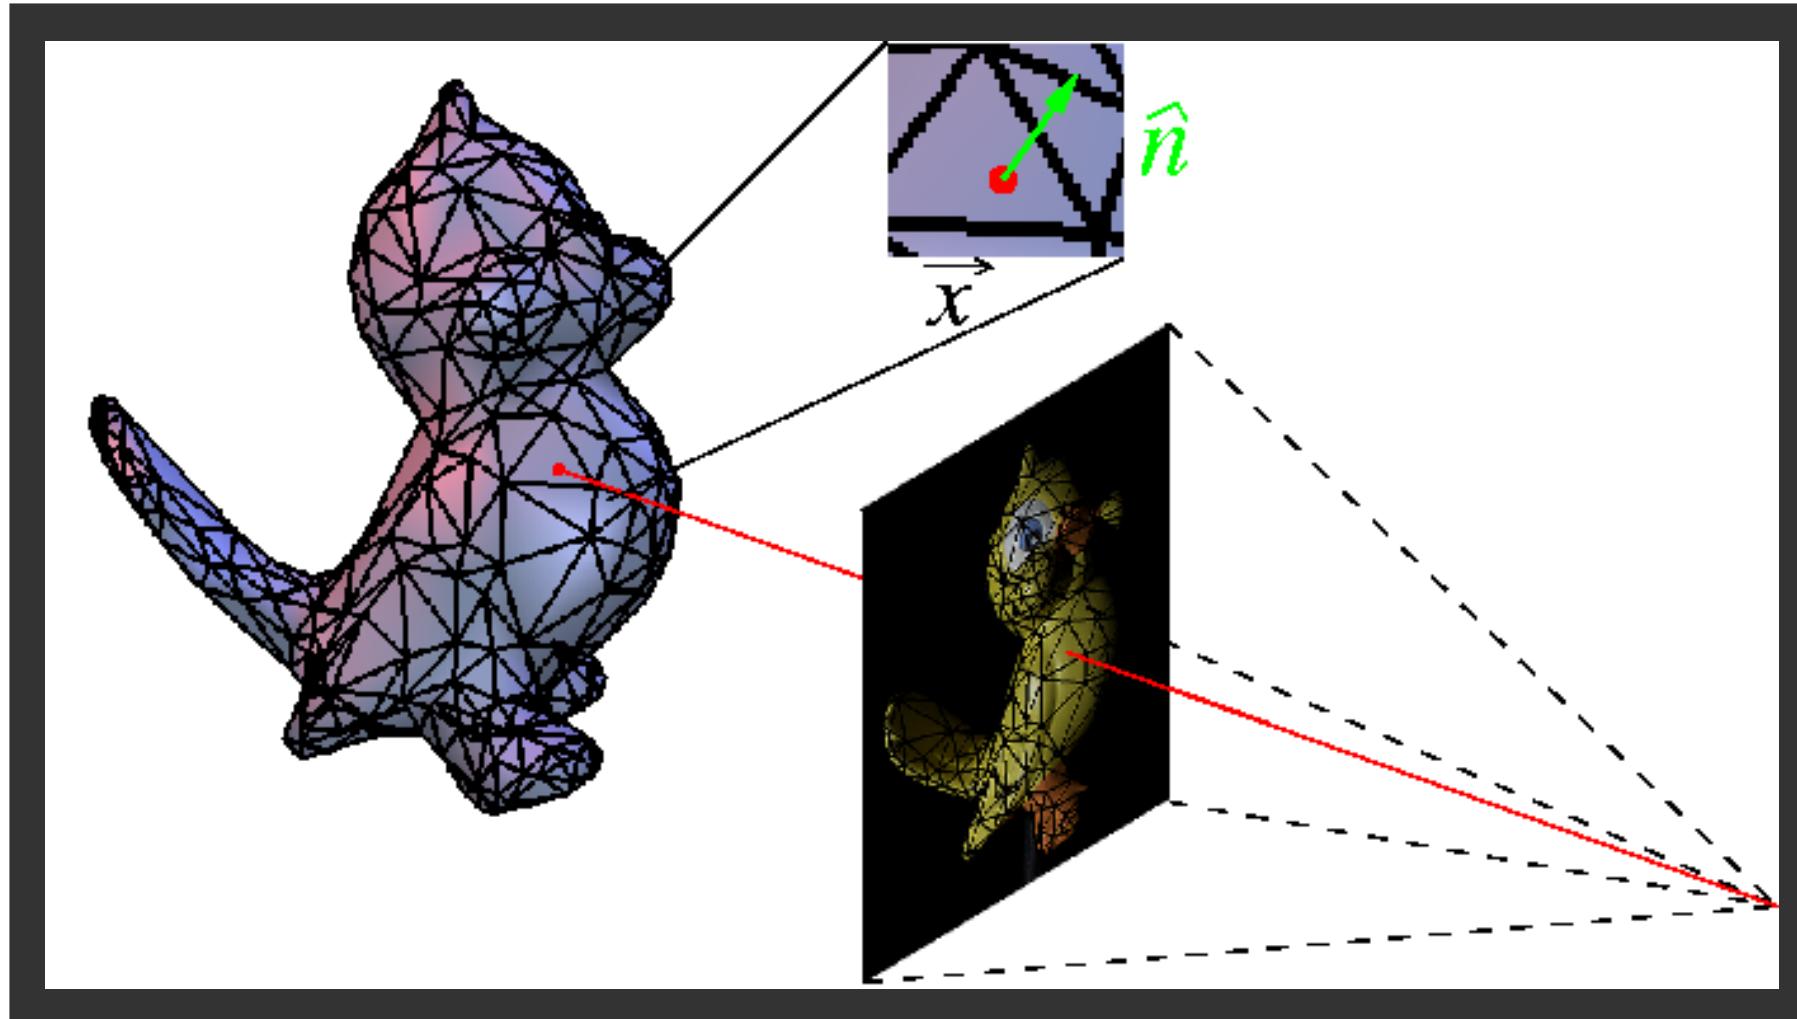
Texturing with real images

- Find the camera settings for each 2D image.





Resampling



- Image to surface mapping is given by projective mapping from camera position



Merged Texture





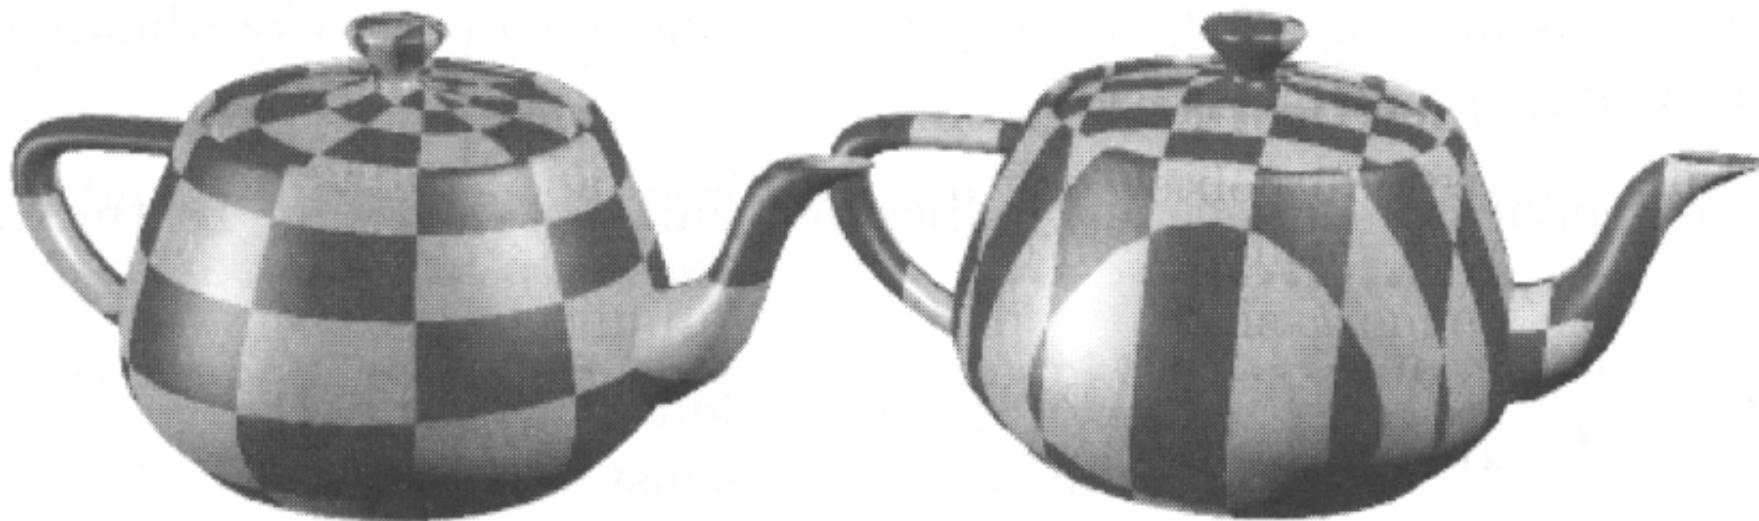
Texture Parameterization

How to map the texture image onto the geometry?



Surface Parameterization

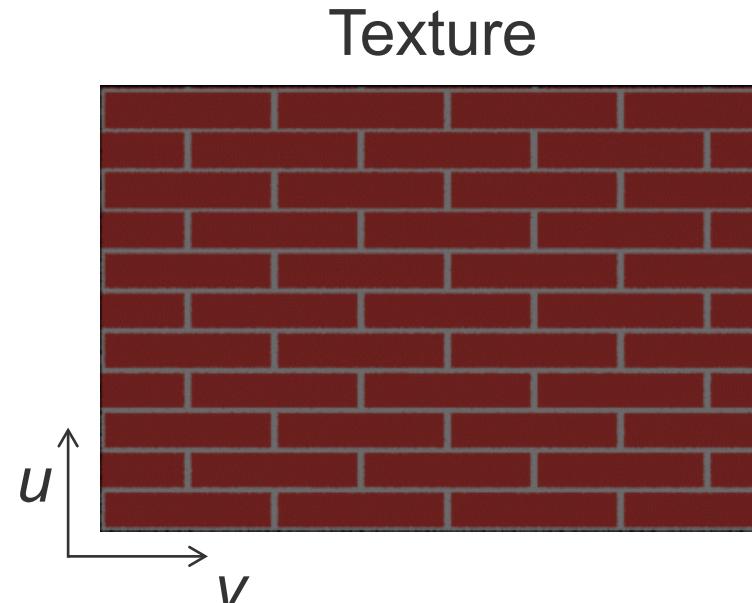
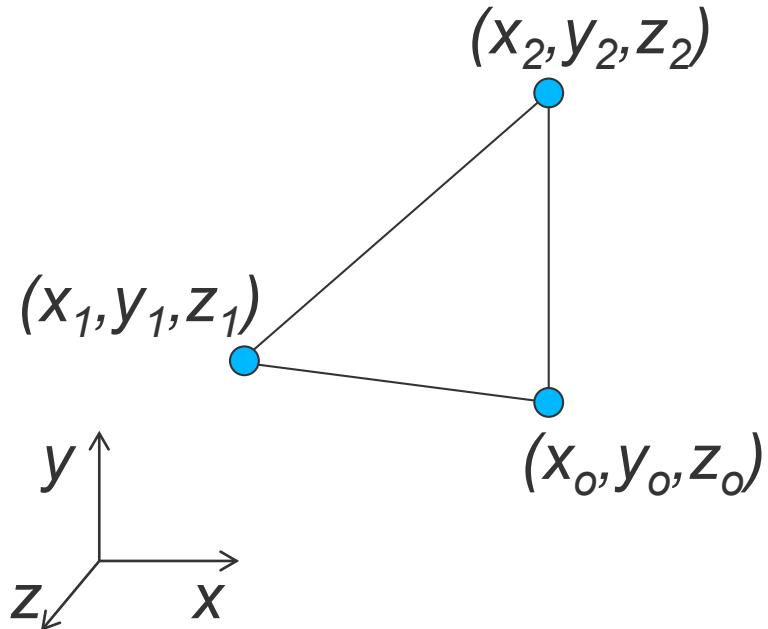
- To apply textures we need 2D coordinates on surfaces
→ Parameterization
- Some objects have a natural parameterization
 - Sphere: spherical coordinates $(\varphi, \theta) = (2\pi u, \pi v)$
 - Cylinder: cylindrical coordinates $(\varphi, z) = (2\pi u, H v)$
 - Parametric surfaces (such as B-spline or Bezier surfaces → later)
- Parameterization is less obvious for
 - Polygons, implicit surfaces, ...





Triangle Parameterization

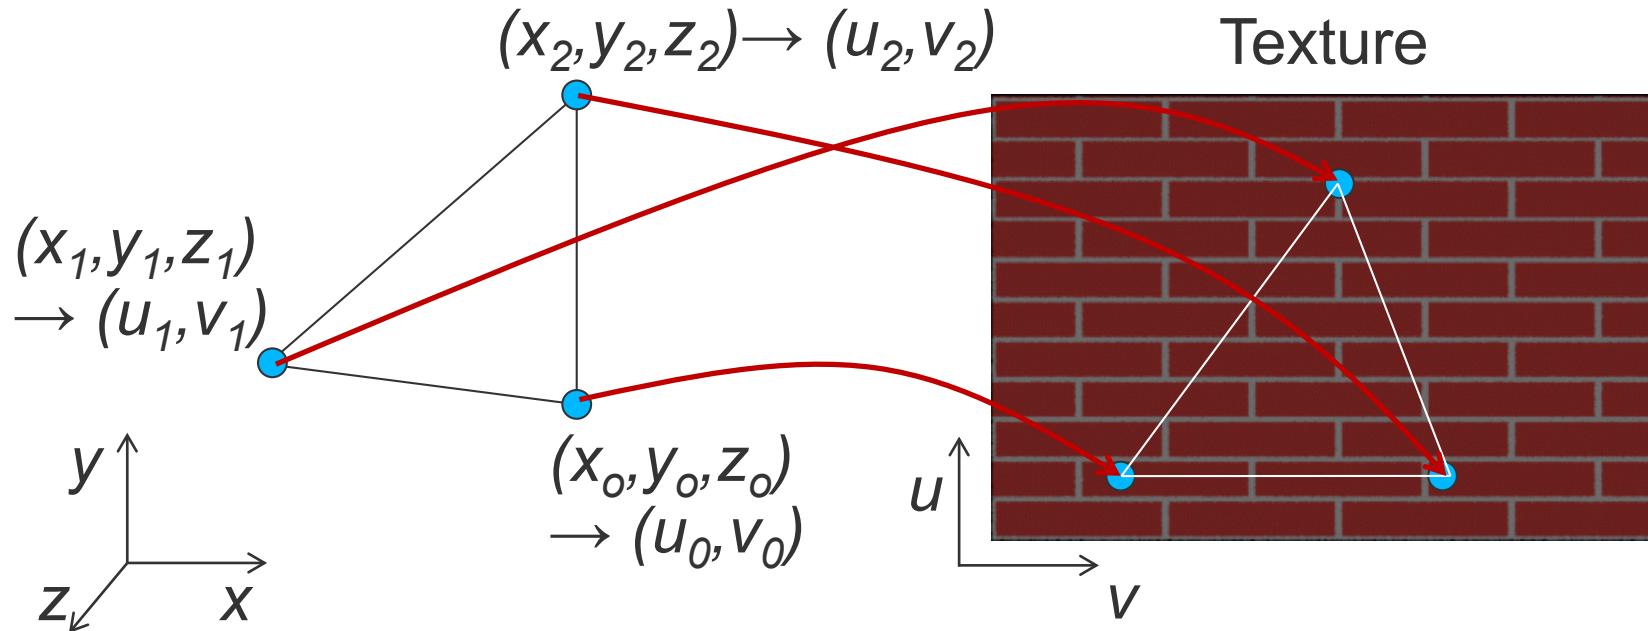
- Triangle is a planar object
 - Has implicit parameterization (e.g. barycentric coordinates)
 - But we need more control: Placement of triangle in texture space
- Assign texture coordinates (u_i, v_i) to each vertex $(x, y, z)_i$





Triangle Parameterization

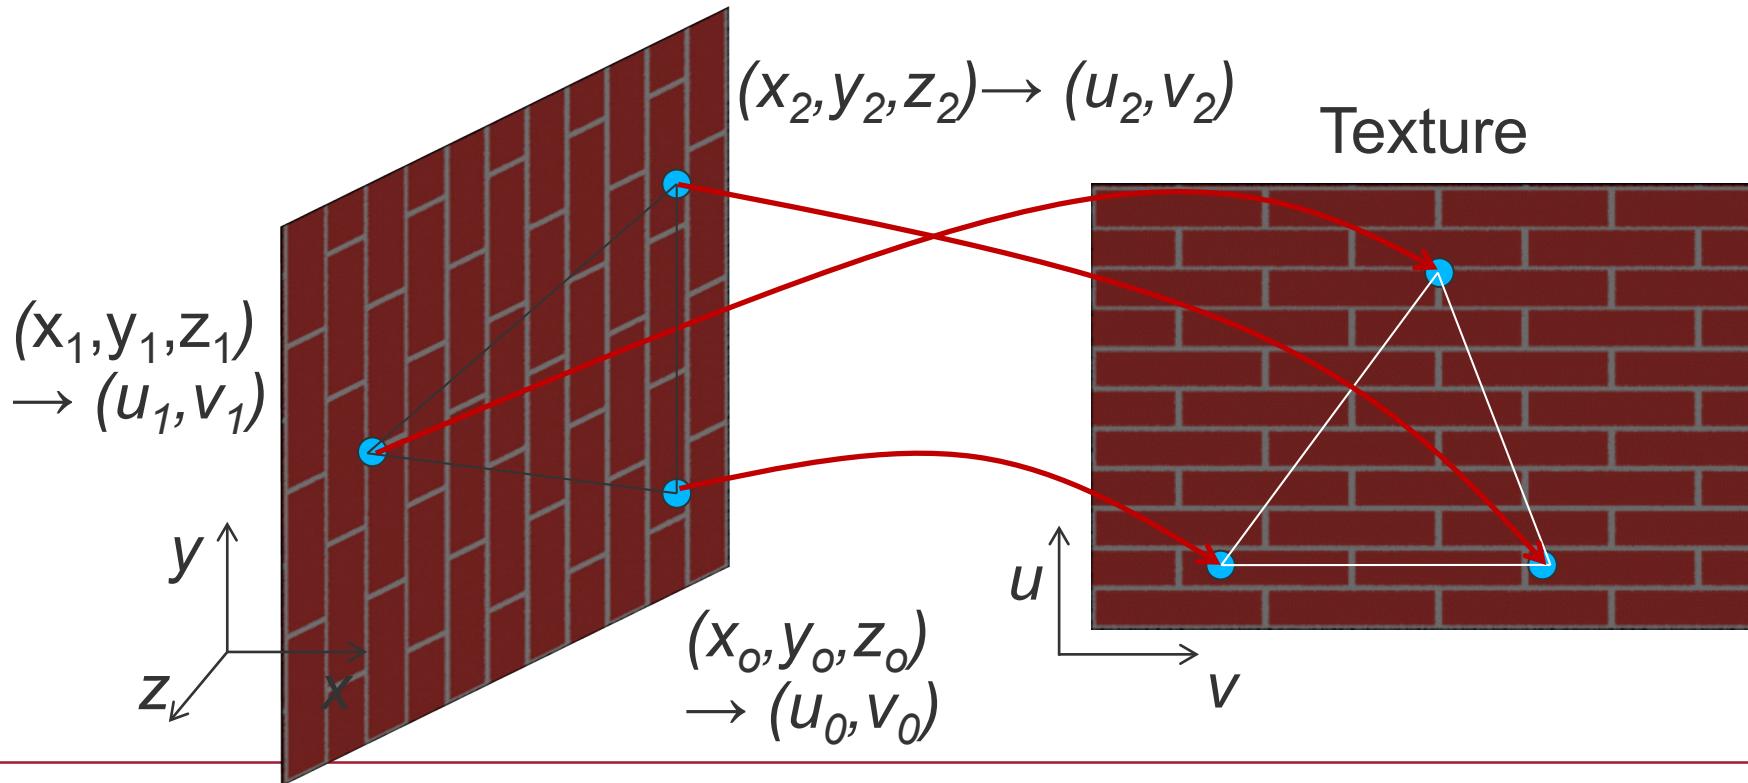
- Triangle is a planar object
 - Has implicit parameterization (e.g. barycentric coordinates)
 - But we need more control: Placement of triangle in texture space
- Assign texture coordinates (u_i, v_i) to each vertex $(x, y, z)_i$





Triangle Parameterization

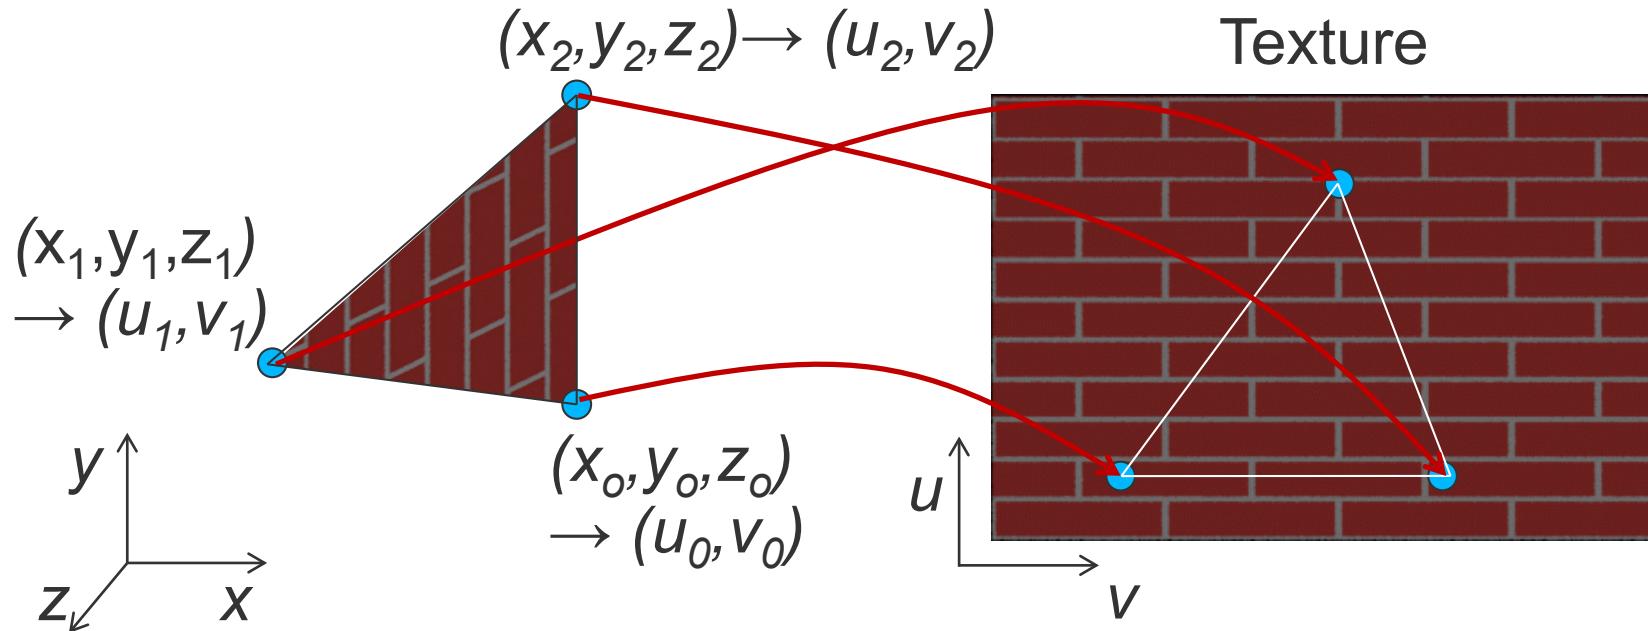
- Triangle is a planar object
 - Has implicit parameterization (e.g. barycentric coordinates)
 - But we need more control: Placement of triangle in texture space
- Assign texture coordinates (u_i, v_i) to each vertex $(x, y, z)_i$





Triangle Parameterization

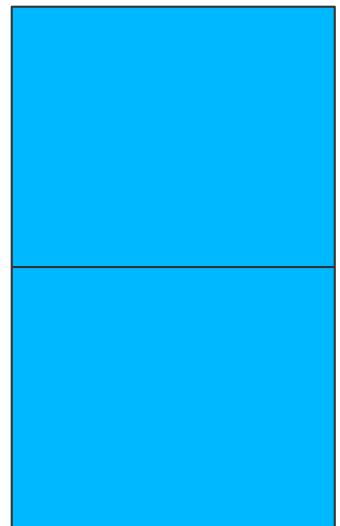
- Triangle is a planar object
 - Has implicit parameterization (e.g. barycentric coordinates)
 - But we need more control: Placement of triangle in texture space
- Assign texture coordinates (u_i, v_i) to each vertex $(x, y, z)_i$



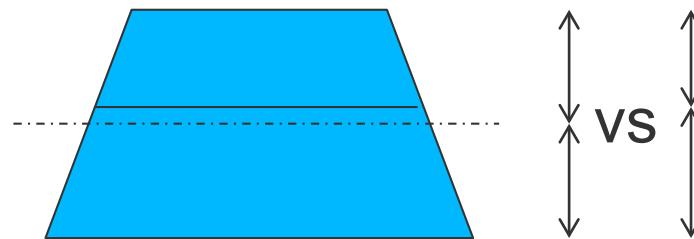


Influence of Perspective

- Due to perspective transformation the front patch is larger than the back patch. The center line of the texture is not at the center between top and bottom: $\sim 1/z$



orthographic



perspective projection

- Need to calculate texture coordinate in 3D, then project!
 - no problem for a ray tracer

Triangle Parameterization

- Direct Perspective mapping
- Assign texture coordinates (u,v) to each vertex (x,y,z)
- Apply viewing projection $(x,y,z) \rightarrow (x,y)$
- Yields full texture transformation (warping) $(u,v) \rightarrow (x,y)$

$$x = \frac{au + bv + c}{gu + hv + i} \quad y = \frac{du + ev + f}{gu + hv + i}$$

- In homogeneous coordinates (by embedding (u,v) as $(u',v',1)$)

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} u' \\ v' \\ q \end{bmatrix} \quad (x, y) = (x'/w, y'/w) \\ (u, v) = (v'/q, v'/q)$$

- Transformation coefficients determined by 3 pairs $(u,v) \rightarrow (x,y)$
 - Three linear equations
 - Invertible iff neither set of points is collinear

Triangle Parameterization II

- Given

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} u' \\ v' \\ q \end{bmatrix}$$

- the inverse transform $(x,y) \rightarrow (u,v)$ is defined as

$$\begin{bmatrix} u' \\ v' \\ q \end{bmatrix} = \begin{bmatrix} A & B & C \\ D & E & F \\ G & H & I \end{bmatrix} \begin{bmatrix} x' \\ y' \\ w \end{bmatrix} \quad \begin{bmatrix} u' \\ v' \\ q \end{bmatrix} = \begin{bmatrix} ei - fh & ch - bi & bf - ce \\ fg - di & ai - cg & cd - af \\ dh - eg & bg - ah & ae - bd \end{bmatrix} \begin{bmatrix} x' \\ y' \\ w \end{bmatrix}$$

- Important to account for perspective transformation!
- Coefficients must be calculated for each triangle
 - Rasterization
 - Incremental bilinear update of (u',v',q) in screen space
 - Using the partial derivatives of the linear function (i.e. constants)
 - vs. Ray tracing
 - Simply interpolate texture coordinates using barycentric coordinates of intersection point



Cylinder Parameterization

- Transformation from texture space to the cylinder parametric representation can be written as:

$$(\theta, h) = (2\pi u, vH)$$

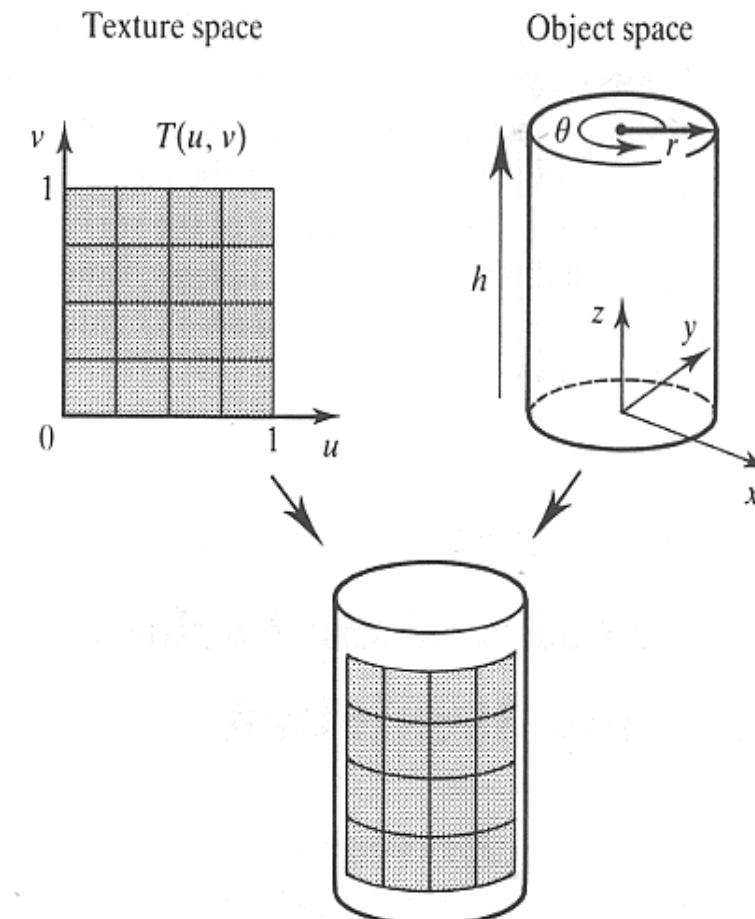
where H is the height of the cylinder.

- The surface coordinates in the Cartesian reference frame can be expressed as:

$$x_o = r \cos \theta$$

$$y_o = r \sin \theta$$

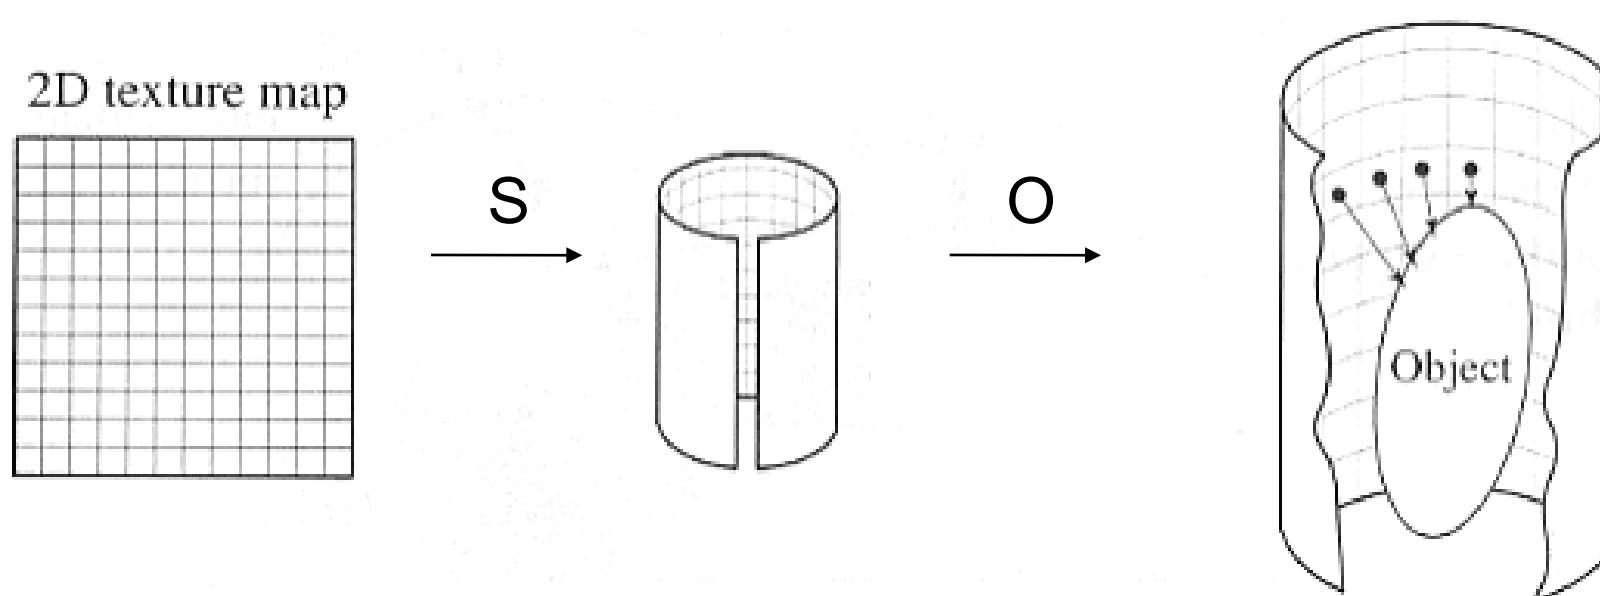
$$z_o = h$$





Two-Stage Mapping

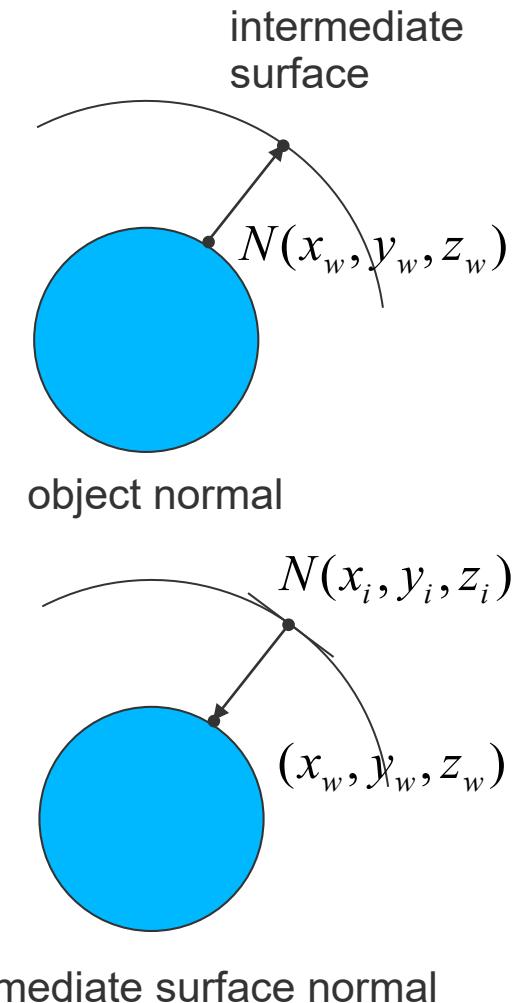
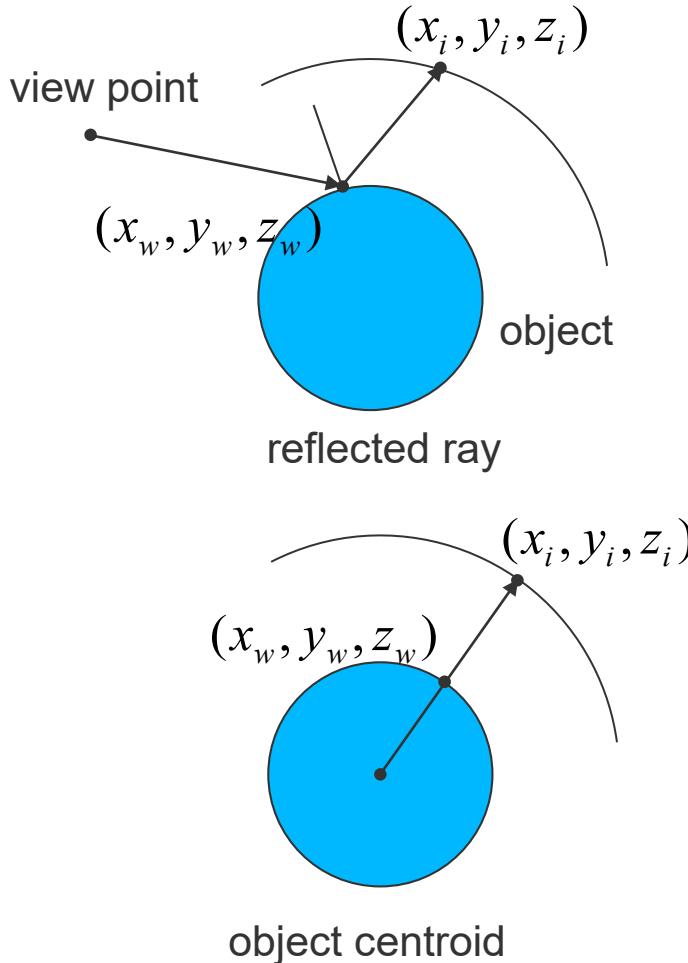
- Inverse Mapping for arbitrary 3D surfaces too complex
- Approximation technique is used:
 - Mapping from 2D texture space to a simple 3D intermediate surface (S mapping)
 - Should be a reasonable approximation of the destination surface
 - E.g.: plane, cylinder, sphere, cube, ...
 - Mapping from the intermediate surface to the destination object surface (O mapping)



O-Mapping

- Determine point on intermediate surface through

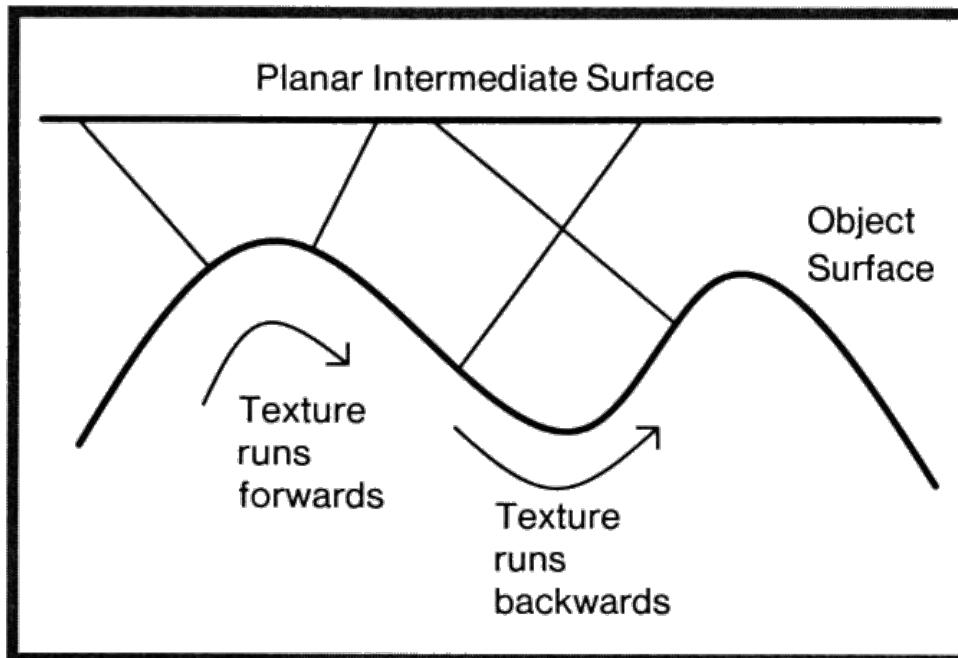
- Reflected view ray
 - Reflection or environment mapping
- Normal mapping
- Line through object centroid
- Shrinkwrapping
 - Forward mapping
 - Normal mapping from intermediate surface





Two-Stage Mapping: Problems

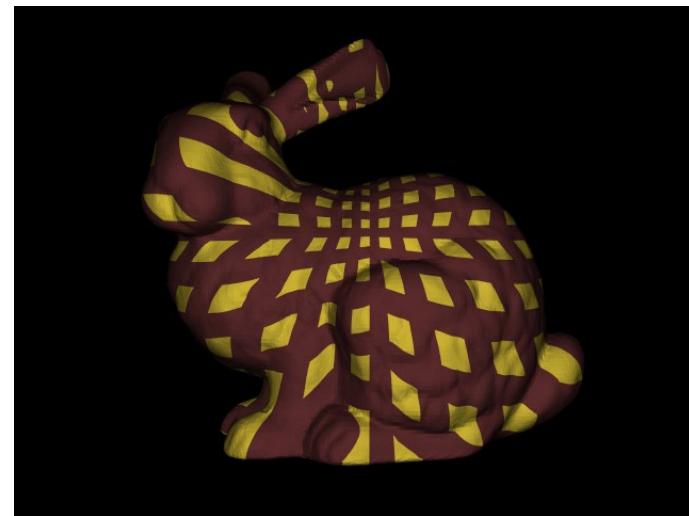
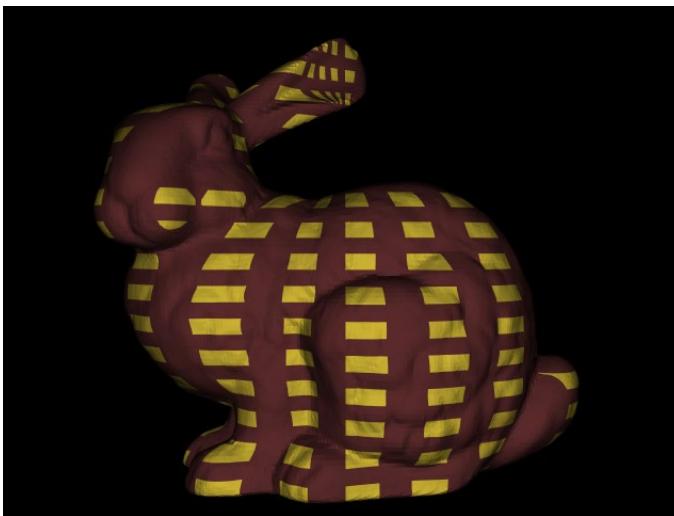
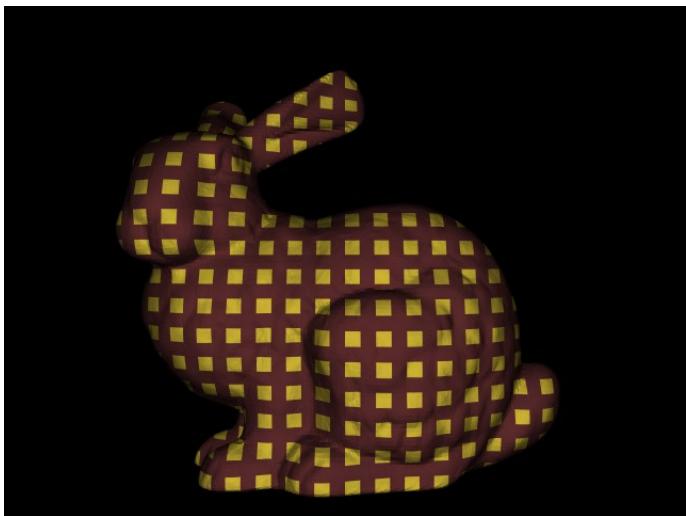
- Problems
 - May introduce undesired texture distortions if the intermediate surface differs too much from the destination surface
 - Still often used in practice because of its simplicity



Surface concavities can cause the texture pattern to reverse if the object normal mapping is used.



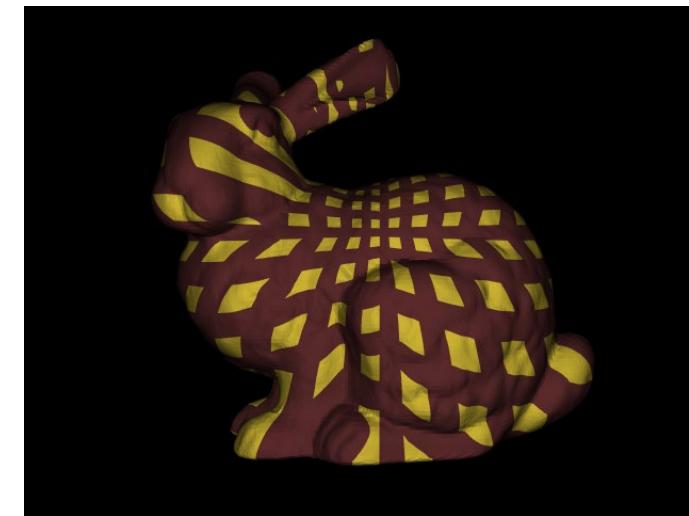
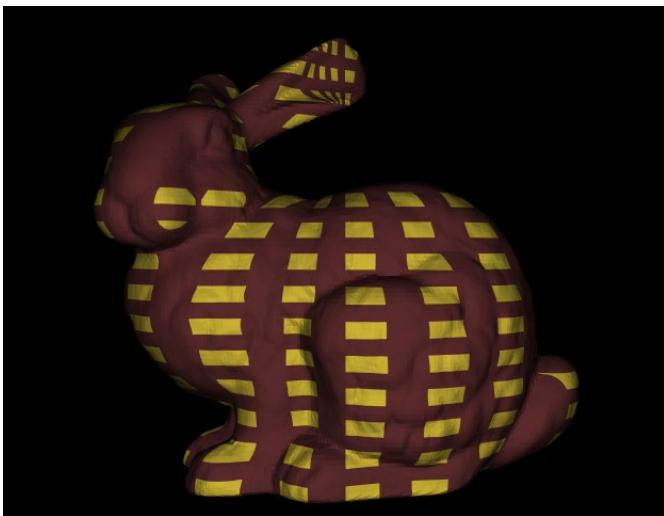
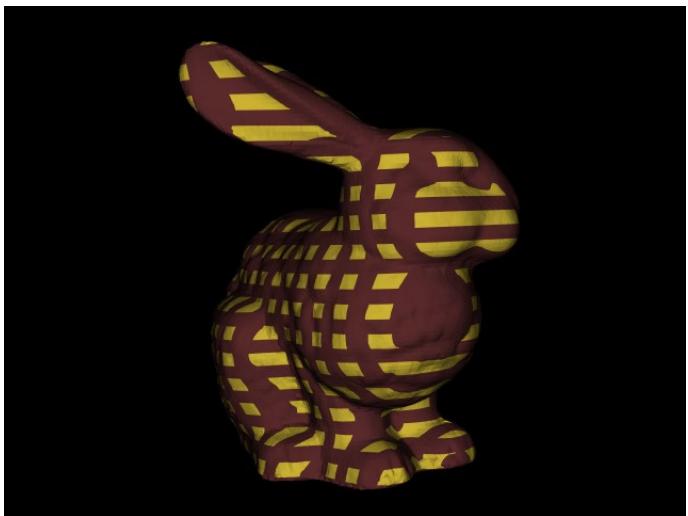
Two-Stage Mapping: Example



- Different intermediate surfaces
- Plane
 - Strong distortion where object surface normal \neq plane normal
- Cylinder
 - Reasonably uniform mapping (symmetry !)
- Sphere
 - Problems with concave regions



Two-Stage Mapping: Example

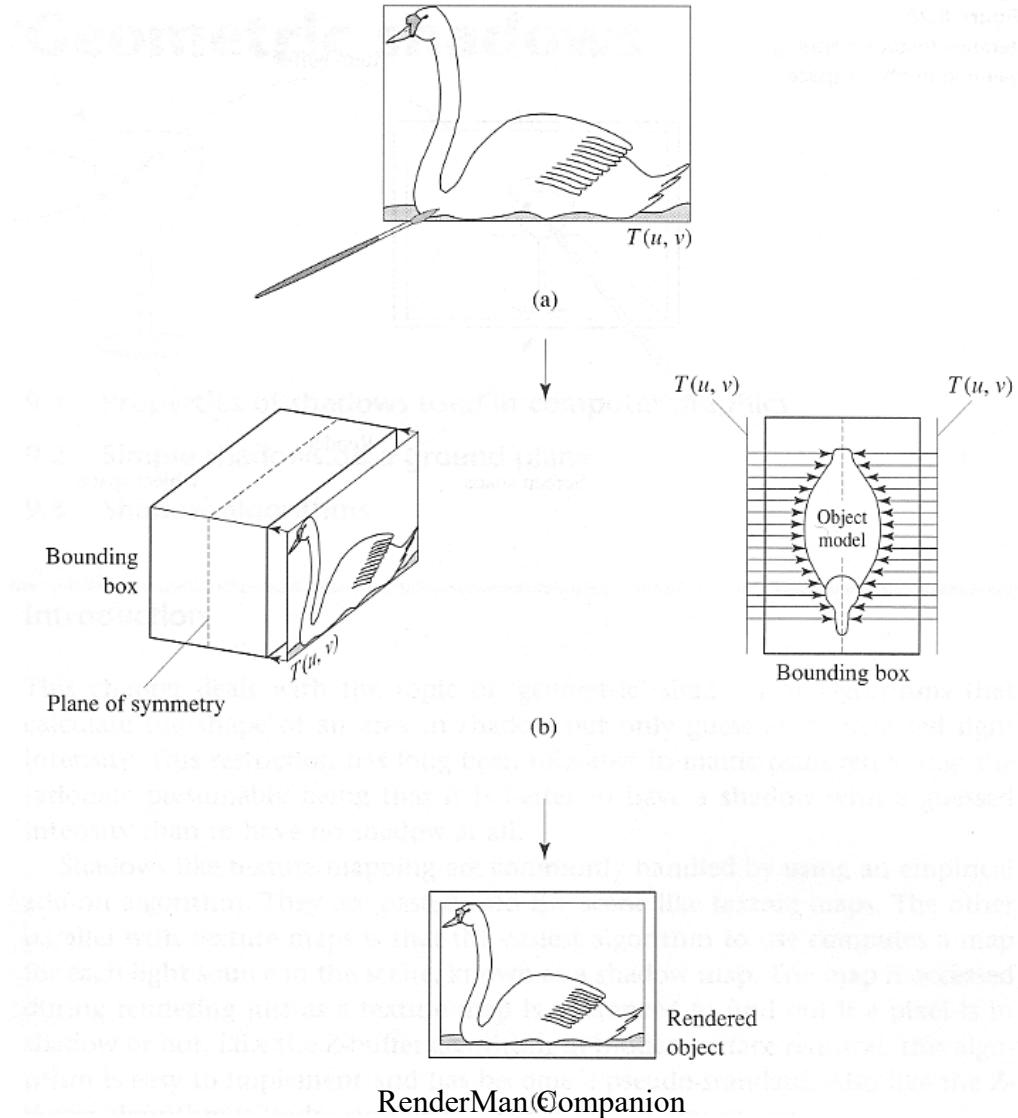


- Different intermediate surfaces
- Plane
 - Strong distortion where object surface normal \neq plane normal
- Cylinder
 - Reasonably uniform mapping (symmetry !)
- Sphere
 - Problems with concave regions



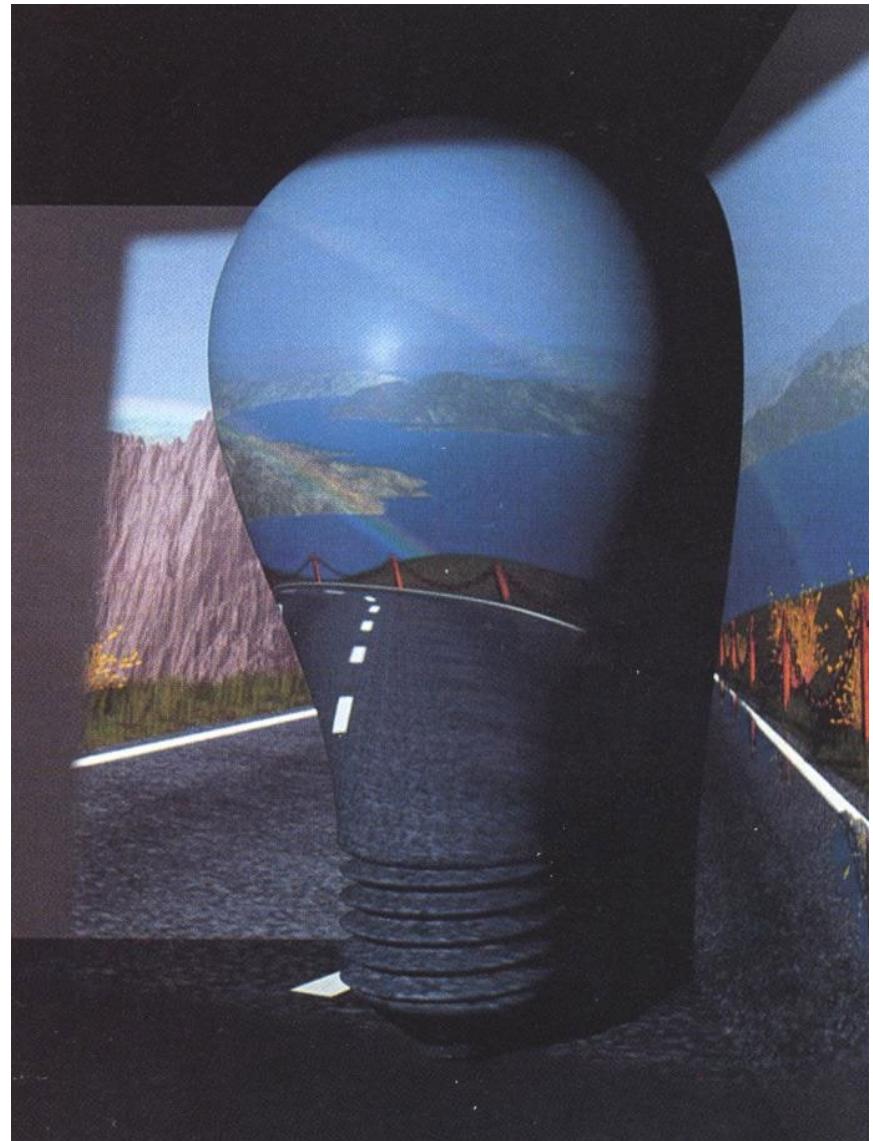
Projective Textures

- Project texture onto object surfaces
 - Slide projector
- Parallel or perspective projection
- Use photographs as textures
- Multiple images
- View-dependent texturing
- Perspective Mapping





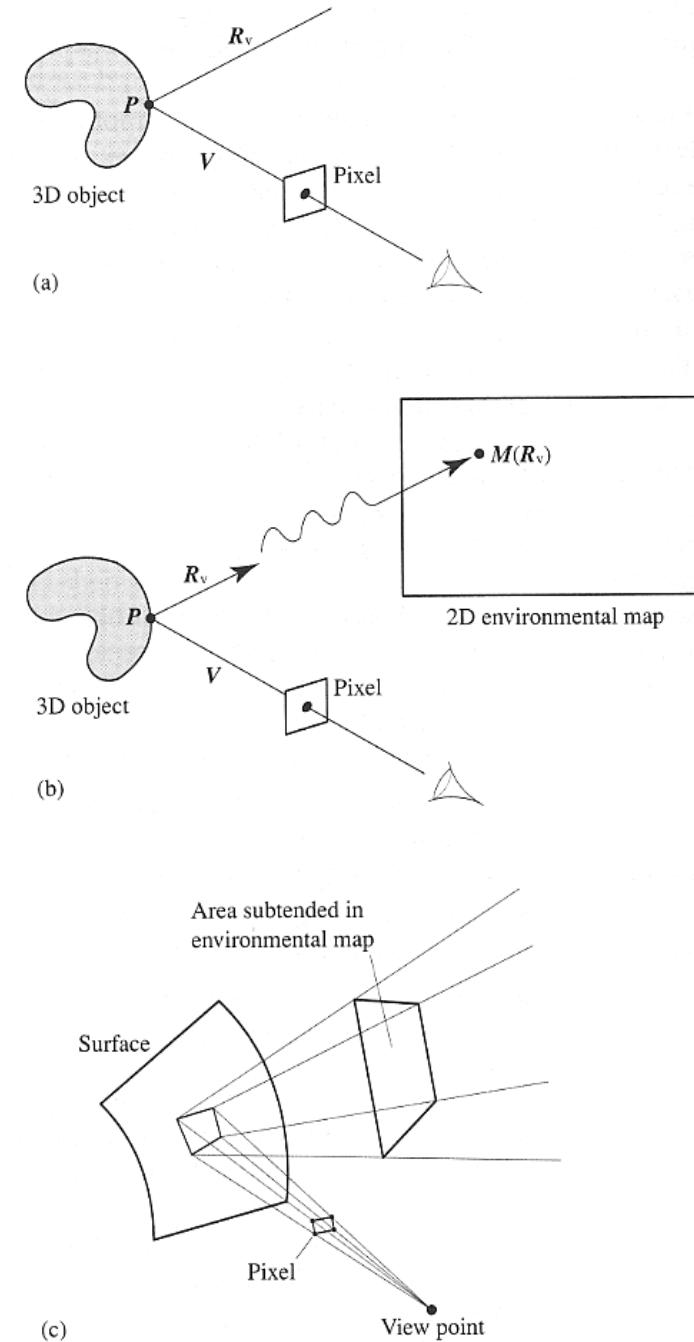
Projective Texturing: Examples





Reflection Mapping

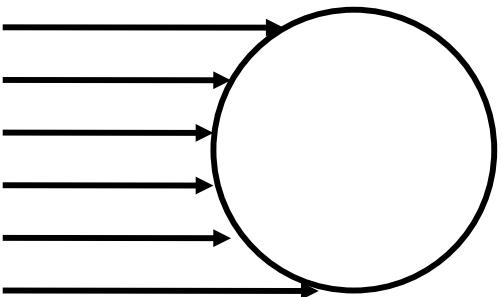
- Also called Environment Mapping
- Mirror reflections
 - Surface curvature: beam tracing
 - Map filtering
- Reflection map parameterization
 - Intermediate surface in 2-stage mapping
 - Often cube, sphere, or double paraboloid
- Assumption: Distant illumination
 - Parallax-free illumination
 - No self-reflections, distortion of near objects
- Option: Separate map per object
 - Often necessary to be reasonable accurate
 - Reflections of other objects
 - Maps must be recomputed after changes





Reflection Map Acquisition

- Generating spherical maps (original 1982/83)
 - i.e. photo of a reflecting sphere (gazing ball)

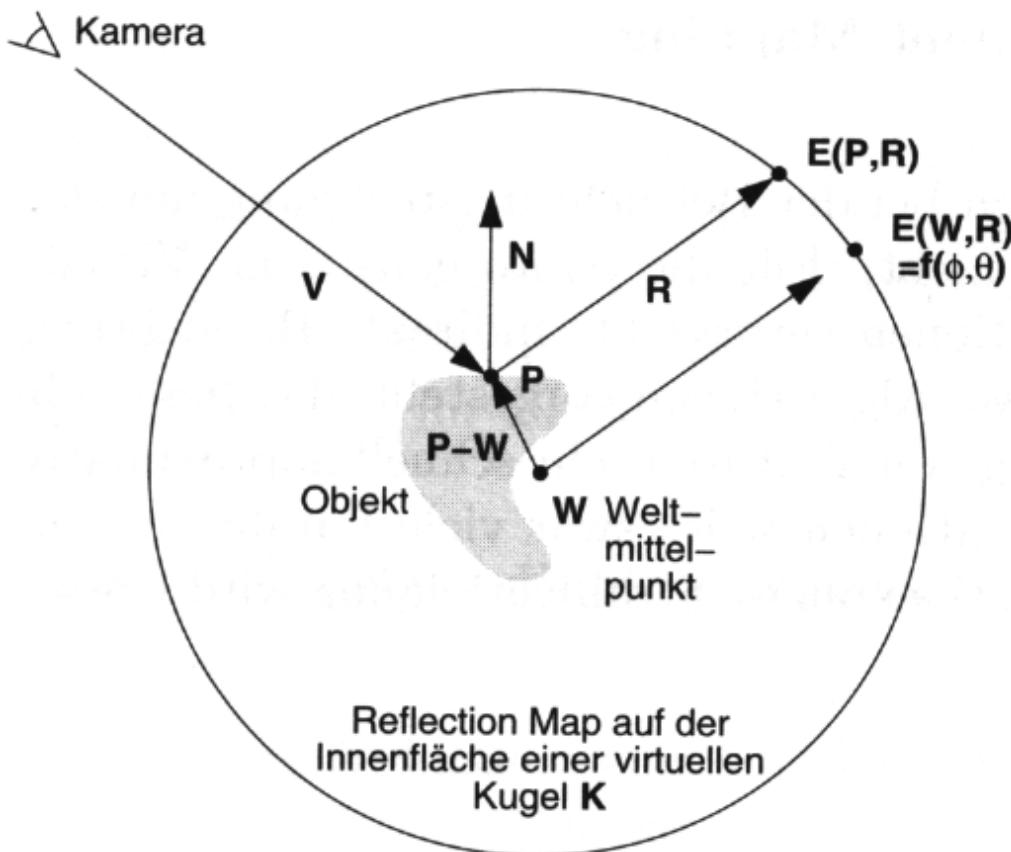


Peter Chou



Reflection Map Rendering

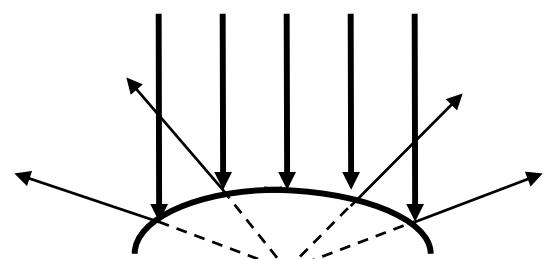
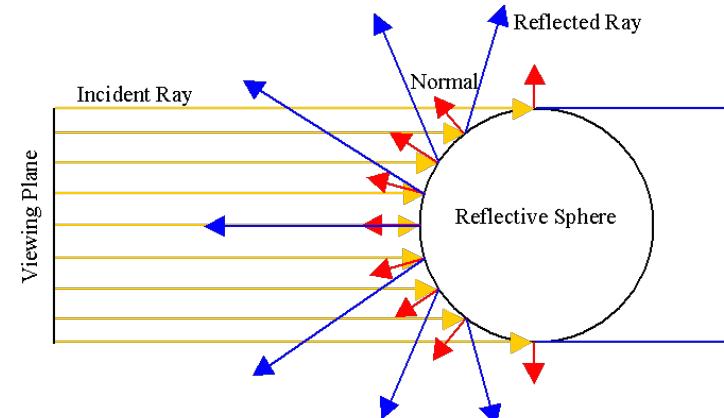
- Spherical parameterization
- O-mapping using reflected view ray intersection





Reflection Map Parameterization

- Spherical mapping
 - Single image
 - Bad utilization of the image area
 - Bad scanning on the edge
 - Artifacts, if map and image do not have the same direction
- Double parabolic mapping
 - Subdivide in 2 images (facing and back facing side)
 - Less bias on the edge
 - Arbitrarily reusable
 - Supported by OpenGL extensions

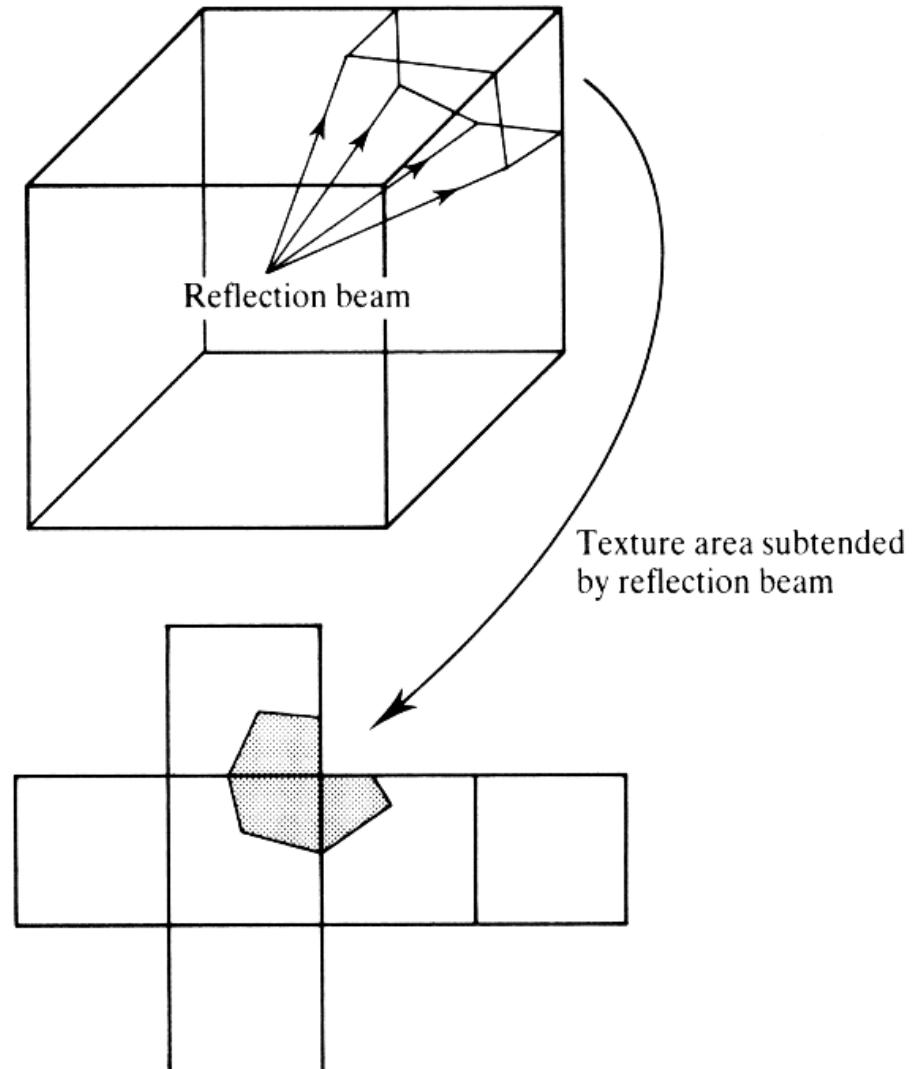
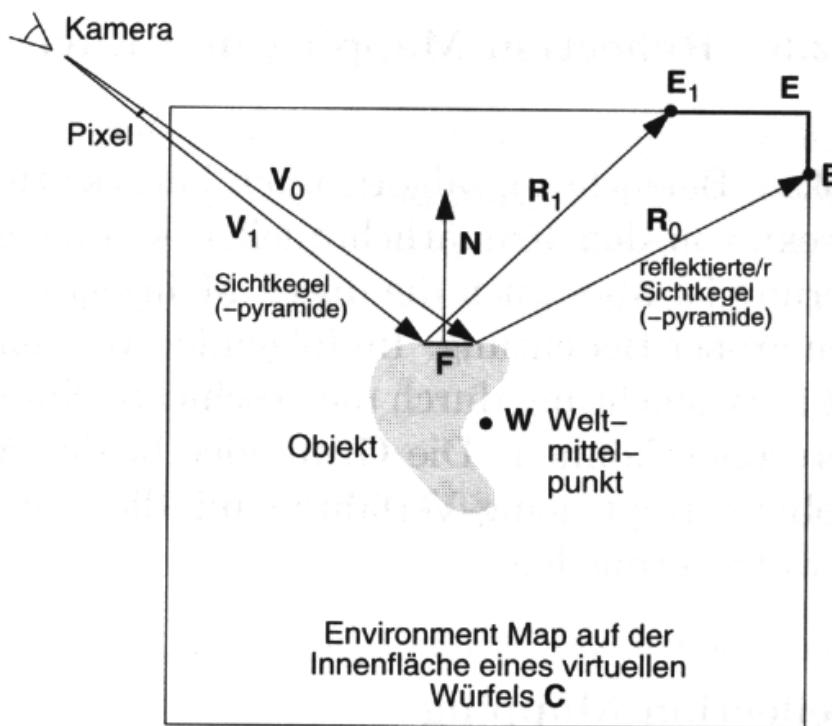


$$f(x, y) = \frac{1}{2} - \frac{1}{2}(x^2 + y^2)$$



Reflection Map Parameterization

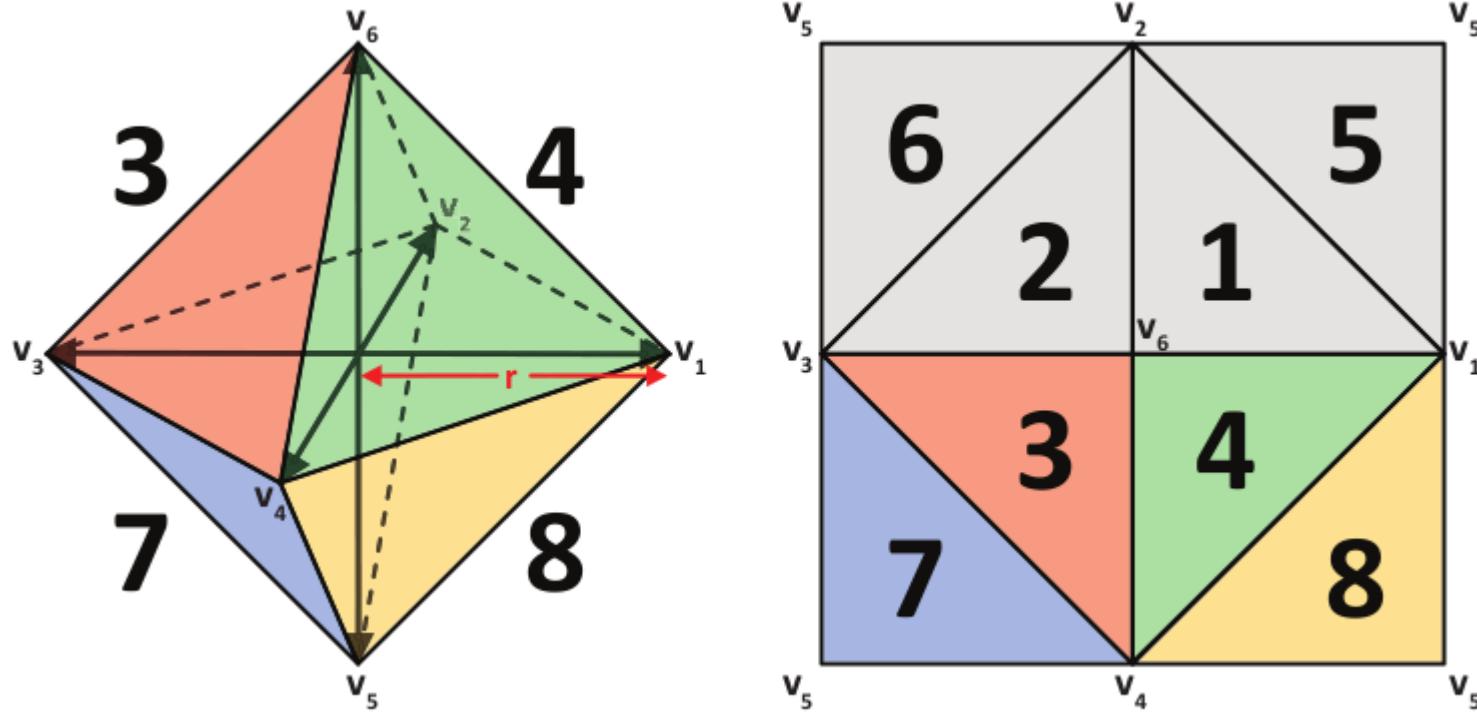
- Cubical environment map, cube map, box map
 - Enclose object in cube
 - Images on faces are easy to compute
 - Poorer filtering at edges
 - Support in OpenGL





Reflection Map Parameterization

- double pyramid [Hoppe et al. 2003]





Reflection Mapping

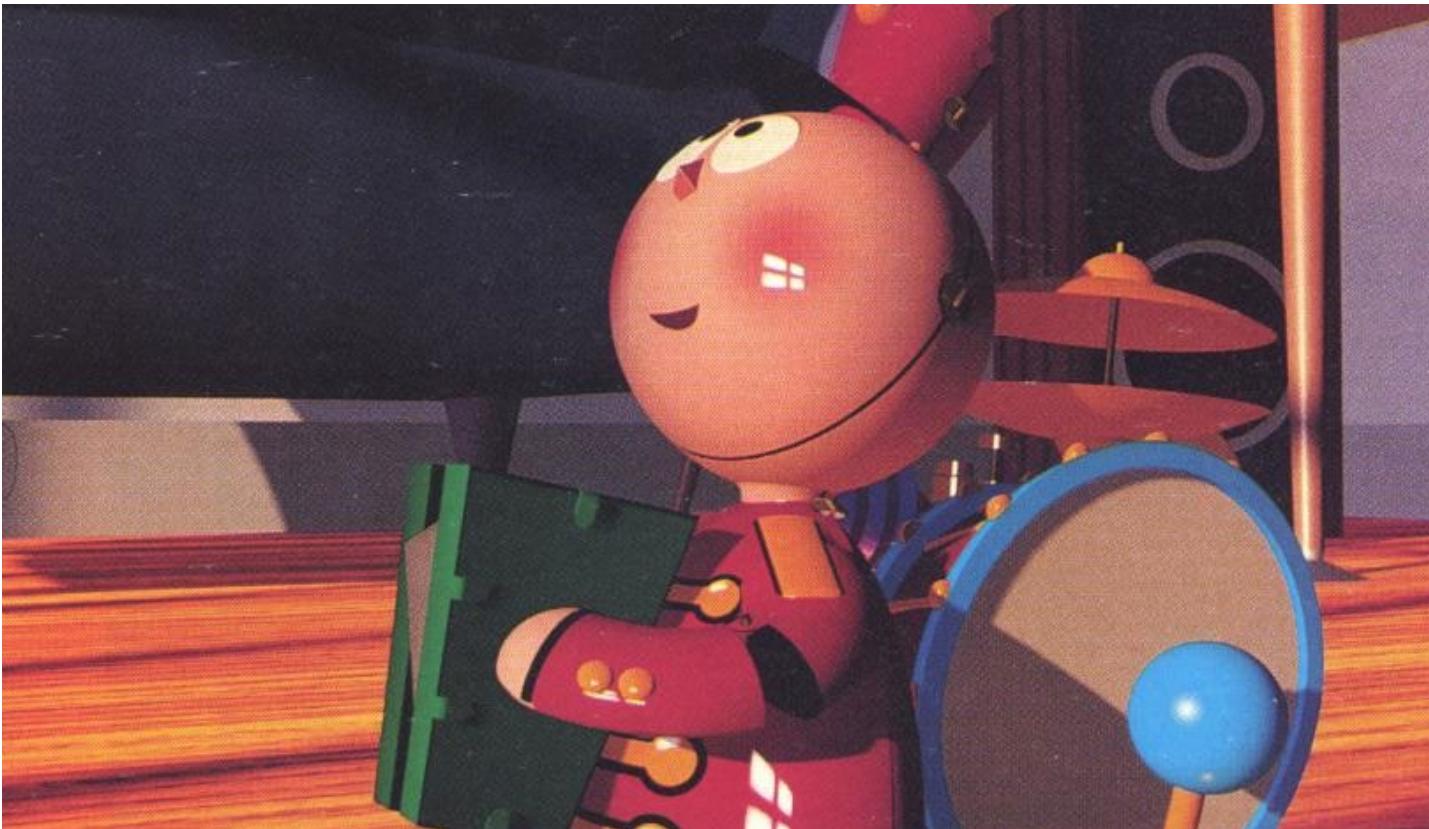


Terminator II motion picture



Reflection Mapping Example II

- Reflection mapping with Phong reflection
 - Two maps: diffuse & specular
 - Diffuse: index by surface normal
 - Specular: indexed by reflected view vector

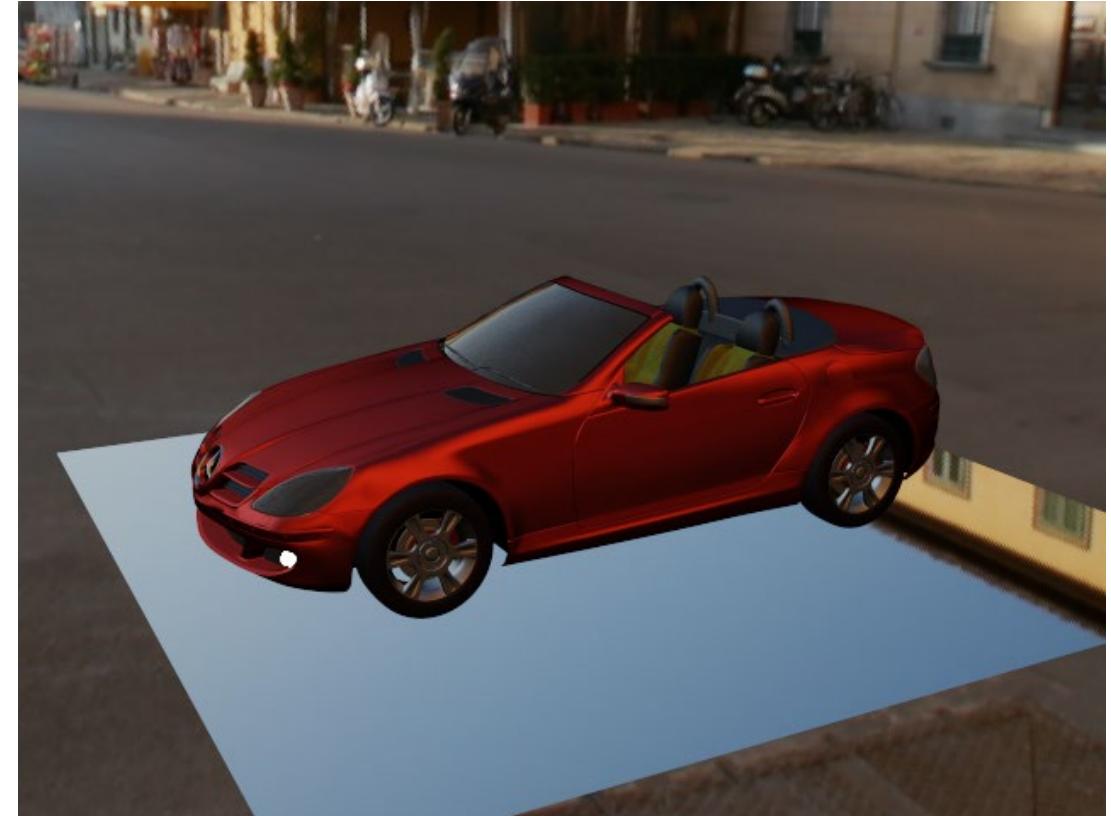


RenderMan
Companion



Ray Tracing vs. Reflection Mapping

- Differences ?





Recursive Ray Tracing

- How to fake it with reflection mapping?

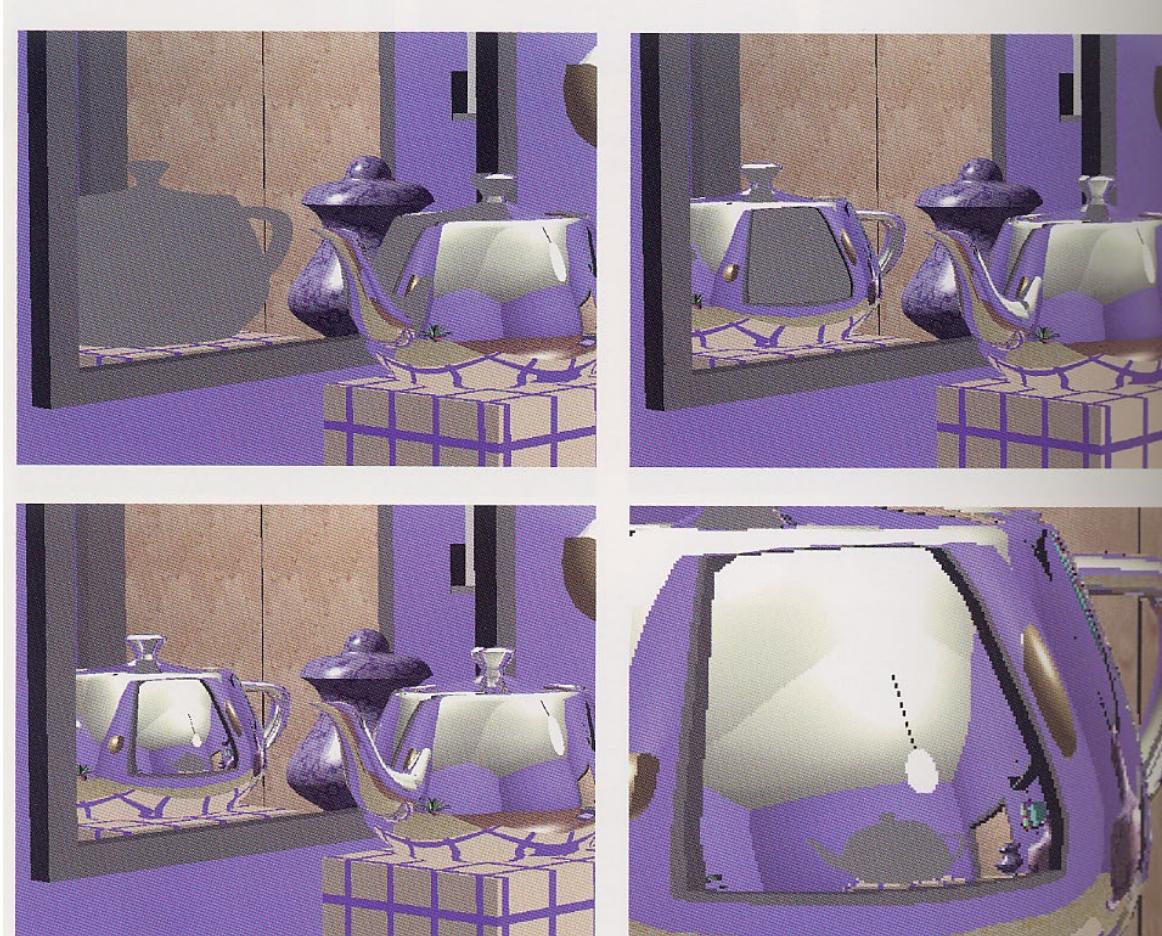


Figure 18.11

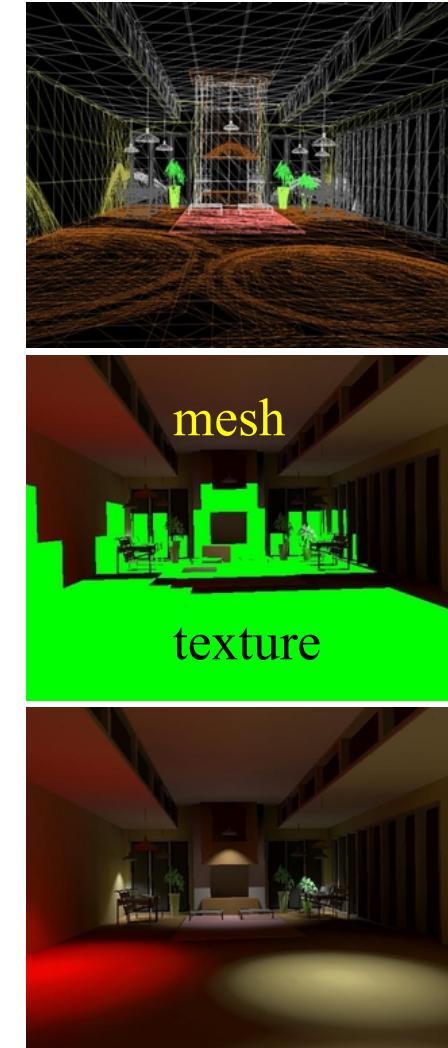
A recursive depth demonstration. The trace terminates at depth 2, 3, 4 and 5 (zoom image) respectively. 'Unassigned' pixels are coloured grey. Bad aliasing as a function of recursive depth (the light cable) is apparent.



Light Maps

- Light maps (i.e. in Quake)
 - Pre-calculated illumination (local irradiance)
 - Often very low resolution
 - Multiplication of irradiance with base texture
 - Diffuse reflectance only
 - Provides surface radiosity
 - View-independent
 - Animated light maps
 - Animated shadows, moving light spots etc.

$$\begin{array}{ccc}
 \text{Reflectance} & \times & \text{Irradiance} \\
 \text{Radiosity}
 \end{array}$$



Representing radiosity
in a mesh or texture



Bump Mapping

- Modulation of the normal vector
 - Surface normals changed only
 - Influences shading only
 - No self-shadowing, contour is not altered



Base
Model



Bump
Mapping



Displacement
Mapping

Image courtesy of www.chromesphere.com

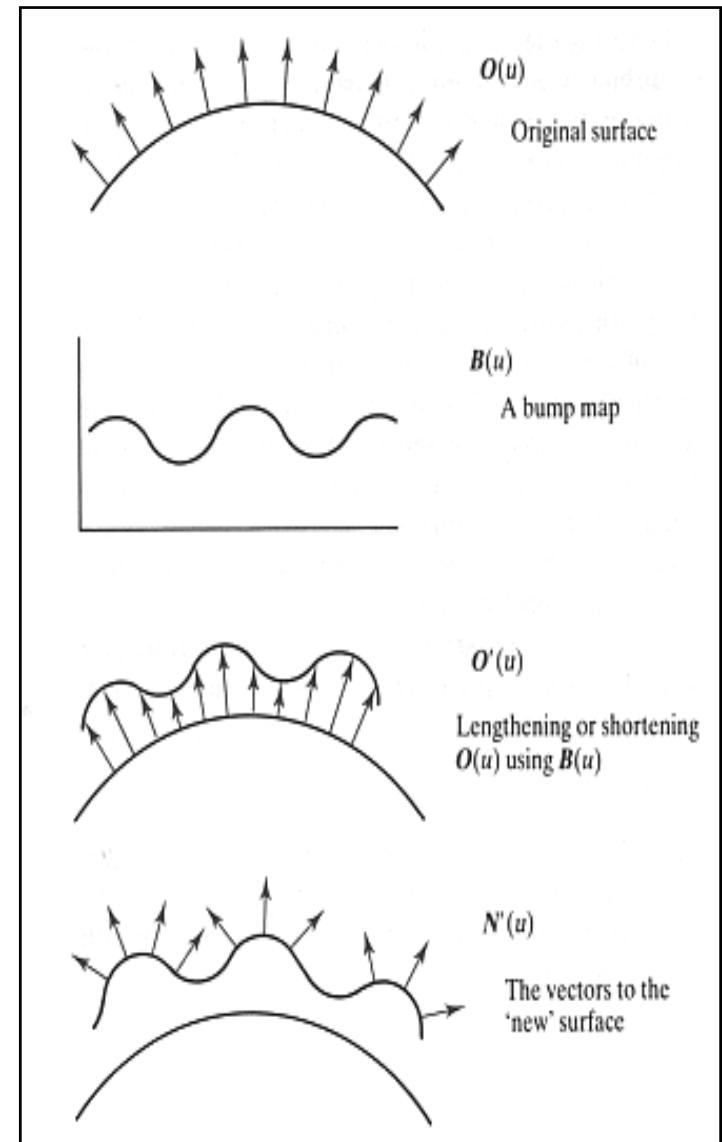


Bump Mapping

- Original surface $O(u, v)$
 - Surface normals are known
- Bump map $B(u, v) \in R$
 - Surface is offset in normal direction according to bump map intensity
 - New normal directions $N'(u, v)$ are calculated based on *virtually displaced surface* $O'(u, v)$
 - Originals surface is rendered with new normals $N'(u, v)$



Grey-valued texture used for bump height





Bump Mapping

$$\mathbf{O}'(u, v) = \mathbf{O}(u, v) + \mathbf{B}(u, v) \frac{\mathbf{N}}{|\mathbf{N}|}$$

Now differentiating this equation gives:

$$\mathbf{O}'_u = \mathbf{O}_u + \mathbf{B}_u \frac{\mathbf{N}}{|\mathbf{N}|} + \mathbf{B} \left(\frac{\mathbf{N}}{|\mathbf{N}|} \right)_u$$

$$\mathbf{O}'_v = \mathbf{O}_v + \mathbf{B}_v \frac{\mathbf{N}}{|\mathbf{N}|} + \mathbf{B} \left(\frac{\mathbf{N}}{|\mathbf{N}|} \right)_v$$

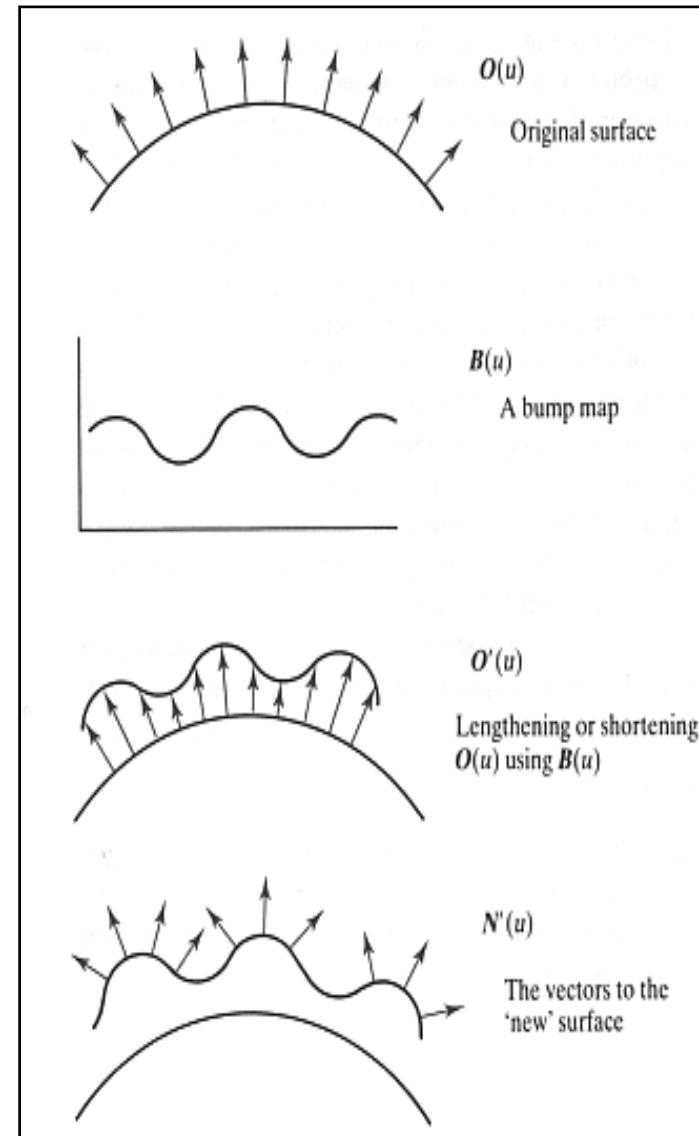
If \mathbf{B} is small (that is, the bump map displacement function is small compared with its spatial extent) the last term in each equation can be ignored and

$$\begin{aligned} \mathbf{N}'(u, v) &= \mathbf{O}_u \times \mathbf{O}_v + \mathbf{B}_u \left(\frac{\mathbf{N}}{|\mathbf{N}|} \times \mathbf{O}_v \right) + \mathbf{B}_v \left(\mathbf{O}_u \times \frac{\mathbf{N}}{|\mathbf{N}|} \right) \\ &\quad + \mathbf{B}_u \mathbf{B}_v \left(\frac{\mathbf{N} \times \mathbf{N}}{|\mathbf{N}|^2} \right) \end{aligned}$$

The first term is the normal to the surface and the last term is zero, giving:

$$\mathbf{D} = \mathbf{B}_u (\mathbf{N} \times \mathbf{O}_v) - \mathbf{B}_v (\mathbf{N} \times \mathbf{O}_u)$$

$$\mathbf{N}' = \mathbf{N} + \mathbf{D}$$





Solid Textures

- Color (and transparency) per voxel
- Evaluated on visible surface

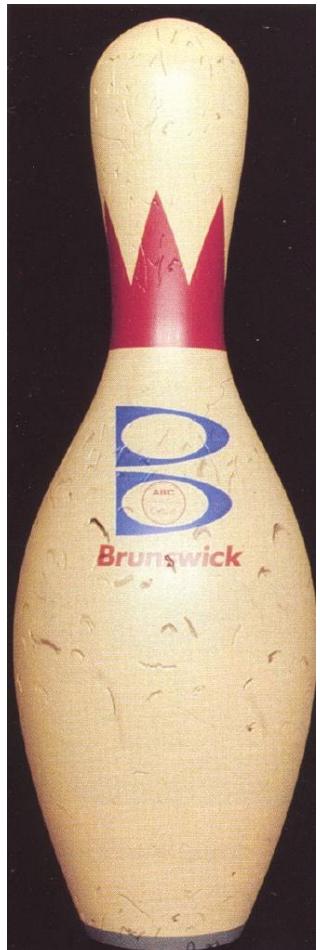


[Kopf et al. SIGGRAPH 2007]

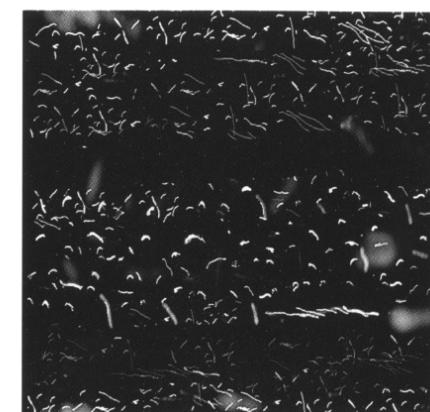


Texture Examples

- Complex optical effects
 - Combination of multiple textures



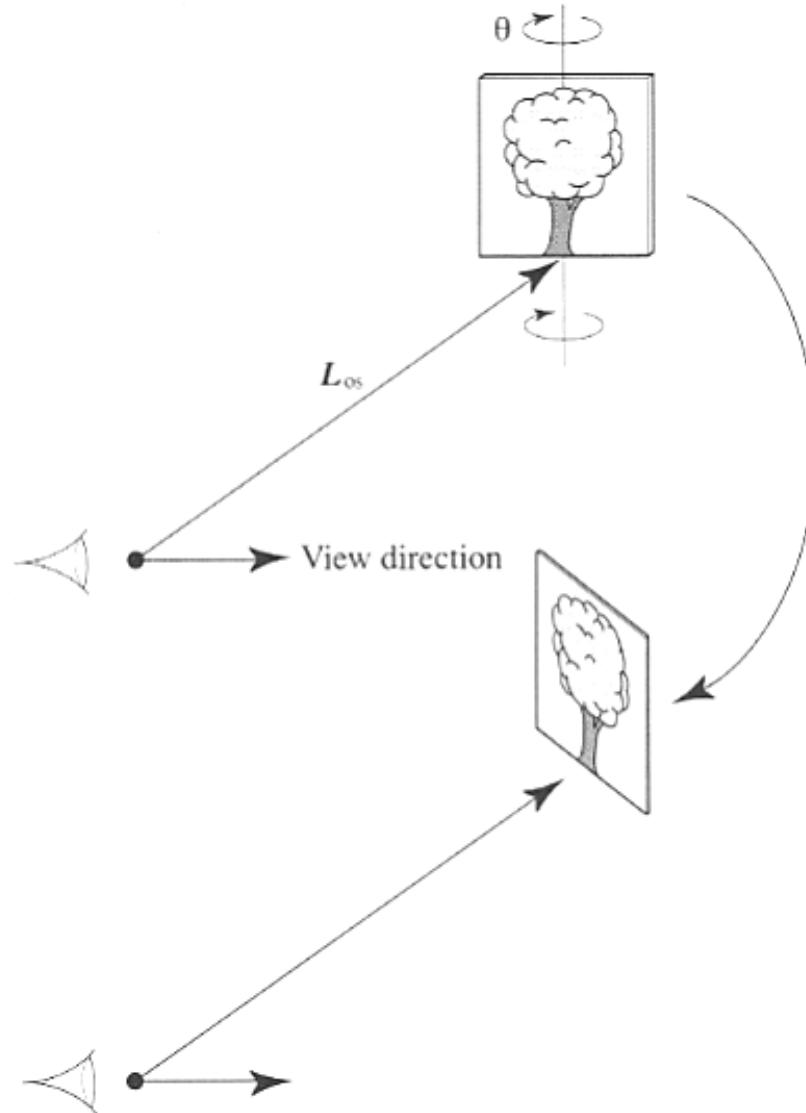
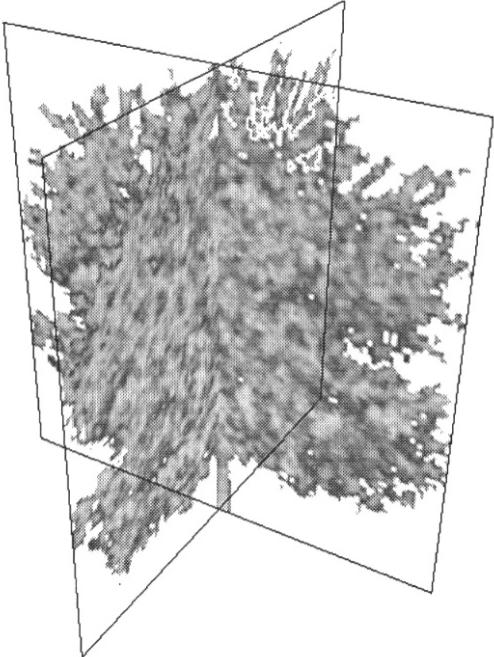
RenderMan Companion





Billboards

- Single textured polygons
 - Often with transparency texture
- Rotates, always facing viewer
- Used for rendering distant objects
- Best results if approximately radially or spherically symmetric



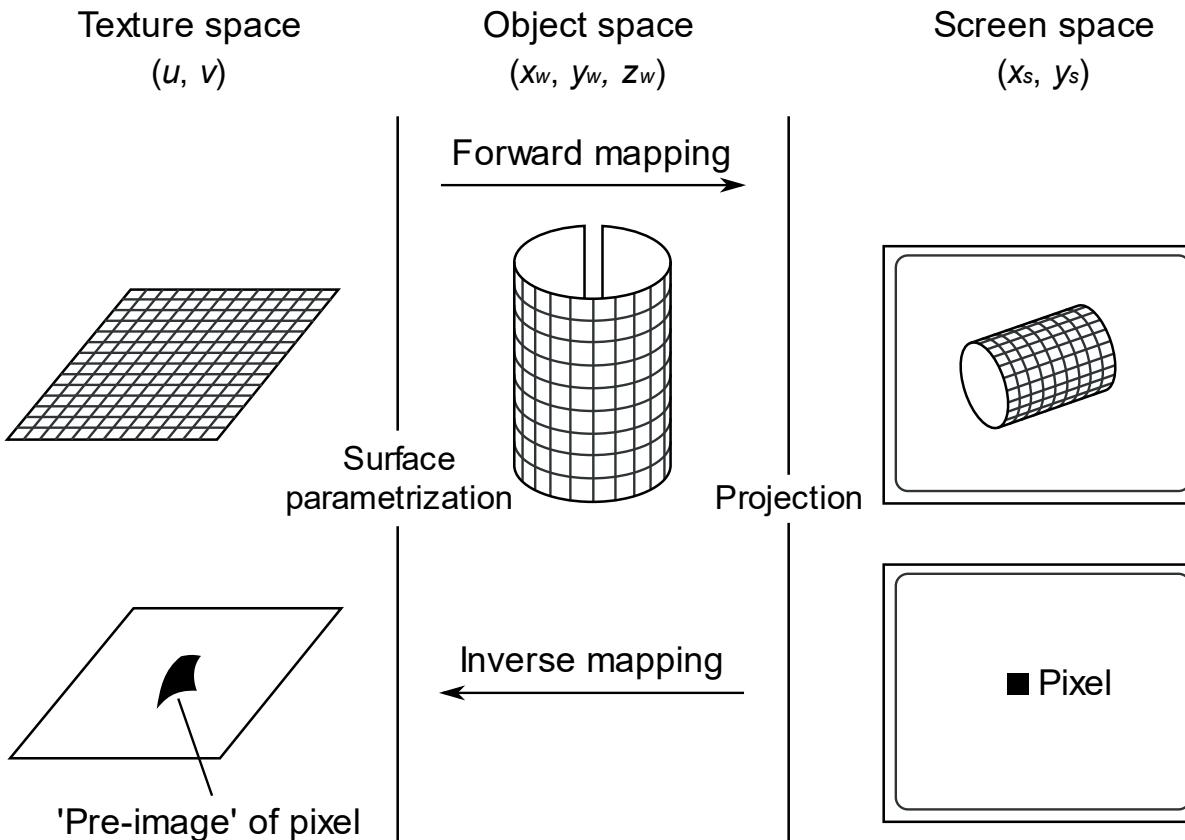


Texture Filtering

How to deal with varying texel size due to projection?



2D Texture Mapping



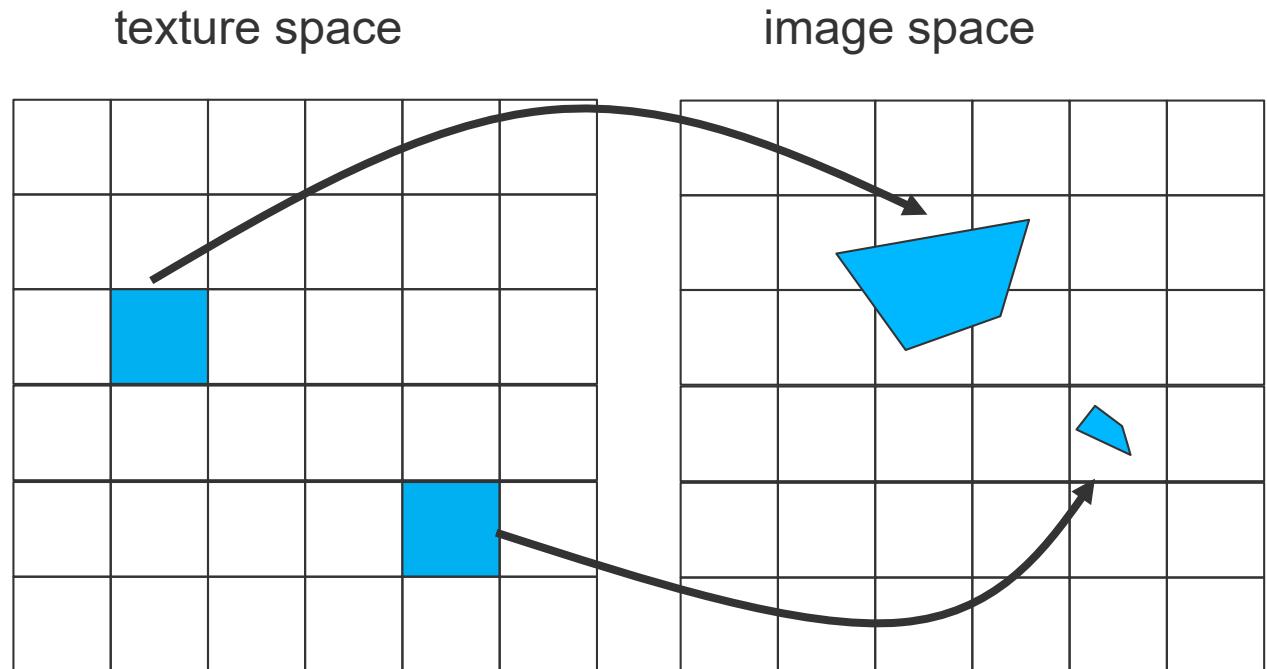
- Forward mapping
 - Object surface parameterization
 - Projective transformation
- Inverse mapping
 - Find corresponding pre-image/footprint of each pixel in texture
 - Integrate over pre-image

Forward Mapping

- Maps each texel to its position in the image
- Uniform sampling of texture space does not guarantee uniform sampling in screen space
- Possibly used if
 - The texture-to-screen mapping is difficult to invert
 - The texture image does not fit into memory

Texture scanning:

```
for v
  for u
    compute x(u,v) and y(u,v)
    copy TEX[u,v] to SCR[x,y]
(or in general
  rasterize image of TEX[u,v])
```



Inverse Mapping

- Requires inverting the mapping transformation
- Preferable when the mapping is readily invertible and the texture image fits into memory
- The most common mapping method
 - for each pixel in screen space, the pre-image of the pixel in texture space is found and its area is integrated over

Screen scanning:

for y

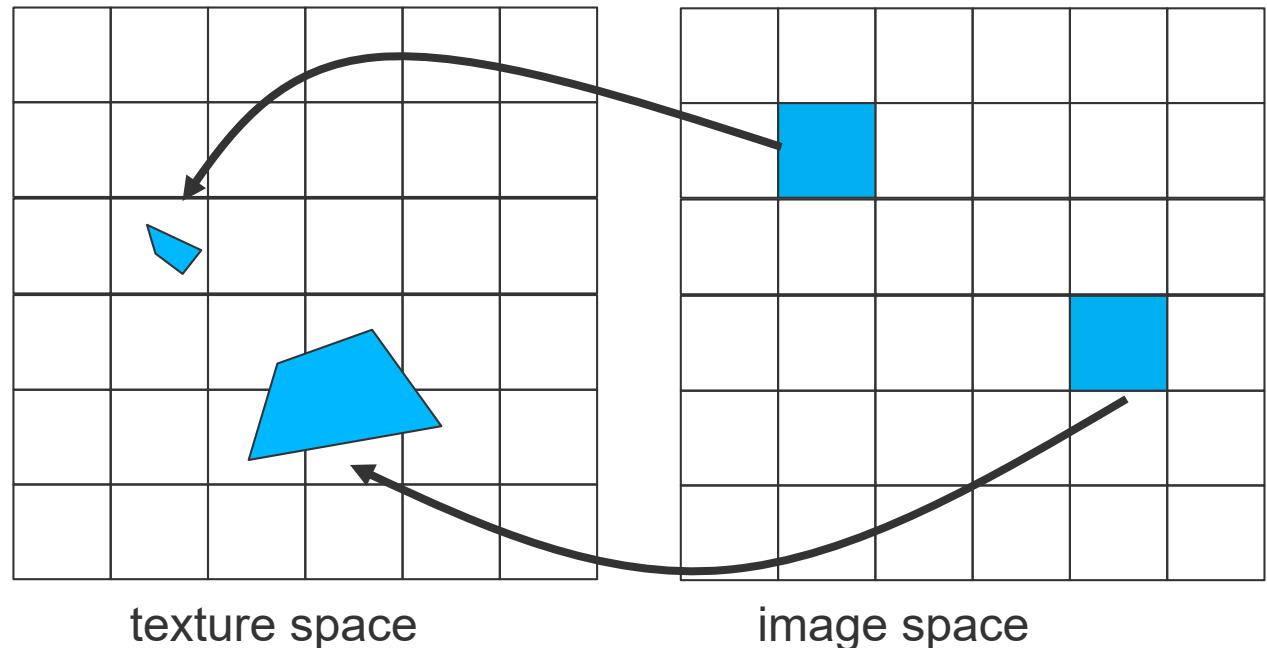
 for x

 compute $u(x,y)$ and $v(x,y)$

 copy $\text{TEX}[u,v]$ to $\text{SCR}[x,y]$

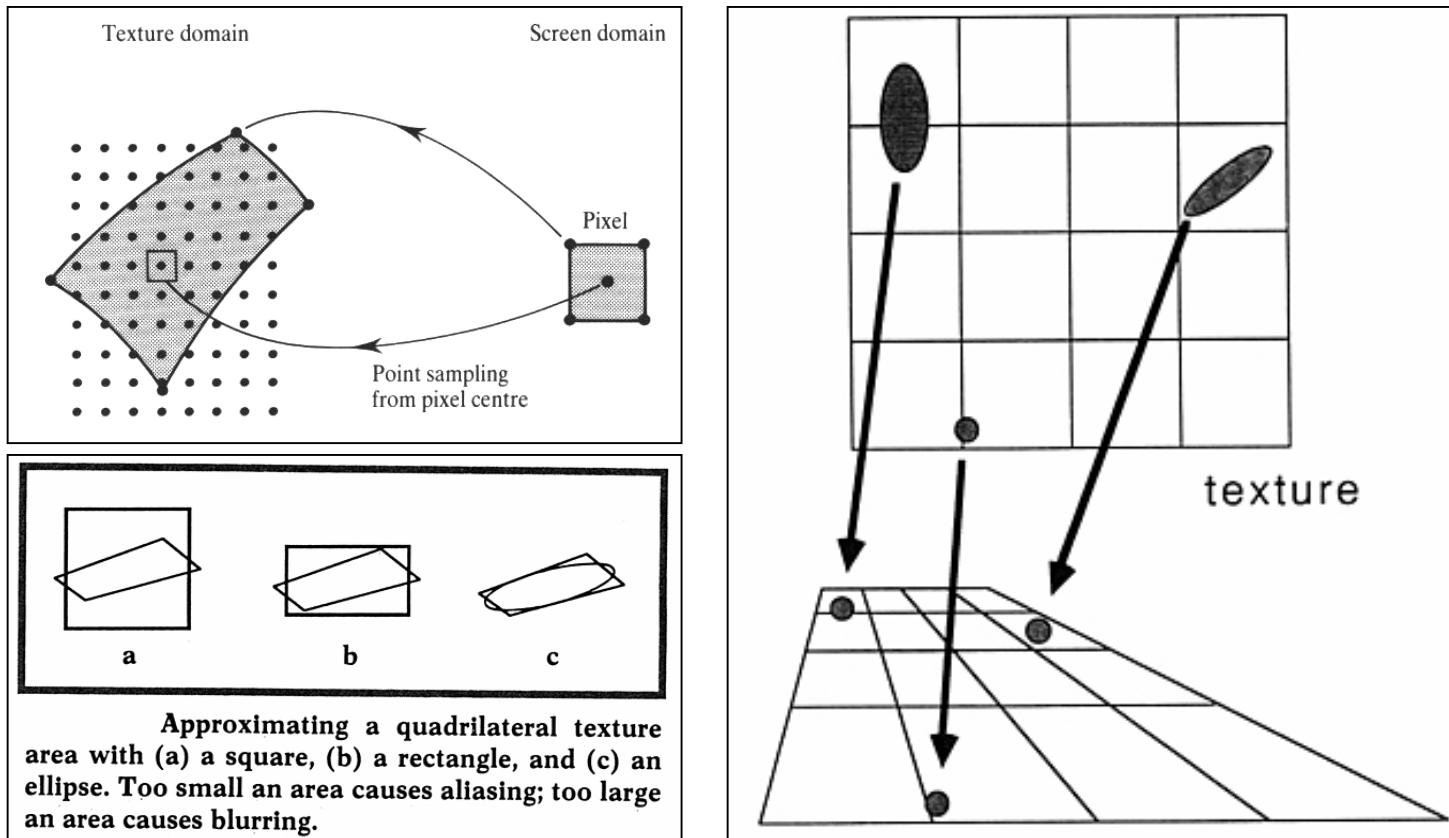
(or in general

 integrate over image of $\text{SCR}[u,v]$)



Pixel Pre-Image in Texture Space

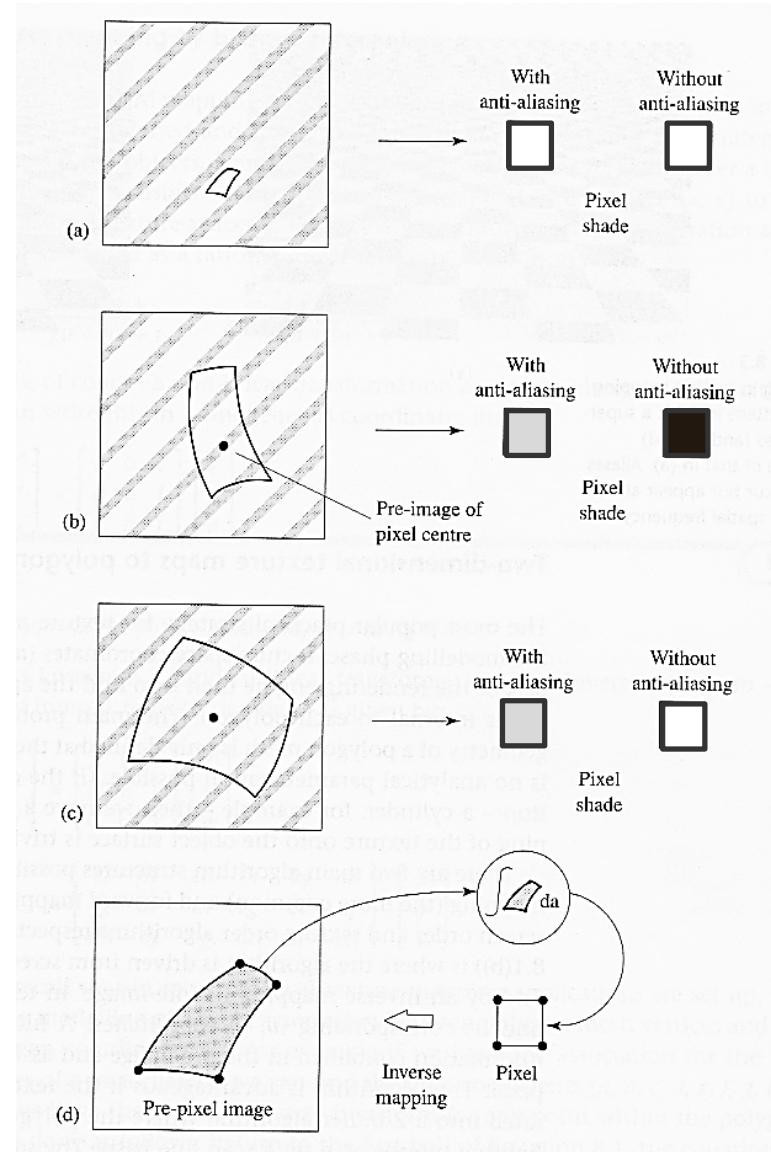
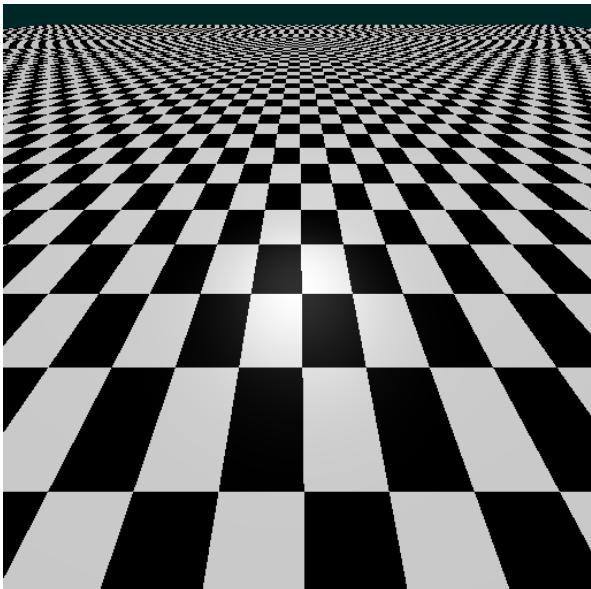
- A square screen pixel that intersects a curved surface has a curvilinear quadrilateral pre-image in texture space. Most methods approximate the true mapping by a quadrilateral or parallelogram. Or they take multiple samples within a pixel. If pixels are instead regarded as circles, their pre-images are ellipses.





Inverse Mapping: Filtering

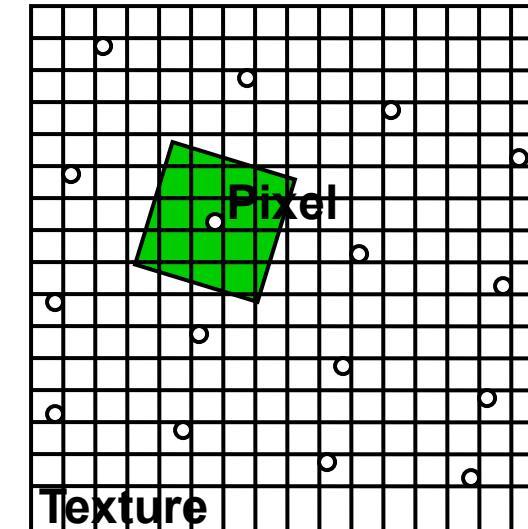
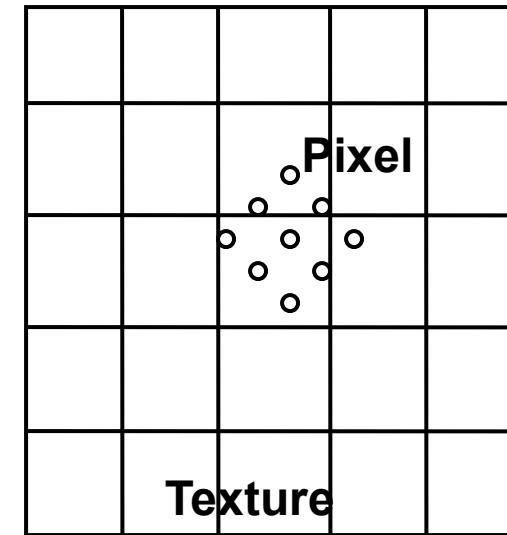
- Integration of Pre-image
 - Integration over pixel footprint in texture space
- Aliasing
 - Texture insufficiently sampled
 - Incorrect pixel values
 - “Randomly” changing pixels when moving





Filtering

- Magnification
 - Map few texels onto many pixels
 - Nearest:
 - Take the nearest texel
 - Bilinear interpolation:
 - Interpolation between 4 nearest texels
 - Need fractional accuracy of coordinates
- Minification
 - Map many texels to one pixel
 - Aliasing:
 - Reconstructing high-frequency signals with low level frequency sampling
 - Filtering
 - Averaging over (many) associated texels
 - Computationally expensive



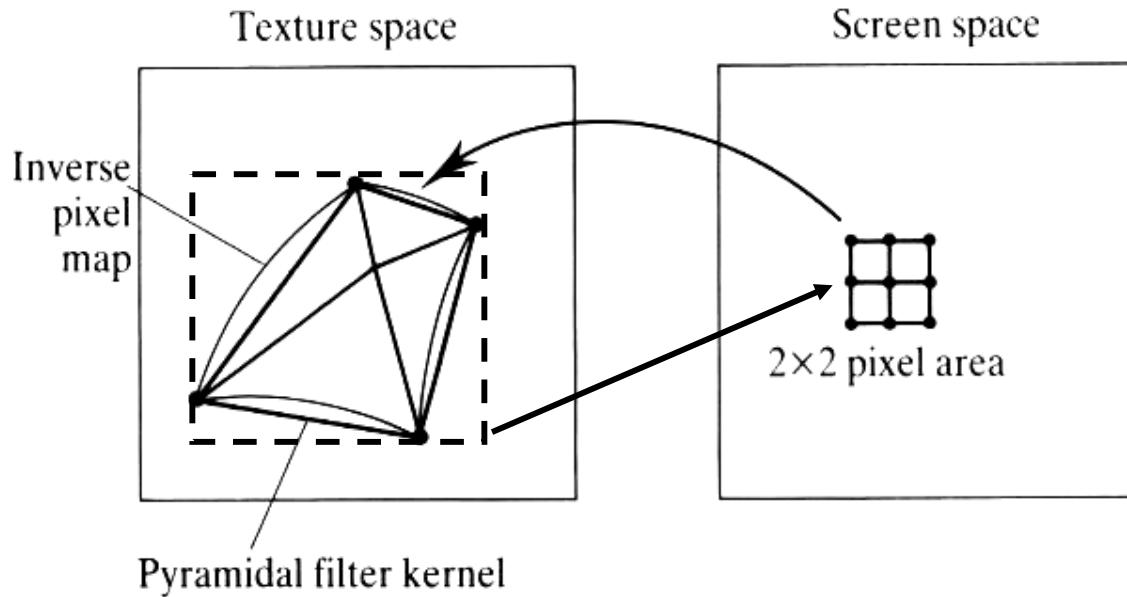


Filtering – Texture Minification

- Space-variant filtering
 - Mapping from texture space (u,v) to screen space (x,y) not affine
 - Filtering changes with position
- Space variant filtering methods
 - **Direct convolution**
 - Numerically compute the Integral
 - **Pre-filtering**
 - Precompute the integral for certain regions -- more efficient
 - Approximate footprint with regions

Direct Convolution

- Convolution in texture space
 - Texels weighted according to distance from pixel center (e.g. pyramidal filter kernel)

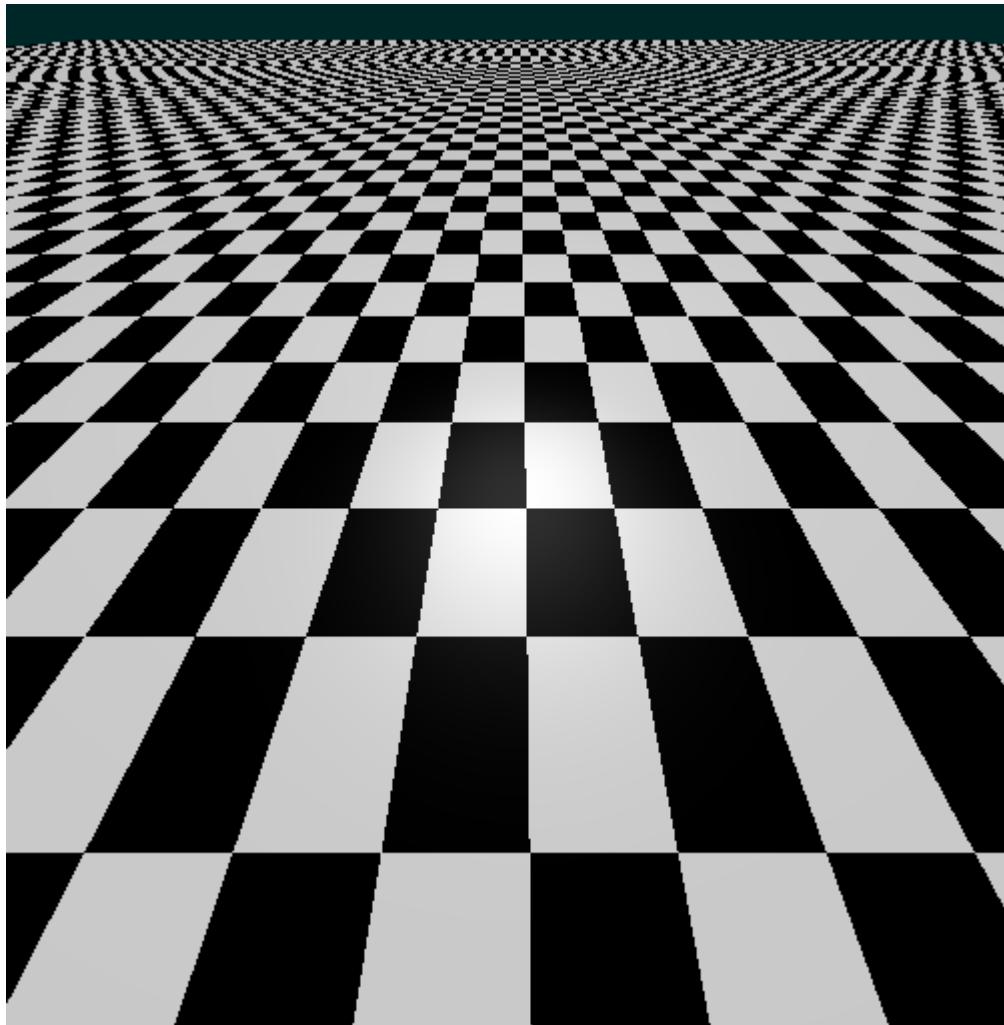


- Convolution in image space
 - 1 Center the filter function on the pixel (in image space) and find its bounding rectangle.
 - 2 Transform the rectangle to the texture space, where it is a quadrilateral. The sides of this rectangle are assumed to be straight. Find a bounding rectangle for this quadrilateral.
 - 3 Map all pixels inside the texture space rectangle to screen space.
 - 4 Form a weighted average of the mapped texture pixels using a two-dimensional lookup table indexed by each sample's location within the pixel.



Aliasing – Sampling too sparse

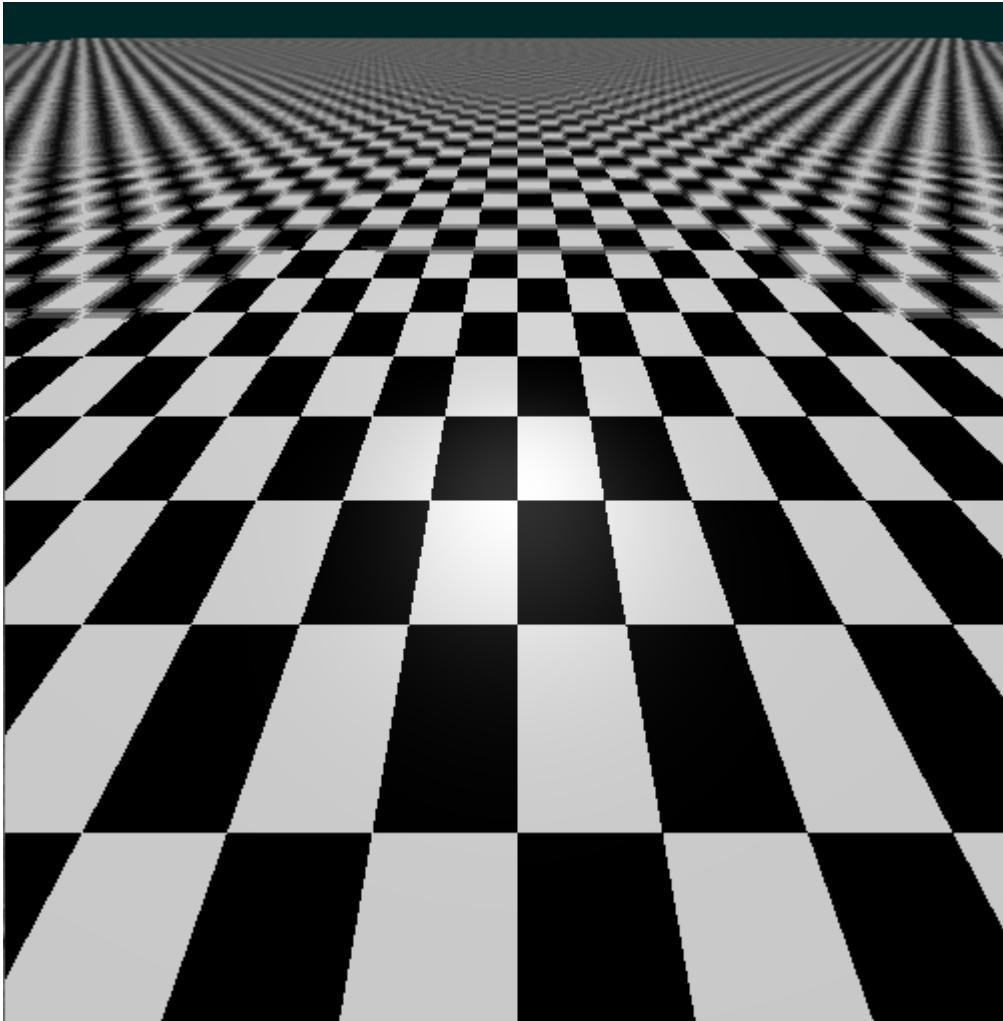
- checker board gets distorted





EWA Filtering – Anti-Aliasing

- elliptical filtering plus Gaussian





Accelerated Filtering



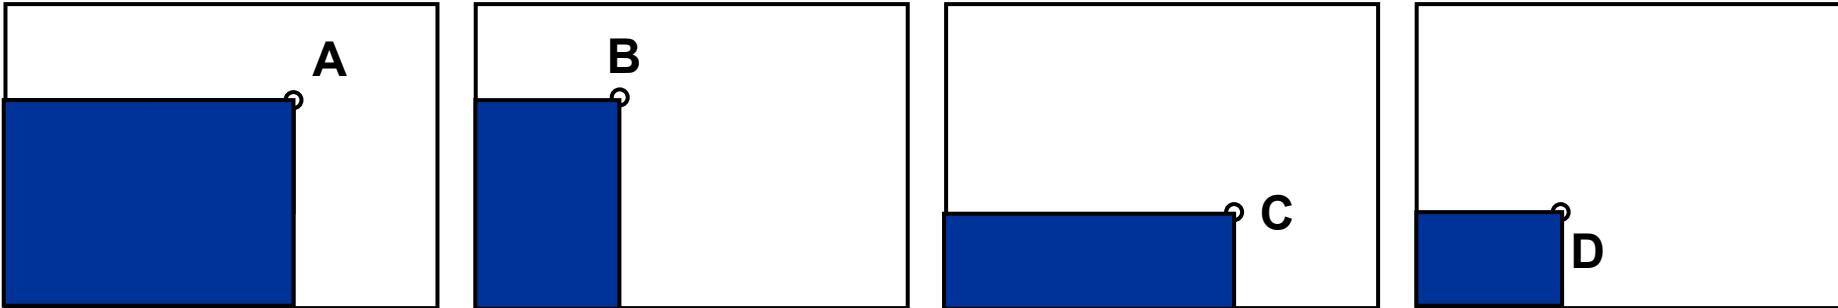
Filtering – Texture Minification

- Direct convolution methods are slow
 - A pixel pre-image can be arbitrarily large along silhouettes or at the horizon of a textured plane
 - Horizon pixels can require averaging over thousands of texture pixels
 - Texture filtering cost grows in proportion to projected texture area
- Speed up
 - The texture can be prefiltered so that during rendering only a few samples will be accessed for each screen pixel
- Two data structures are commonly used for prefiltering:
 - Integrated arrays (summed area tables)
 - Image pyramids (mipmaps) Space-variant filtering



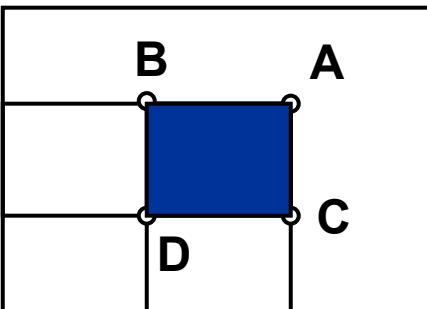
Summed Area Tables

- Per texel, store sum from $(0, 0)$ to (u, v)



- Many bits per texel (sum !)
- Evaluation of 2D integrals in constant time!

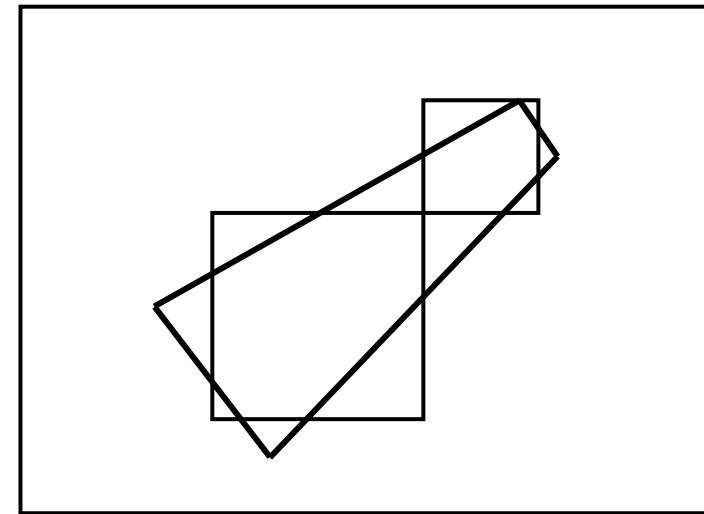
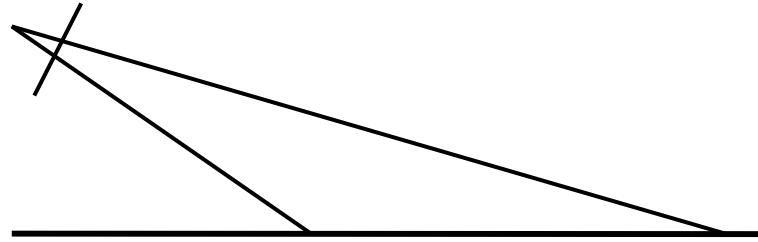
$$\int_{Bx}^{Ax} \int_{Cy}^{Ay} I(x, y) dx dy = A - B - C + D$$





Integrated Arrays

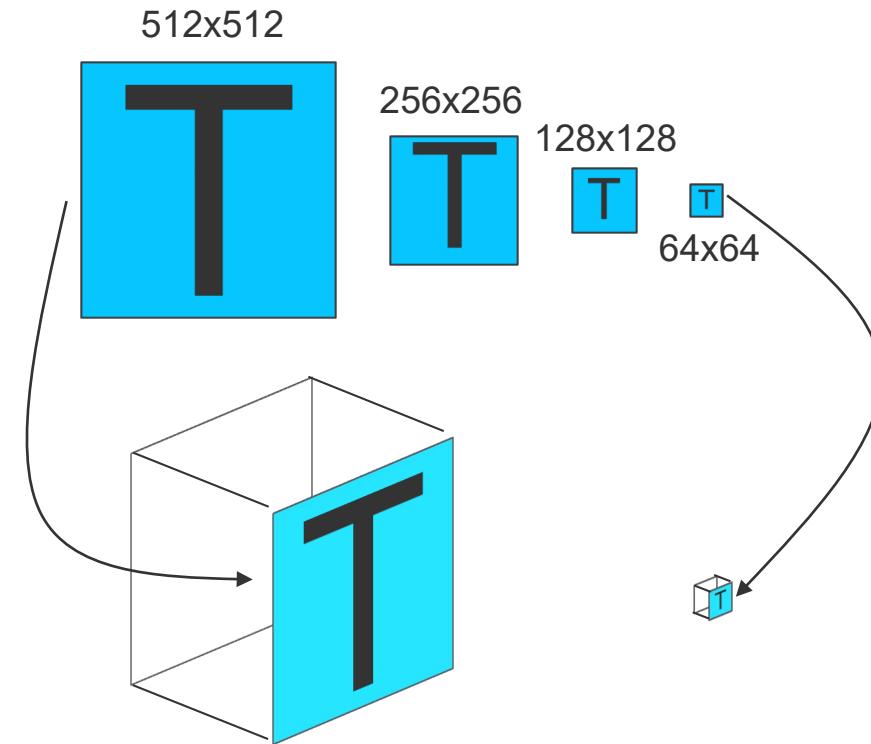
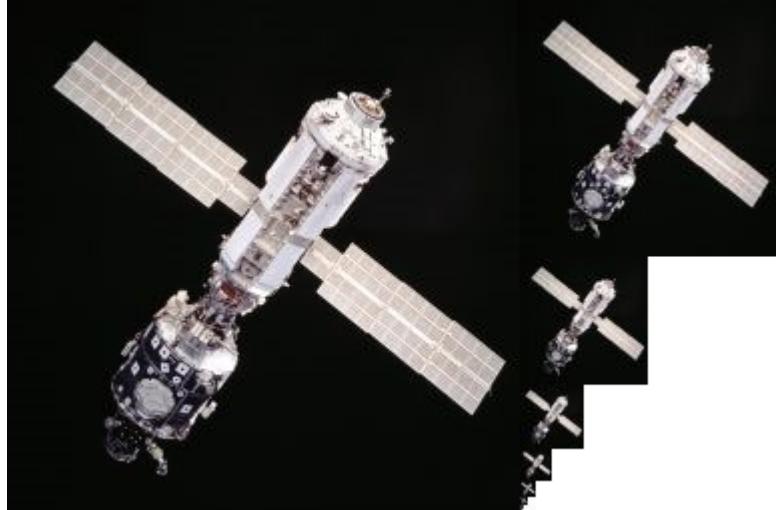
- Footprint assembly
 - Good for space variant filtering
 - e.g. inclined view of terrain
 - Approximation of the pixel area by rectangular texel-regions
 - The more footprints the better accuracy
 - Often fixed number of texels because of economical reasons





MipMapping

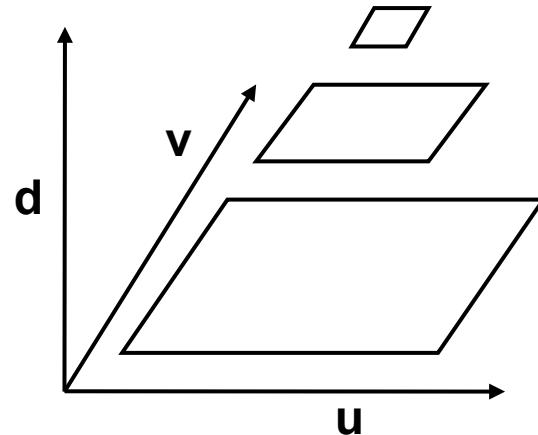
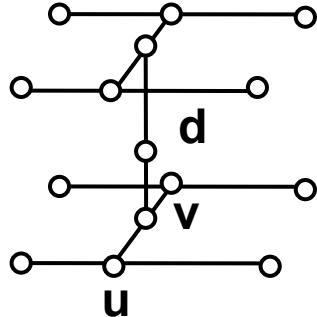
- Texture available in multiple resolutions
 - Pre-processing step
- Rendering: select appropriate texture resolution
 - Selection is usually per pixel !!
 - Texel size(n) < extent of pixel footprint < texel size($n+1$)



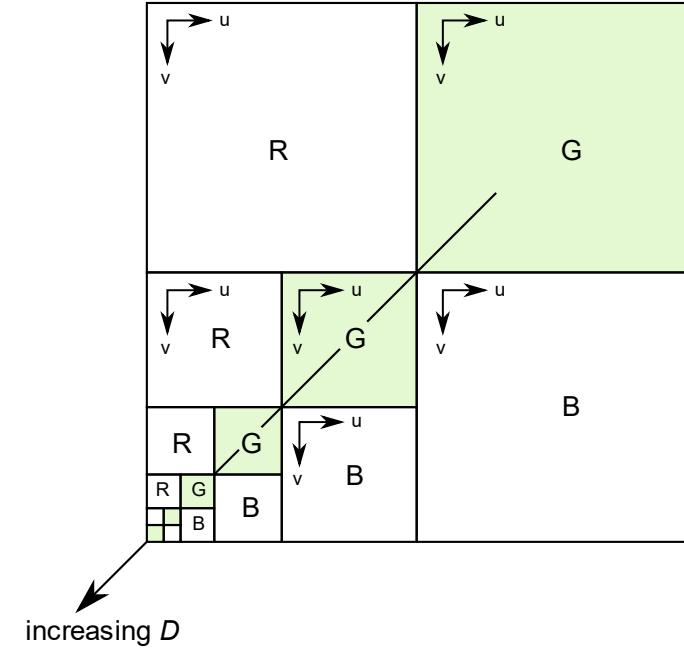
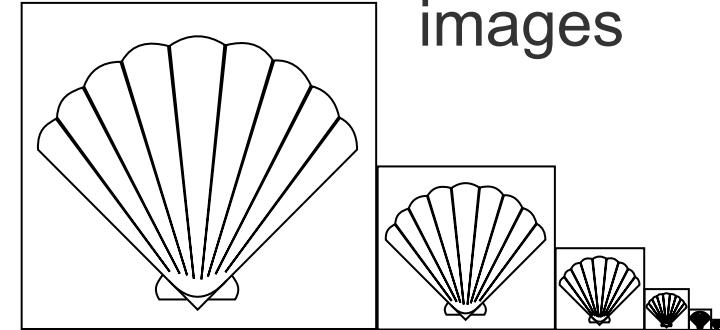


MipMapping II

- Multum In Parvo (MIP): much in little
- Hierarchical resolution pyramid
 - Repeated averaging over 2×2 texels
- Rectangular arrangement (RGB)
- Reconstruction
 - Tri-linear interpolation of 8 nearest texels



original - prefiltered
images





MipMap Example

- RGB → 4/3 the size





MipMaps

- Why is MipMapping sometimes faster?
 - Bottleneck is memory bandwidth
 - Using of texture caches
 - Texture minification required for far geometry
- No MipMap
 - „Random“ access to texture
 - Always 4 new texels
- MipMap
 - Next pixel at about one texel distance (1:1 mapping)
 - Access to 8 texels at a time, but
 - Most texels are still in the cache



Procedural Textures

How to describe a natural 2D distribution of samples?



Texture Maps vs. Procedural Textures

- Texture maps (photos, simulations, videos, ...)
 - Simple acquisition
 - Illumination „frozen“ during acquisition
 - Limited resolution, aliasing
 - High memory requirements
 - Mapping issues
- Procedural textures
 - Non-trivial programming
 - Flexibility & parametric control
 - Unlimited resolution
 - Anti-aliasing possible
 - Low memory requirements
 - Low-cost visual complexity
 - Can adapt to arbitrary geometry



Ken Perlin



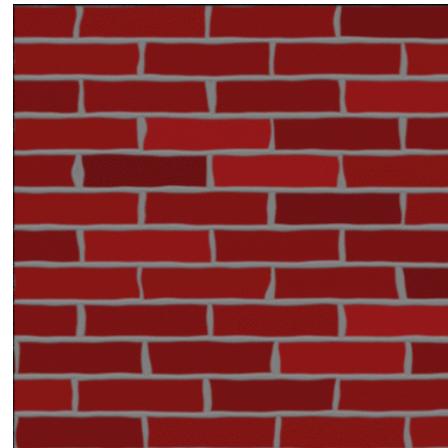
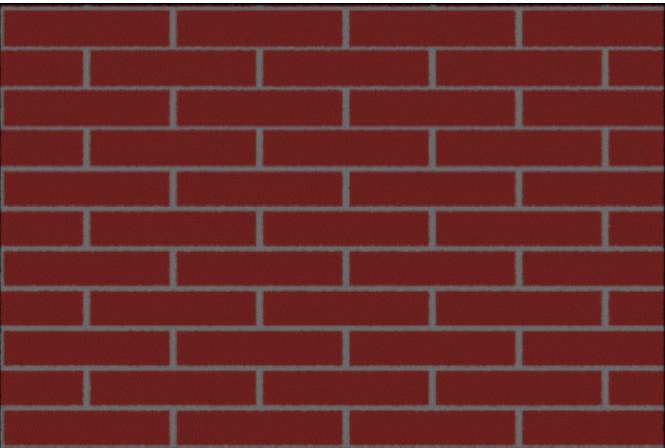
Procedural Textures

- Function of some shading parameter, e.g.
 - world space, texture coordinates, ...
- Texturing: evaluation of function on object surface
 - Ray tracing: At intersection point with surface
- Observation: Textures of natural objects
 - Similarity between patches at different locations
 - Repetitiveness, coherence (e.g. skin of a tiger)
 - Similarity on different resolution scales
 - Self-similarity
 - But never completely identical
 - Additional disturbances, turbulence, noise
- Goal: Generic procedural texture function
 - Mimics statistical properties of natural textures
 - Purely empirical approach
 - Looks convincing, but has nothing to do with material's physics

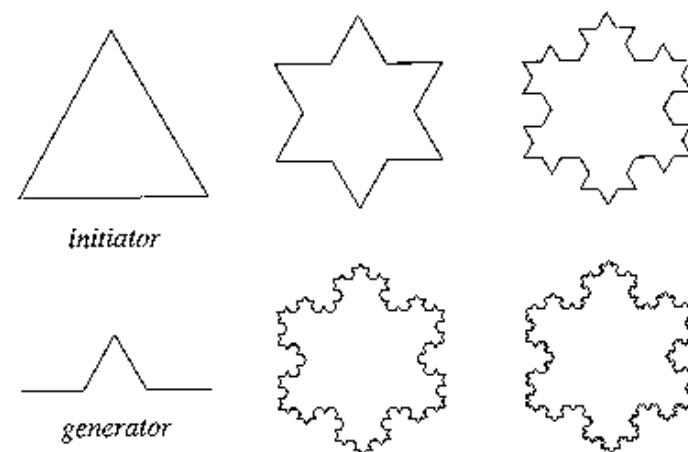


Texture Examples

- Translational similarity

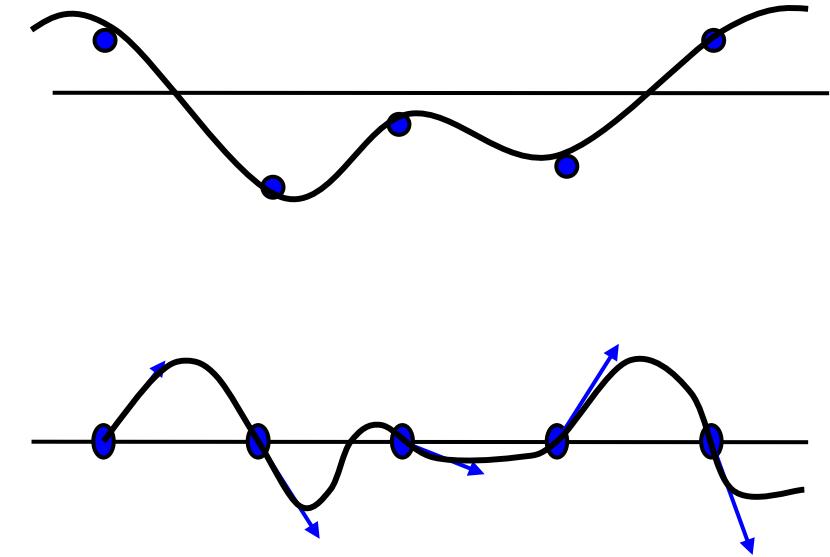


- Similarity on different scales



3D / Solid Noise: Perlin Noise

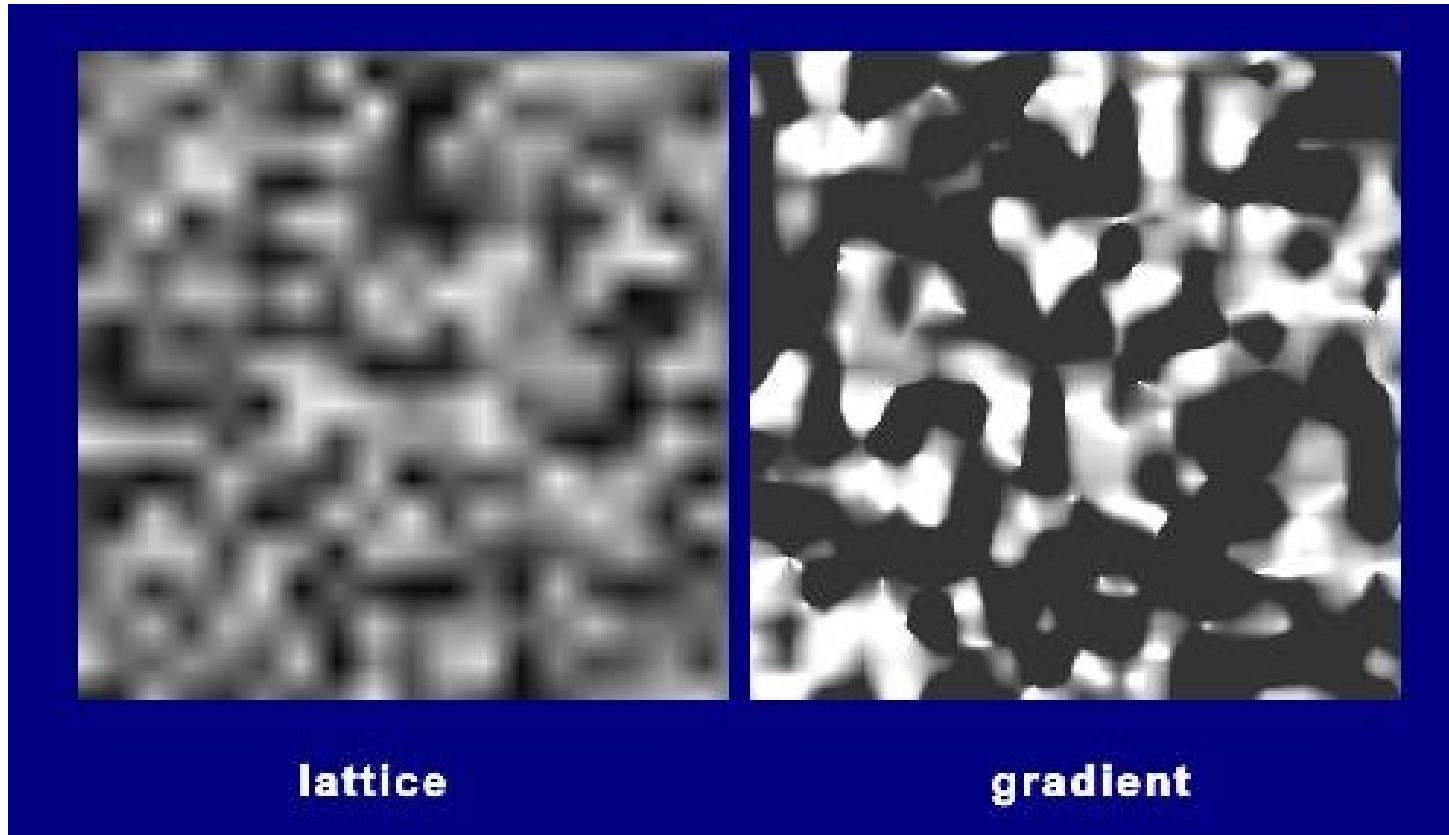
- $\text{Noise}(x,y,z)$
 - Statistical invariance under rotation
 - Statistical invariance under translation
 - Roughly one specific frequency
- Integer lattice (i,j,k)
 - Value noise: Random number at lattice
 - Look-up table or hashing function into hash map
 - Gradient lattice noise
 - Random (hashed) gradient vectors
 - Fixed fundamental frequency of ~ 1 Hz over lattice
- Evaluation at (x,y,z)
 - Tri-linear interpolation
 - Cubic interpolation (Hermite spline → later)
- Unlimited domain due to lattice and hashing
- Also see
 - http://freespace.virgin.net/hugo.elias/models/m_perlin.htm





Gradient vs. Value Noise

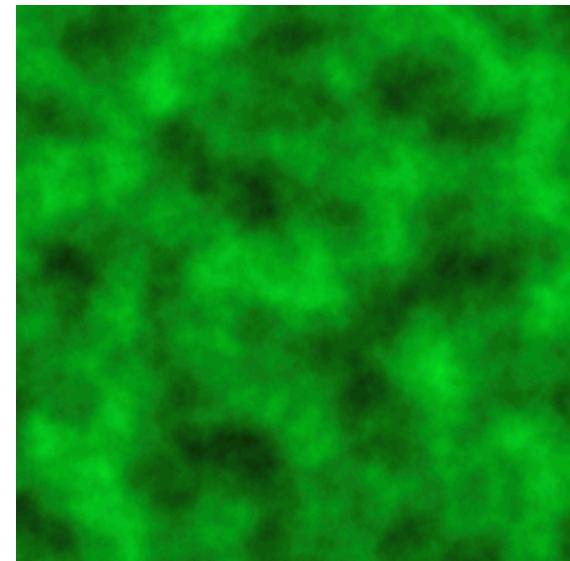
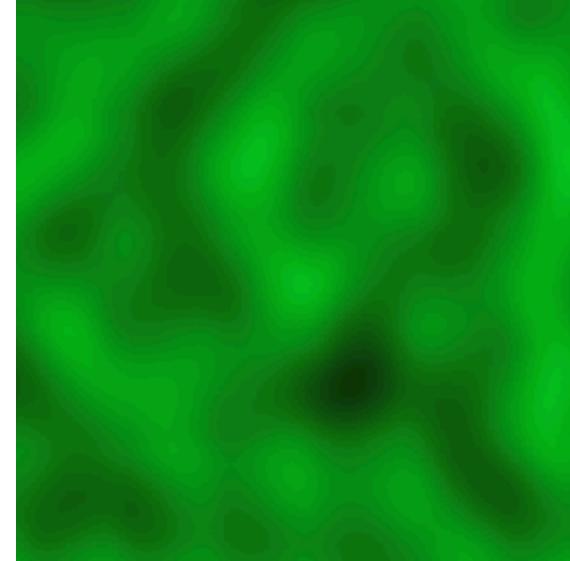
- Gradient noise better than value noise
 - Less regularity artifacts
 - More high frequencies in noise spectrum
 - Even tri-linear interpolation produces good results





Turbulence Function

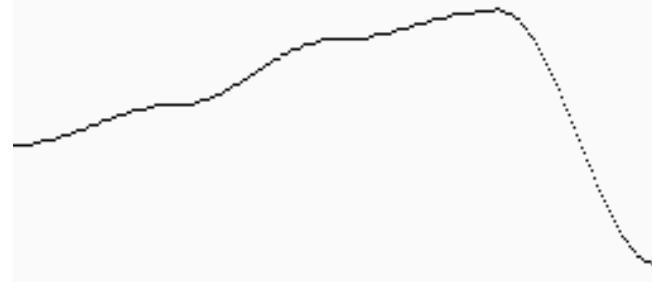
- Noise function
 - “White” frequency spectrum
- Natural textures
 - Decreasing power spectrum towards high frequencies
- Turbulence from noise
 - $\text{Turbulence}(x) = \sum_{i=0}^k \text{abs}(\text{noise}(2^i x) / p^i)$
 - persistence p typically $p=2$
 - Summation truncation
 - $1/2^{k+1} < \text{size of one pixel}$ (band limit)
 - 1. Term: $\text{noise}(x)$
 - 2. Term: $\text{noise}(2x)/2$
 - ...
 - Power spectrum: $1/f$
 - (Brownian motion has $1/f^2$)



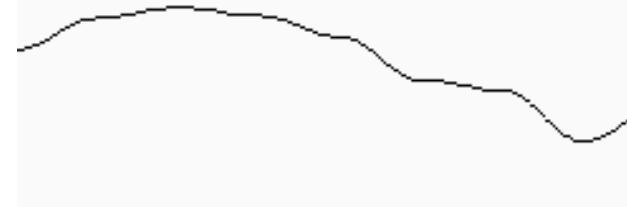


Synthesis of Turbulence (1D)

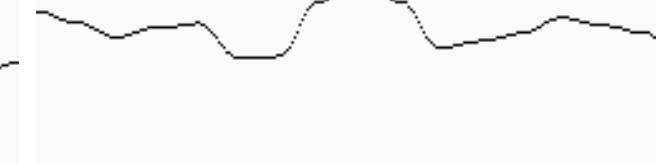
Amplitude : 128
frequency : 4



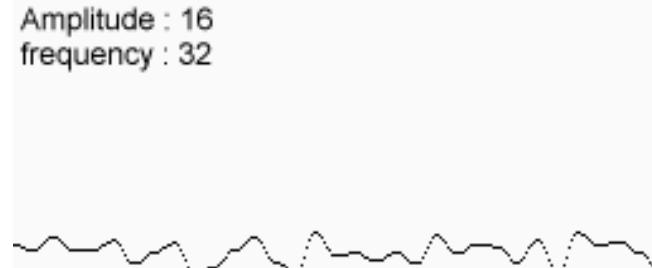
Amplitude : 64
frequency : 8



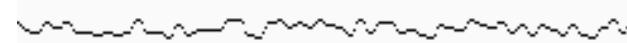
Amplitude : 32
frequency : 16



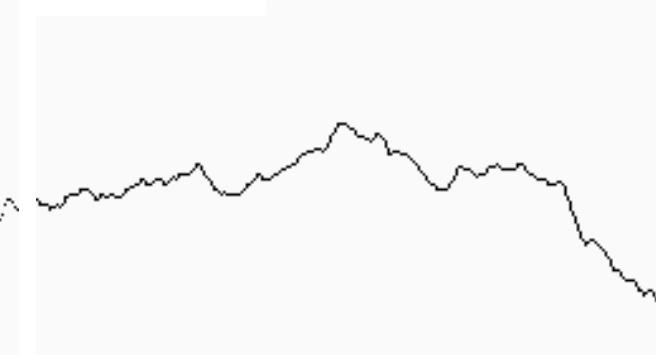
Amplitude : 16
frequency : 32



Amplitude : 8
frequency : 64

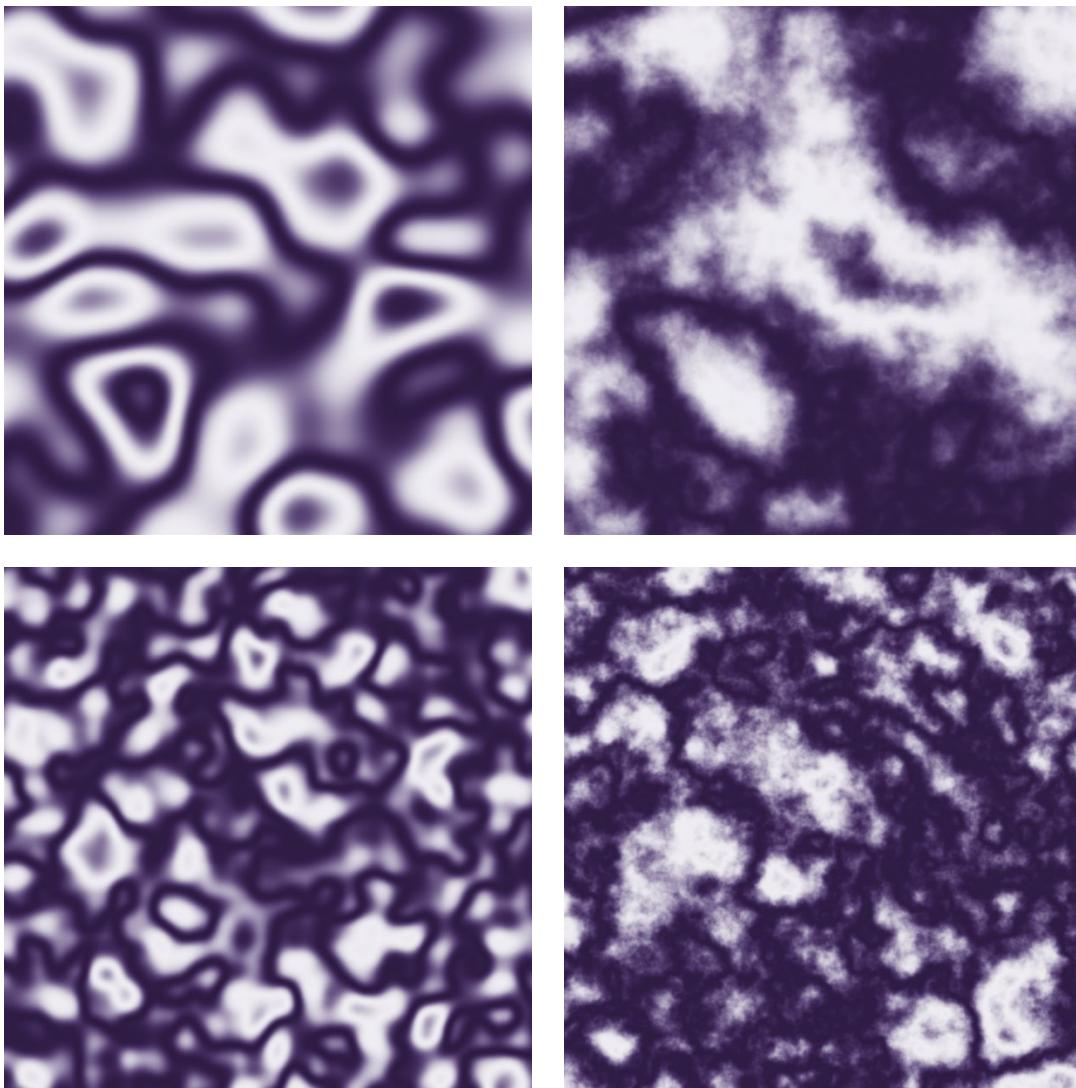


Sum of Noise Functions = (Perlin Noise)



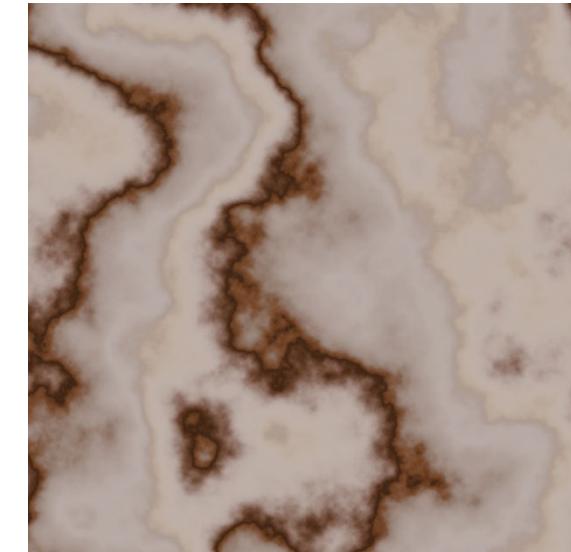


Synthesis of Turbulence (2D)



Example: Marble Texture Function

- Overall structure: alternating layers of white and colored marble
 - $f_{\text{marble}}(x,y,z) := \text{marble_color}(\sin(x))$
 - `marble_color` : transfer function (see lower left)
- Realistic appearance: simulated turbulence
 - $f_{\text{marble}}(x,y,z) := \text{marble_color}(\sin(x + \text{turbulence}(x,y,z)))$
- Moving object: turbulence function also transformed





Further Procedural Texturing Applications

- Bark
 - Turbulated sawtooth function
 - Bump mapping
- Clouds
 - White blobs
 - Turbulated transparency along edge
 - Transparency mapping
- Animation
 - Vary procedural texture function's parameters over time



Fractal Landscapes

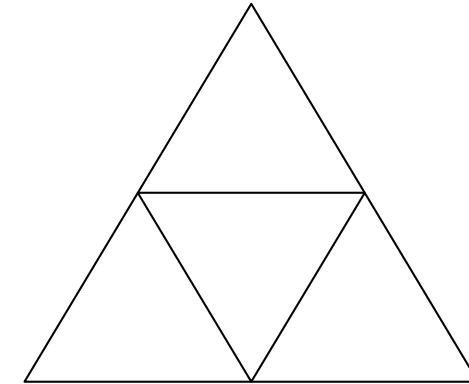
- Procedural generation of geometry
- Complex geometry at virtually no memory cost
 - Can be difficult to ray trace !





Fractal Landscapes

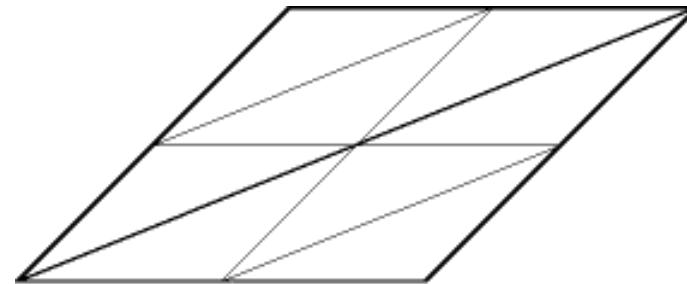
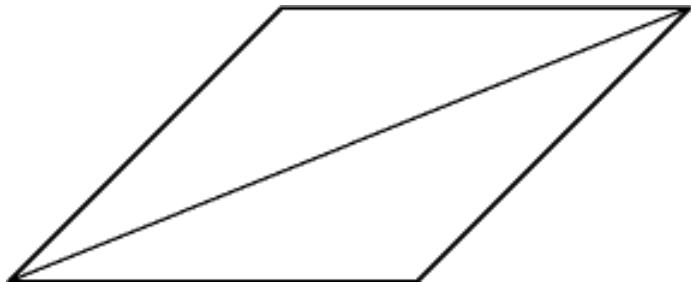
- Coarse triangle mesh approximation
- 1:4 triangle subdivision
 - Vertex insertion at edge-midpoints
- New vertex perturbation
 - Random displacement along normal
 - Scale of perturbation depends on subdivision level
 - Decreasing power spectrum
 - Parameter models surface roughness
- Recursive subdivision
 - Level of detail (LOD) determined by # subdivisions
- All done inside renderer
 - LOD generated locally when/where needed (bounding box test)
 - Minimal I/O cost (coarse mesh only)



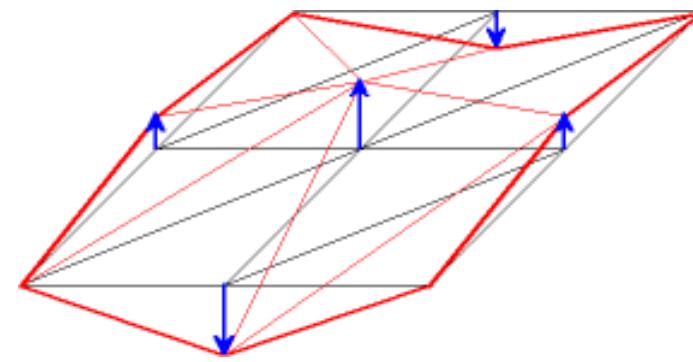
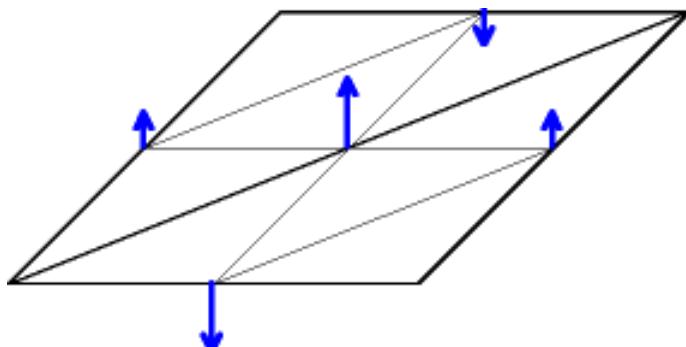


Fractal Landscapes

- Triangle subdivision
 - Insert new vertices at edge midpoints
 - 1:4 triangle subdivision



- Vertex displacement
 - Along original triangle normal

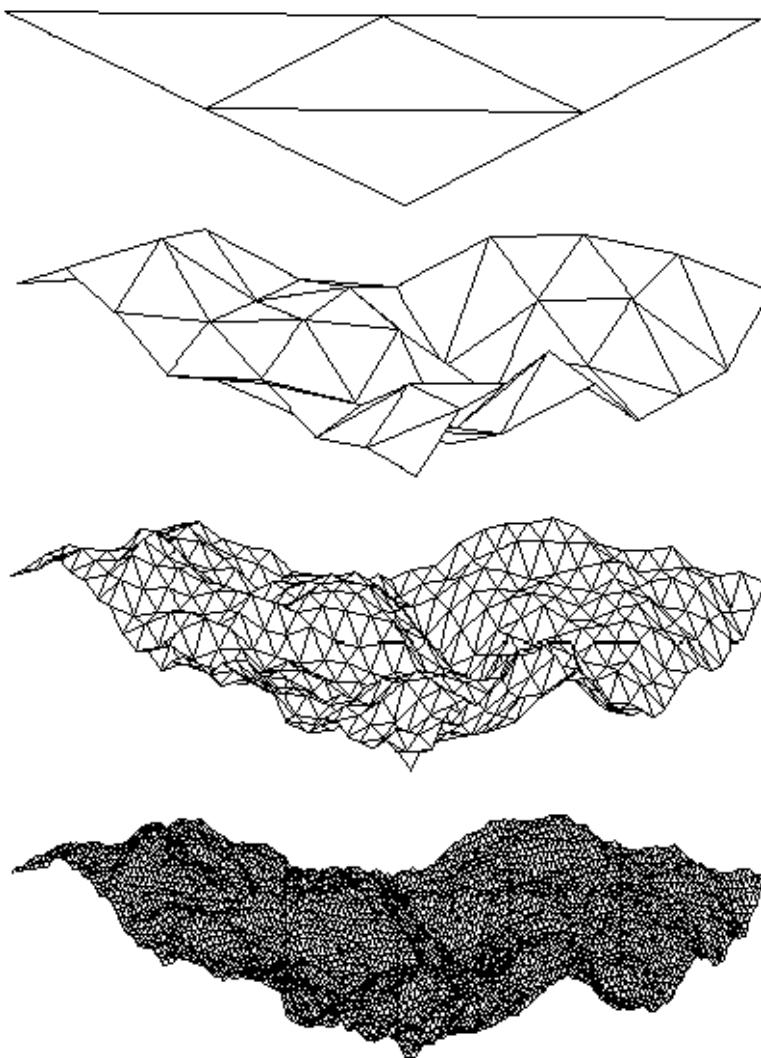
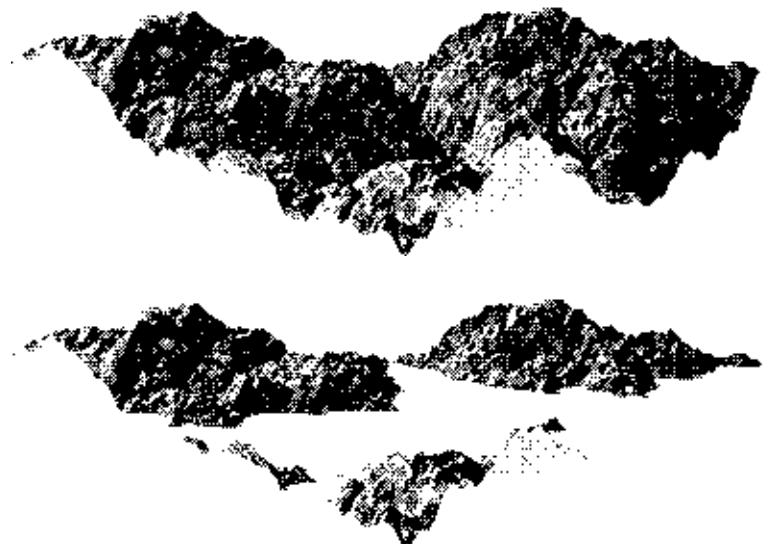


Courtesy <http://www.uni-paderborn.de/SFB376/projects/a2/zBufferMerging/>



Fractal Landscape Generation

- Base mesh
- Repeated subdivision & vertex displacement
- Shading
- + Water surface
- + Fog
- + ...



Courtesy <http://www.uwp.edu/academic/computer.science/morris.csci/CS.320/Week.11/Ch11b.www/Ch11b.html>



Fractal Landscape Ray Tracing

- Fractal terrain generated on-the-fly
- Problem: where is the ray-surface interaction ?
 - Triangle mesh not a-priori known
- Solution: bounding boxes
 - Maximum possible bounding box around each triangle
 - Decreasing displacement amplitude: finite bounding box
- Algorithm
 - Intersect ray with bounding box
 - Subdivide corresponding triangle
 - Compute bounding boxes of 4 new triangles
 - Test against 4 new bounding boxes
 - Iterate until termination criterion fulfilled (LOD / pixel size)



Questions

- What spaces are involved in texture mapping?
- What is forward and inverse mapping?
- Why do we need filtering in texture mapping?
- How is a mip map organized?
- What are bump mapping and reflection mapping?



Wrap-Up

- Texture Parameterization
 - mapping vertices to texture coordinates
 - problem of perspective transformation
- Procedural Textures
 - Noise
 - Turbulence
 - Fractal Landscapes
- Texture Filteringung
 - Minification
 - Magnification
 - EWA
 - Summed Area Tables
 - MipMaps