



# Computer Graphics (Graphische Datenverarbeitung)

## - Image Processing -

WS 2021/2022



# Corona

---

- Regular random lookup of the 3G certificates
- Contact tracing: We need to know who is in the class room
  - New ILIAS group for every lecture slot
  - Register via ILIAS or this QR code (only if you are present in this room)





# Aliasing

- In Fourier space

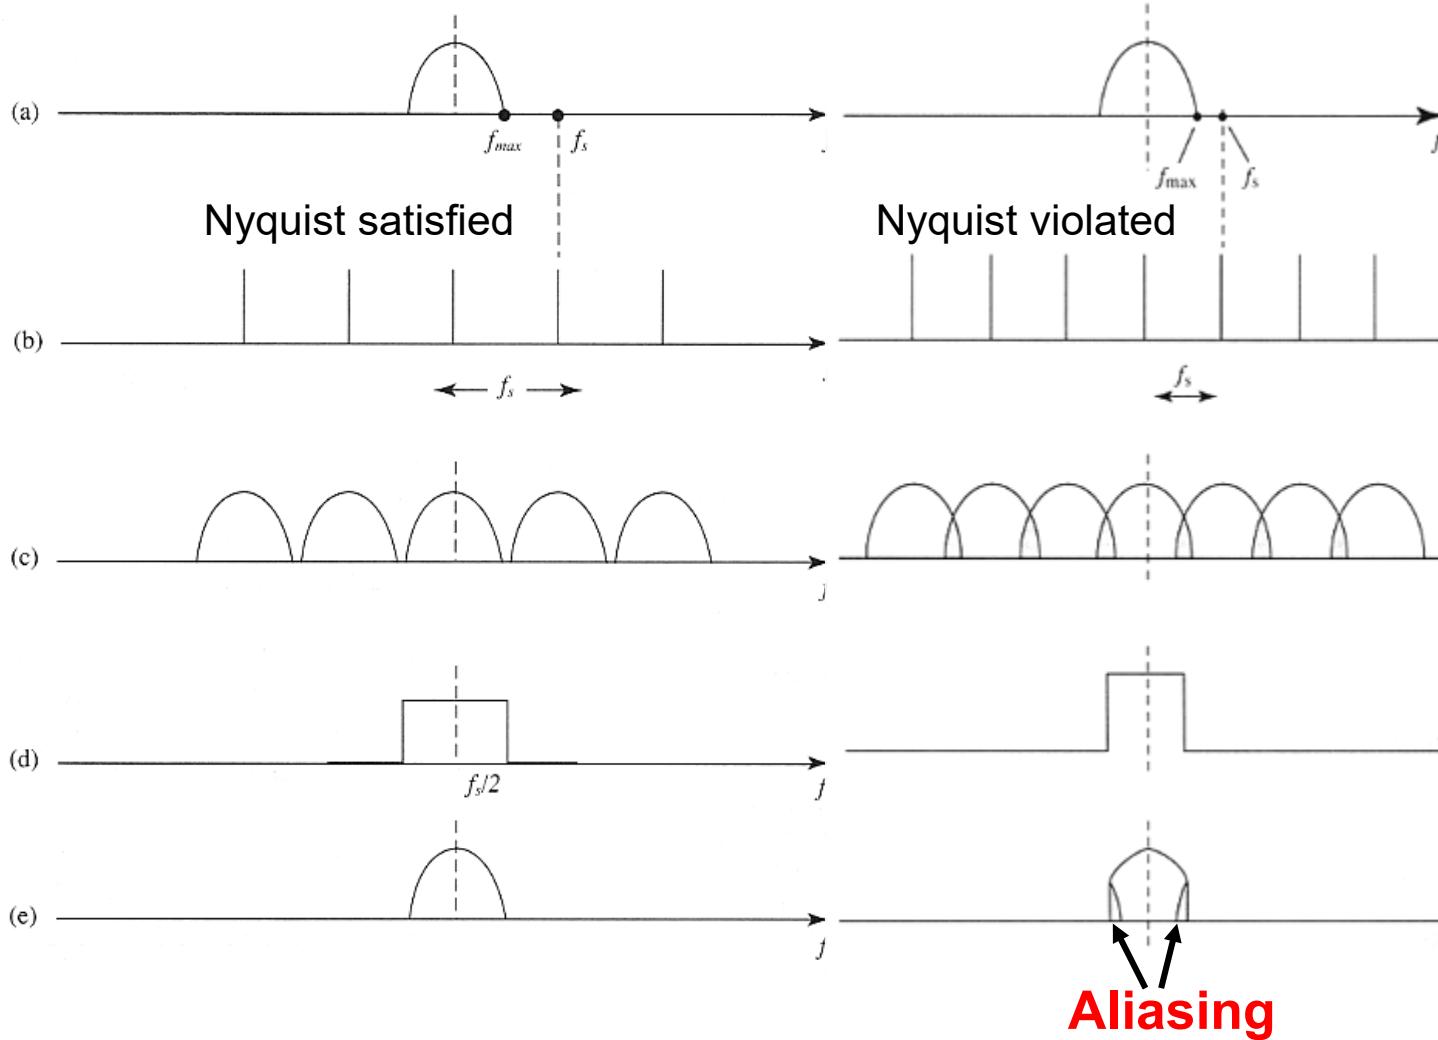
- Original spectrum

- Sampling comb

- Resulting spectrum

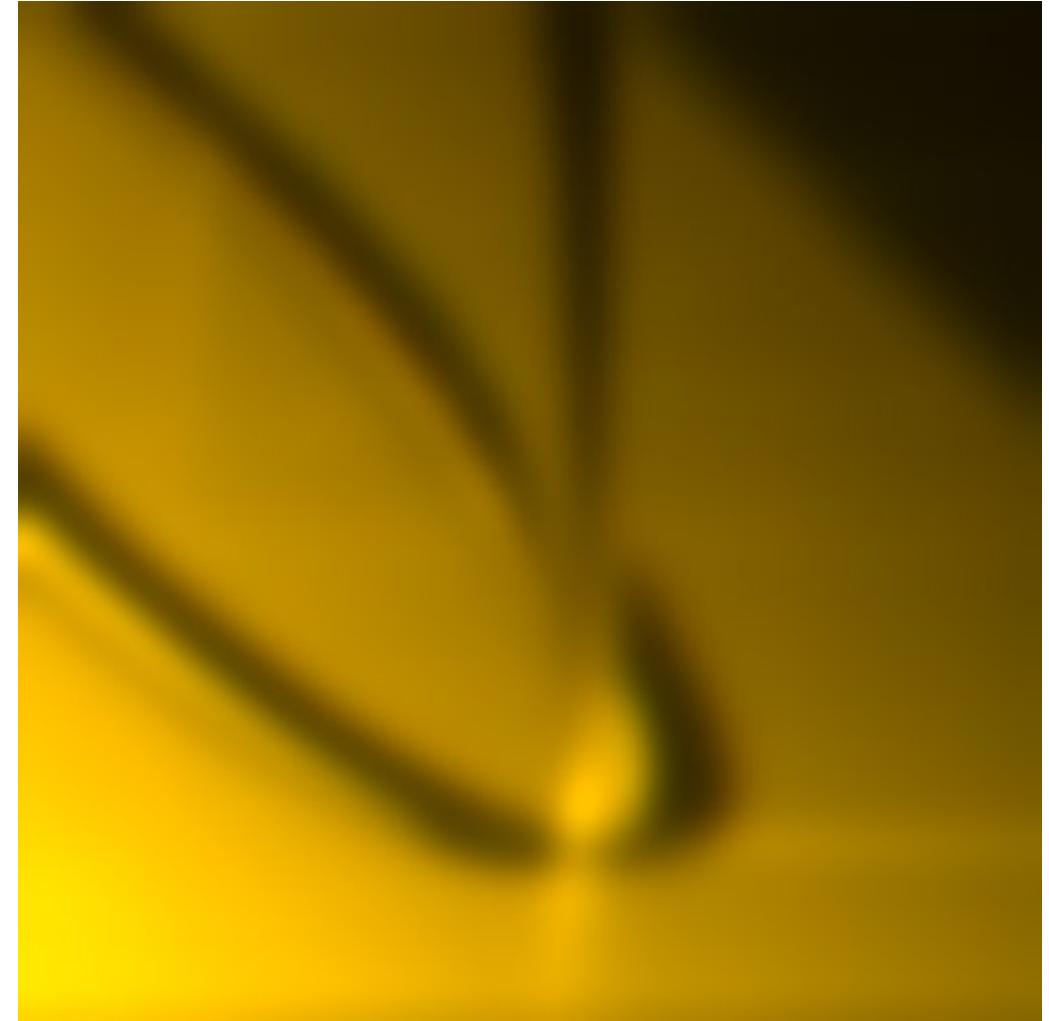
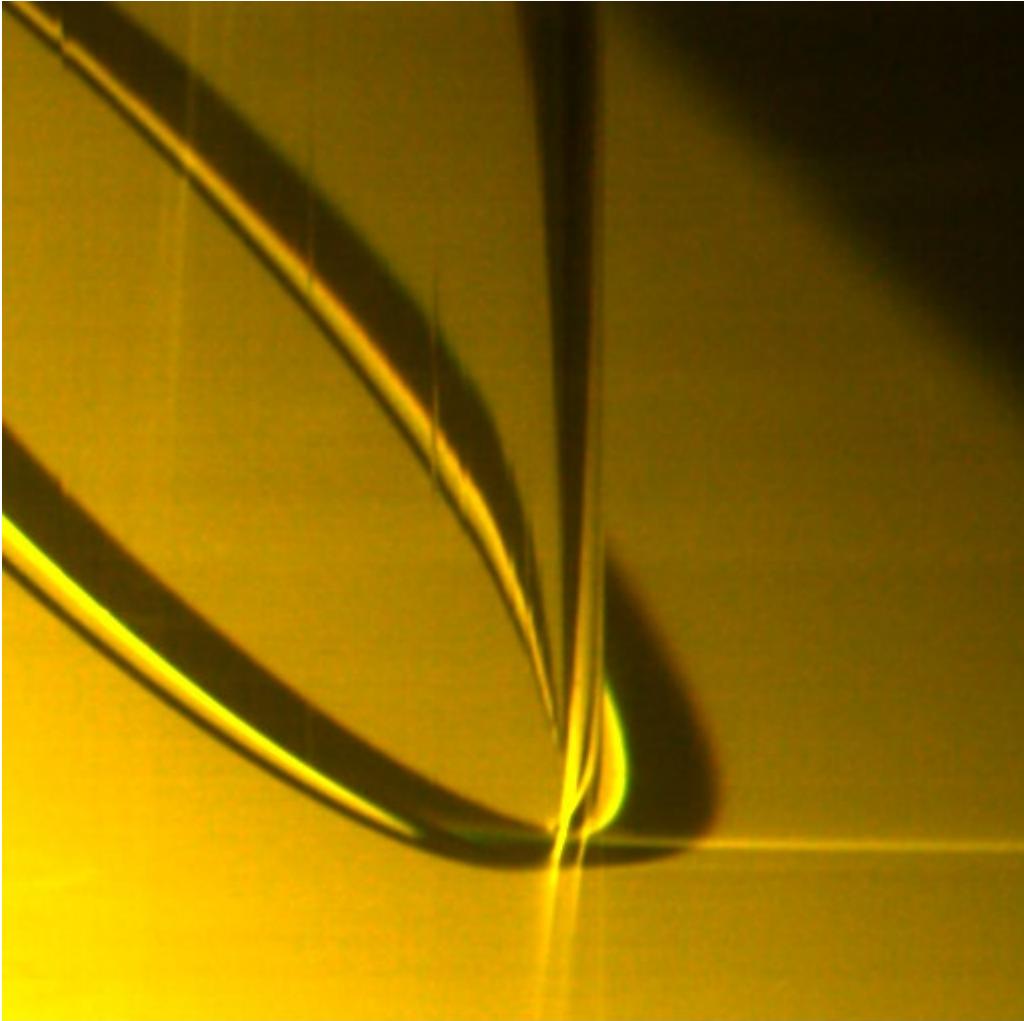
- Reconstruction Filter

- Reconstructed spectrum





# What kind of filter operation?





---

# Which Filters do you know?



# Overview

---

- Last lecture
  - Sampling
  - Antialiasing & supersampling
- Now
  - Image Filters
  - Wavelets
- Next
  - Human Visual System



# What you should learn today

---

- Definition of standard filter kernels
- Properties of Wavelets
- Edge-Avoiding Wavelets
- When to apply which filter
- How to compute a filtered image

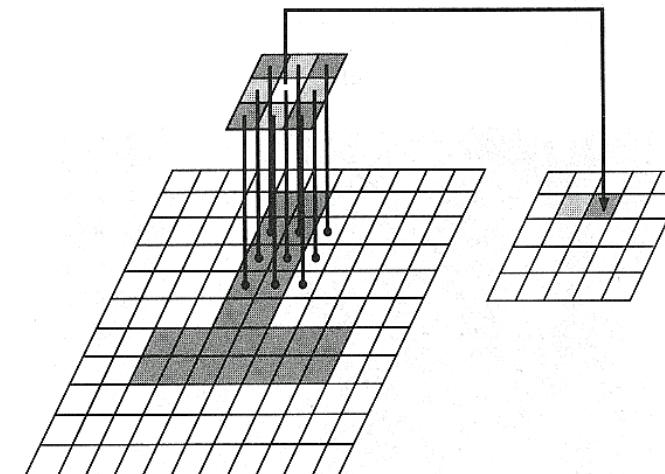
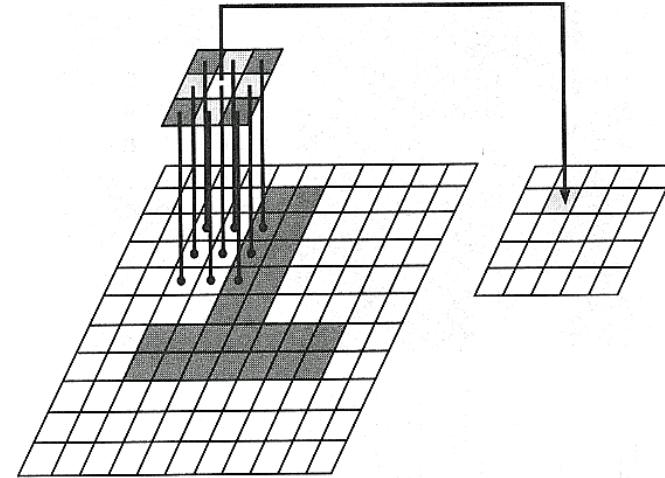


# Image Filters



# Convolution and Filtering

- Technical Realization
  - In image domain
  - Pixel mask with weights
  - OpenGL: Convolution extension
- Problems (e.g. sinc)
  - Large filter support
    - Large mask
    - A lot of computation
  - Negative weights
    - Negative light?





# Gaussian Filter

---

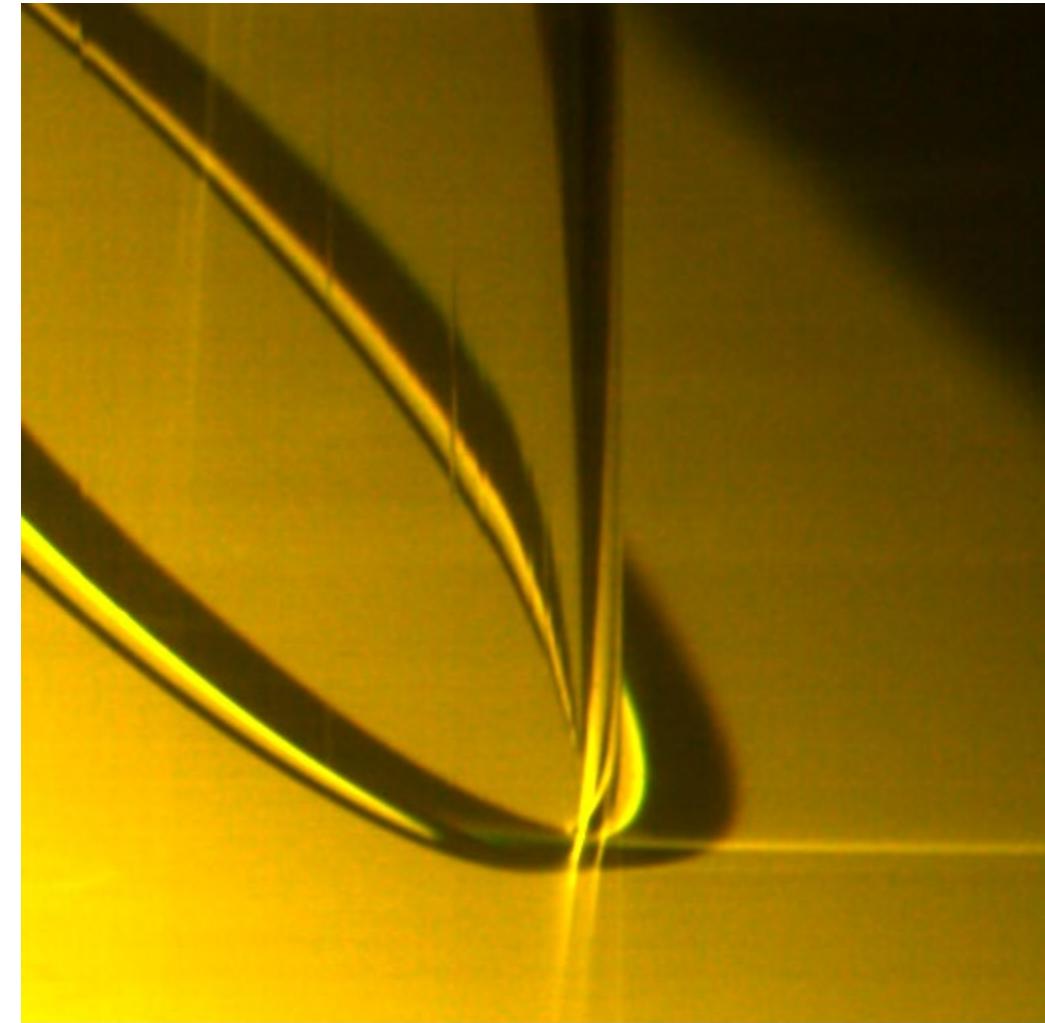
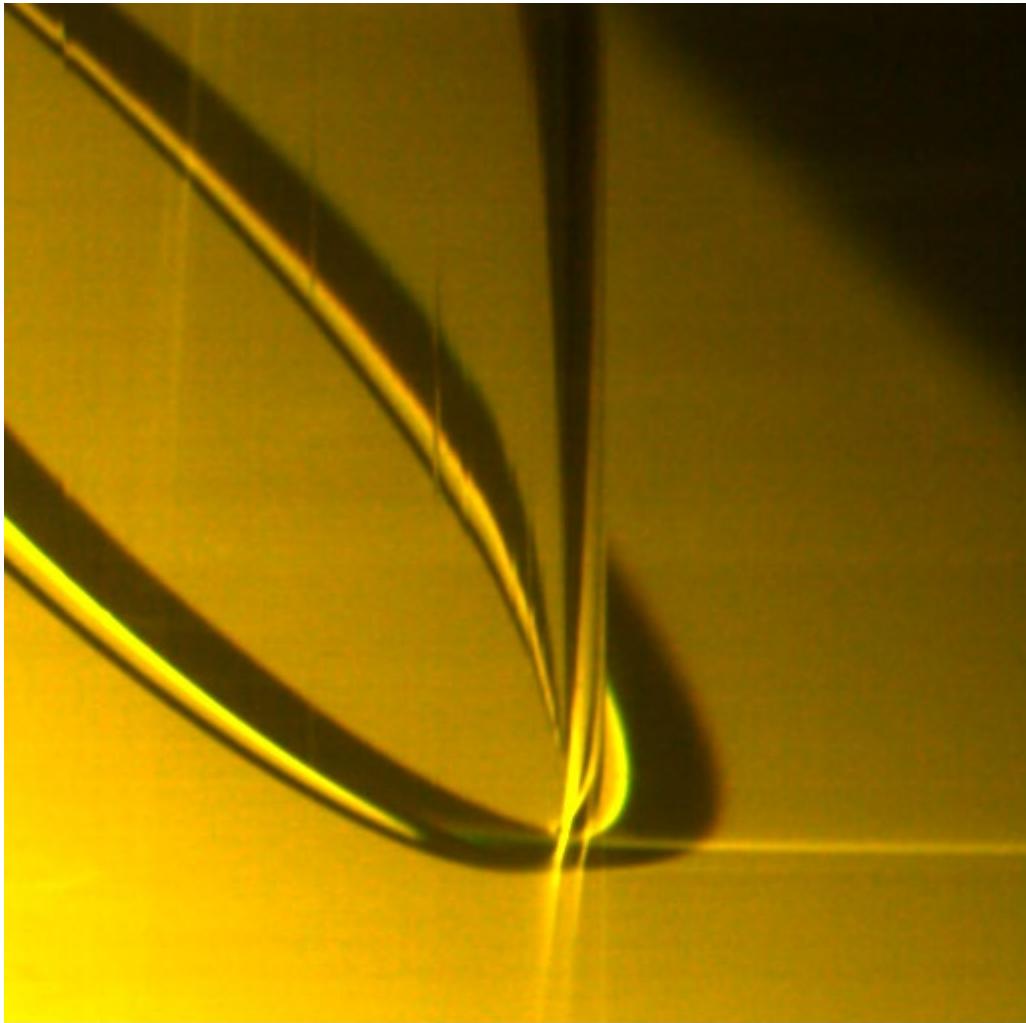
- Convolution with a 2D Gaussian

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

- variance  $\sigma$  is defining the “blur”
- Gaussian approximates a low-pass filter
- It smoothes the image

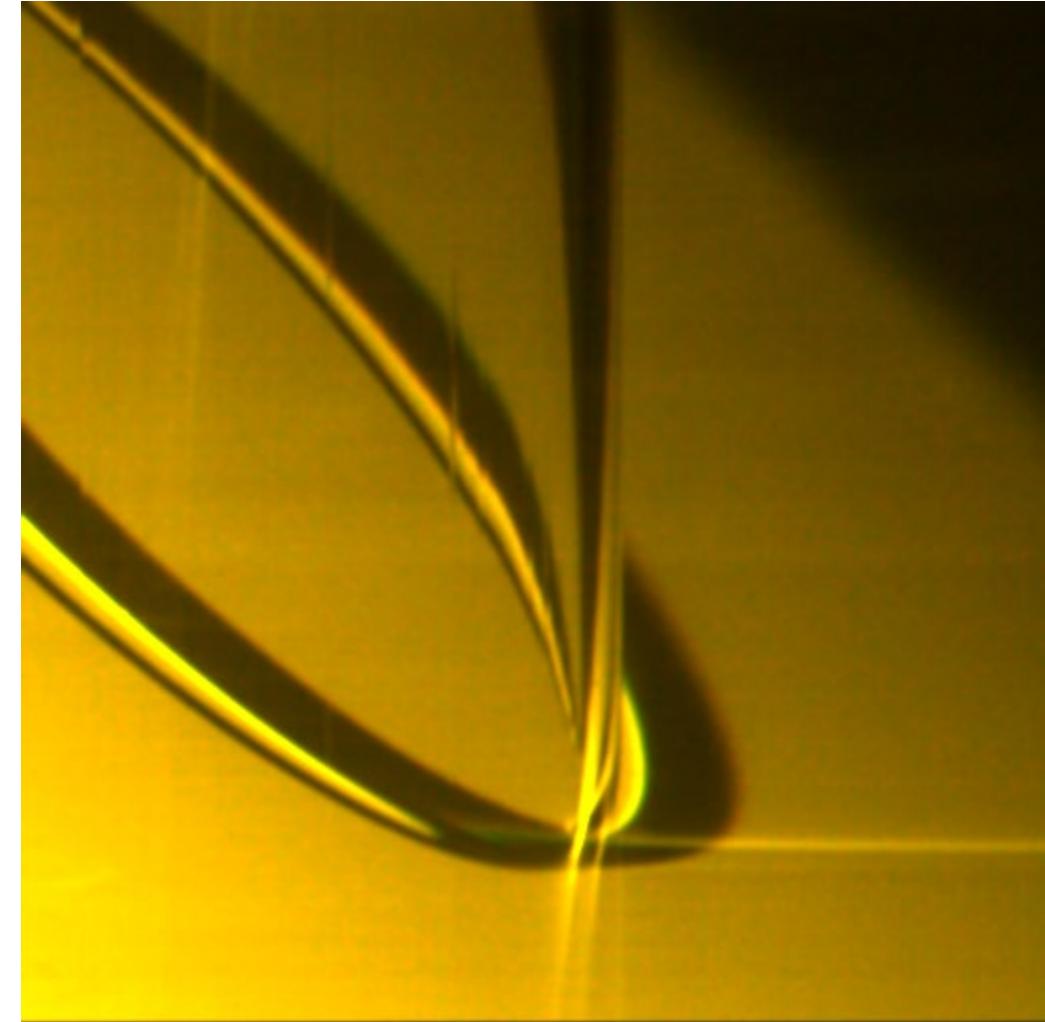
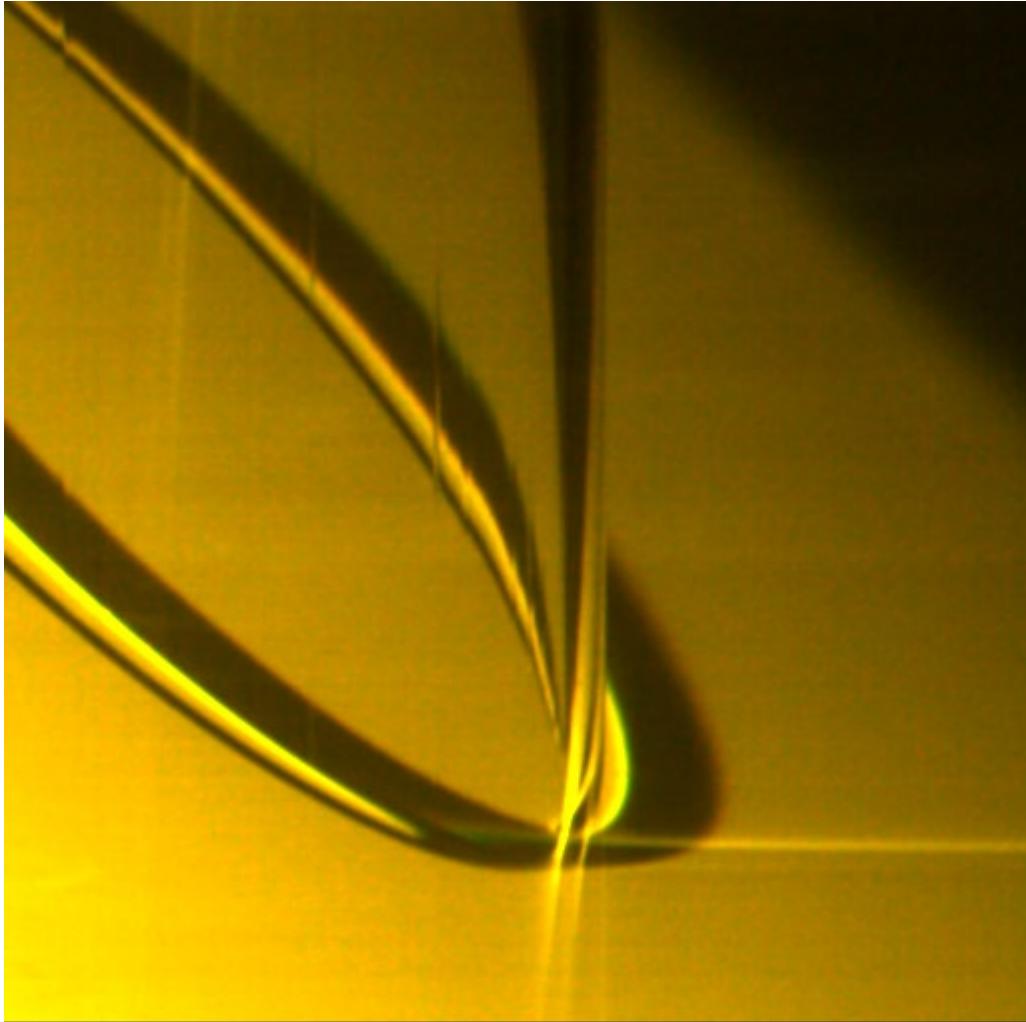


# Gaussian – $\sigma=0.2$



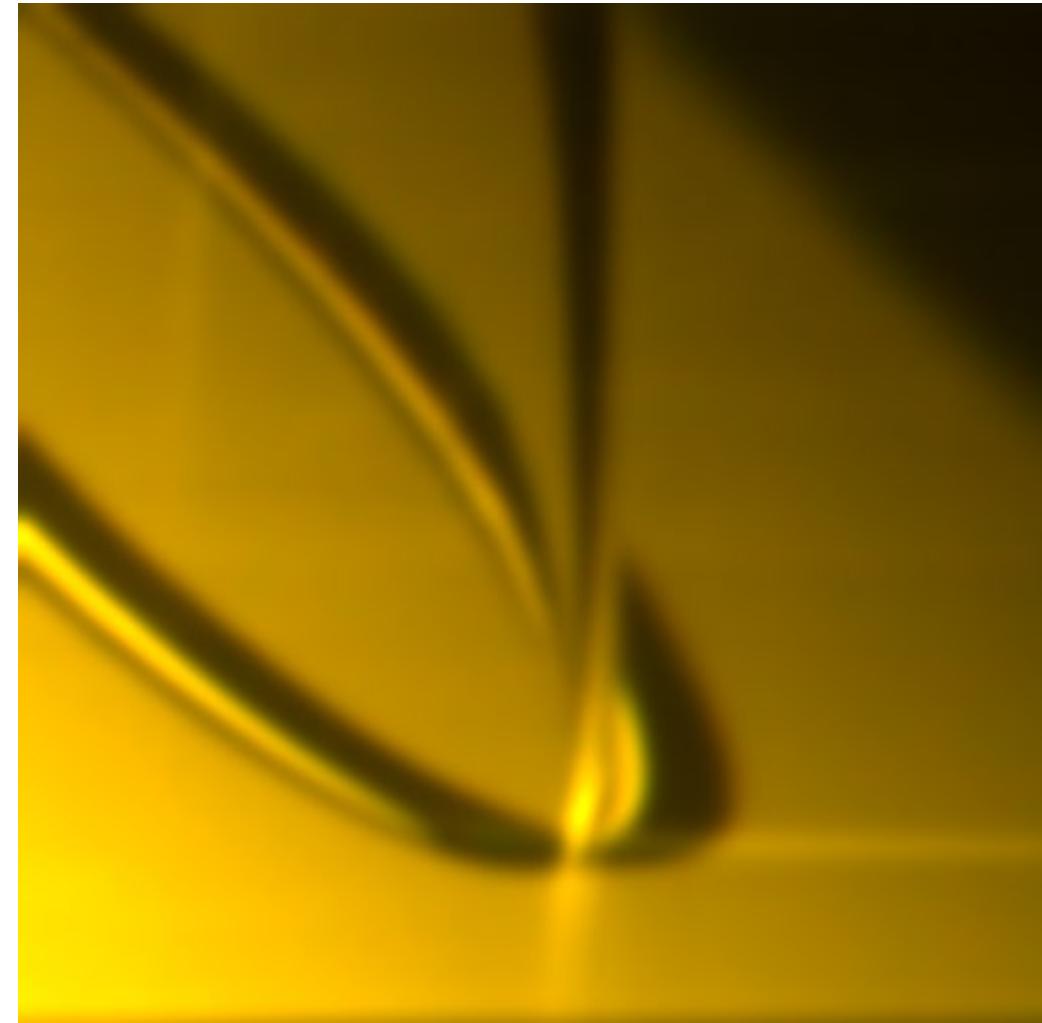
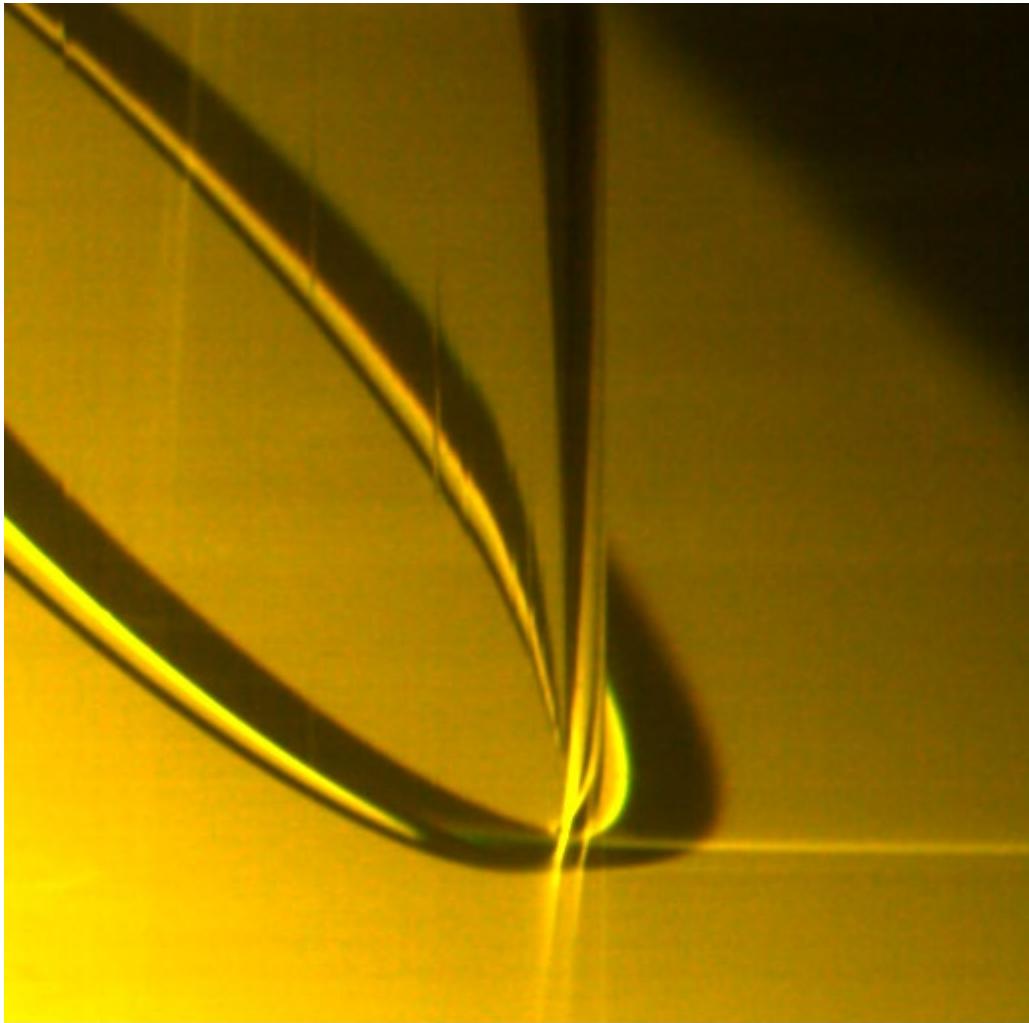


# Gaussian – $\sigma=0.8$



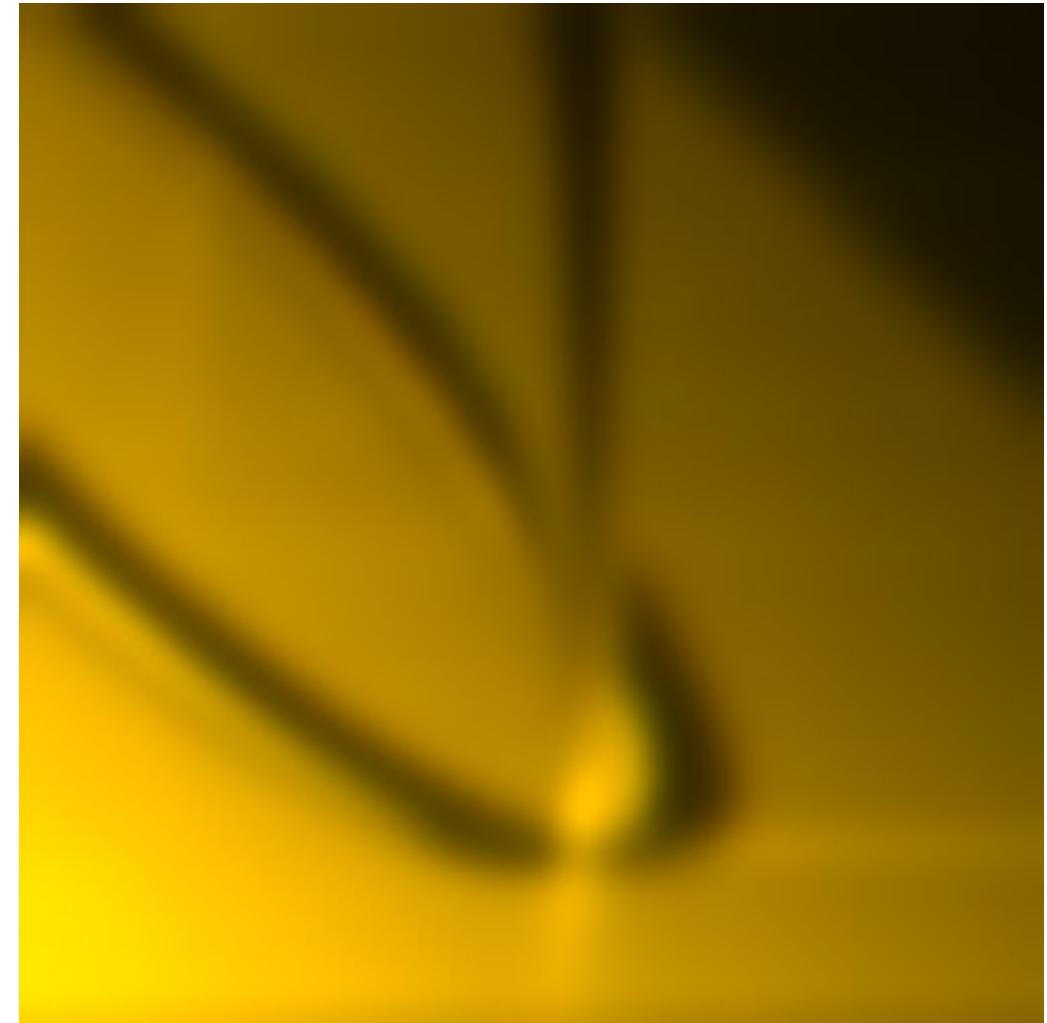
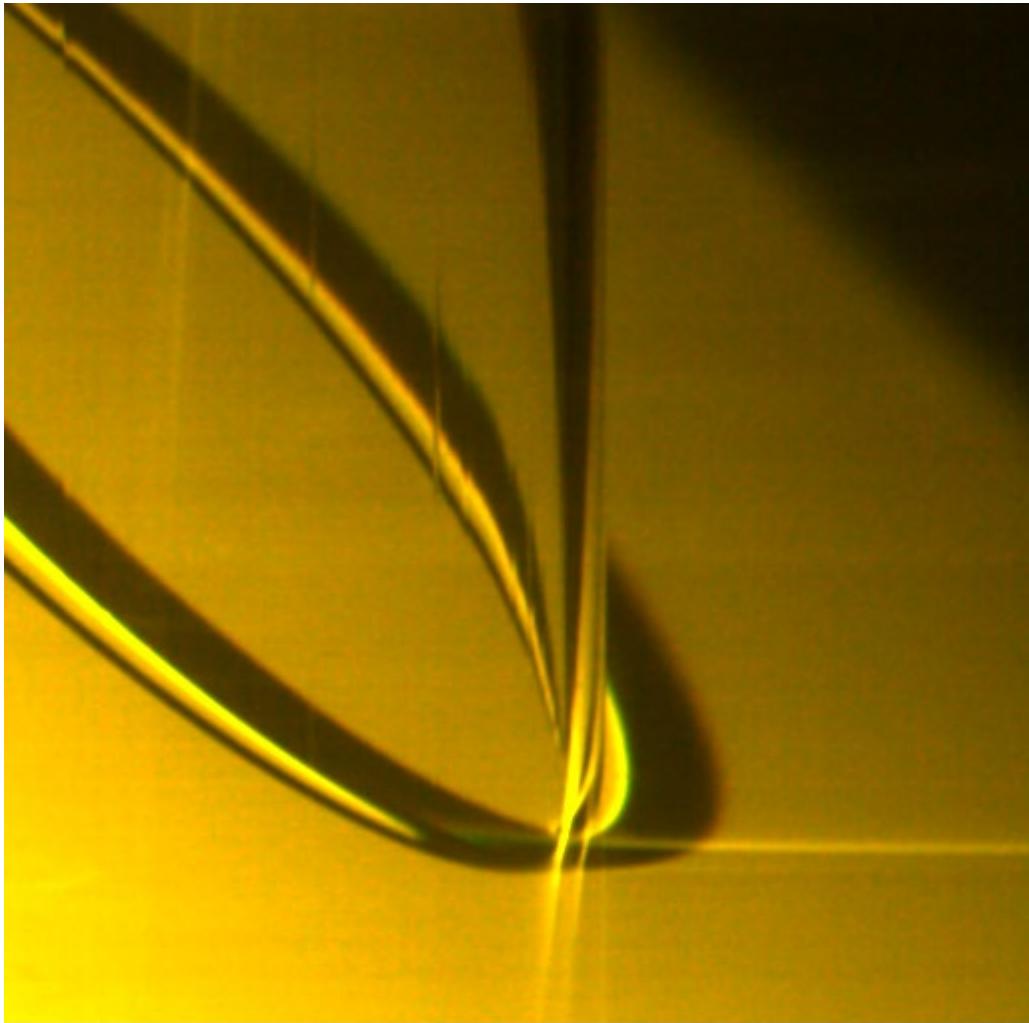


# Gaussian – $\sigma=5$





# Gaussian – $\sigma=10$



# Gaussian - Implementation

- 2D filter mask  $\sigma = 0.84$
- Width of the mask should follow  $3\sigma$  rule
  - at  $3\sigma$  the value of the Gaussian is less than 0.1%.

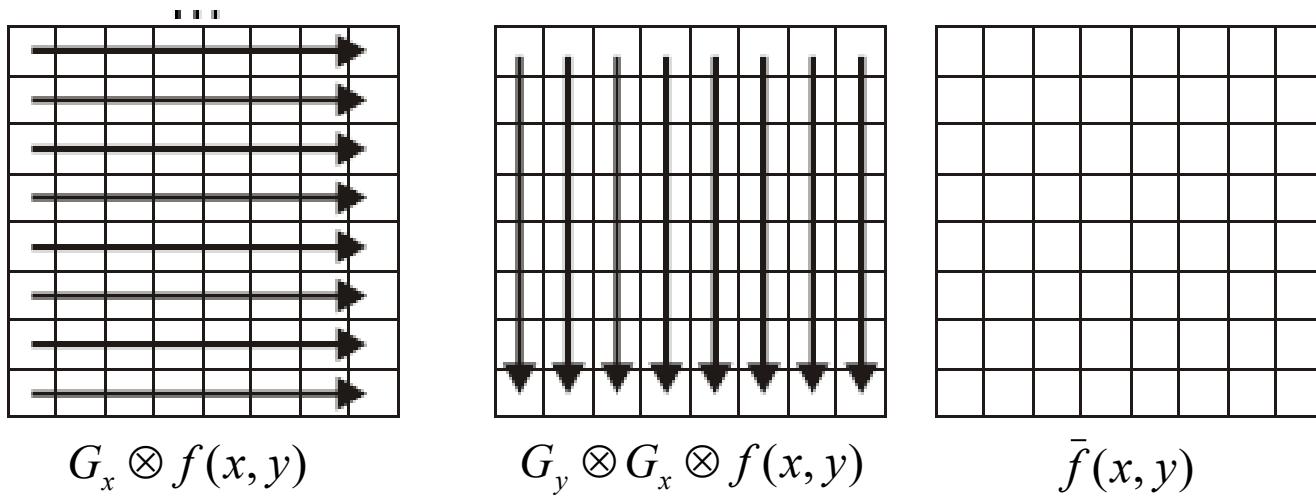
0.00000067	0.00002292	<b>0.00019117</b>	0.00038771	<b>0.00019117</b>	0.00002292	0.00000067
0.00002292	0.00078633	0.00655965	0.01330373	0.00655965	0.00078633	0.00002292
<b>0.00019117</b>	0.00655965	0.05472157	0.11098164	0.05472157	0.00655965	<b>0.00019117</b>
0.00038771	0.01330373	0.11098164	<b>0.22508352</b>	0.11098164	0.01330373	0.00038771
<b>0.00019117</b>	0.00655965	0.05472157	0.11098164	0.05472157	0.00655965	<b>0.00019117</b>
0.00002292	0.00078633	0.00655965	0.01330373	0.00655965	0.00078633	0.00002292
0.00000067	0.00002292	<b>0.00019117</b>	0.00038771	<b>0.00019117</b>	0.00002292	0.00000067

# Gaussian - Implementation

- Gaussian filter is separable

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\bar{f}(x, y) = G_{(x,y)} \otimes f(x, y) = G_y \otimes (G_x \otimes f(x, y))$$





# Unsharp Masking

---

- Sharpen the image by subtracting the low frequency (unsharpened or blurred) component

$$U(x, y) = G_{(x, y)} \otimes f(x, y)$$

$$\bar{f}(x, y) = \frac{c}{2c-1} f(x, y) - \frac{(1-c)}{2c-1} U(x, y)$$



# Unsharp Masking

s1



- s2



=



sharpened



# Sobel Operator

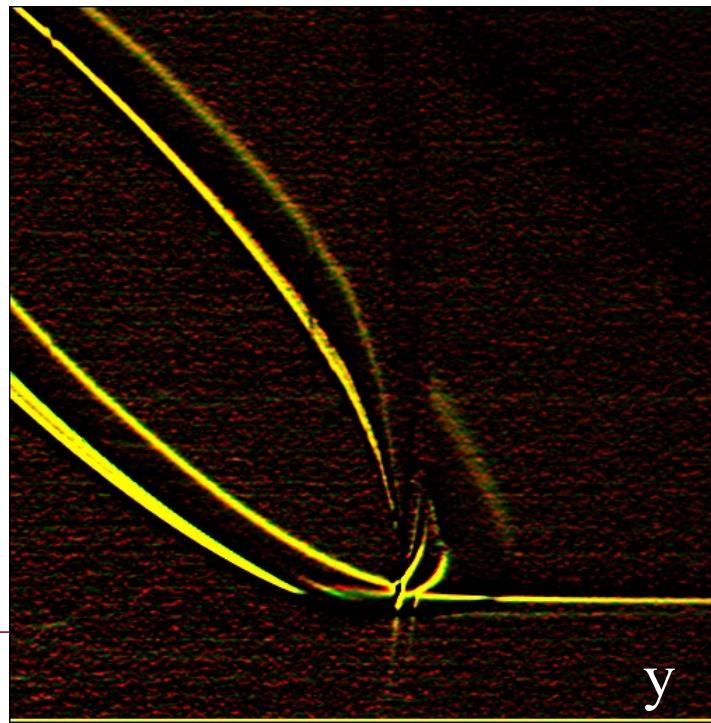
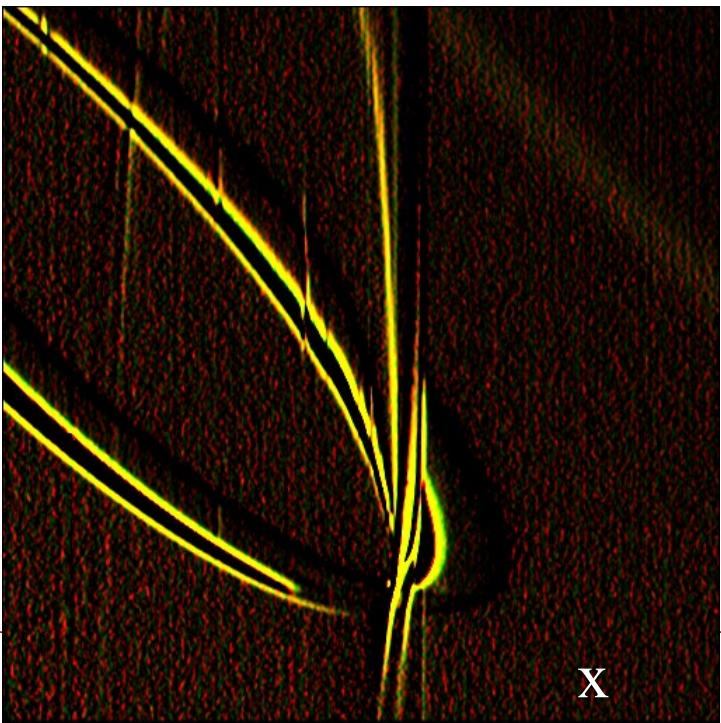
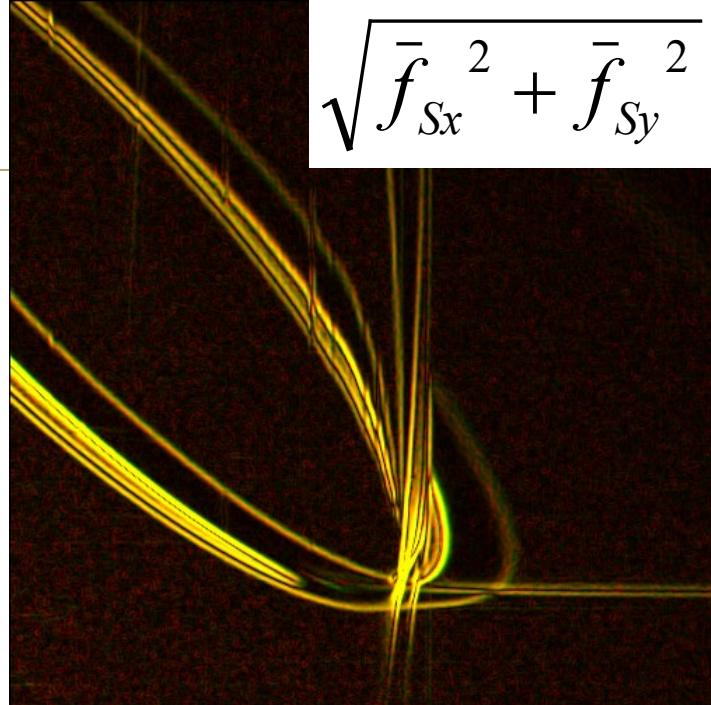
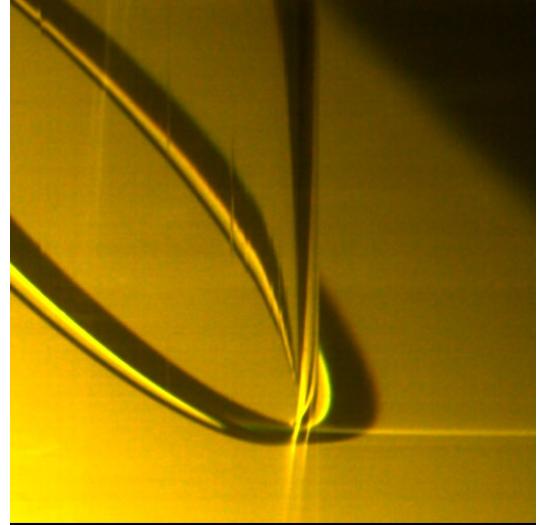
- A simple edge detection filter for horizontal and vertical edges

$$S_x = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \quad S_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

- Arbitrary direction (high-pass filter) approximated by

$$\sqrt{\bar{f}_{Sx}^2 + \bar{f}_{Sy}^2}$$

# Sobel Operator





# Non-linear Filters: Median

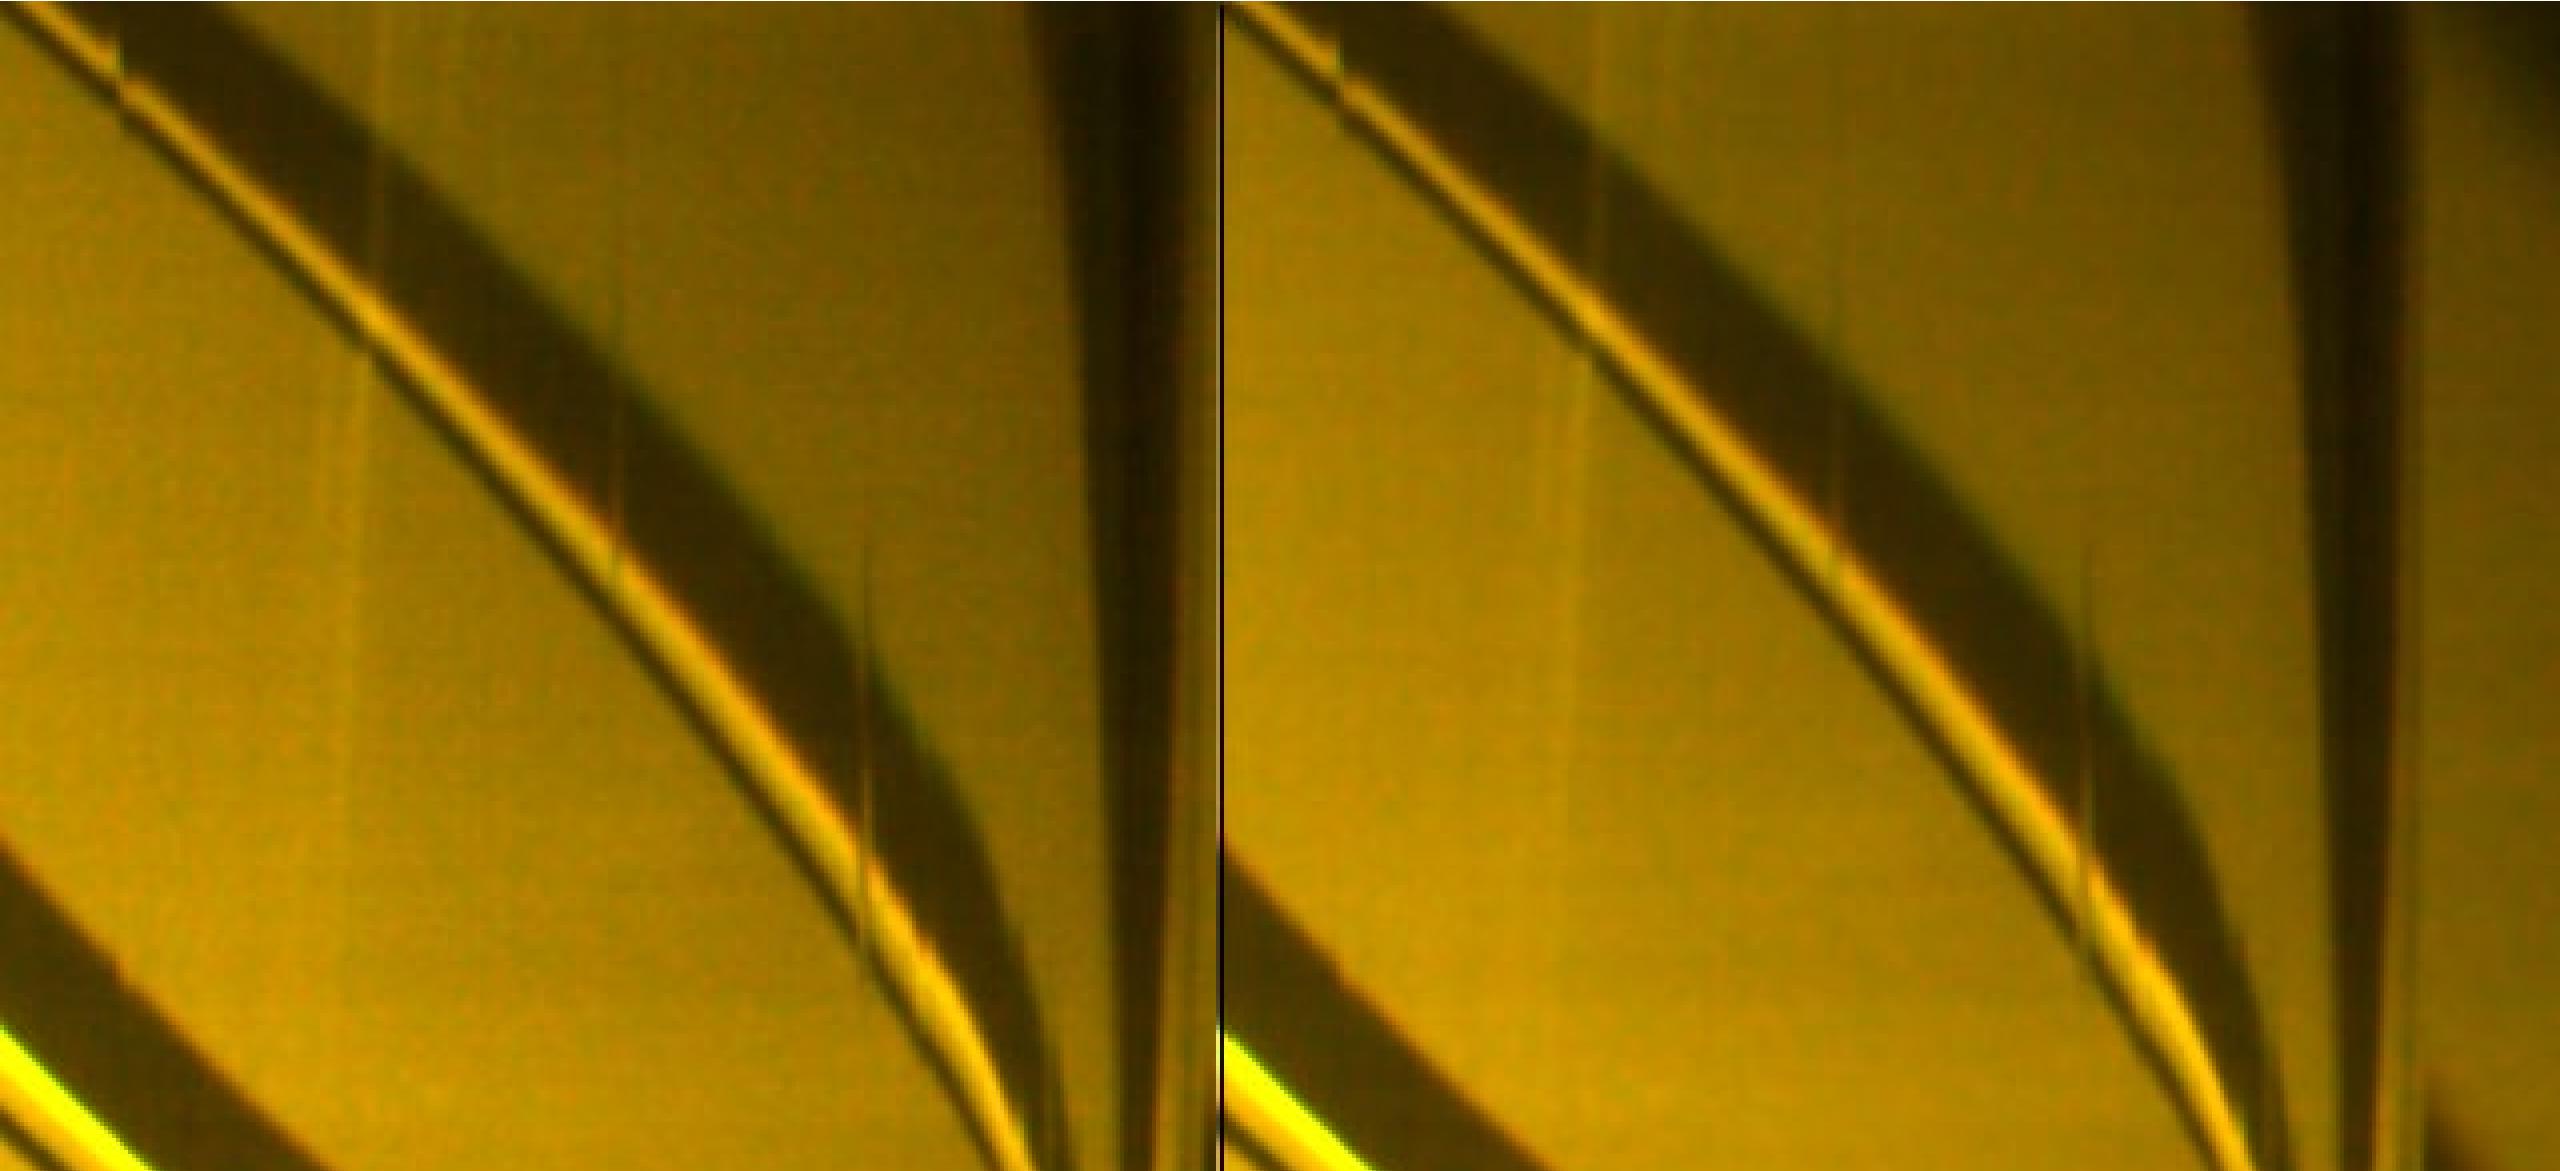
---

- Take all samples in under the filter mask, sort them
- Store back the median, element at position ( $window^*window/2$ )
- Nicely removes outliers and “sticky pixels”
  
- Non-separable but lots and lots of accelerations



# Median 3x3

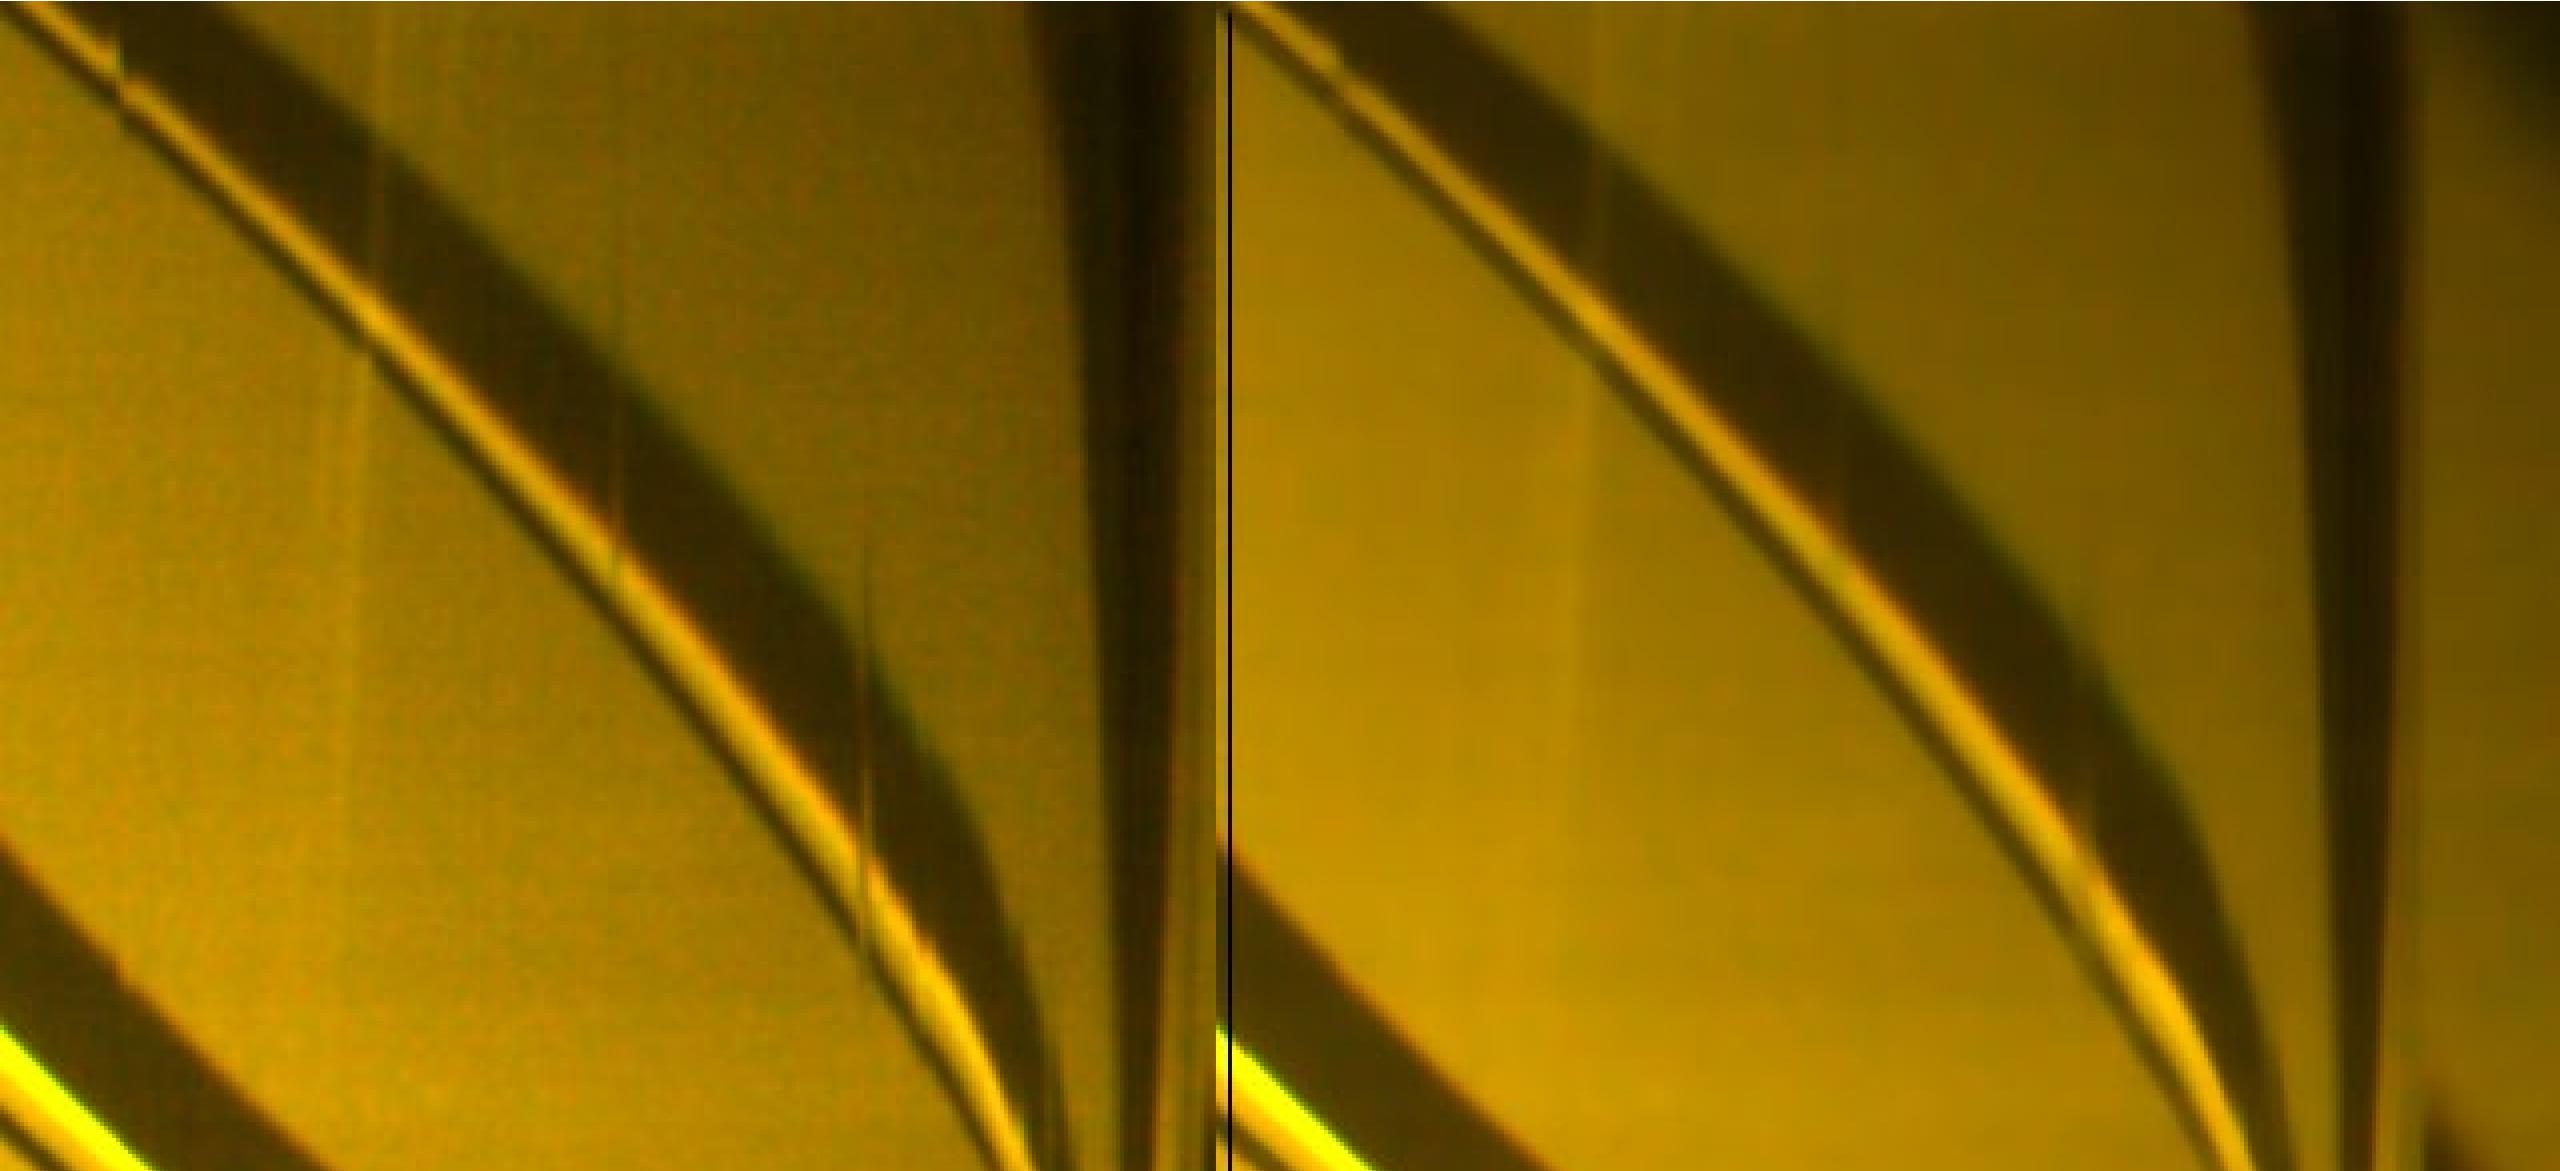
---





# Median 7x7

---





# Bilateral Filter

# Bilateral Filtering

---

- Gaussian filter kernel - spatial distance

$$g(p_1, p_2) = e^{-\frac{1}{2} \left( \frac{d(p_1 - p_2)}{\sigma_g} \right)^2}$$

- Bilateral filter kernel – spatial and value distance

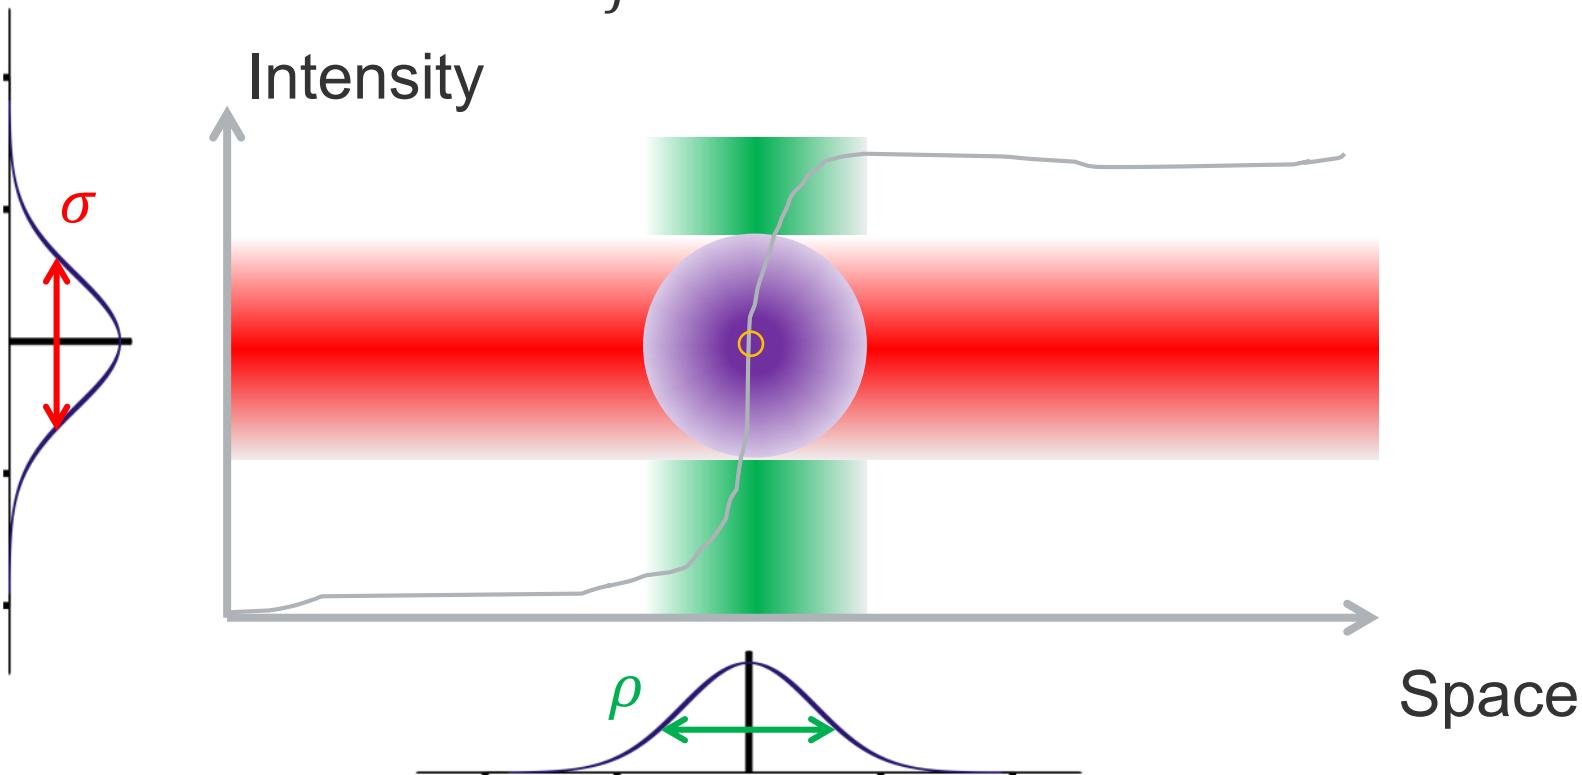
$$b(p_1, p_2) = g(p_1, p_2) \cdot e^{-\frac{1}{2} \left( \frac{d(f(p_1) - f(p_2))}{\sigma_b} \right)^2}$$

- prevents blurring across edges
- for two points  $p_1, p_2$  under the filter mask and
- some distance function  $d$



# Bilateral Filtering

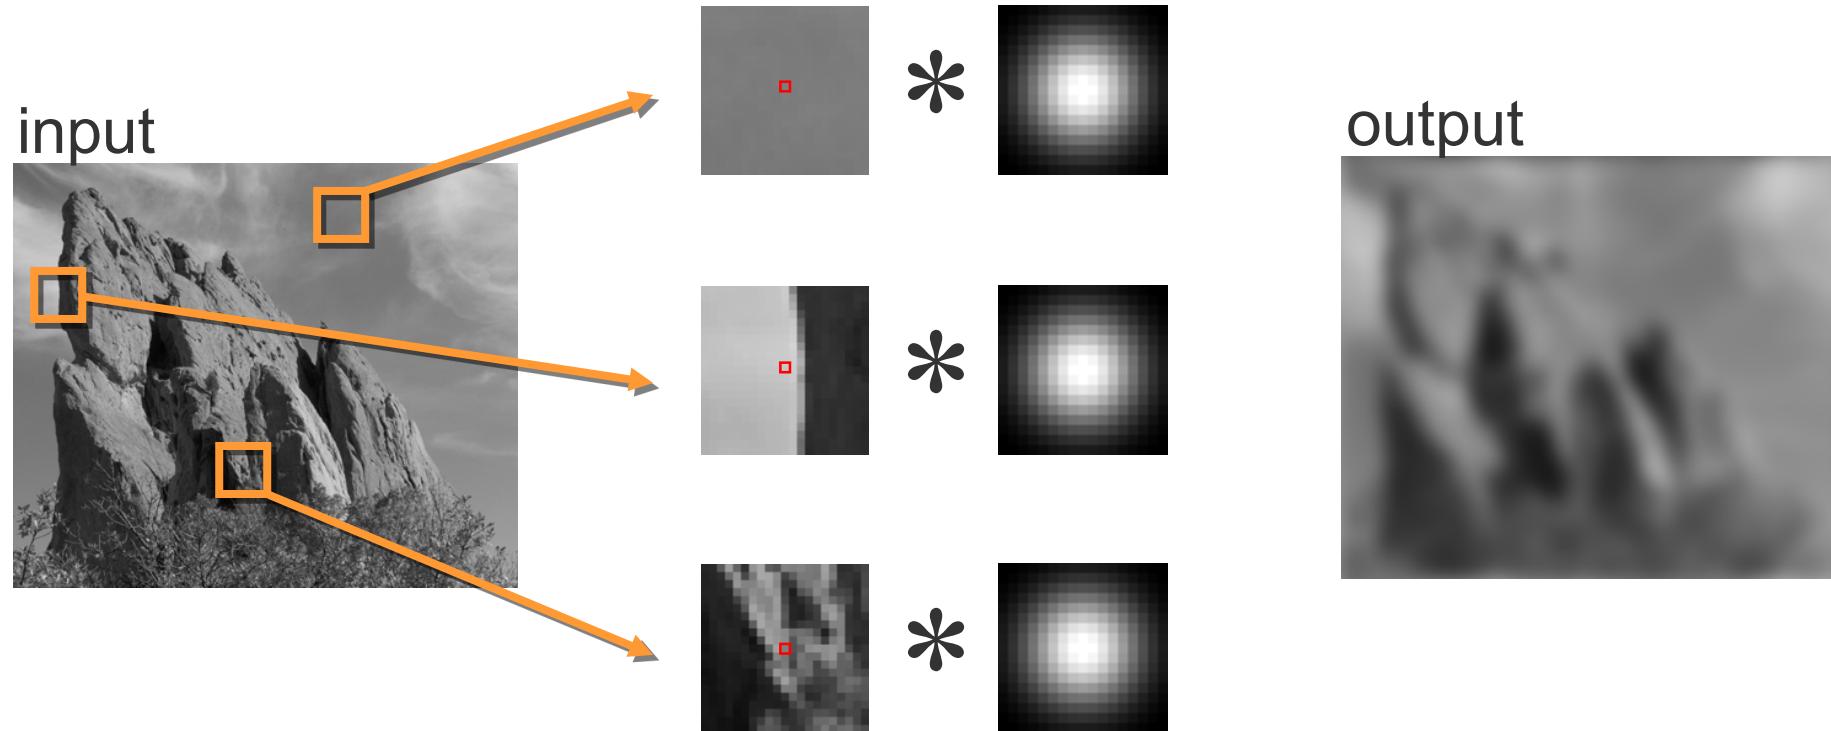
$$\hat{x}(i) = \frac{1}{C_i} \sum_j y(j) e^{-\frac{\|i-j\|^2}{2\rho^2}}$$



Smith and Brady  
(1997)

# Gaussian Smoothing

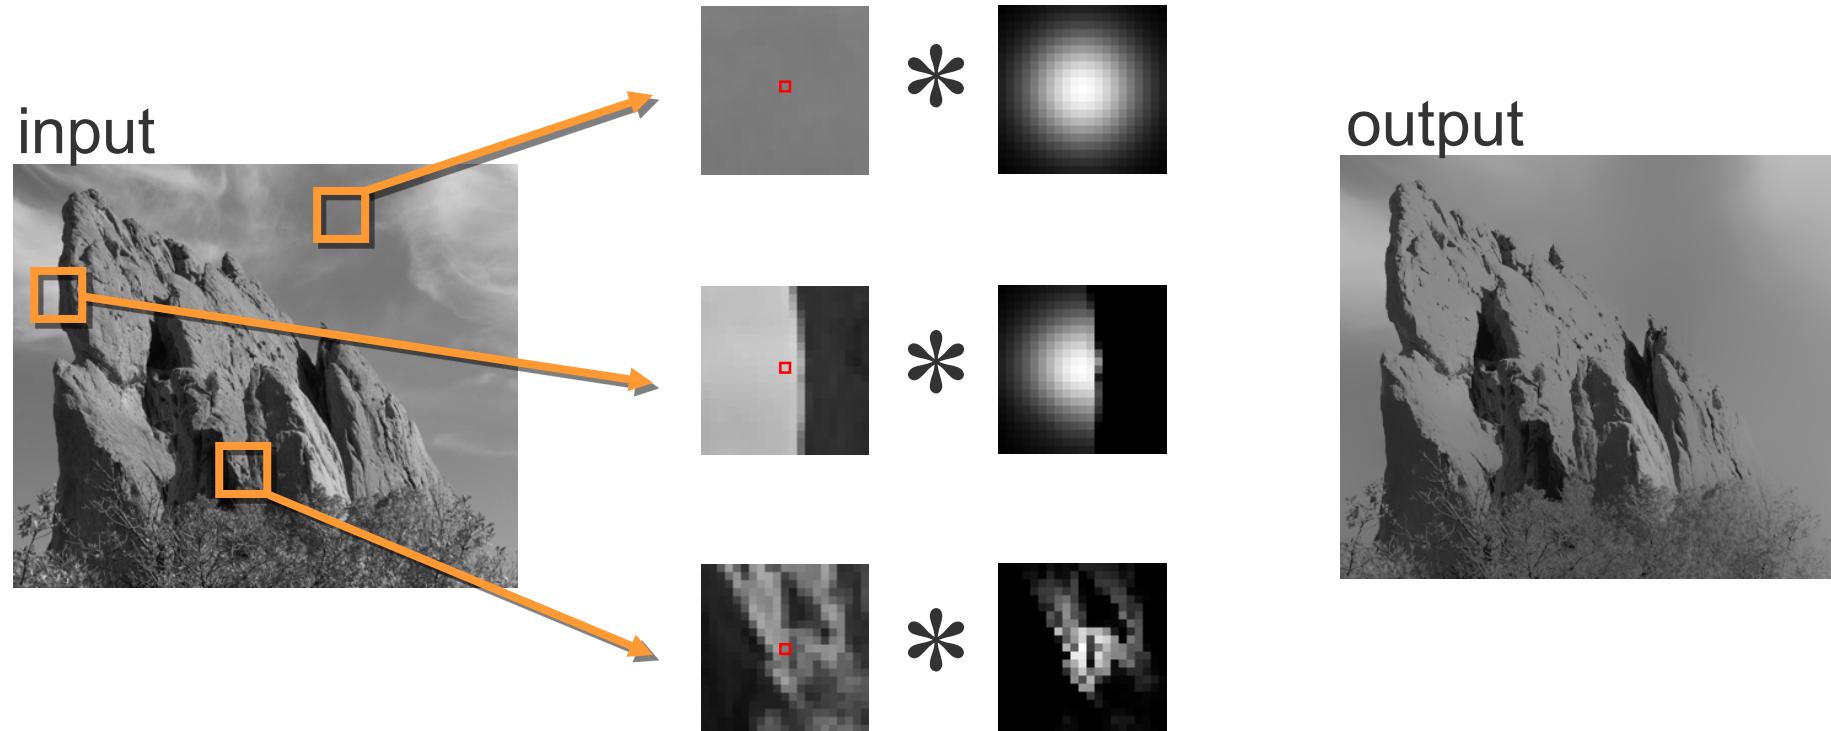
Slides taken from Sylvain Paris,  
Siggraph 2007



Same Gaussian kernel everywhere  
Averages across edges  $\Rightarrow$  blur

# Bilateral Filtering

Slides taken from Sylvain Paris,  
Siggraph 2007



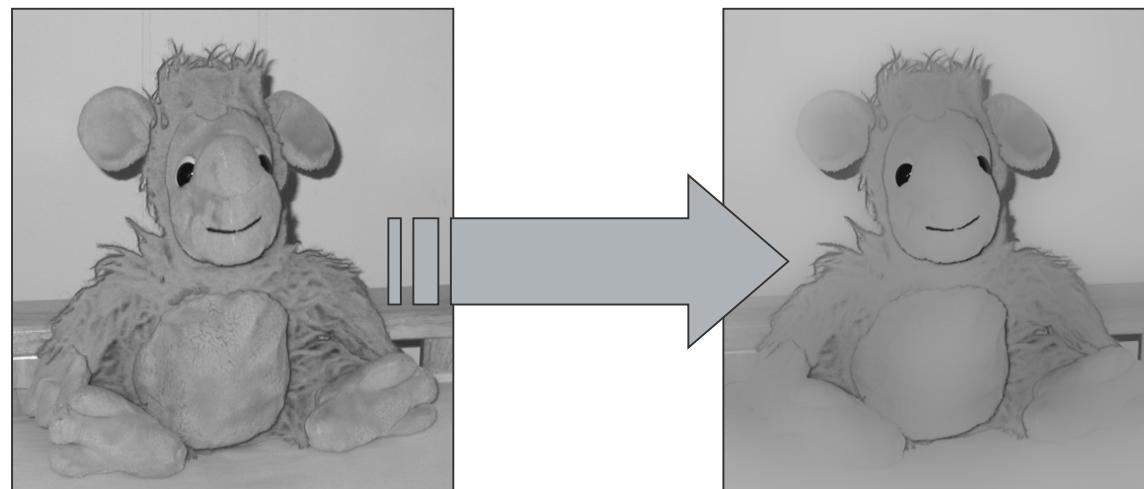
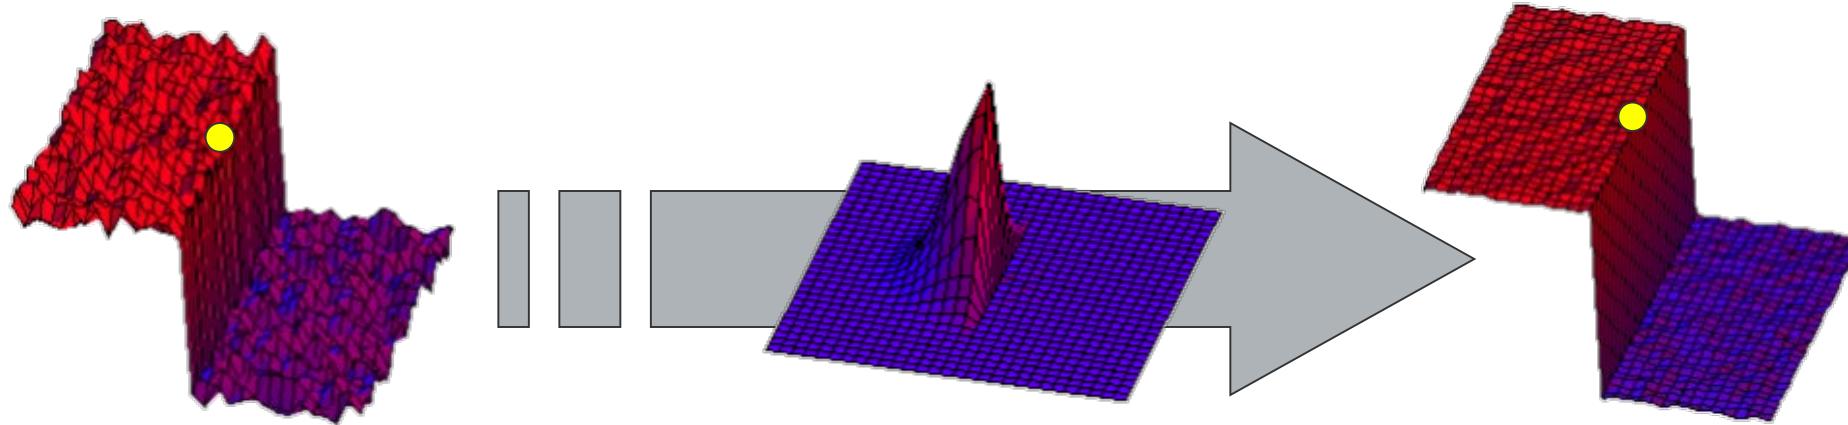
Kernel shape depends on image content  
Avoids averaging across edges



# Large-scale Layer

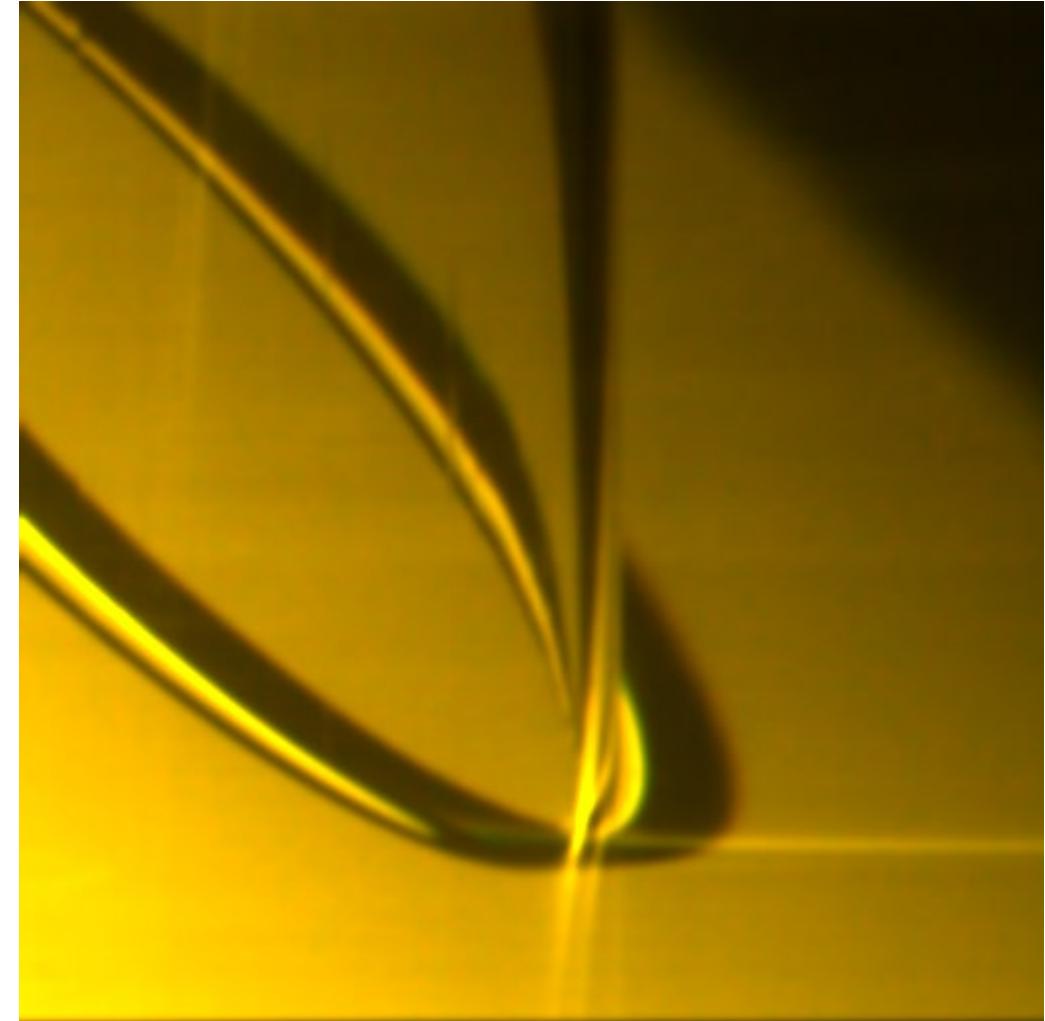
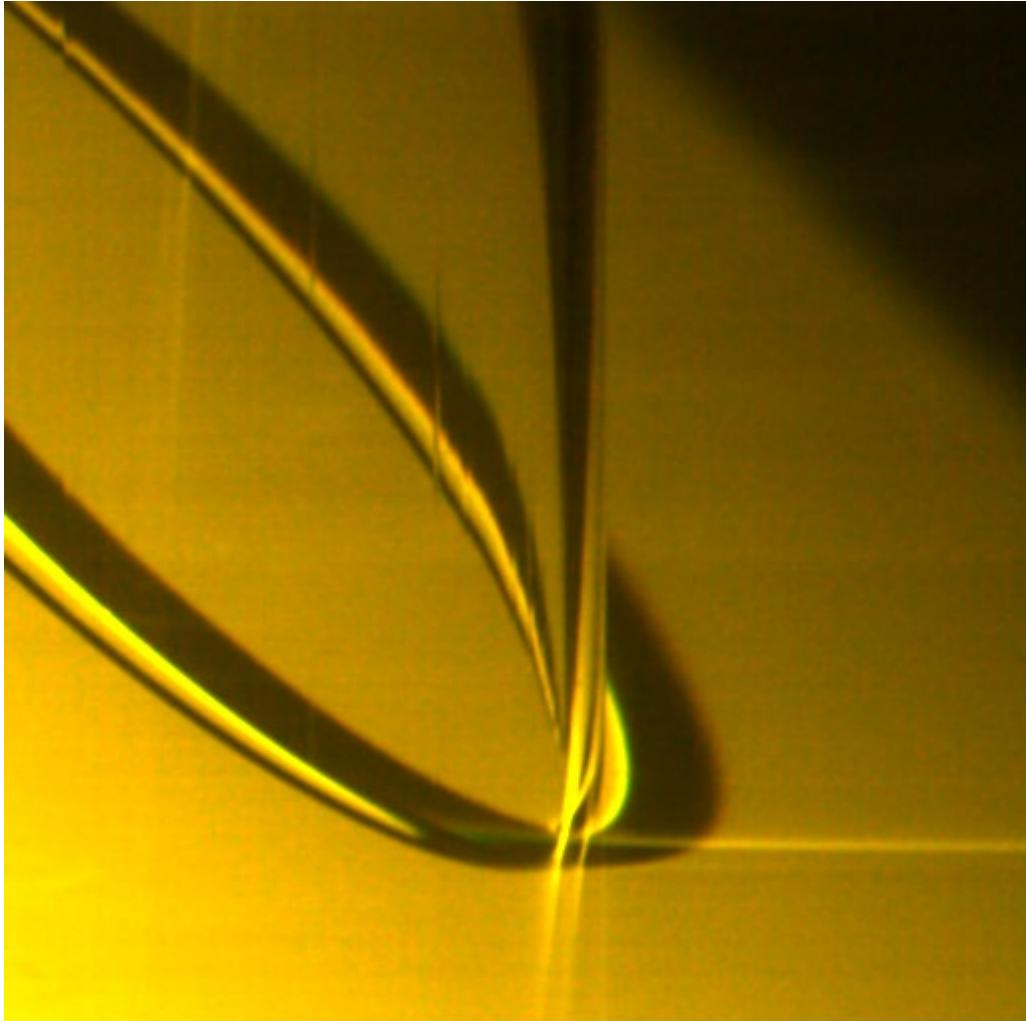
- Bilateral filter – edge preserving filter

Smith and Brady 1997; Tomasi and Manducci 1998; Durand et al. 2002



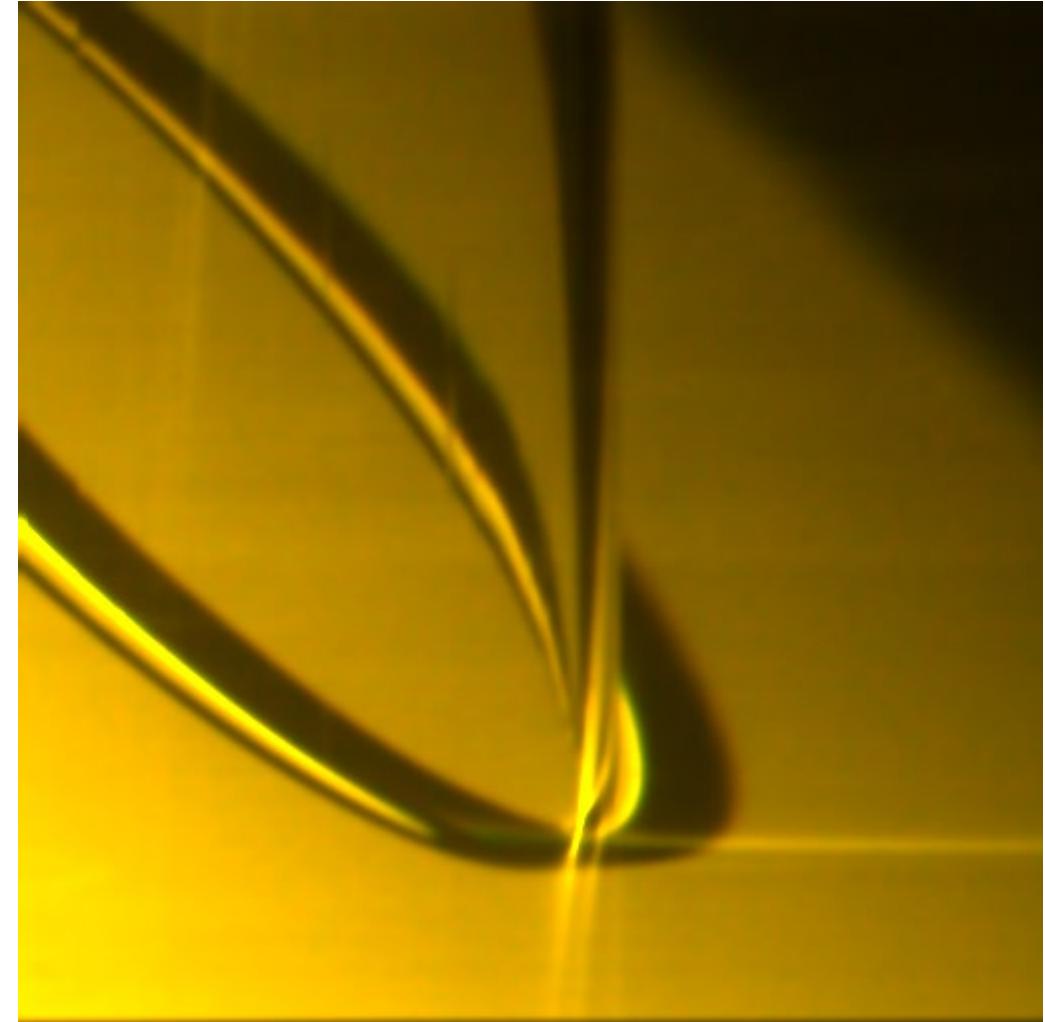
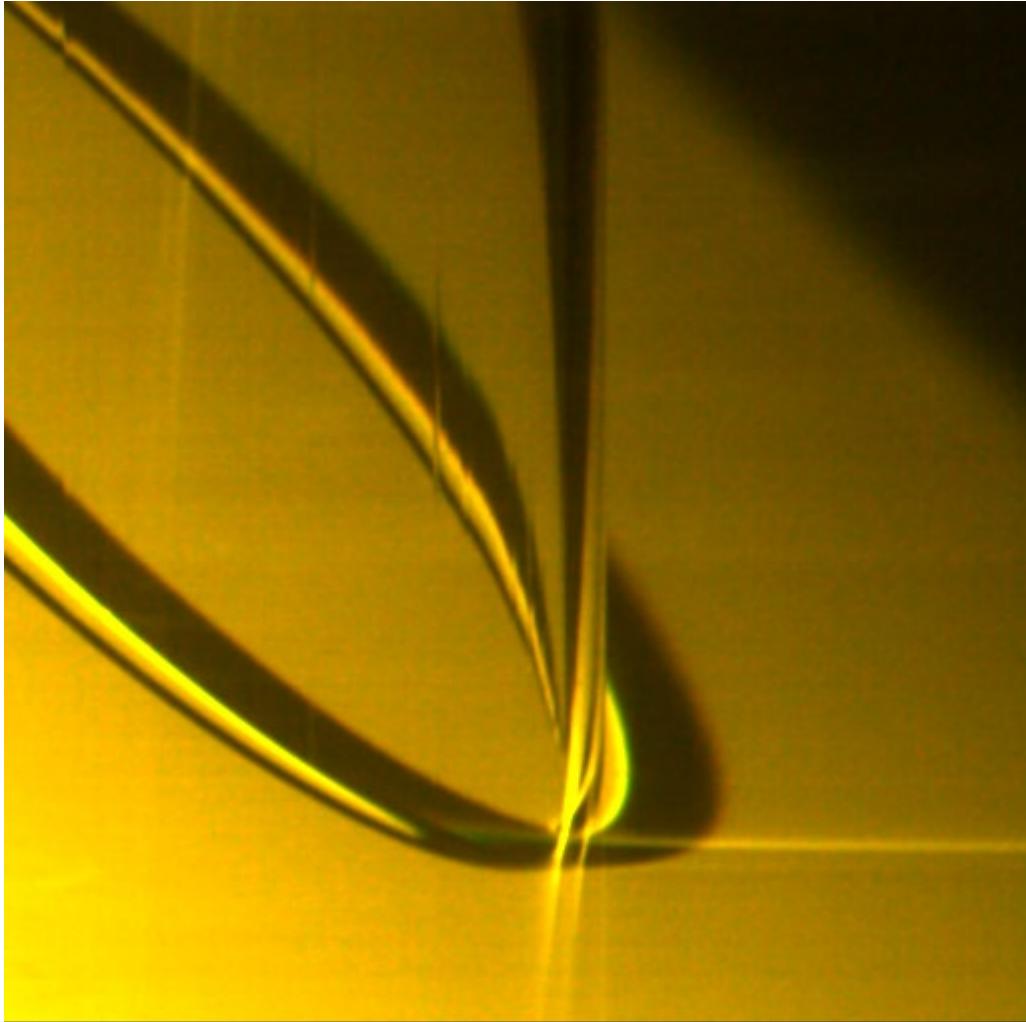


# Bilateral Filter – $\sigma_v=0.1$



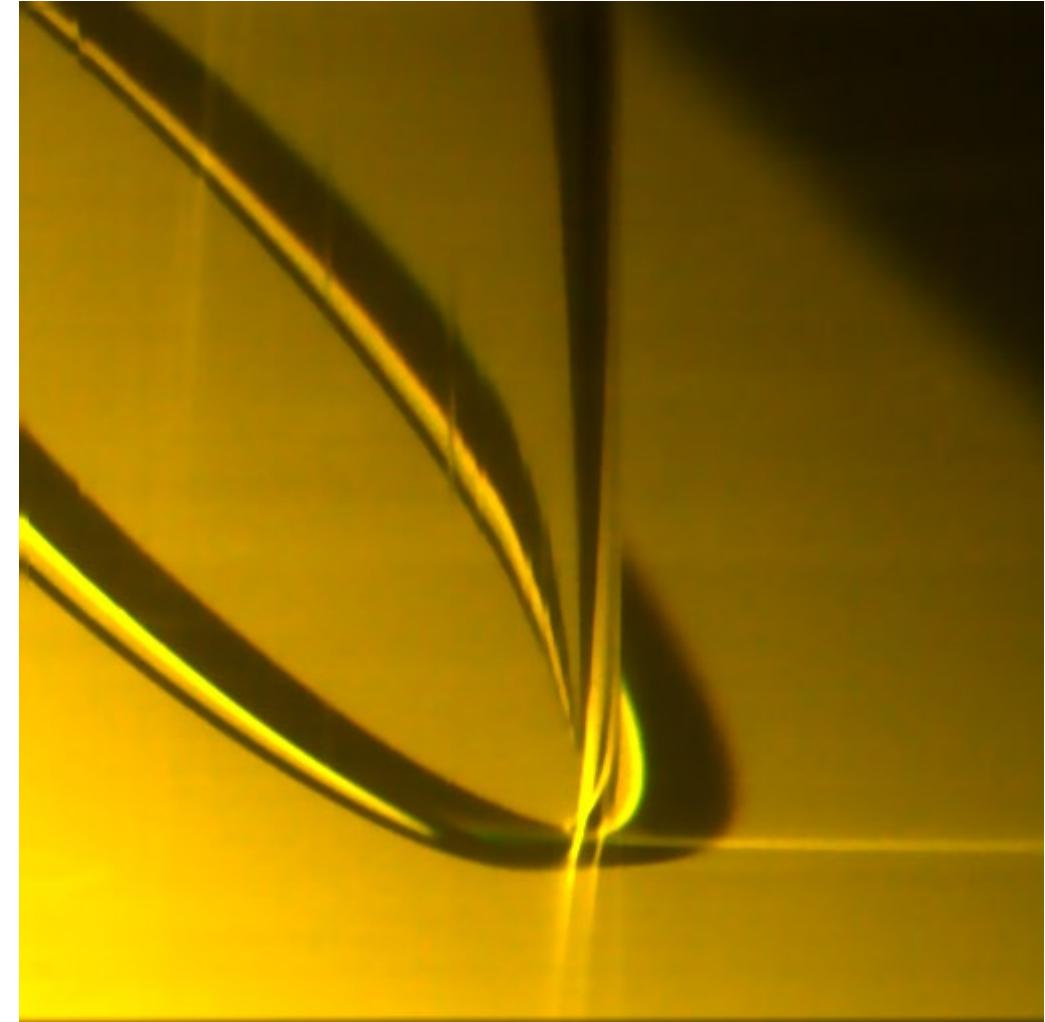
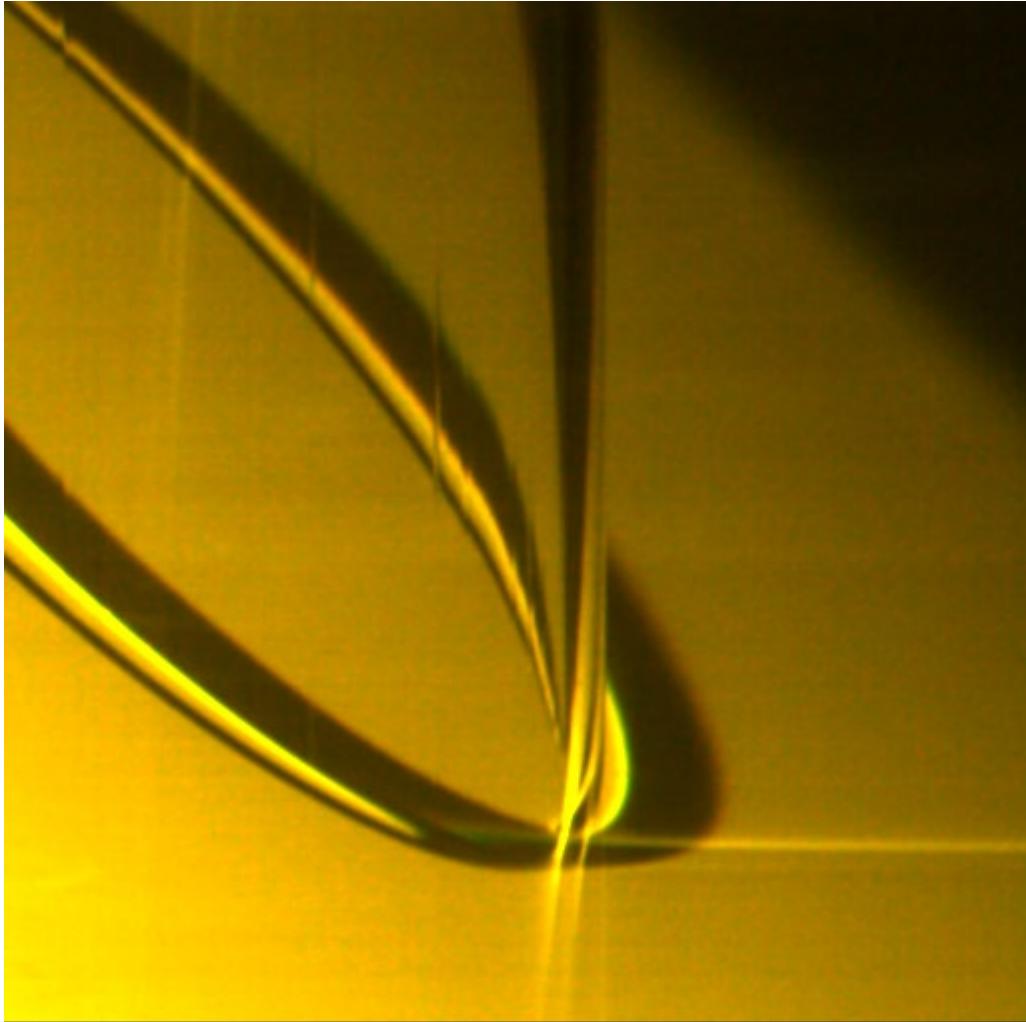


# Bilateral Filter – $\sigma_v=0.01$



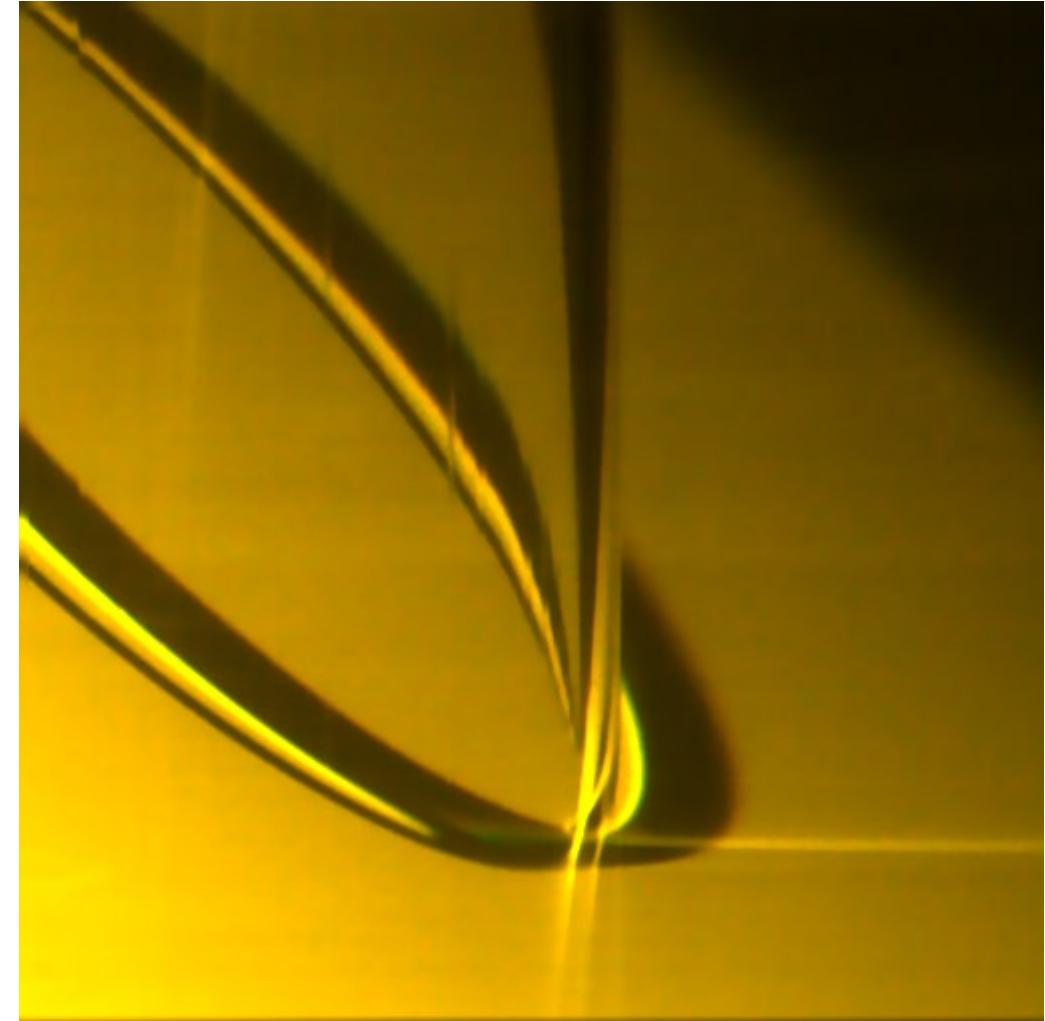
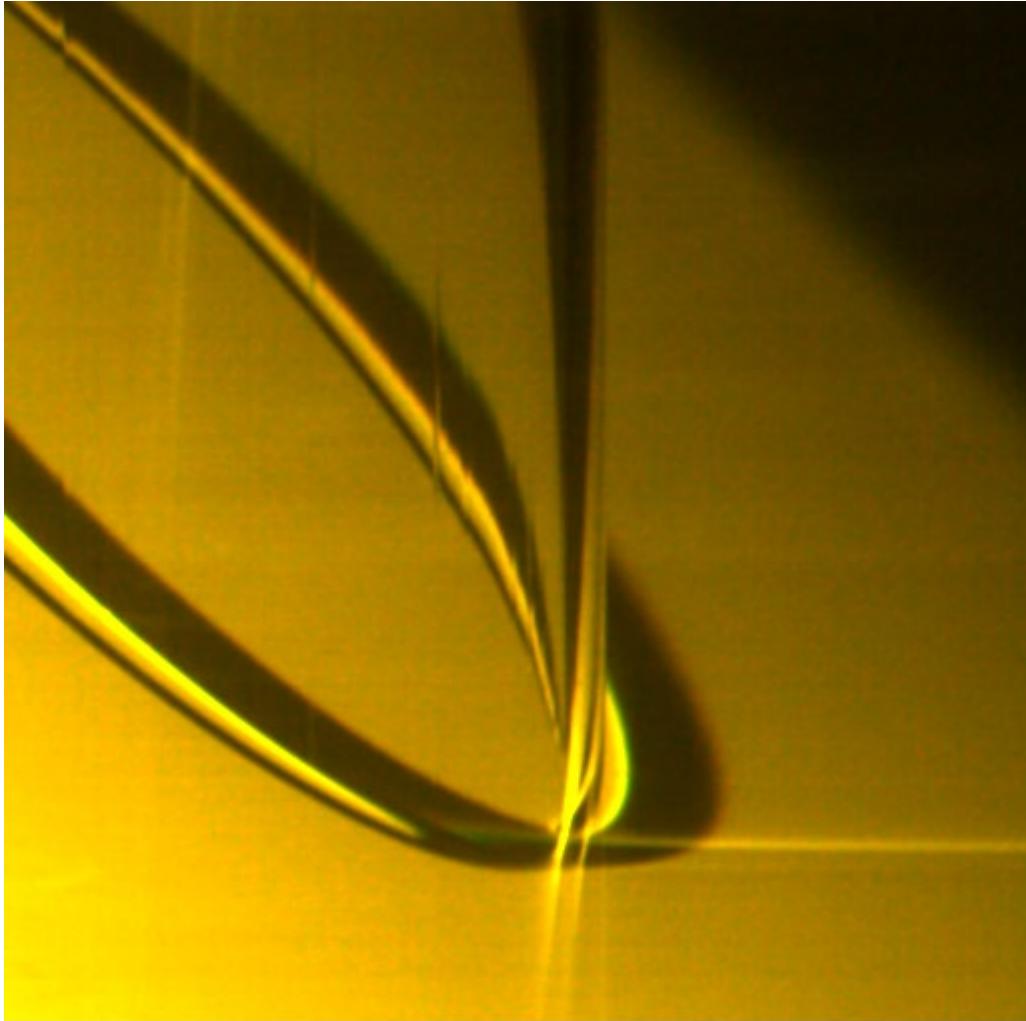


# Bilateral Filter – $\sigma_v=0.001$





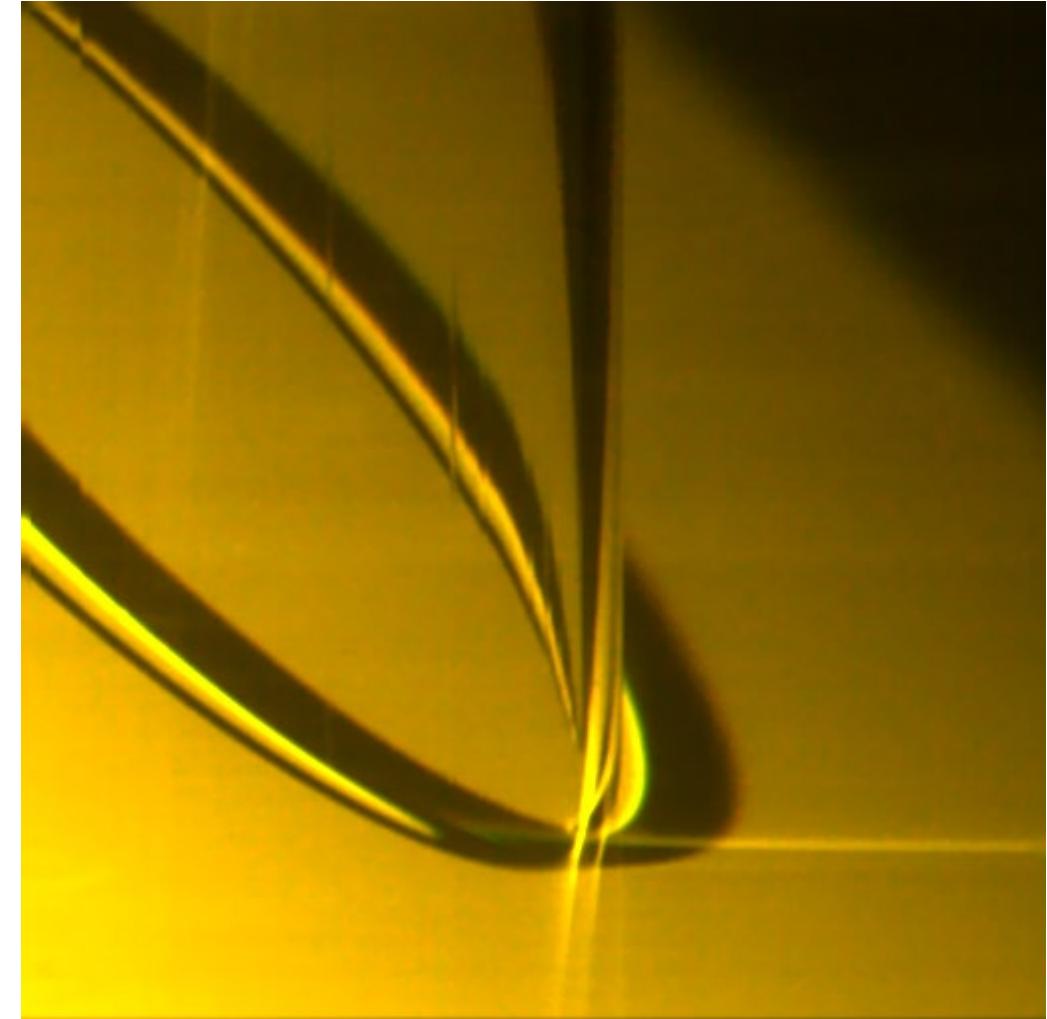
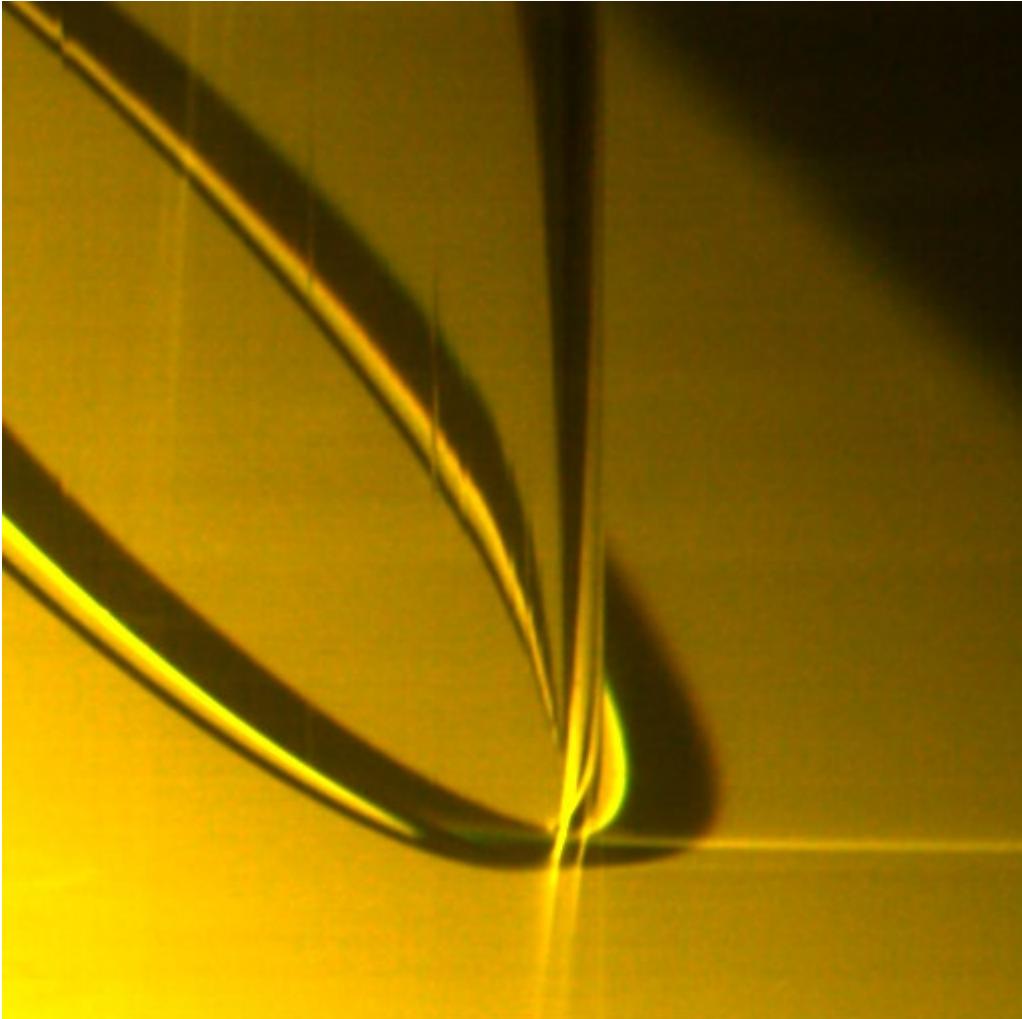
# Bilateral Filter – $\sigma_v=0.001$





# Bilateral Filter – $\sigma_v=0.0001$

Removes the noise but keeps the structure.





# Flash-No-Flash [Eisemann&Durand04]

- [the following slides are from Eisemann's SIGGRAPH talk]

Available light:

- + nice lighting
- noise/blurriness
- color





# Flash

---

Flash:

- + details
- + color

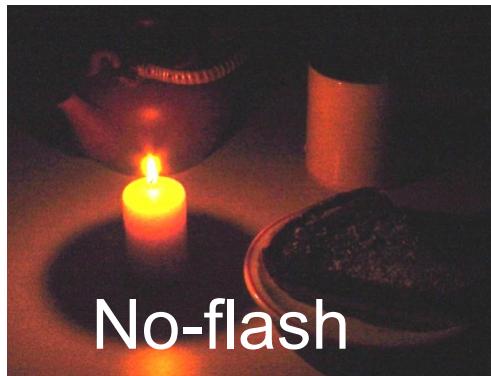
- flat/artificial
- flash shadows
- red eyes





# Combined

- Use no-flash image relight flash image

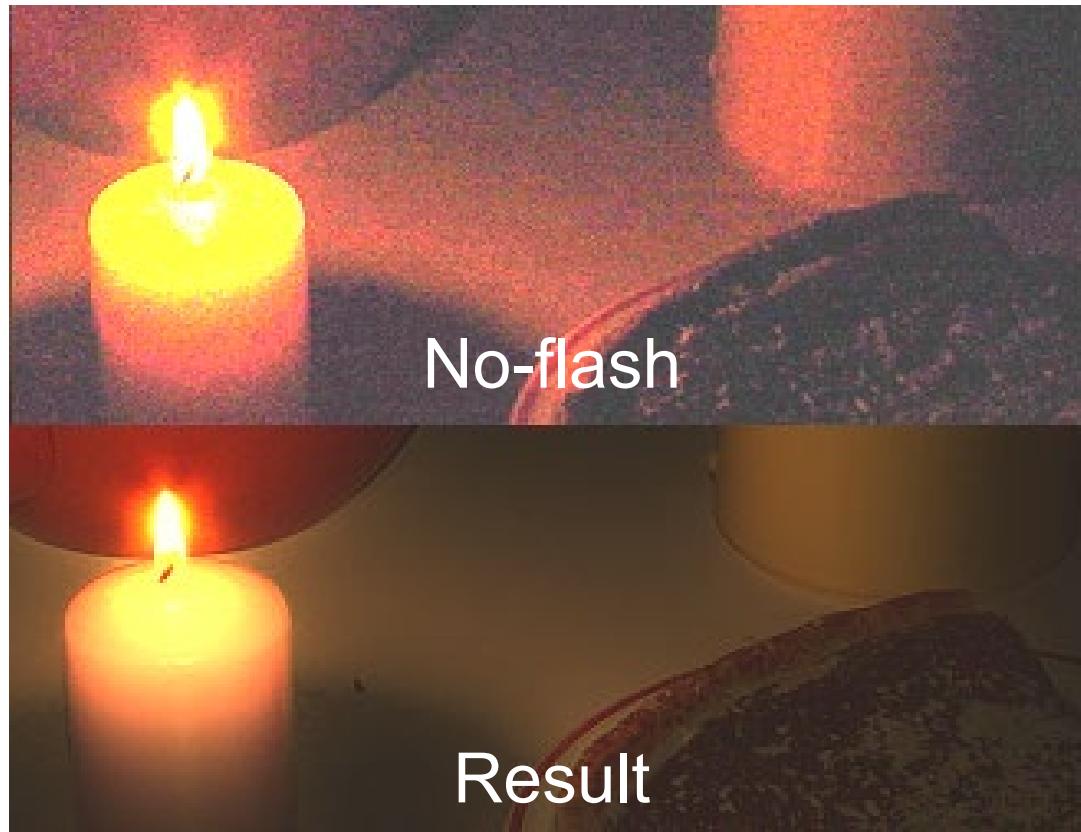




# Introduction

---

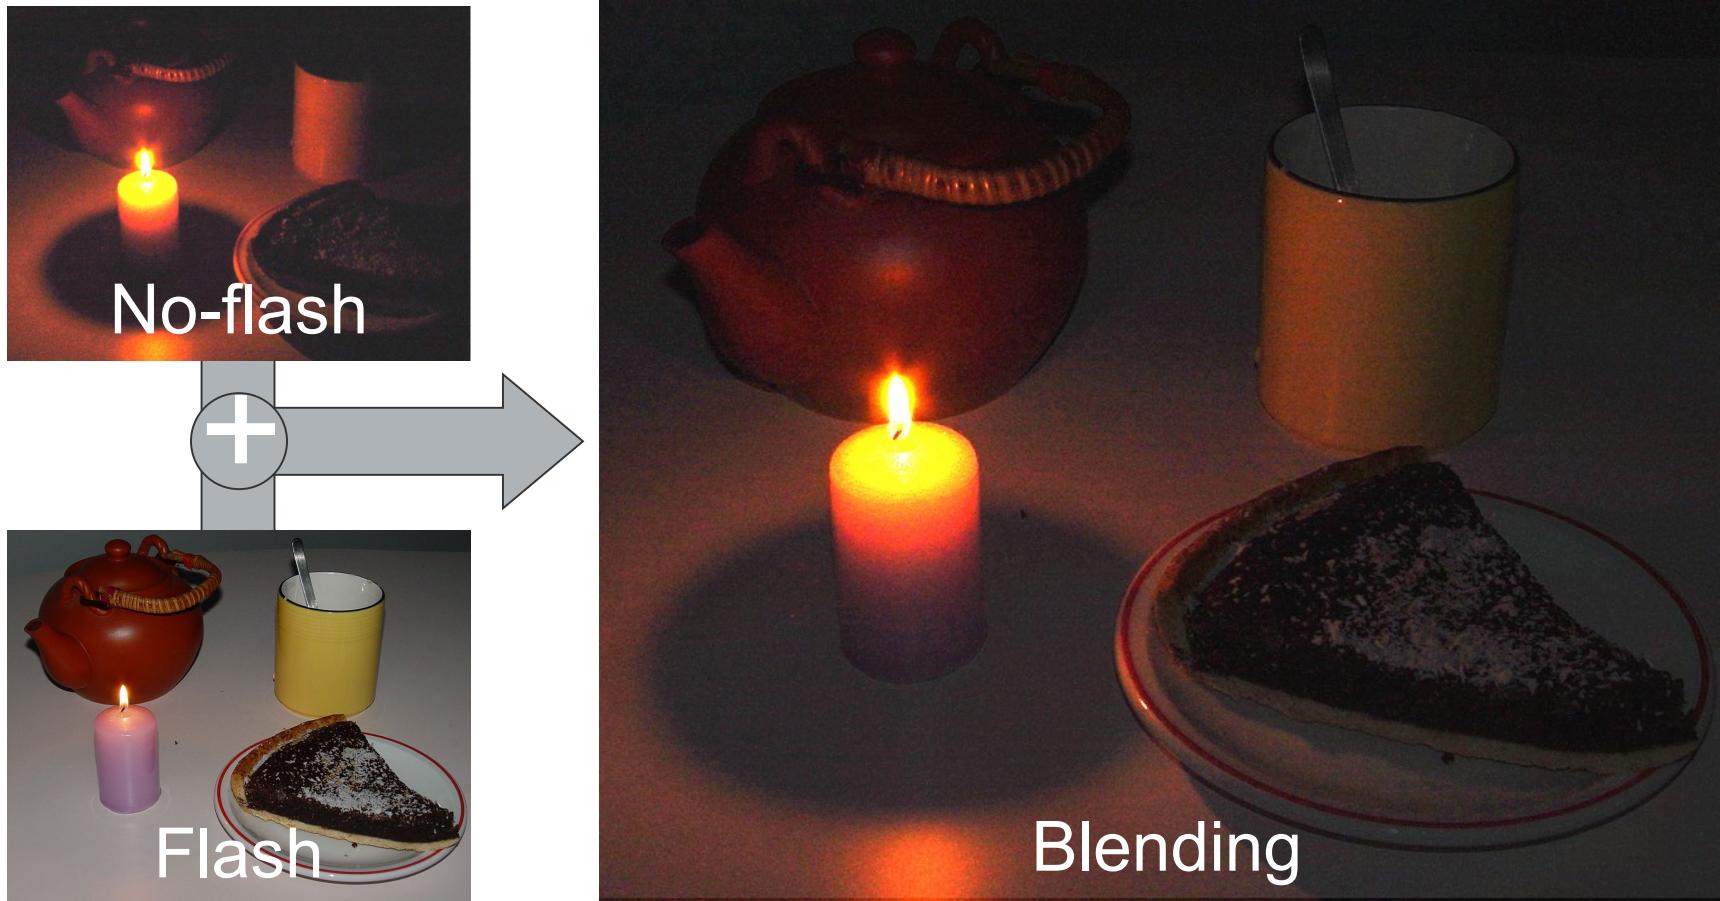
- Use no-flash image relight flash image





# Introduction

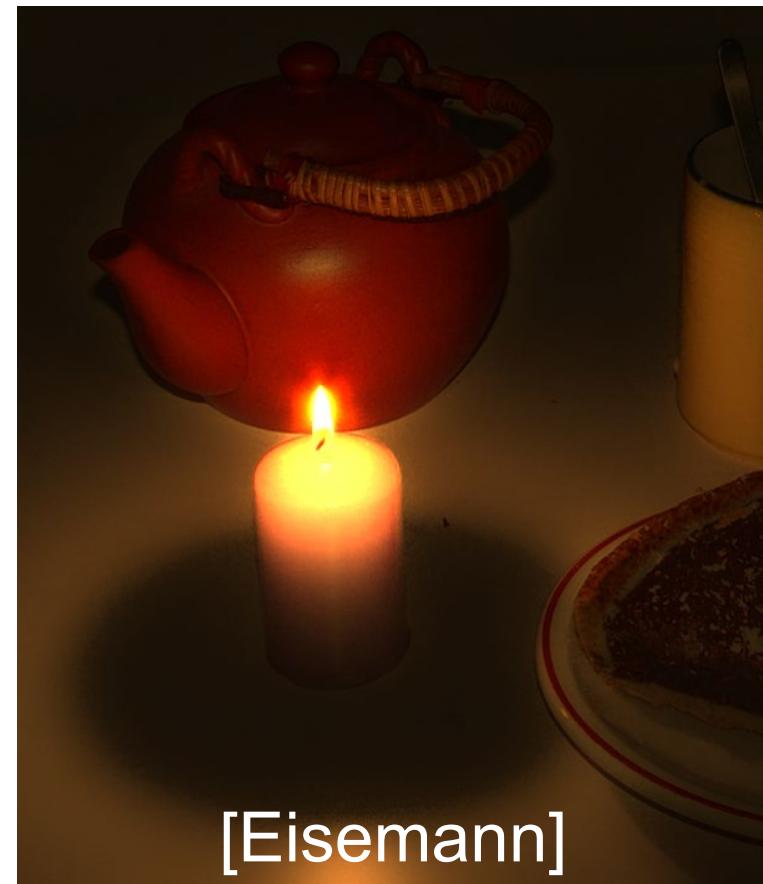
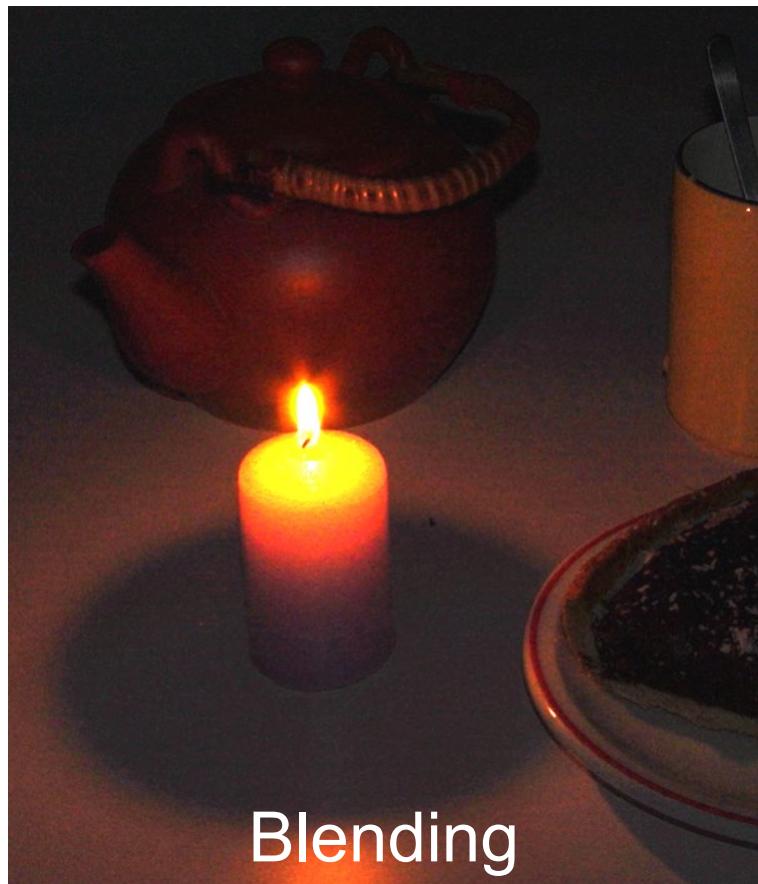
- One approach: Blend the two photos





# Introduction

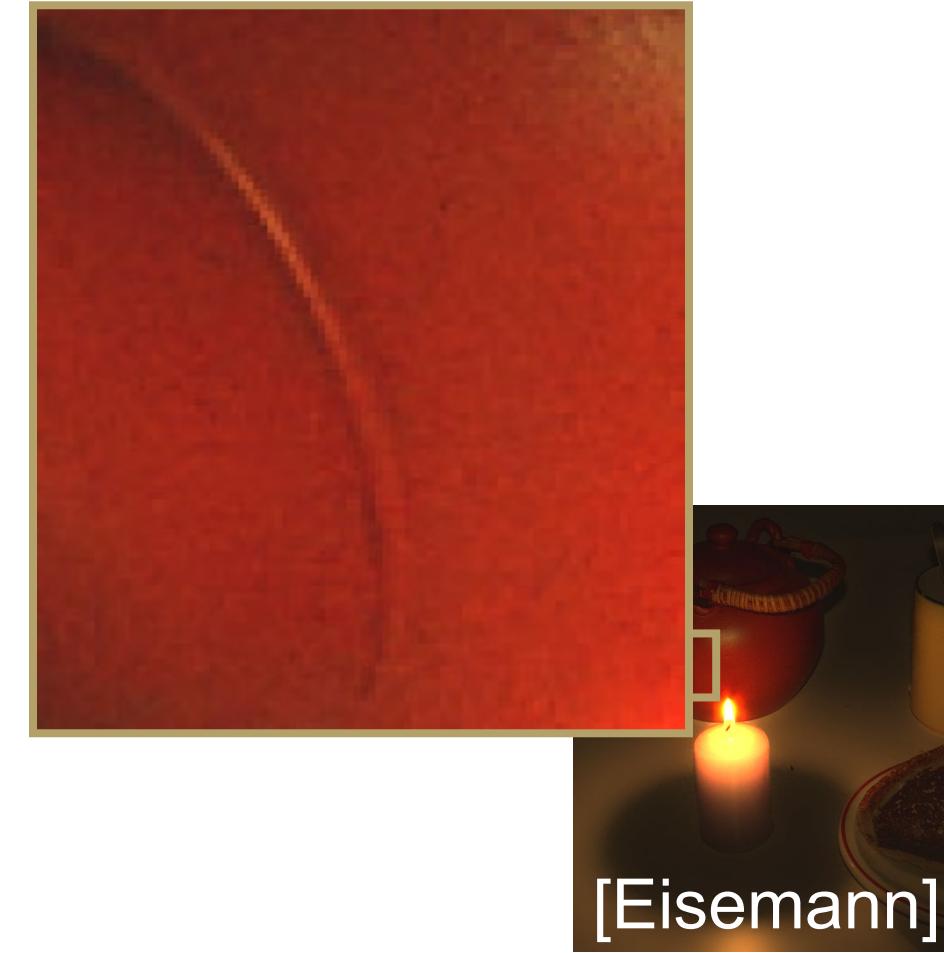
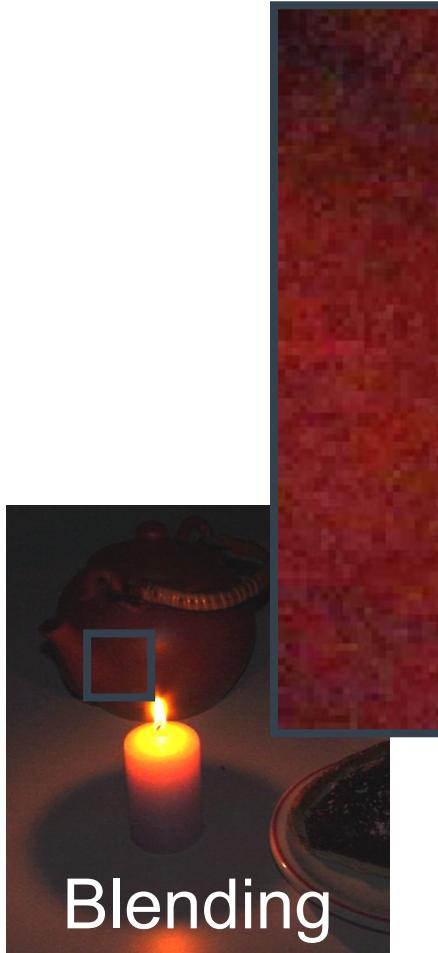
- One approach: Blend the two photos





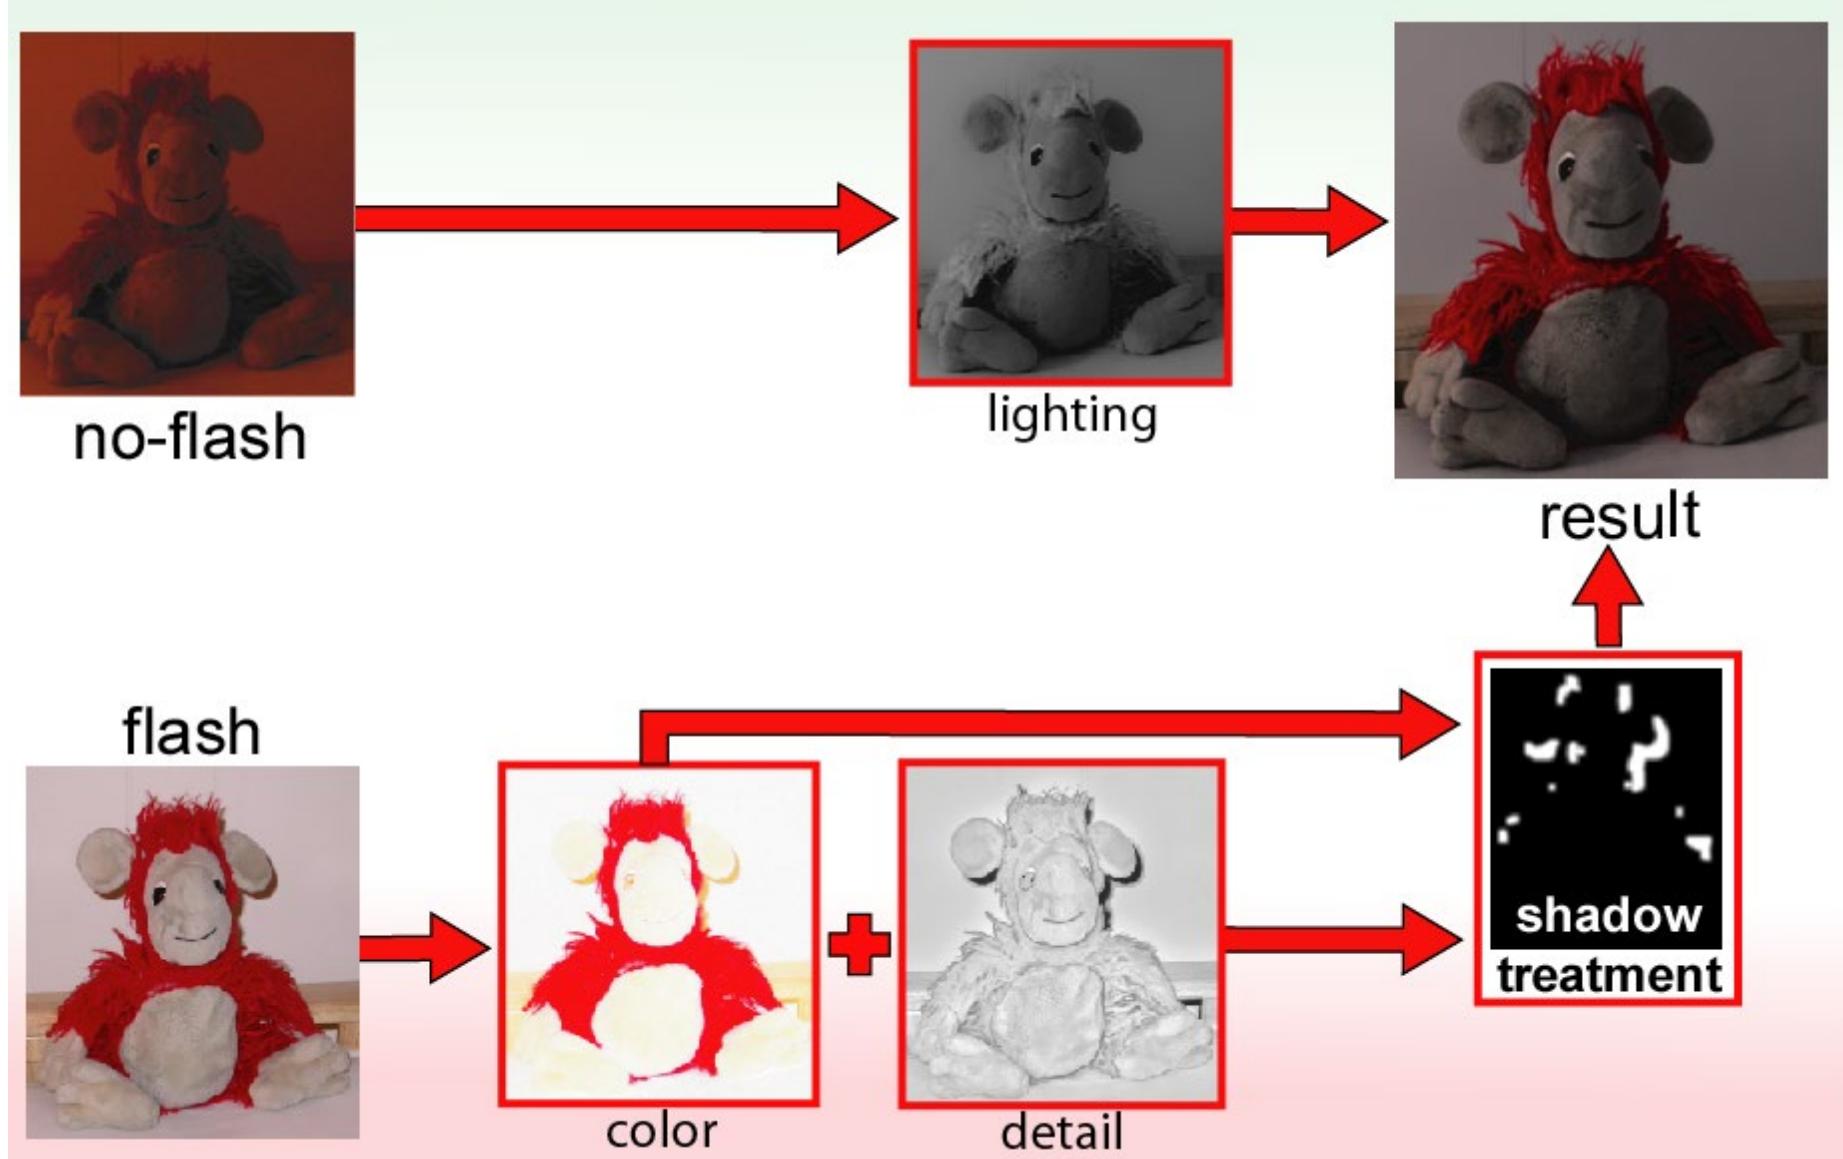
# Introduction

- One approach: Blend the two photos



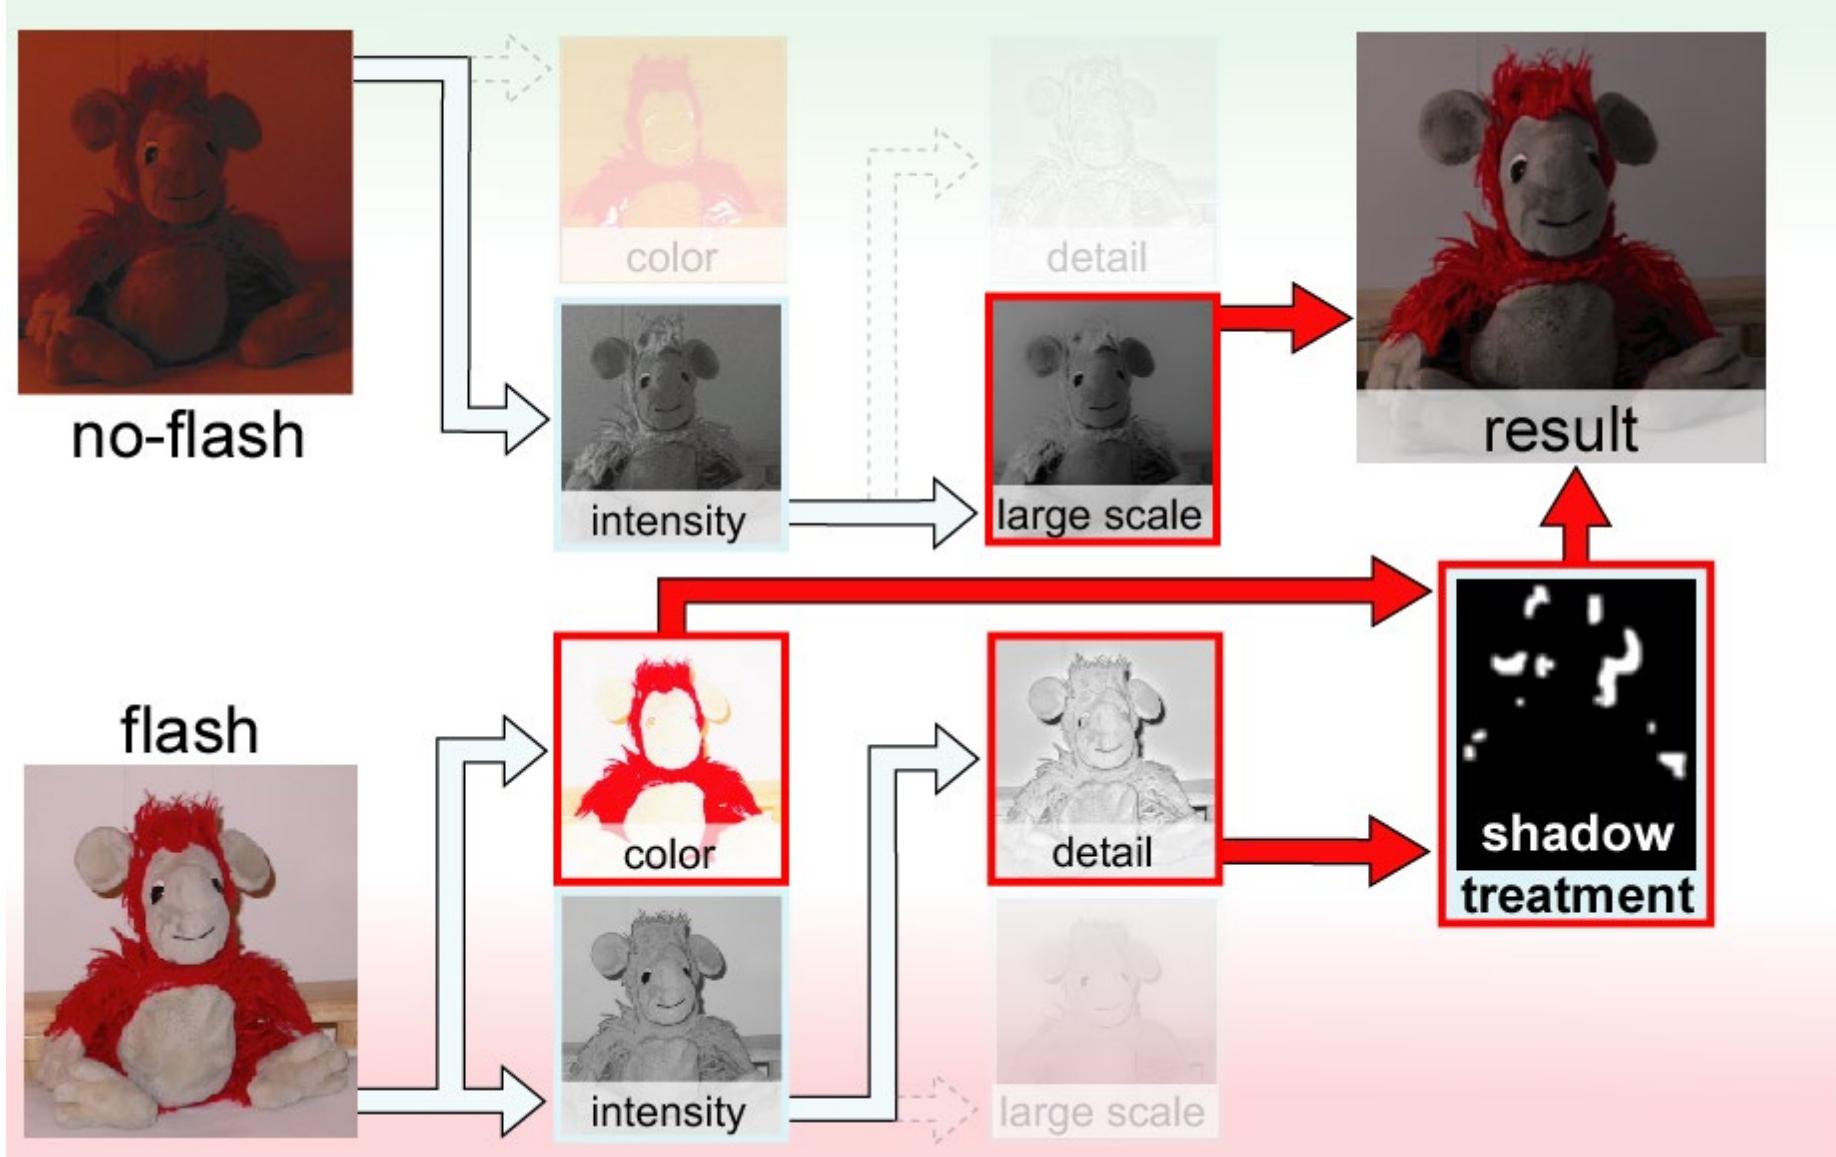


# Main Idea



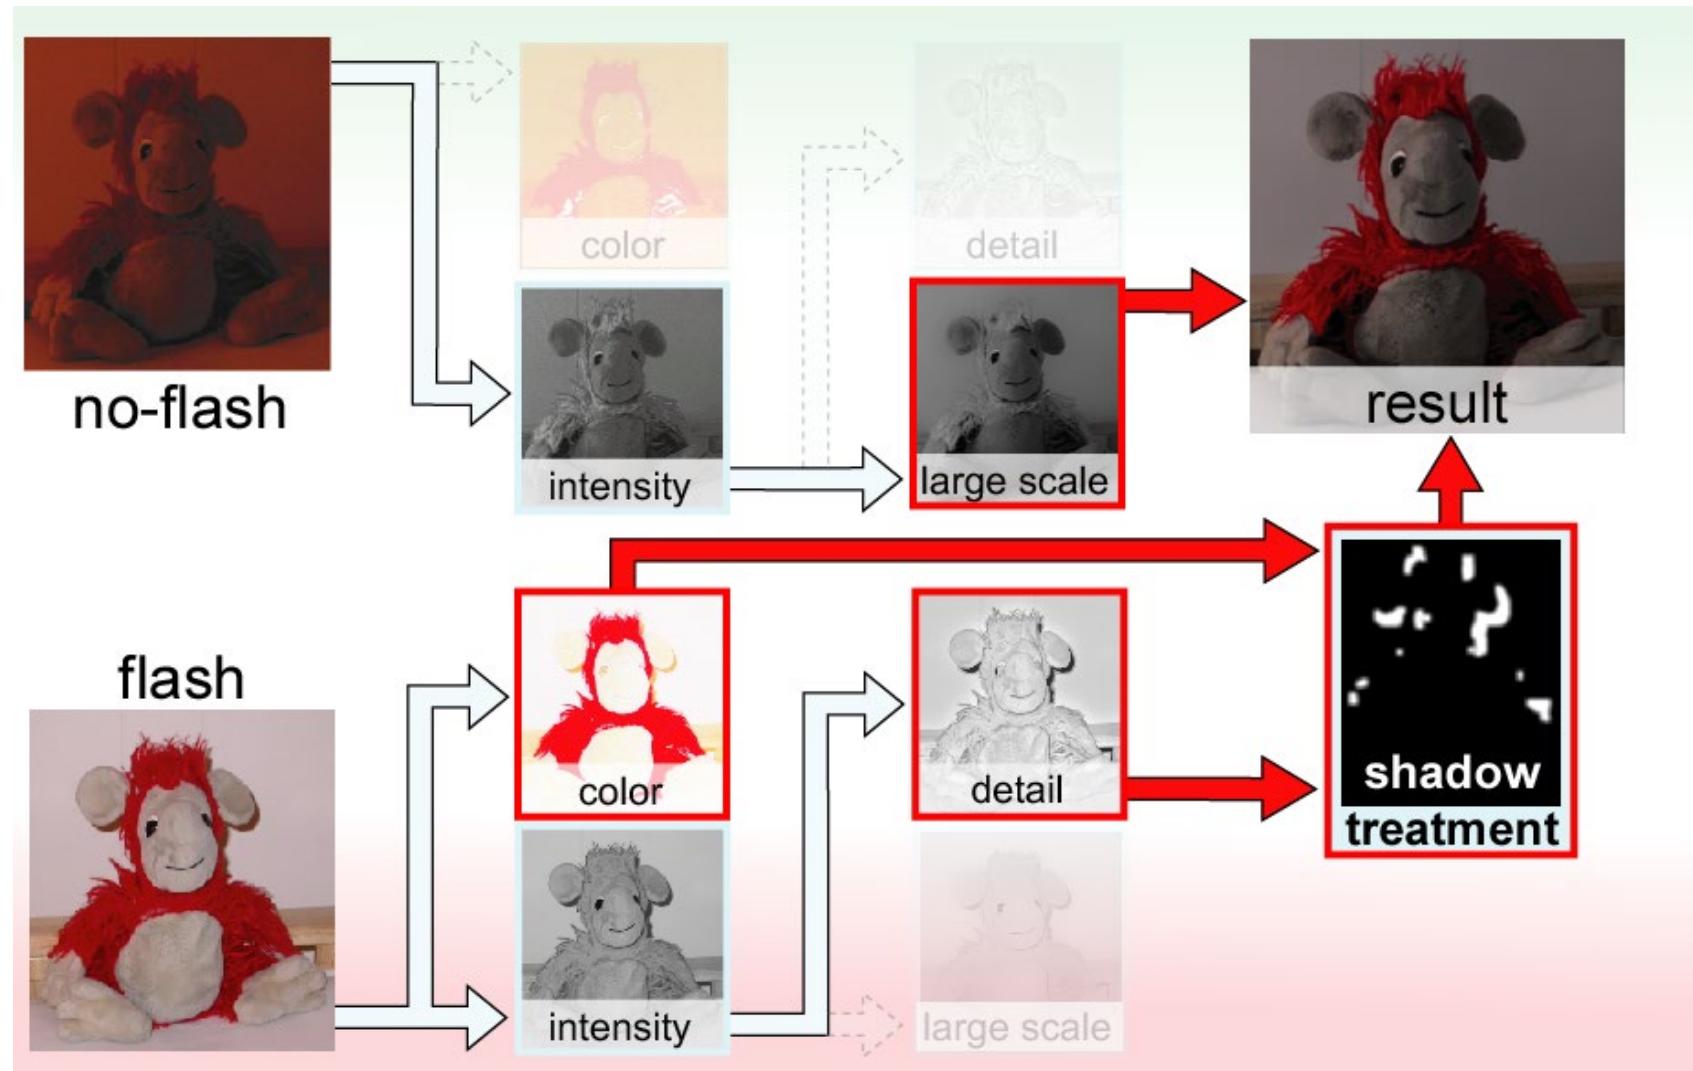


# Structure





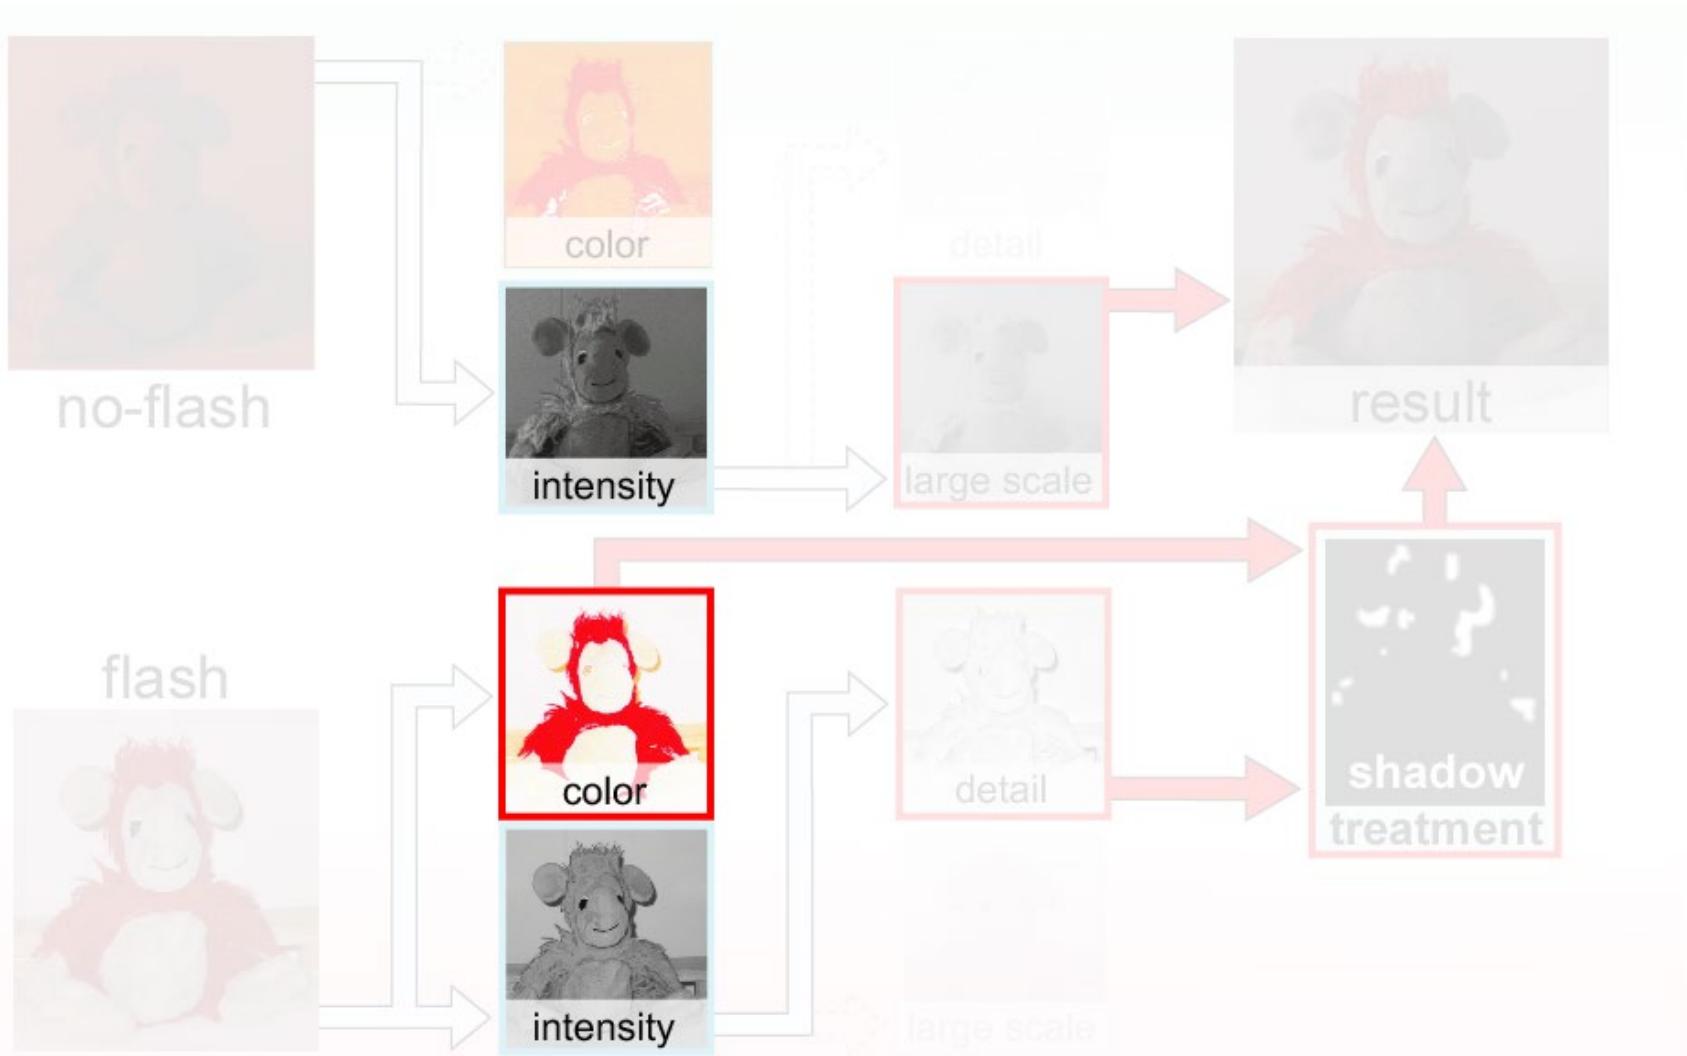
# Structure





# Intensity-Color Decomposition

## Decomposition





# Decomposition

- Color / Intensity:



=



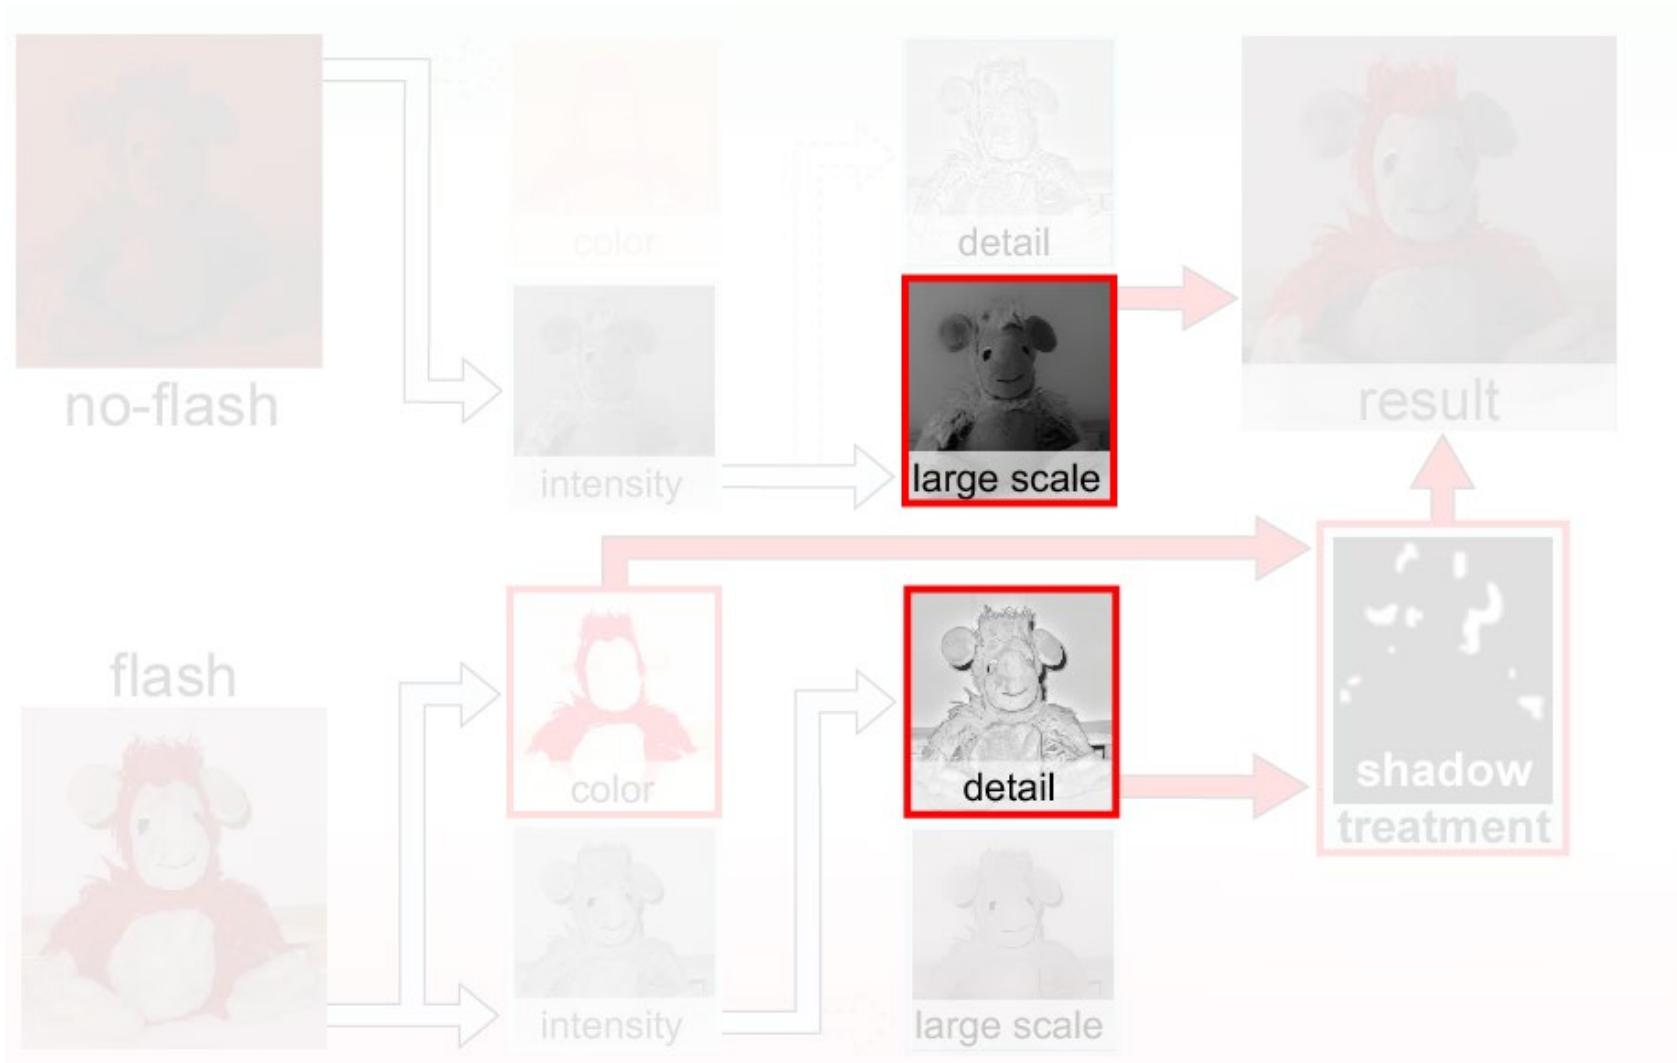
\*





# Frequency Decoupling

## Decoupling





# Decoupling

- Lighting : Large-scale variation
- Texture : Small-scale variation



Lighting

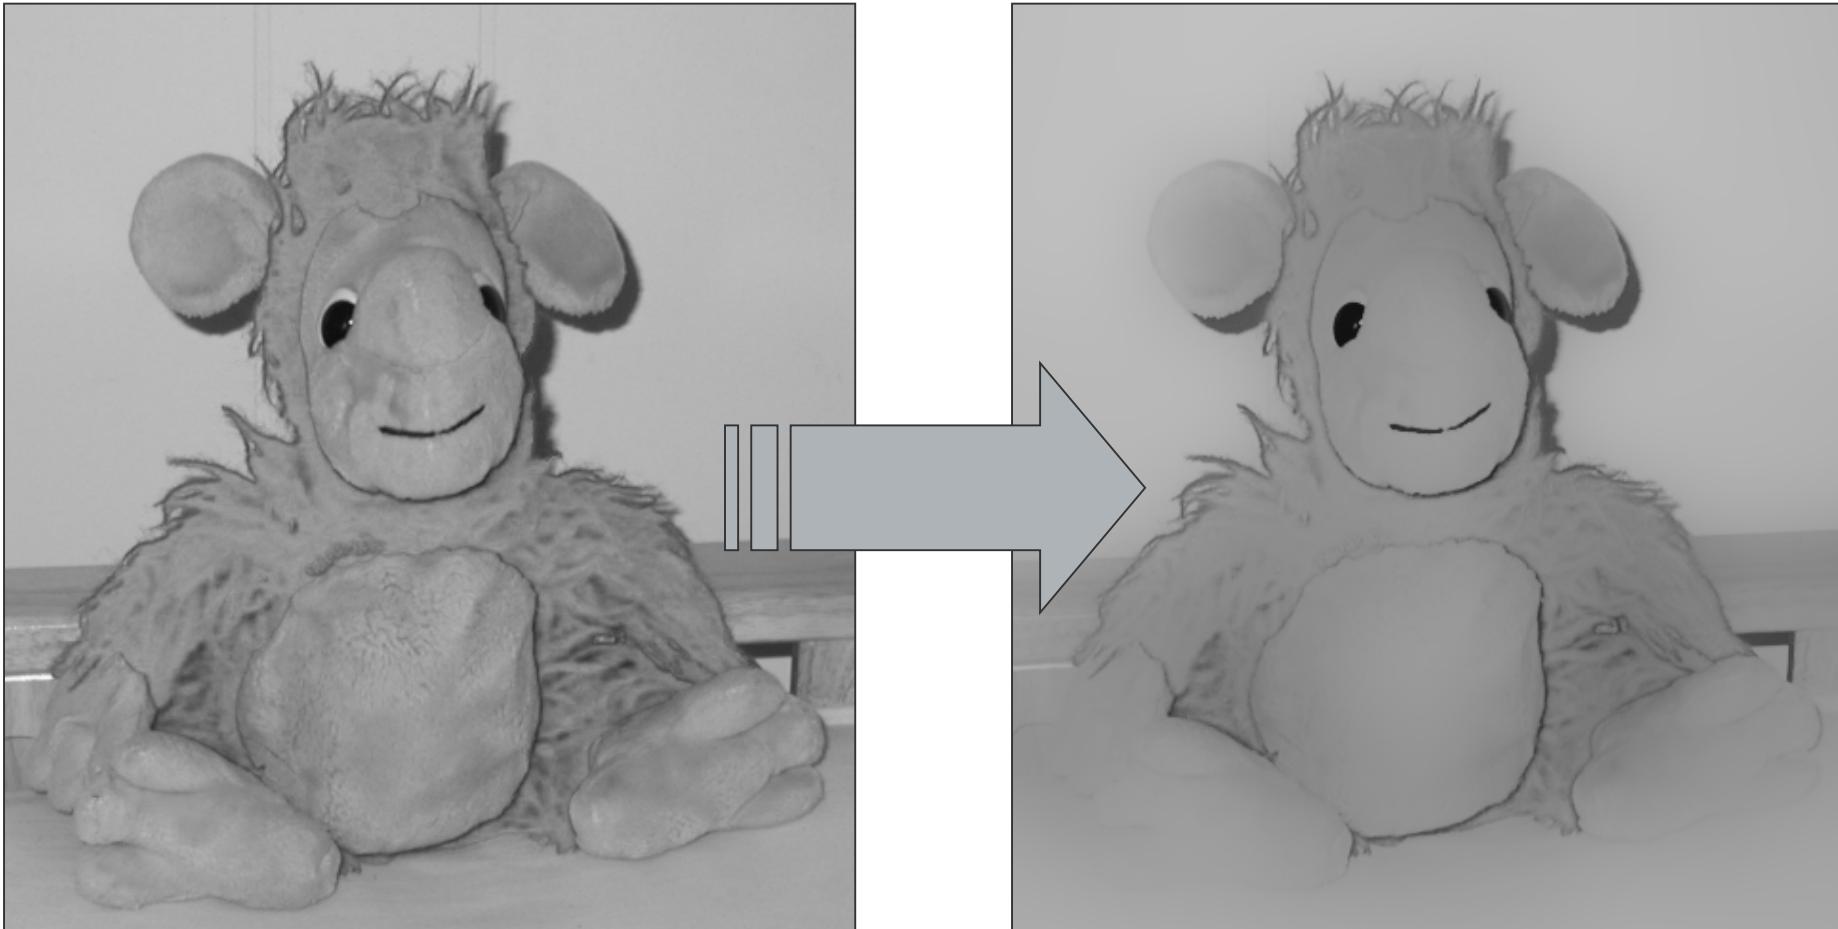


Texture



# Large-scale Layer

- Bilateral filter





# Detail Layer



/



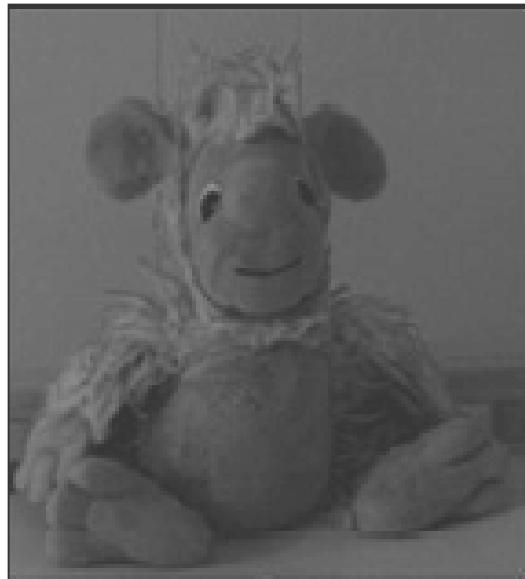
=



Recombination: Large scale \* Detail = Intensity



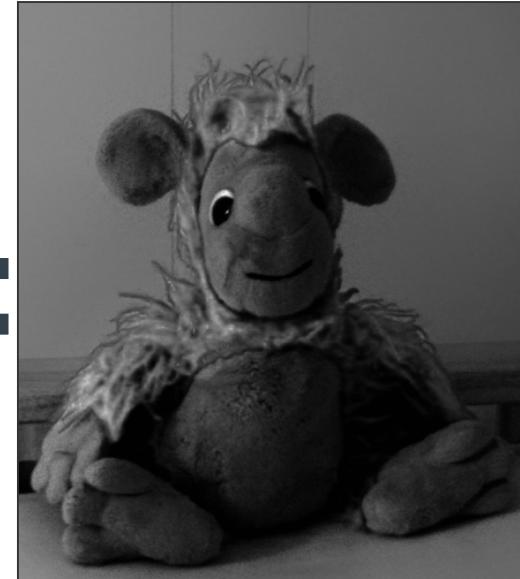
# Recombination



\*



=



Recombination: Large scale \* Detail = Intensity



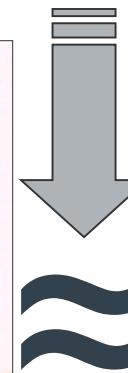
# Recombination



\*



shadows



Result

Recombination: Intensity \* Color = Original



# Results



No-flash



Flash



Result



---

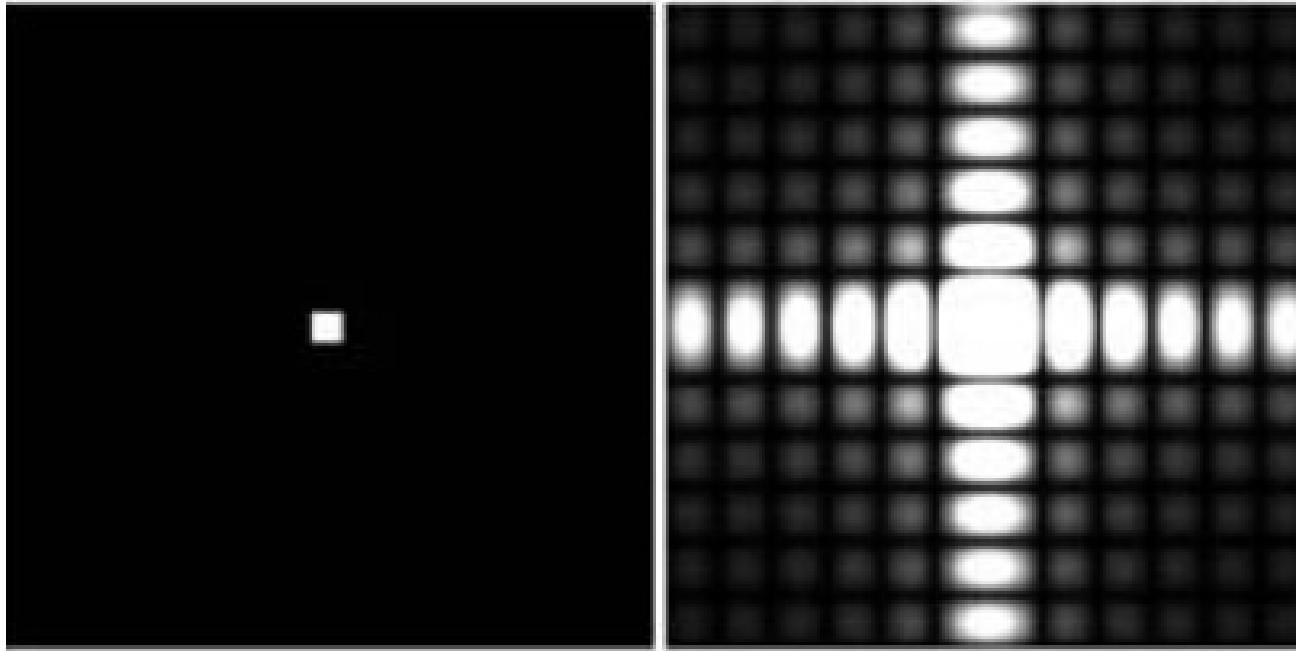
# Wavelets

---



# Motivation

- Fourier or DCT perform a global transformation into the frequency domain
- Image areas where there is low/high frequency content are not distinguished



- This example a single spot in the spatial domain influences the entire frequency spectrum



## Motivation (2)

---

Disadvantage of Fourier/DCT:

Basis function with global support in spatial domain

- Not very efficient in representing the signal
- Loss of any spatial information
- Operations in the frequency domain cannot be confined to specific region in spatial domain



# Motivation (3)

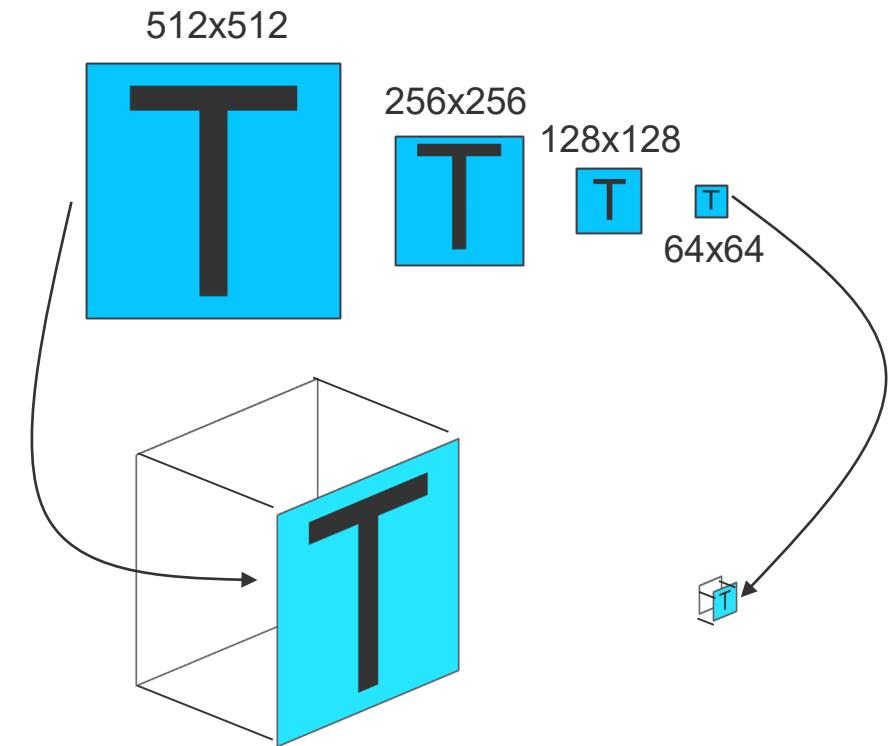
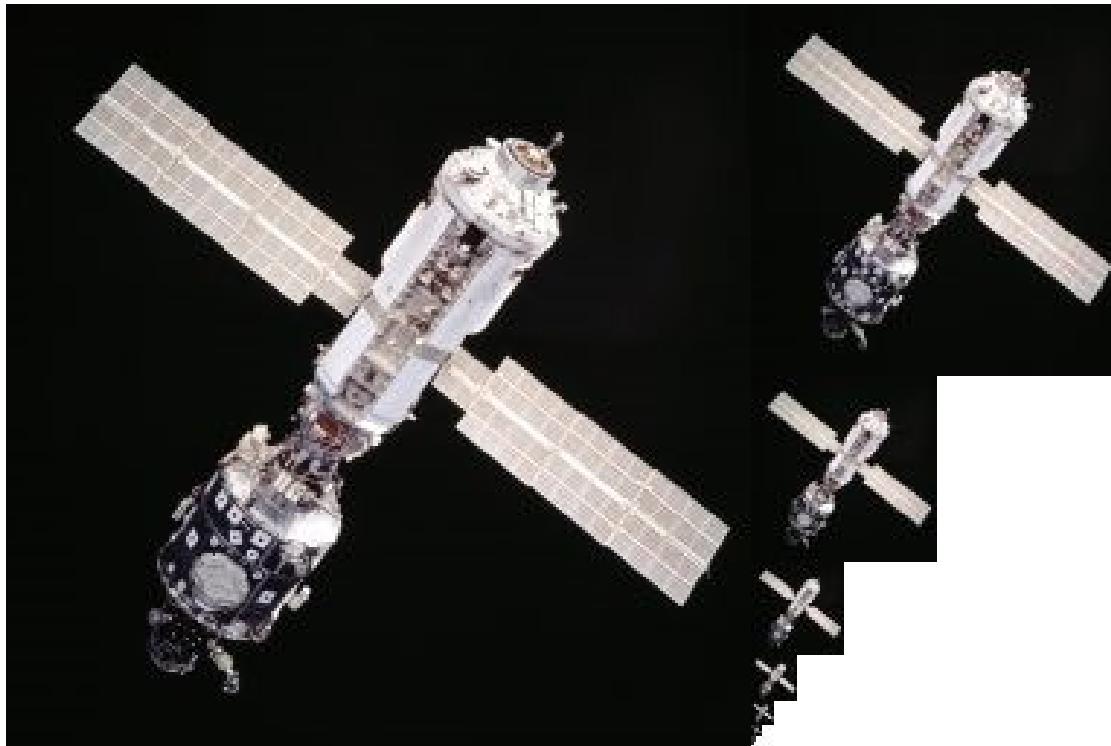
---

- Fourier/DCT:
  - Represents the signal/image as linear combination of basis functions with varying frequency
- Wavelets:
  - Representation as linear combination of basis function with varying frequency (similar to Fourier/DCT)
  - Wavelet basis function with ***spatially localized support***



# MipMapping

- Texture available in multiple resolutions
  - Pre-processing step
- Rendering: select appropriate texture resolution
  - Selection is usually per pixel !!
  - Texel size( $n$ ) < extent of pixel footprint < texel size( $n+1$ )



# Haar Wavelet Basis

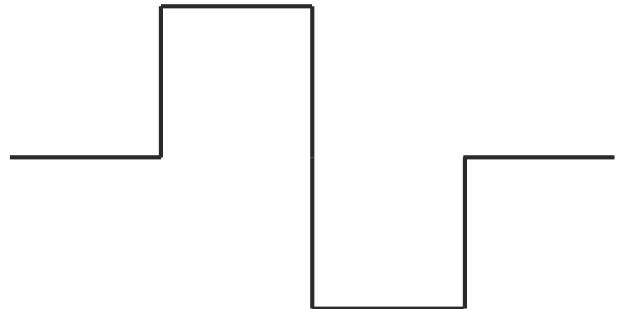
- mother scaling function

$$\phi(x) := \begin{cases} 1, & \text{for } 0 \leq x < 1 \\ 0, & \text{otherwise} \end{cases}$$



- mother wavelet function

$$\psi(x) := \begin{cases} 1, & \text{for } 0 \leq x < 1/2 \\ -1, & \text{for } 1/2 \leq x < 1 \\ 0, & \text{otherwise.} \end{cases}$$



# Haar Wavelet Basis

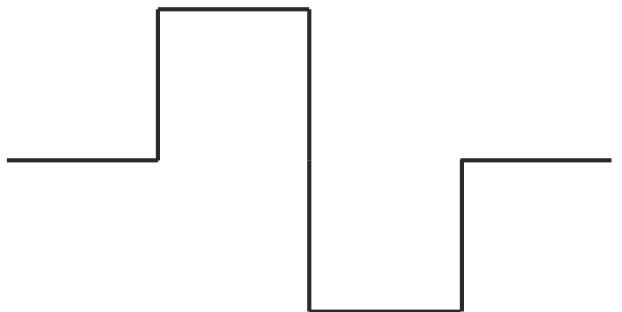
- mother scaling function

$$\phi(x) := \begin{cases} 1, & \text{for } 0 \leq x < 1 \\ 0, & \text{otherwise} \end{cases}$$



- mother wavelet function

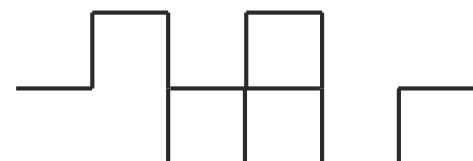
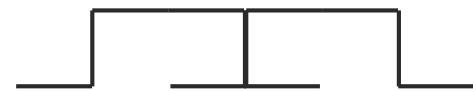
$$\psi(x) := \begin{cases} 1, & \text{for } 0 \leq x < 1/2 \\ -1, & \text{for } 1/2 \leq x < 1 \\ 0, & \text{otherwise.} \end{cases}$$



- Normalized scaling and wavelet basis functions

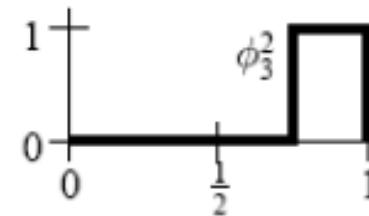
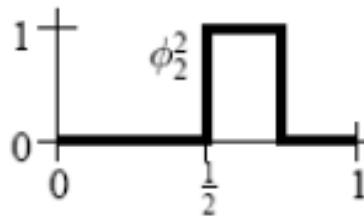
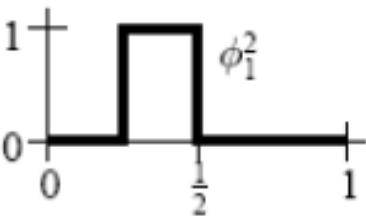
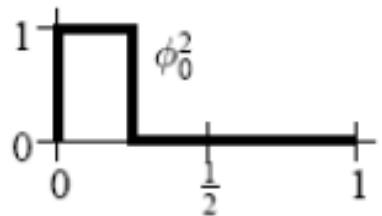
$$\phi_t^l(x) := 2^{l/2} \phi(2^l x - t)$$

$$\psi_t^l(x) := 2^{l/2} \psi(2^l x - t)$$

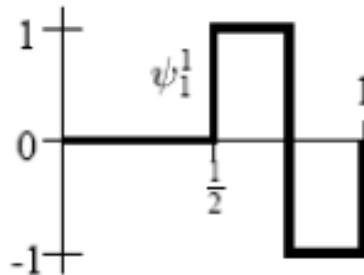
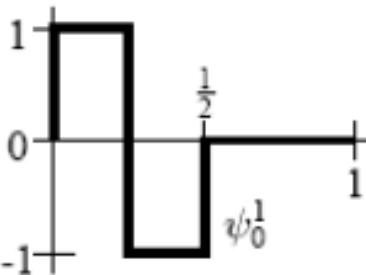




# Haar Basis Functions



Scaling Functions for  $V^2$



Wavelets for  $W^1$



# Wavelet Representation

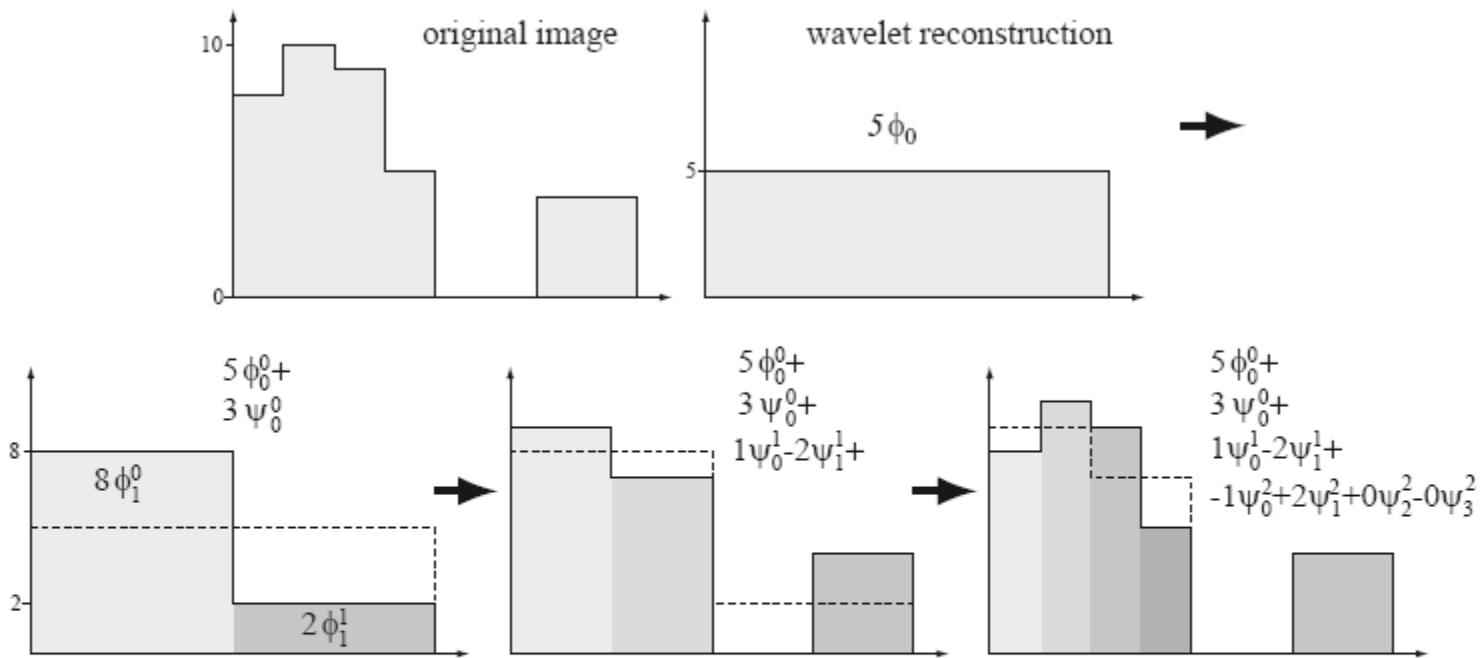


Figure 3: The 1D image (upper, left) is:  $[8, 10, 9, 5, 0, 0, 4, 4]$ , and its unnormalized (used here because it is simpler to display) Haar representation is:  $[5, 3, 1, -2, -1, 2, 0, 0]$ . The image is then reconstructed one level at a time as follows:  $[5] \rightarrow [5 + 3, 5 - 3] = [8, 2] \rightarrow [8 + 1, 8 - 1, 2 - 2, 2 + 2] = [9, 7, 0, 4]$  and so on.

# Wavelet Representation

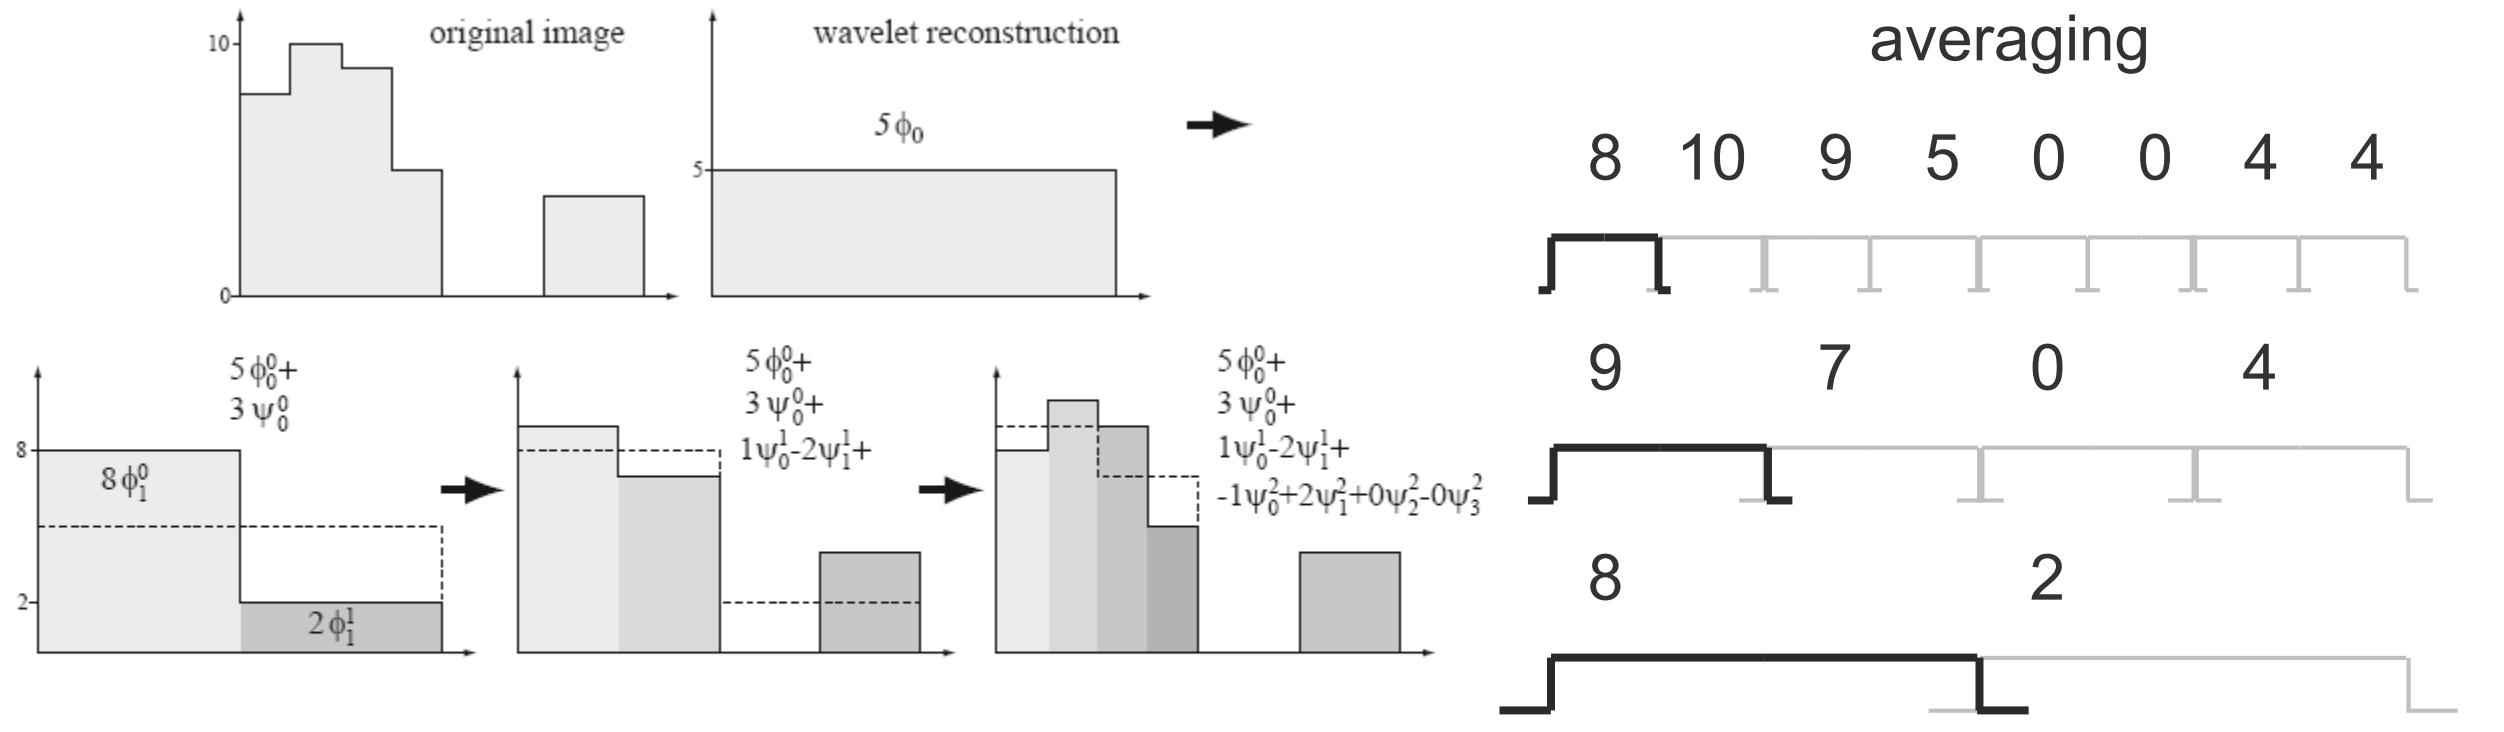


Figure 3: The 1D image (upper, left) is: [8, 10, 9, 5, 0, 0, 4, 4], and its unnormalized (used here because it is simpler to display) Haar representation is: [5, 3, 1, -2, -1, 2, 0, 0]. The image is then reconstructed one level at a time as follows:  $[5] \rightarrow [5 + 3, 5 - 3] = [8, 2] \rightarrow [8 + 1, 8 - 1, 2 - 2, 2 + 2] = [9, 7, 0, 4]$  and so on.



# Wavelet Representation

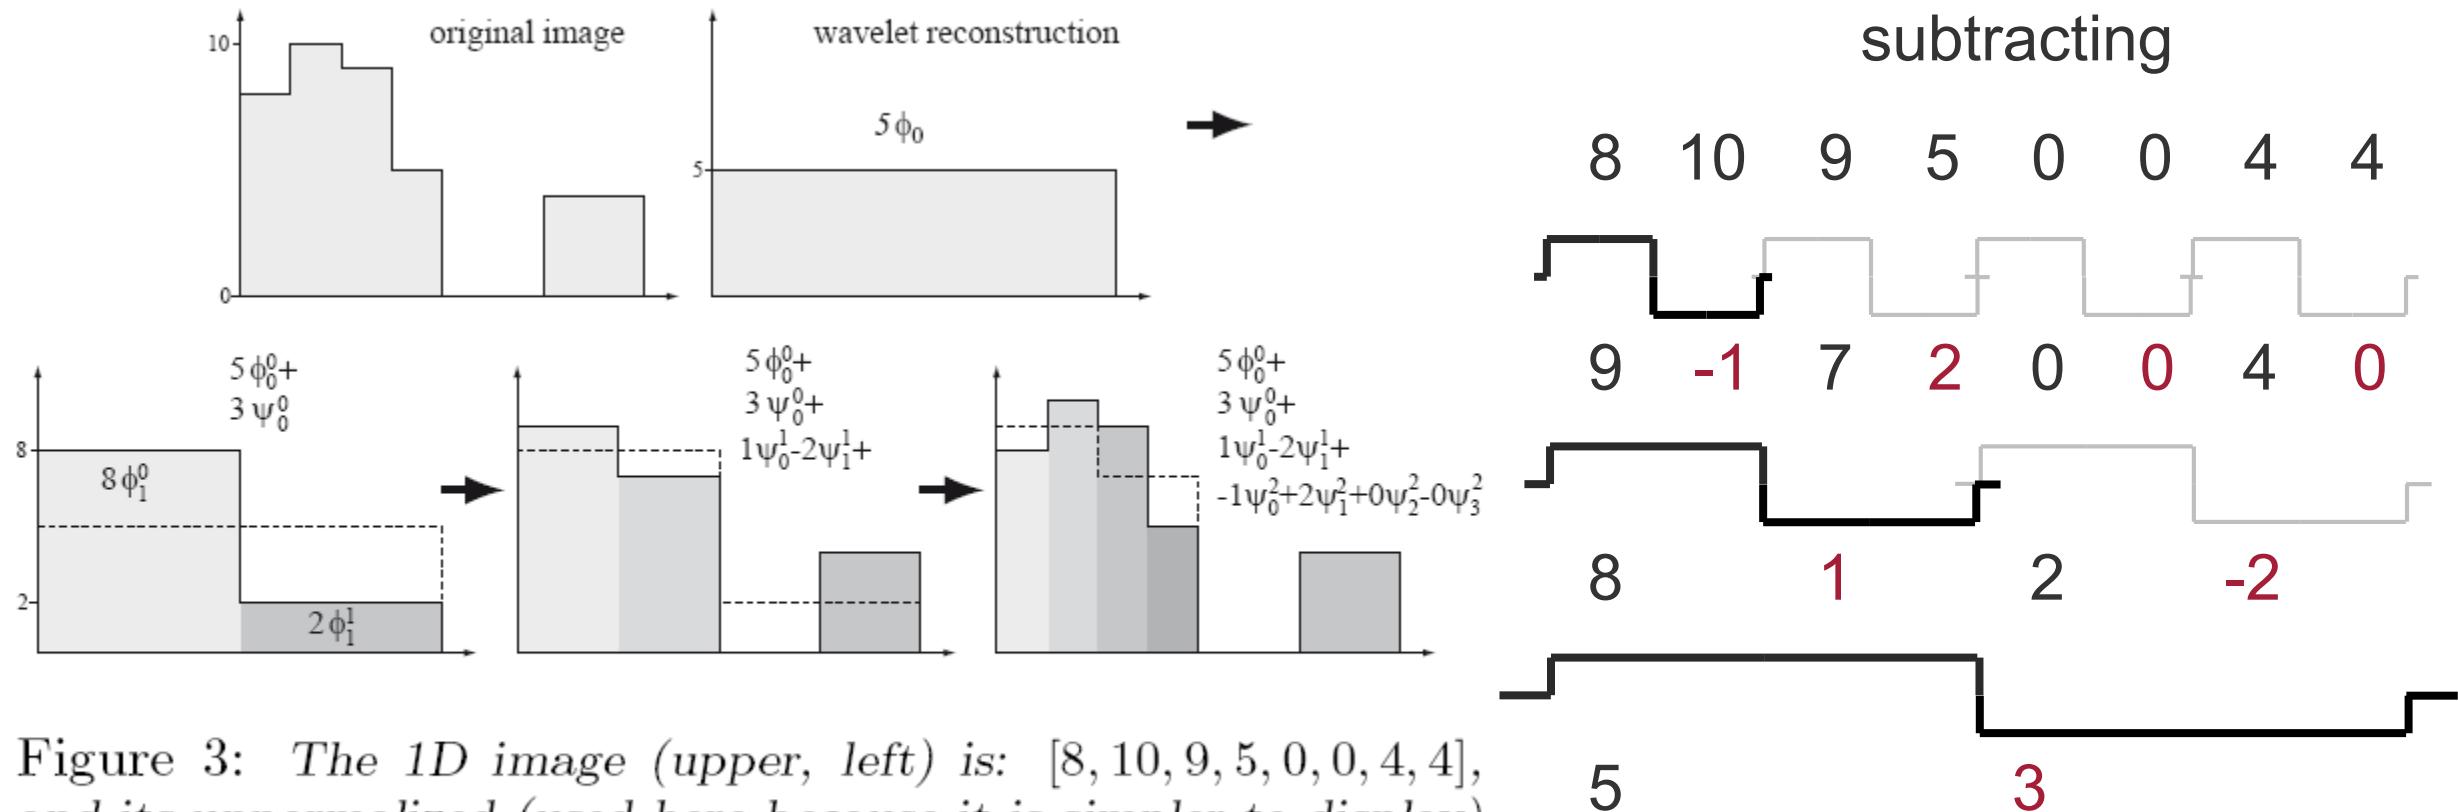


Figure 3: The 1D image (upper, left) is: [8, 10, 9, 5, 0, 0, 4, 4], and its unnormalized (used here because it is simpler to display) Haar representation is: [5, 3, 1, -2, -1, 2, 0, 0]. The image is then reconstructed one level at a time as follows:  $[5] \rightarrow [5 + 3, 5 - 3] = [8, 2] \rightarrow [8 + 1, 8 - 1, 2 - 2, 2 + 2] = [9, 7, 0, 4]$  and so on.

# Wavelet Representation – Reconstruction

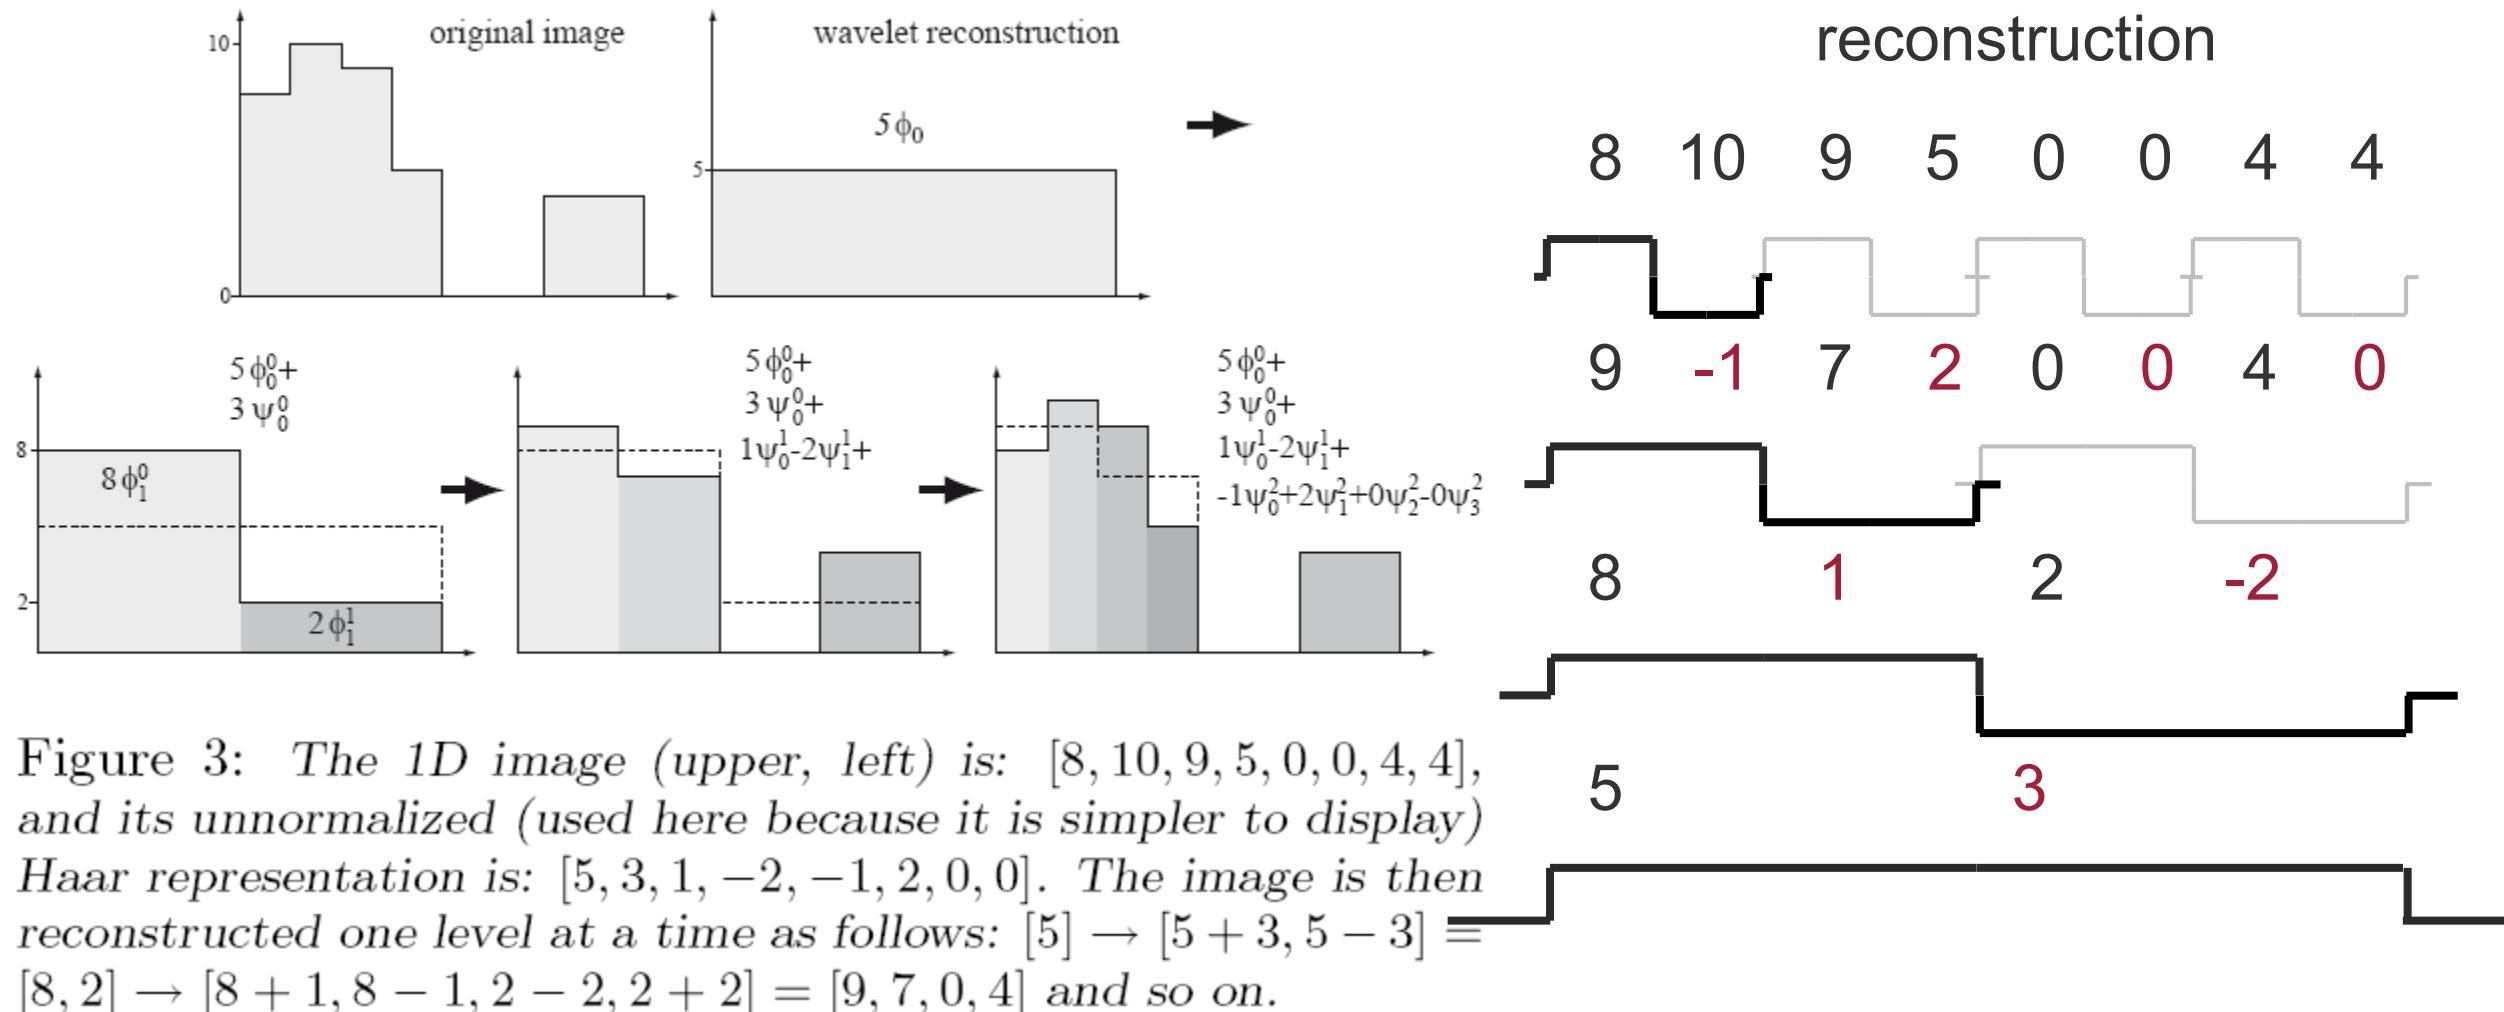


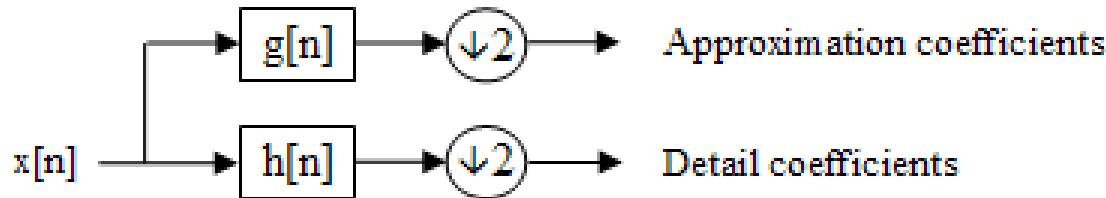
Figure 3: The 1D image (upper, left) is: [8, 10, 9, 5, 0, 0, 4, 4], and its unnormalized (used here because it is simpler to display) Haar representation is: [5, 3, 1, -2, -1, 2, 0, 0]. The image is then reconstructed one level at a time as follows:  $[5] \rightarrow [5 + 3, 5 - 3] = [8, 2] \rightarrow [8 + 1, 8 - 1, 2 - 2, 2 + 2] = [9, 7, 0, 4]$  and so on.

# Fast Wavelet Transform

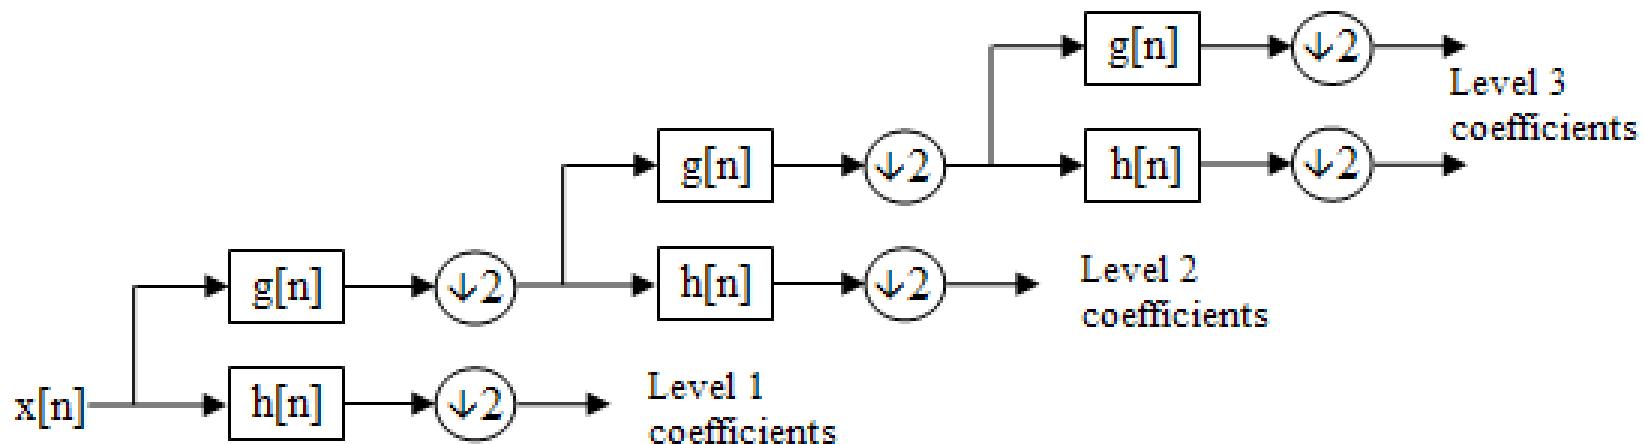
- Downsampling operator e.g.
- selecting every other sample
- high-pass and low-pass filter



- Single level transform

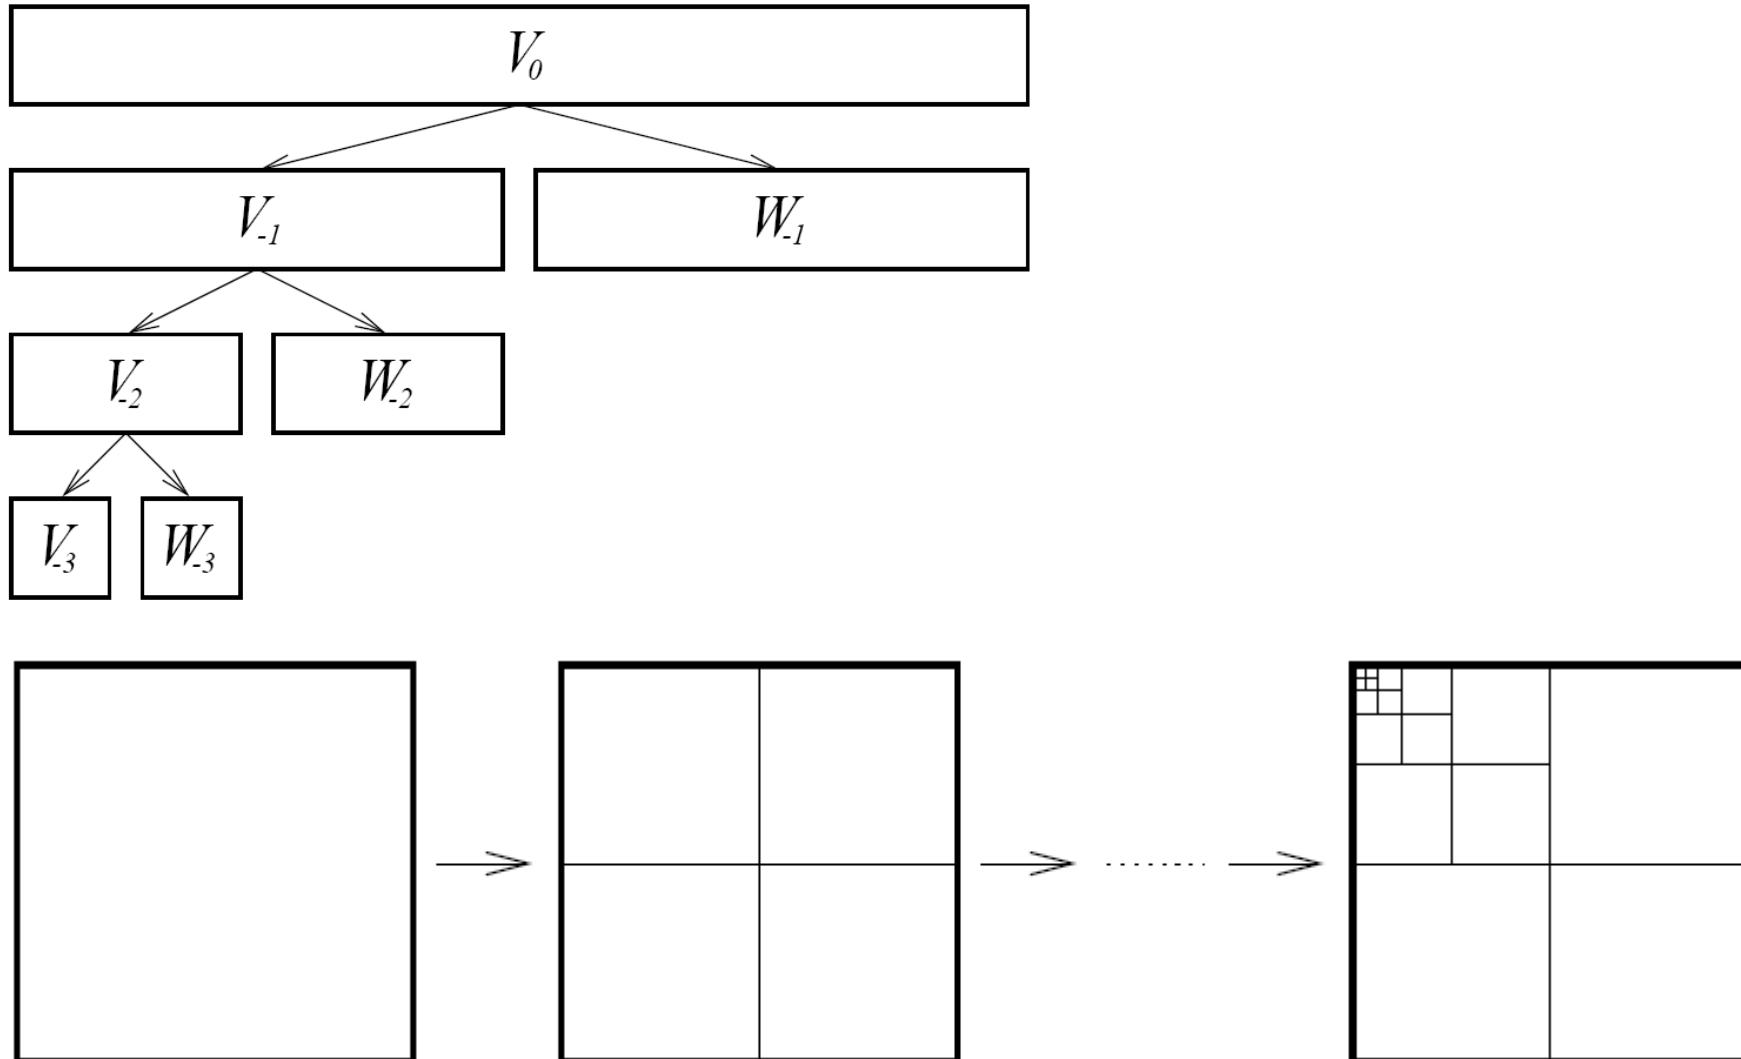


- Multi Resolution Analysis



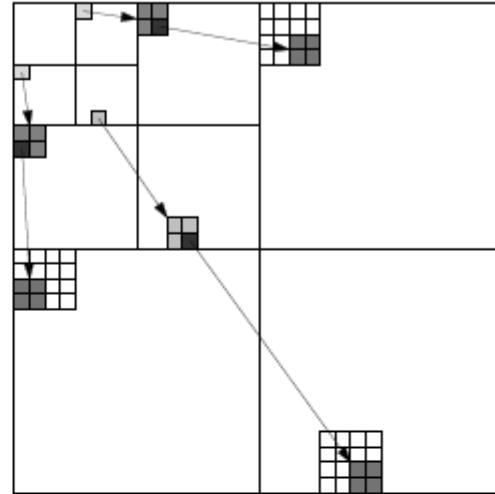
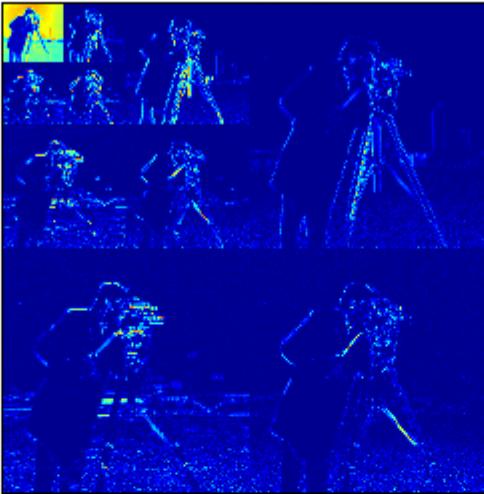


# Coefficients





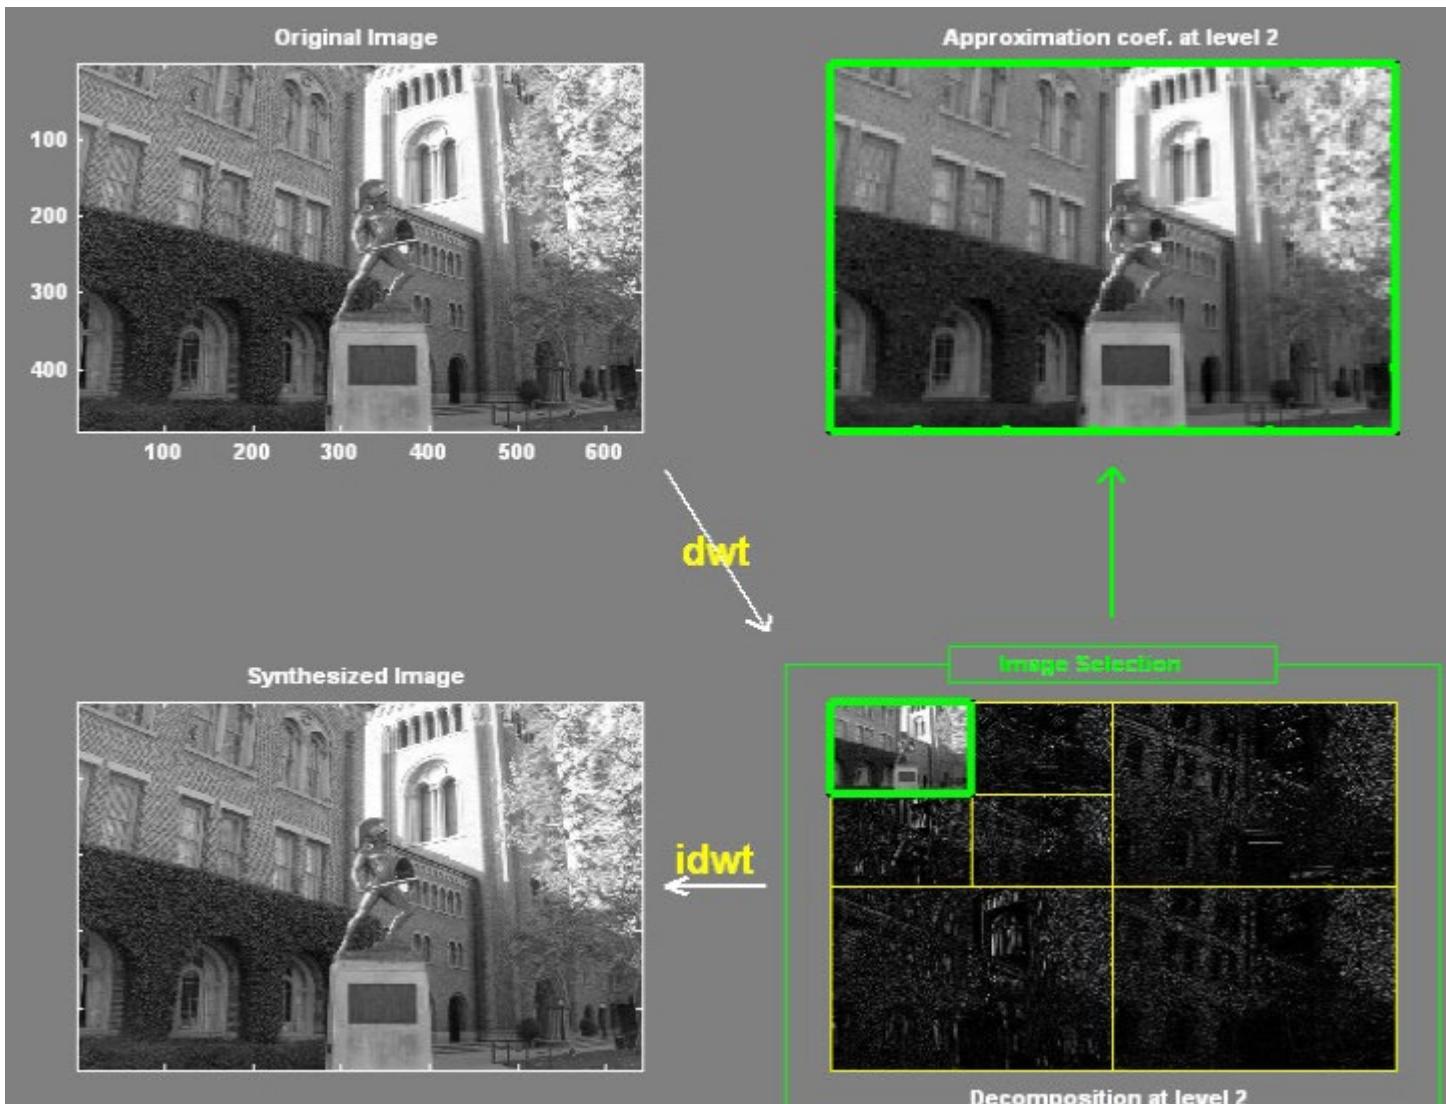
# 2D Wavelets



- Sparse structure: few large coeffs, many small coeffs
- Basis for JPEG2000 image compression standard
- Wavelet approximations: smooths regions great, edges much sharper
- *Fundamentally better than DCT for images with edges*

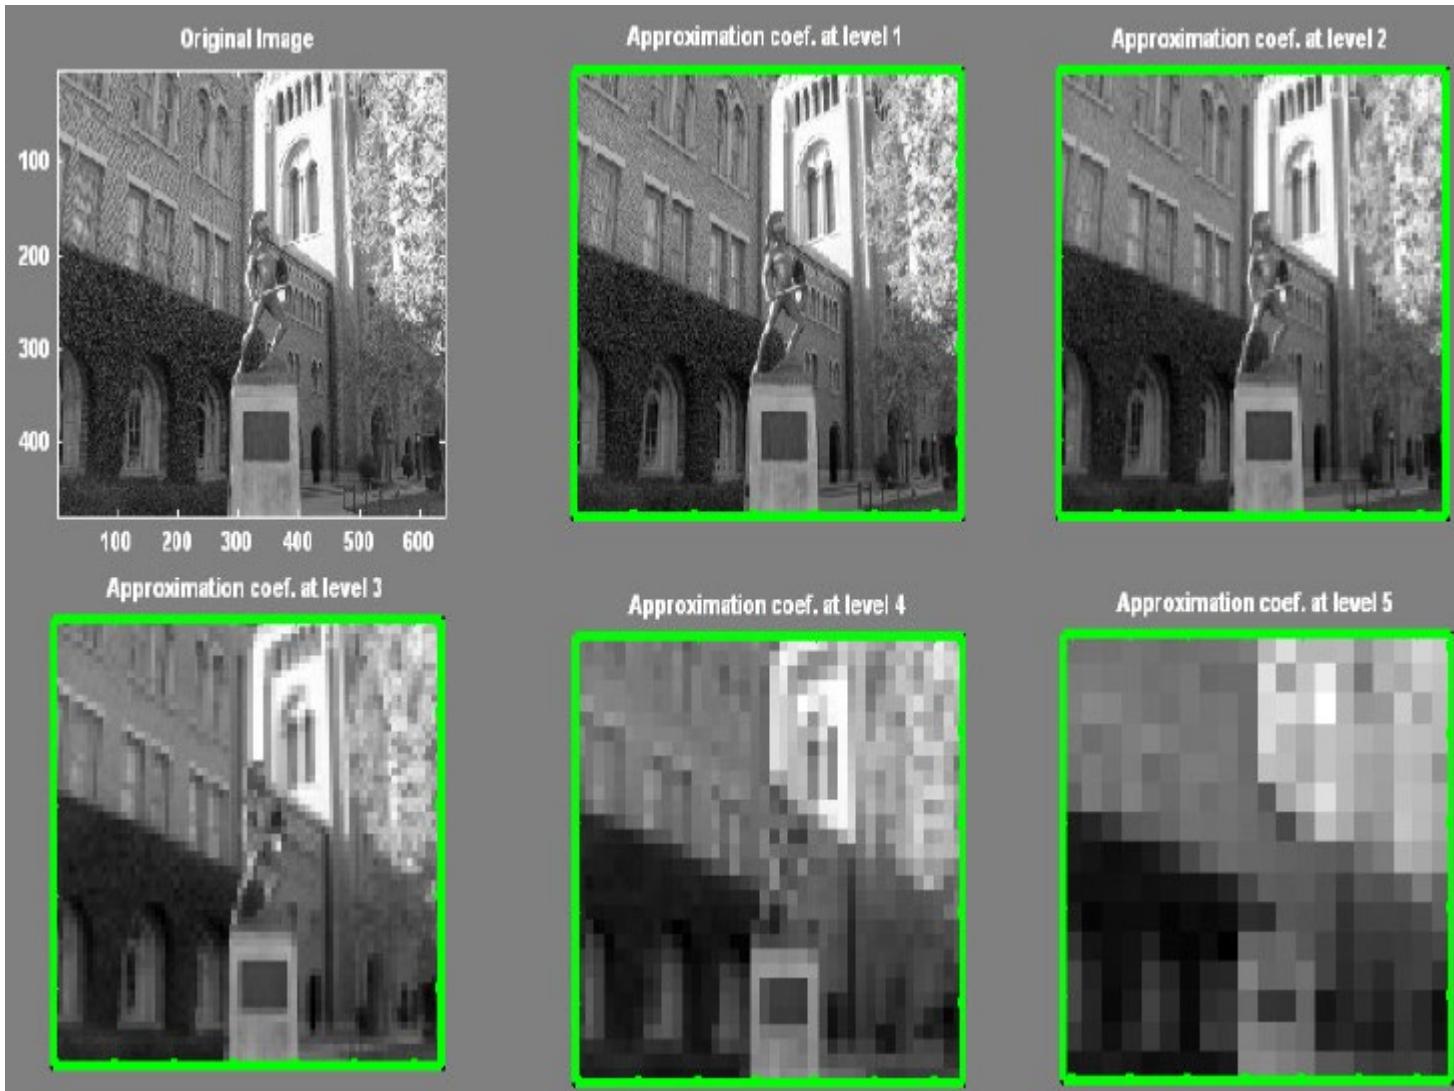


# Example





# Example





# Questions

---

- What happens if I leave out the detail coefficients of the last two levels?
- How many detail coefficients do you need to faithfully represent an image?  
10%, 25%, 50%, 75% or 100%?

# Wavelet Representation

- For representing  $H(x)$  in Haar wavelet basis you need only
  - scaling coefficient for the first level
  - all detail coefficients

$$H(x) = H_{0,0}^0 \phi_0^0 + \sum_l \sum_t H_{t,1}^l \psi_t^l = \sum_i H_i \Psi_i$$

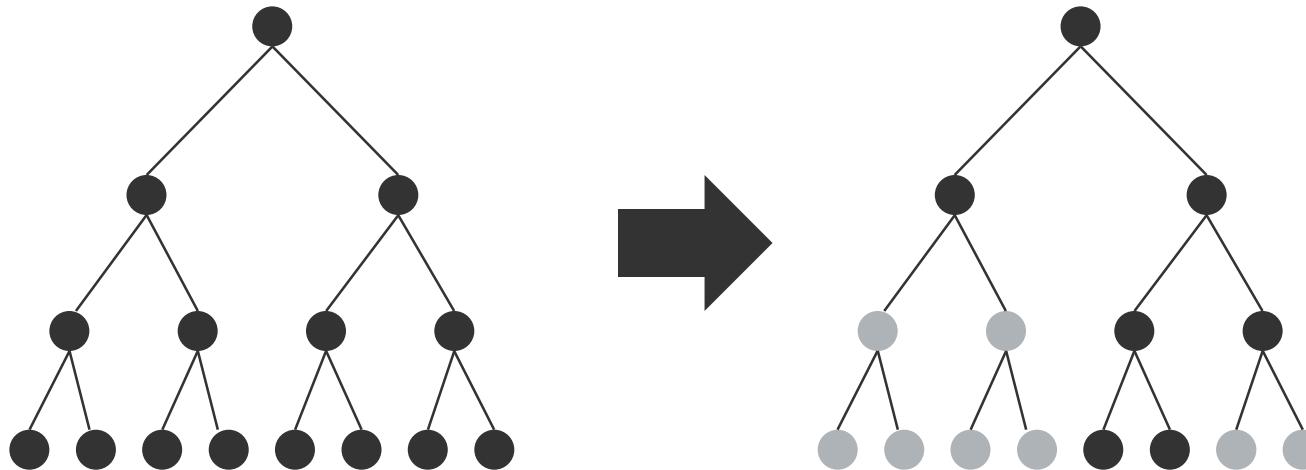
- For a sampled function a lossless representation is achieved if #samples==#wavelet coefficients
- Compression by setting small wavelet coefficients to zero





# Sparse Representation

- First measure entire matrix
- Then compress by determining the sparse tree
  - since all information is given you are on the save side





# Compression - Shrinkage

- The goal of compression is to express an initial set of data using some smaller set of data, either with or without loss of information.

$$f(x) = \sum_{i=1}^m c_i u_i(x) \longrightarrow \tilde{f}(x) = \sum_{i=1}^{\tilde{m}} \tilde{c}_i \tilde{u}_i(x)$$

- For orthonormal basis the square of the *L2 error in this approximation* is

$$\left\| f(x) - \tilde{f}(x) \right\|_2^2 = \sum_{i=\tilde{m}+1}^m (c_i)^2$$

- Compression – set all coefficients  $< t$  to zero!

such that  $\tilde{m} < m$  and  $\left\| \tilde{f}(x) - f(x) \right\| < \varepsilon$  for some norm



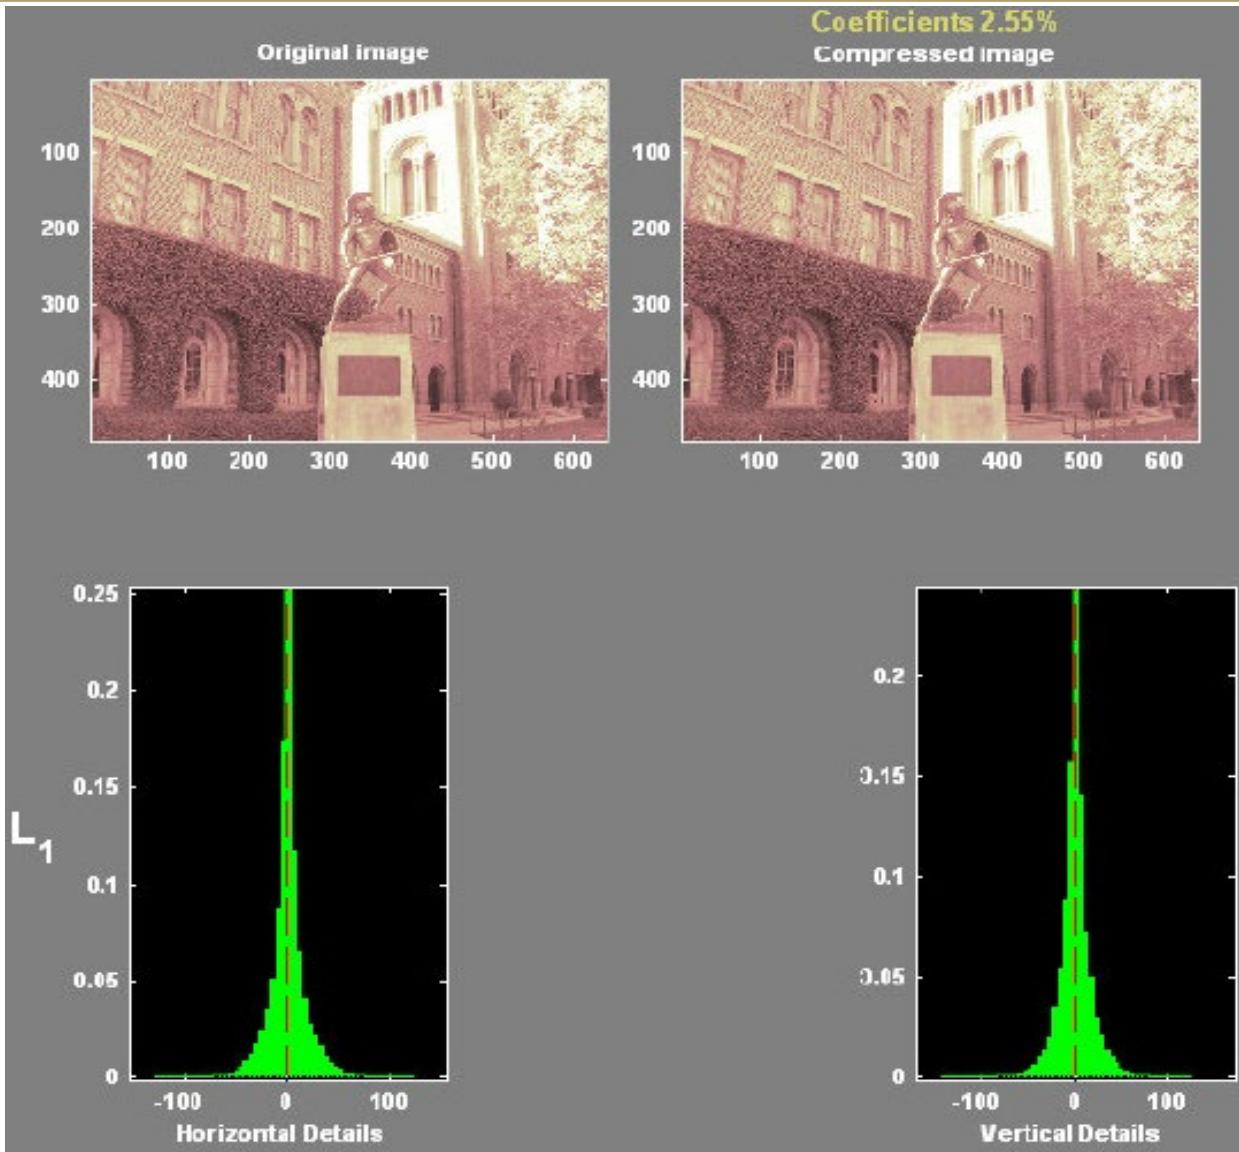
# Wavelet image compression

1. Compute coefficients  $c_1, \dots, c_m$  representing an image in a normalized two-dimensional Haar basis.
2. Sort the coefficients in order of decreasing magnitude to produce the sequence  $c(1), \dots, c(m)$ .
3. Find the smallest  $\tilde{m}$  so that

$$\sum_{i=\tilde{m}+1}^m (c_{\sigma(i)})^2 \leq \varepsilon^2$$



# Natural Signals are often sparse

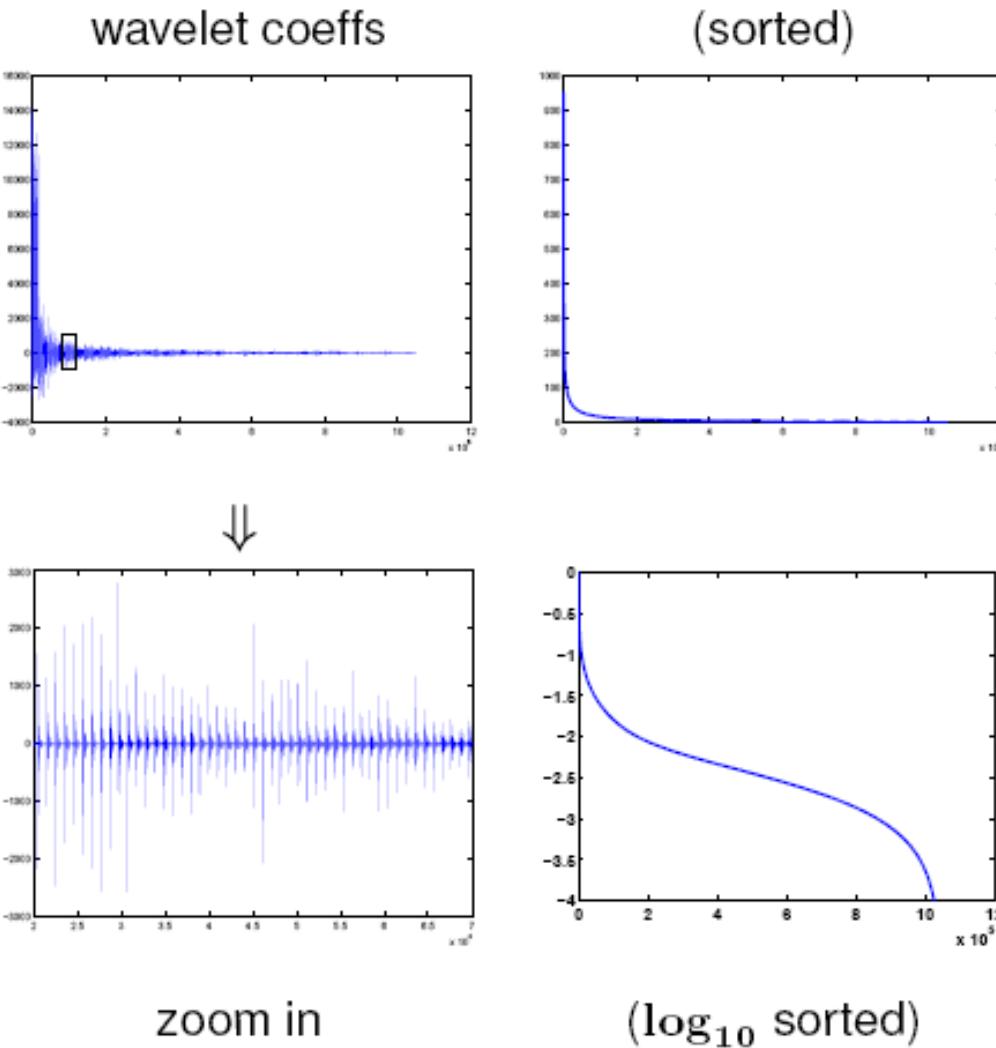




# Sparsity of the Signal



1 megapixel image





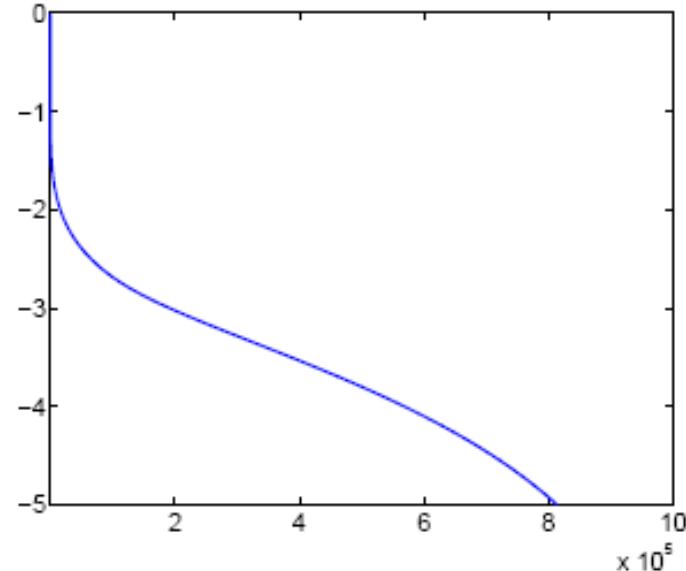
# Decay of Coefficients



1 megapixel image



25k term approx



approx error

- Within 2 digits (in MSE) with  $\approx 2.5\%$  of coeffs
- Original image =  $f$ ,  $K$ -term approximation =  $f_K$

$$\|f - f_K\|_2 \approx .01 \cdot \|f\|_2$$



# Other Types of Wavelets

---

- Mexican Hat
  - Daubechies Wavelets
  - .....
- 
- To better represent edges:
  - Curvelets
  - Ridgelets



# Properties of Wavelets

---

- Multi-resolution
- Locality both in time/space domain and frequency domain
- (bi)Orthogonality
- Smooth basis functions
- Good approximations in  $L^2$
- Fast transform algorithm  $O(n)$
- Stable decomposition
- Sparse coefficients

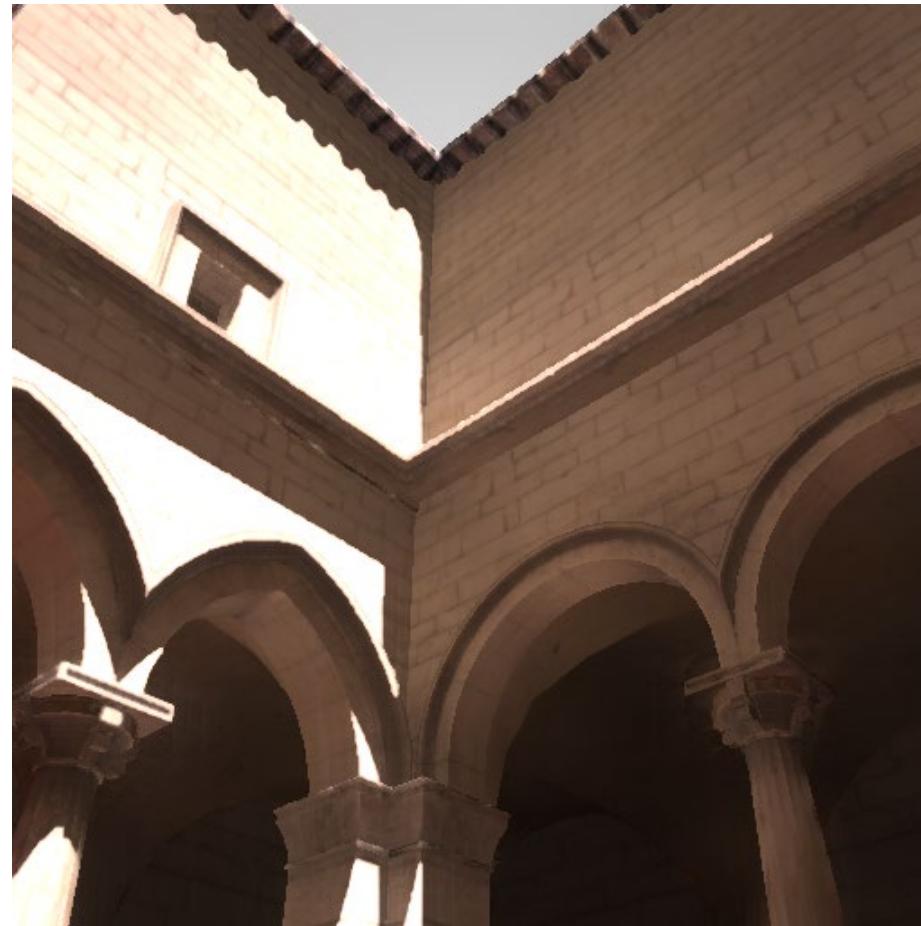


# Edge-Avoiding À-Trous Wavelets

Combining Wavelets and Bilateral Filtering

# Edge-Avoiding À-Trous Wavelet Transform for fast Global Illumination Filtering

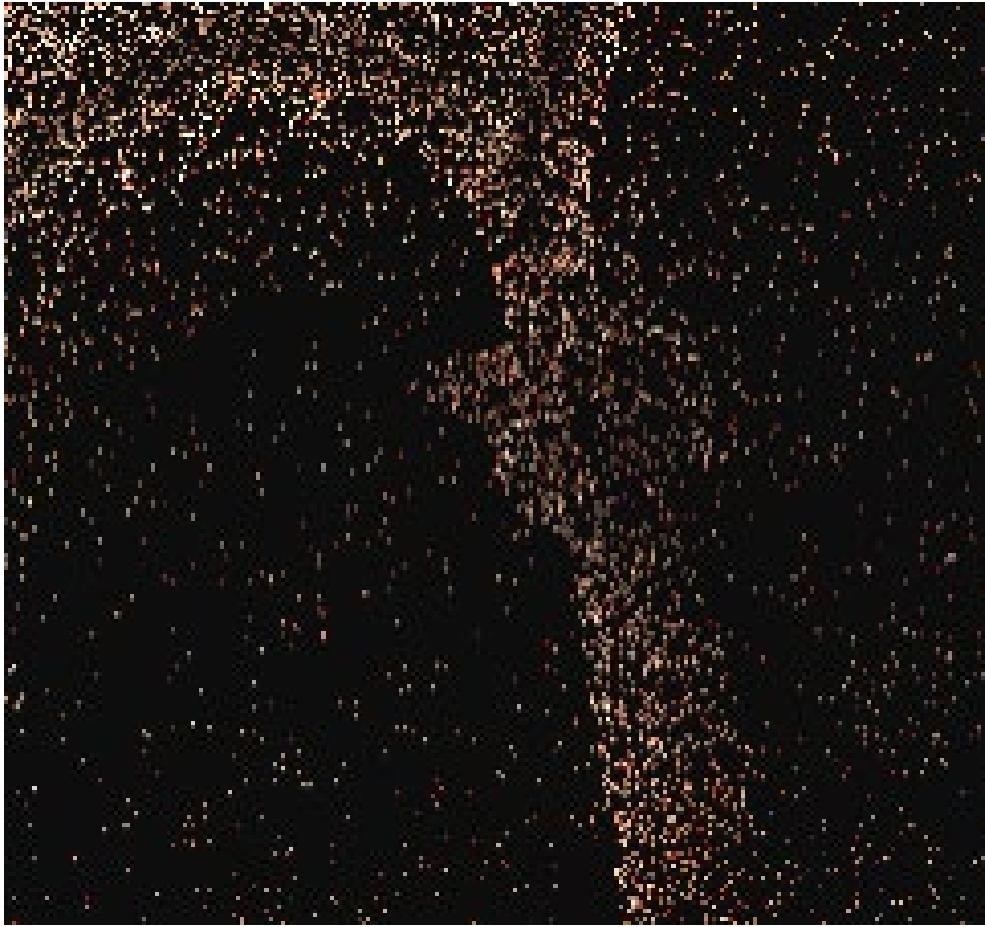
---



[Dammertz, Sewtz, Hanika, Lensch – HPG 2010]

# Edge-Avoiding À-Trous Wavelet Transform for fast Global Illumination Filtering

---



*[Dammertz, Sewtz, Hanika, Lensch – HPG 2010]*

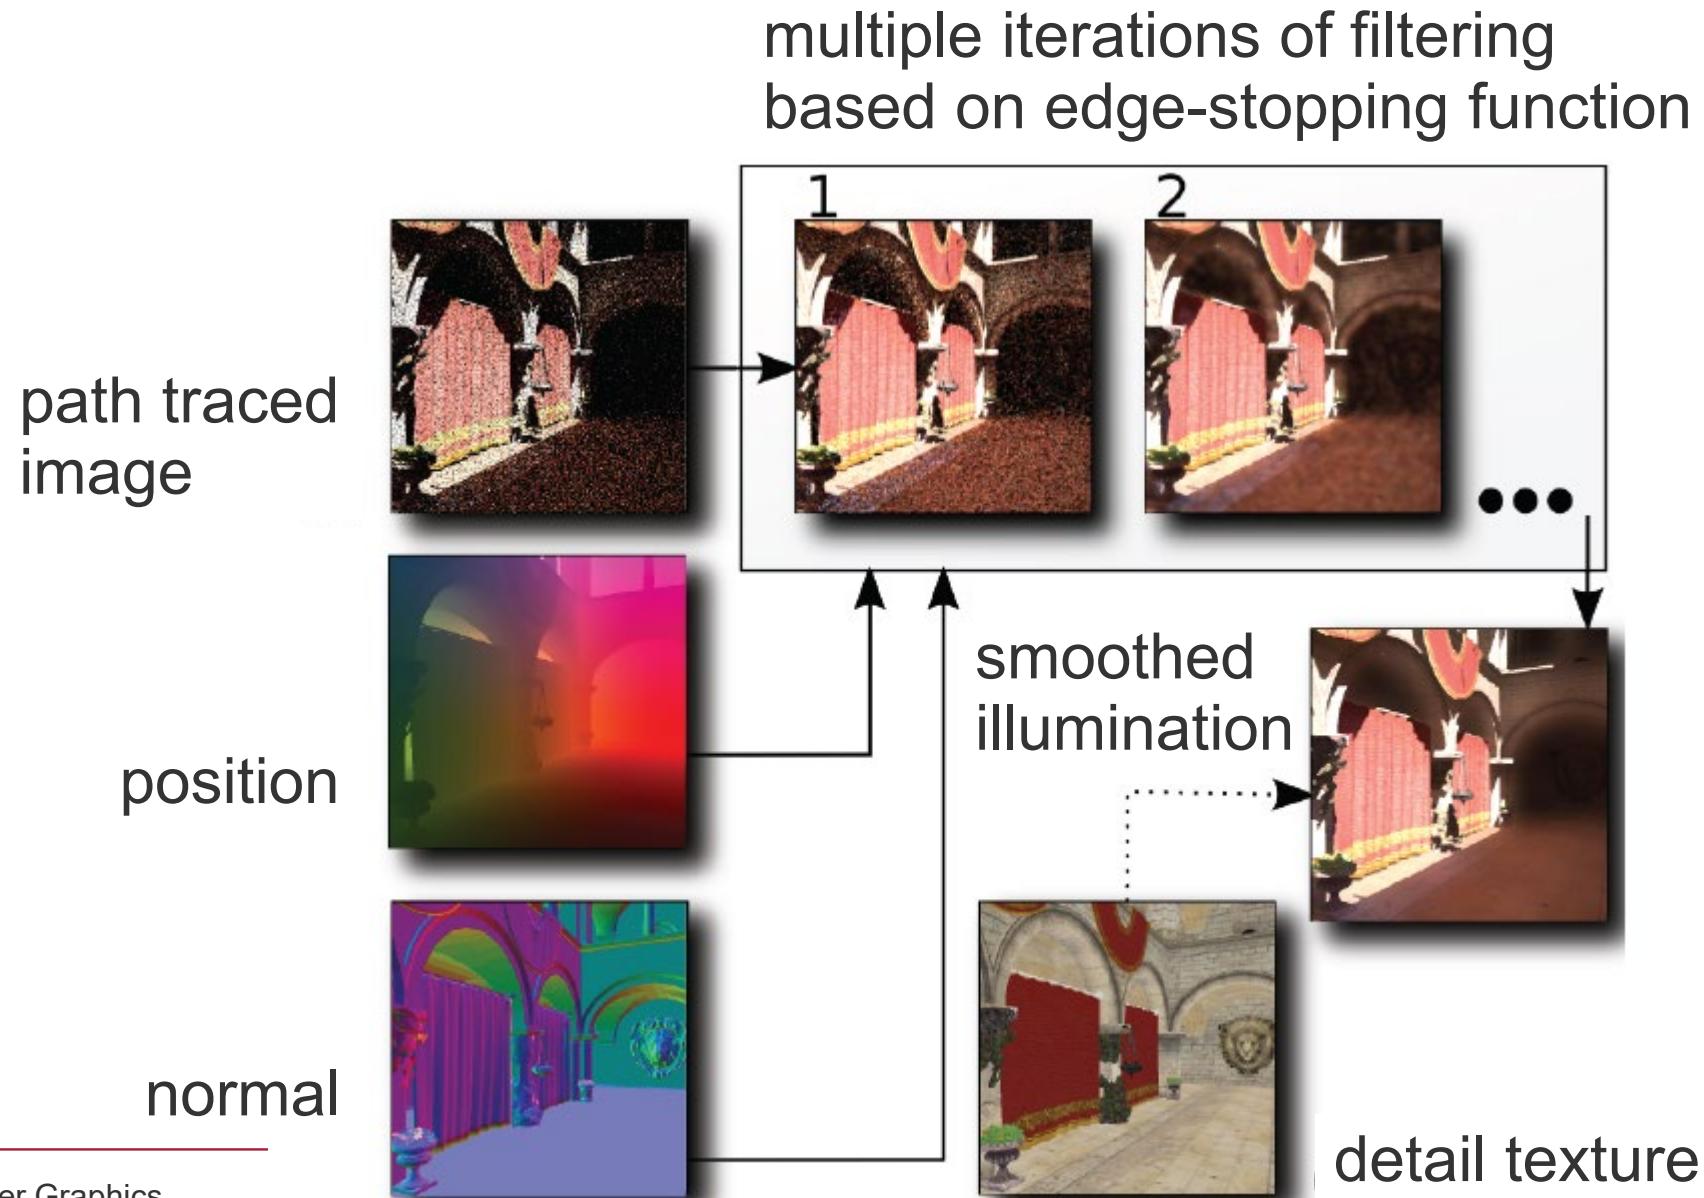
HPG 2010  
paper1029

Edge-Avoiding A-Trous Wavelet Transform  
for fast Global Illumination Filtering

supplemental Material



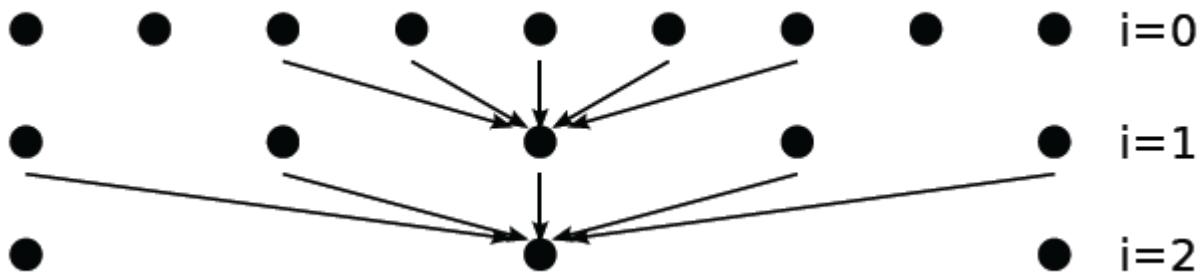
# Pipeline





# Decimating Wavelet Filtering

- Apply the same filter multiple times
  - reduce the image resolution at each iteration

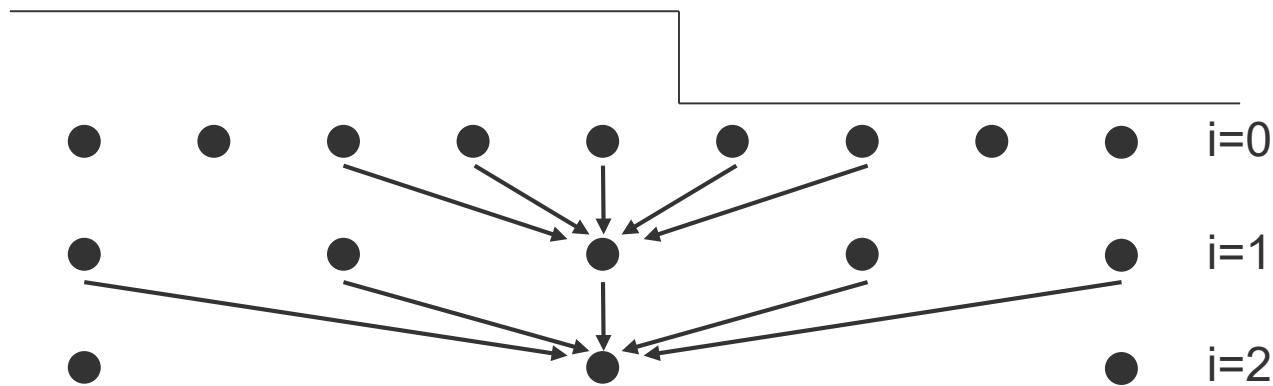


- Benefits:
  - multi-scale representation
  - fast



# Decimating Wavelet Transform

- Apply the same filter multiple times
  - reduce the image resolution at each iteration

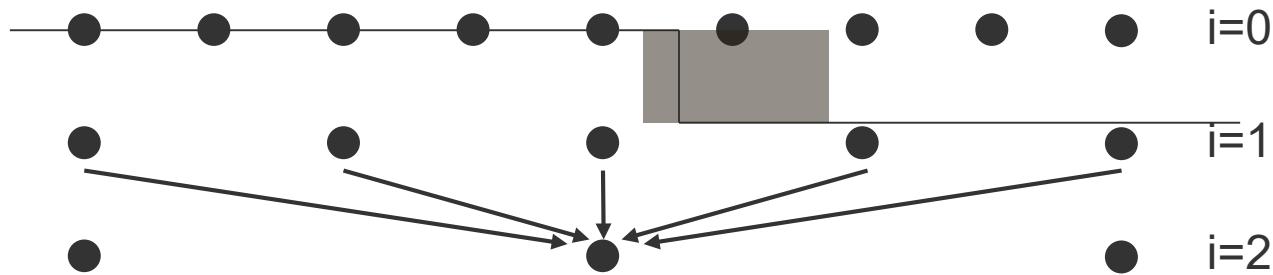


- Benefits:
  - multi-scale representation, fast!
- Drawback: filtered information at few locations only
  - Where is the edge?



# Decimating Wavelet Transform

- Apply the same filter multiple times
  - reduce the image resolution at each iteration

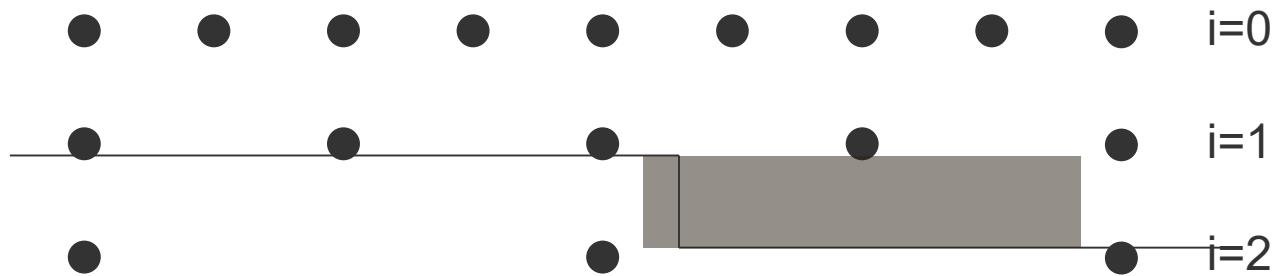


- Benefits:
  - multi-scale representation, fast!
- Drawback: filtered information at few locations only
  - Where is the edge?



# Decimating Wavelet Transform

- Apply the same filter multiple times
  - reduce the image resolution at each iteration

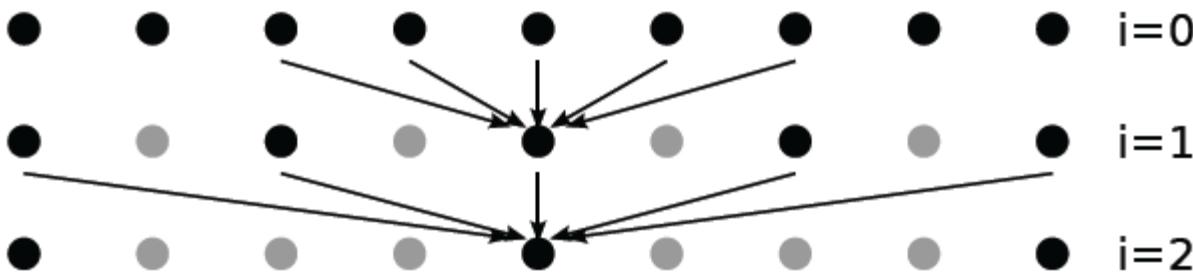


- Benefits:
  - multi-scale representation, fast!
- Drawback: filtered information at few locations only
  - Where is the edge?



# À-Trous Wavelet Filtering

- “With holes”
- At each iteration convolve with a Gaussian
  - double the filter size
  - by introducing more and more holes

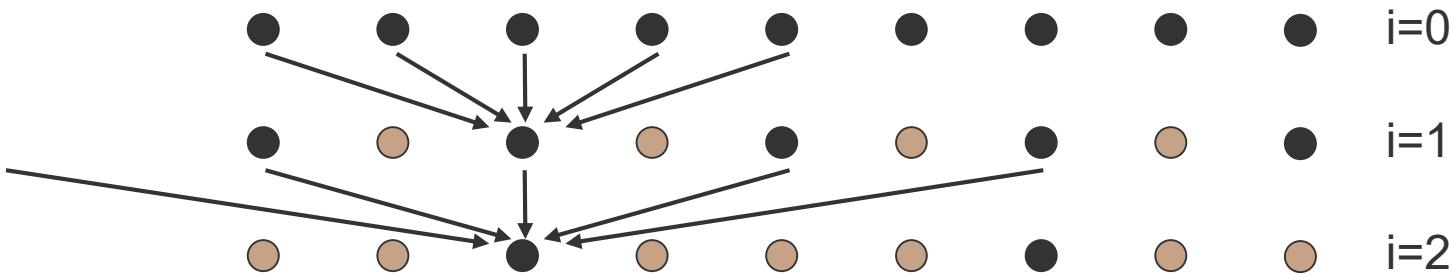


- Benefits:
  - constant effort per iteration (in contrast to undecimated wavelets)
  - filtered information at each pixel (in contrast to decimated wavelets)



# À-Trous Wavelet Transform

- “With holes”
- At each iteration convolve with a Gaussian
  - double the filter size
  - by introducing more and more holes

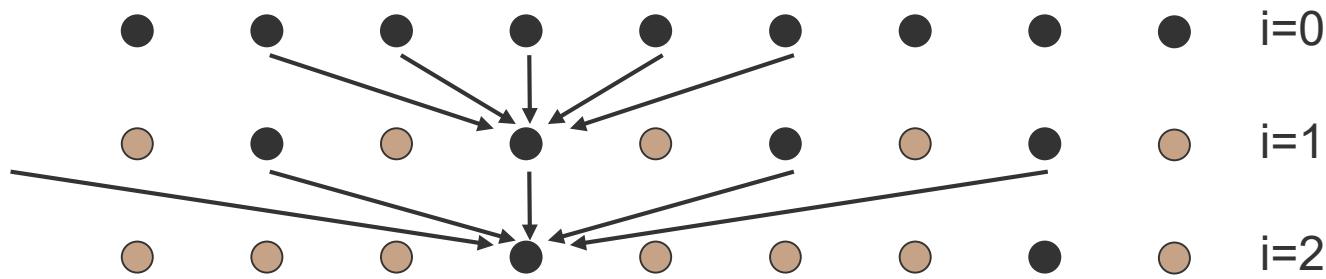


- Benefits:
  - constant effort per iteration (in contrast to undecimated wavelets)
  - filtered information at each pixel (in contrast to decimated wavelets)



# À-Trous Wavelet Transform

- “With holes”
- At each iteration convolve with a Gaussian
  - double the filter size
  - by introducing more and more holes

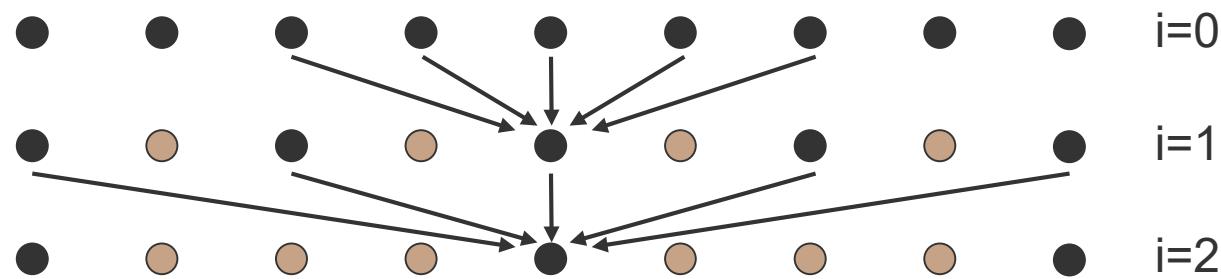


- Benefits:
  - constant effort per iteration (in contrast to undecimated wavelets)
  - filtered information at each pixel (in contrast to decimated wavelets)



# À-Trous Wavelet Transform

- “With holes”
- At each iteration convolve with a Gaussian
  - double the filter size
  - by introducing more and more holes



- Benefits:
  - constant effort per iteration (in contrast to undecimated wavelets)
  - filtered information at each pixel (in contrast to decimated wavelets)

# À-Trous Wavelet Transform

---

1. At level  $i = 0$ , start with input signal  $c_0(p)$ .
2.  $c_{i+1}(p) = c_i(p) * h_i$ , where  $*$  is the discrete convolution.

The distance between the entries in the filter  $h_i$  is  $2^i$ .

3.  $d_i(p) = c_i(p) - c_{i+1}(p)$   
where  $d_i$  are the detail or wavelet coefficients of level
4. Repeat 2 and 3 until  $i = N$  (number of levels to compute)
5.  $\{d_0, d_1, \dots, d_{N-1}, c_N\}$  is the wavelet transform of  $c$ .
6. The reconstruction is given by  $c = c_N + \sum_{i=0}^{N-1} d_i$

# À-Trous Wavelet *Filtering*

---

1. At level  $i = 0$ , start with input signal  $c_0(p)$ .
2.  $c_{i+1}(p) = c_i(p) * h_i$ , where  $*$  is the discrete convolution.

The distance between the entries in the filter  $h_i$  is  $2^i$

3.  $d_i(p) = c_i(p) - c_{i+1}(p)$

where  $d_i$  are the detail or wavelet coefficients of level

4. Repeat 2 and 3 until  $i = N$  (number of levels to compute)
5.  $\{d_0, d_1, \dots, d_{N-1}, c_N\}$  is the wavelet transform of  $c$ .
6. The **filtered** reconstruction is given by

$$c = c_N + \sum_{i=0}^{N-1} \alpha_i d_i$$



# Edge-Stopping Function

- Compute weighted convolution  
(compare to bilateral filter)

$$c_{i+1}(p) = \frac{\sum_{q \in Q} h_i(q) \cdot w(p, q) \cdot c_i(p)}{\sum_{q \in Q} h_i(q) \cdot w(p, q)}$$

$$[c_{i+1}(p) = c_i(p) * h_i]$$

- With weights  $w(p, q) = w_X(p, q)$  with
- $\sigma_X$  can be made large for the first two levels

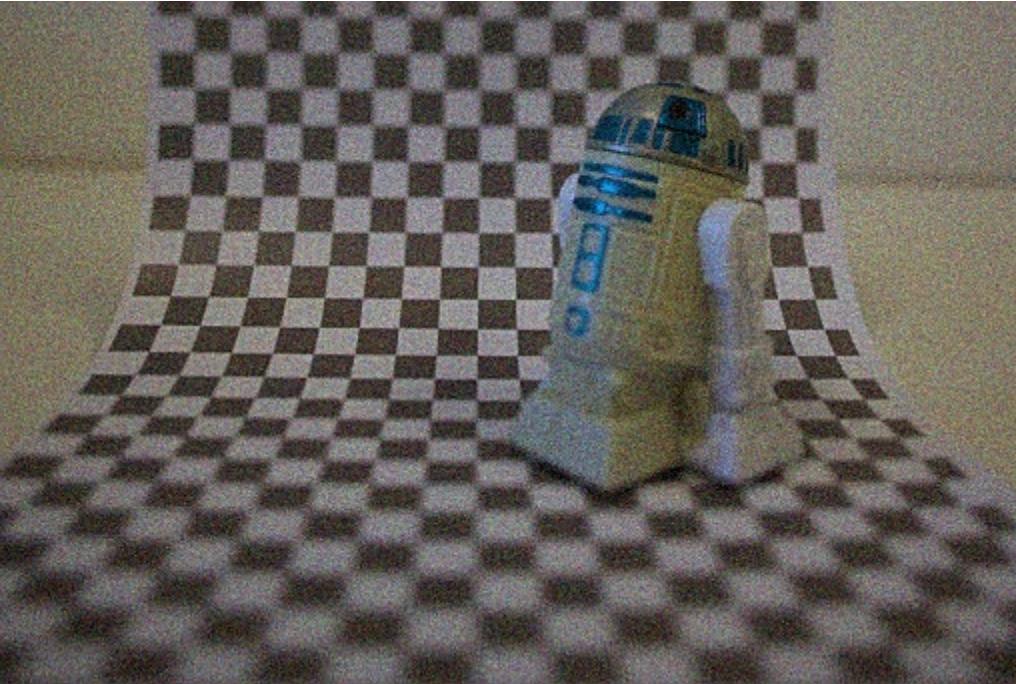
$$w_X(p, q) = e^{\left( -\frac{\|I_p - I_q\|^2}{\sigma_X^2} \right)}$$





# Input

---





# Level 0

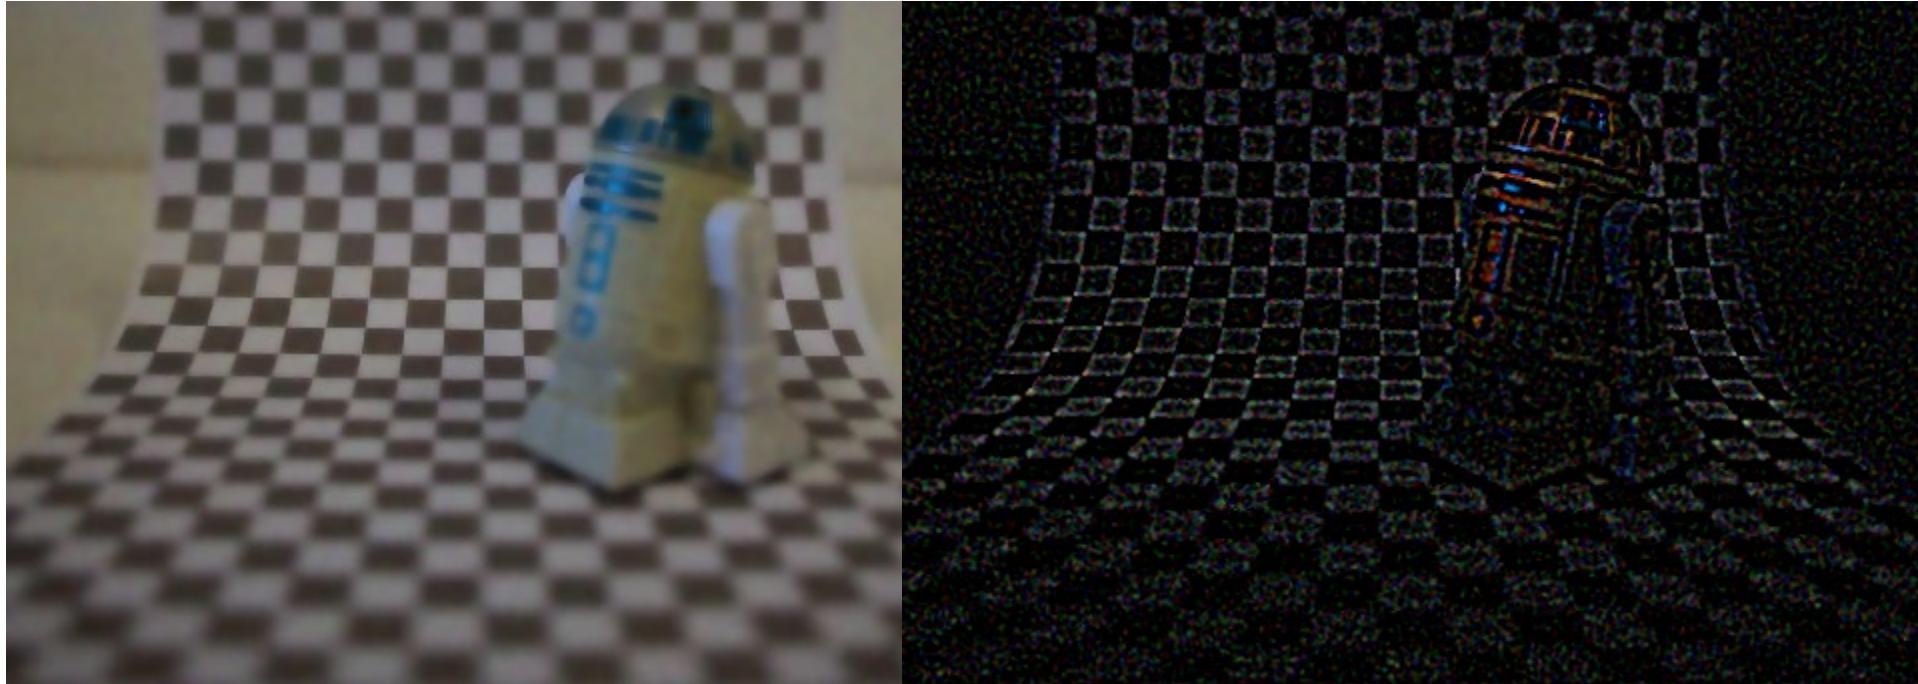
---





# Level 1

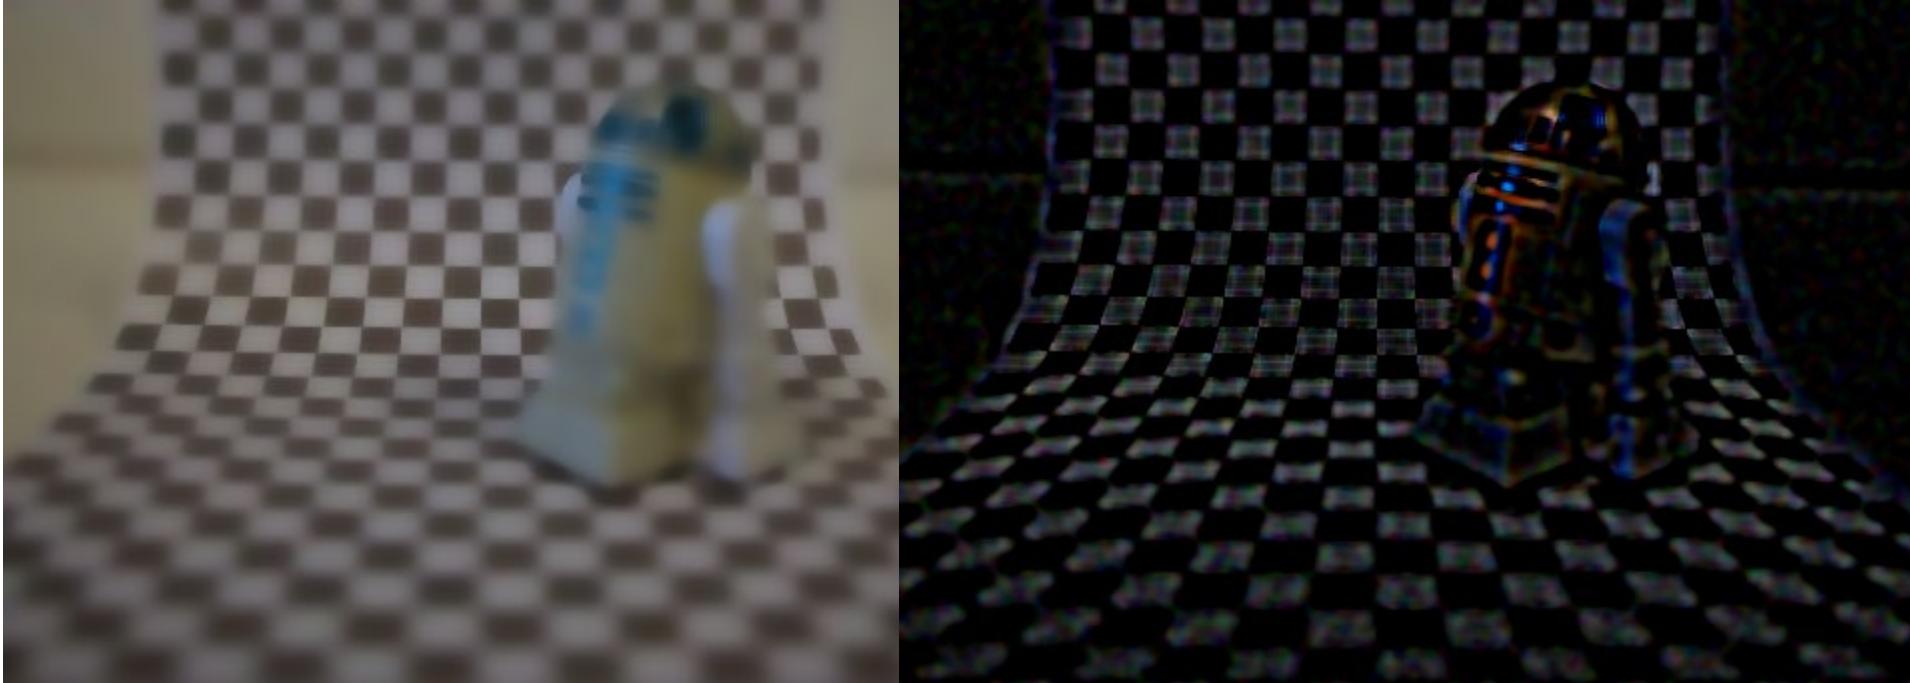
---





# Level 2

---



# Denoising by Shrinkage

---

- Estimate standard deviation of the noise in the image:

$$\sigma_n = \frac{\text{median}(|d_0|)}{0.6745}.$$

- Compute optimal shrinkage threshold based on the signal variance

$$T = \frac{\sigma_{n,i}^2}{\sqrt{\max\{0, \sigma_{y,i}^2 - \sigma_{n,i}^2\}}} \quad ($$

$$\text{with } \sigma_{y,i}^2 = \frac{1}{N} \sum_p d_i(p)^2 \text{ and } \sigma_{n,i} = \sigma_n \cdot 2^{-i}.$$

- Apply shrinkage and contrast boost (if wanted):

$$d'_i = \max\{0, |d_i| - T\} \cdot \text{sign}(d_i)$$

$$\text{and } c_{i-1} = c_i + \beta \cdot d'_i,$$



# Denoising

input:10% noise PSNR 26.274



à-trous PSNR 34.8321



edge-optimized à-trous:PSNR 35.887

# Denoising

---

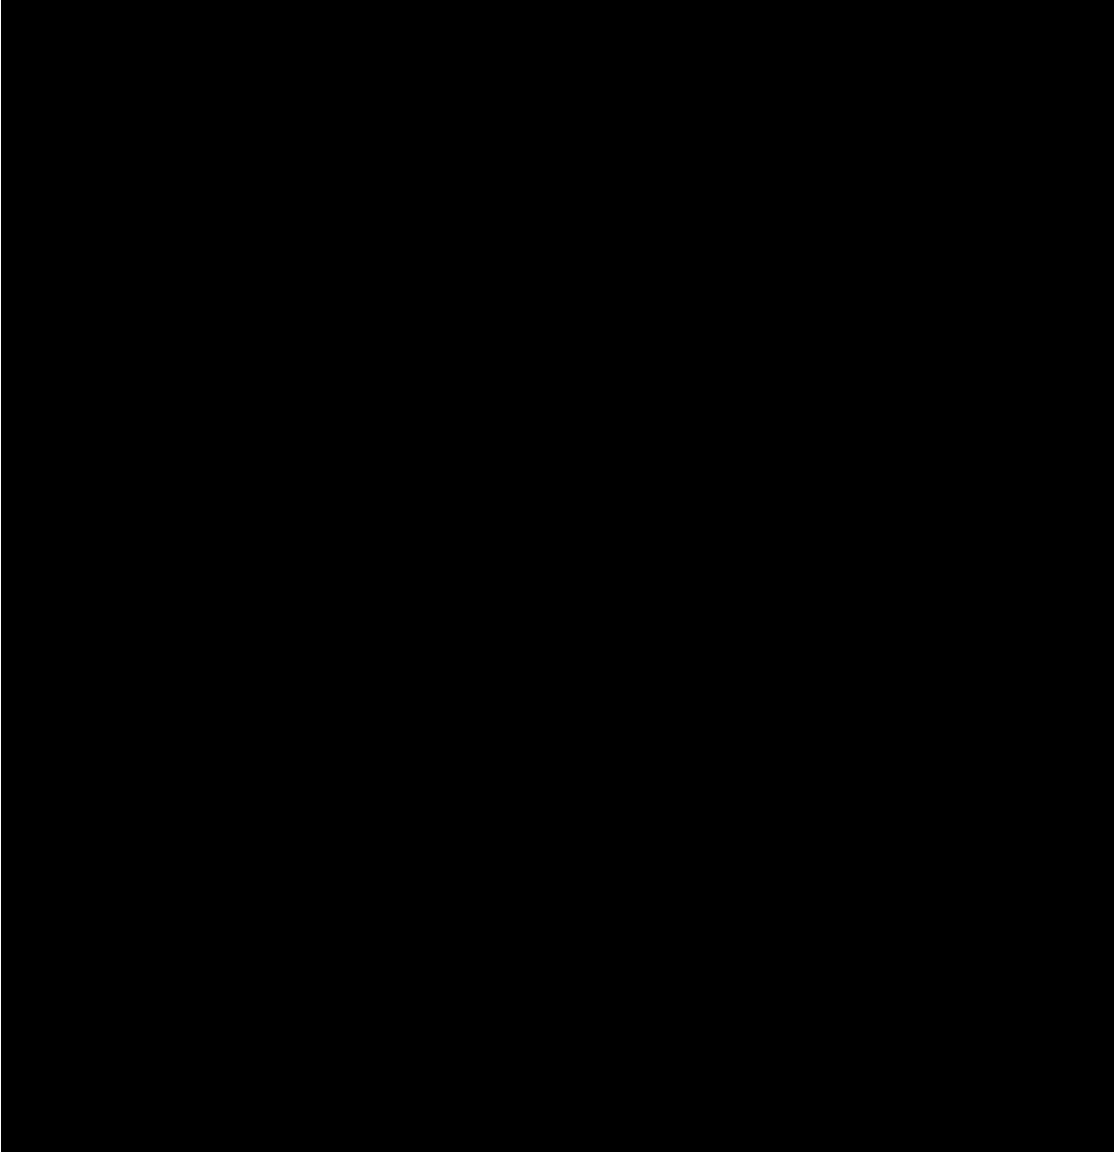
input:10% noise PSNR 26.274





# Contrast Enhancement

---



# Edge-Avoiding À-Trous Wavelets – Contrast Enhancement



# Edge-Avoiding À-Trous Wavelets – Smoothing





# Wrap-Up

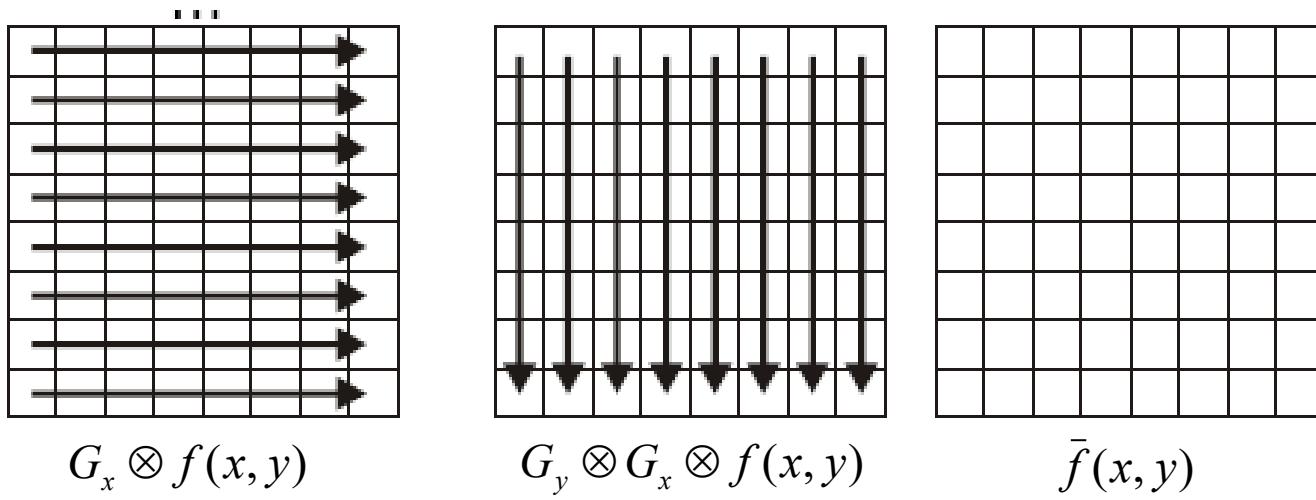
---



# Gaussian - Implementation

- Gaussian filter is separable

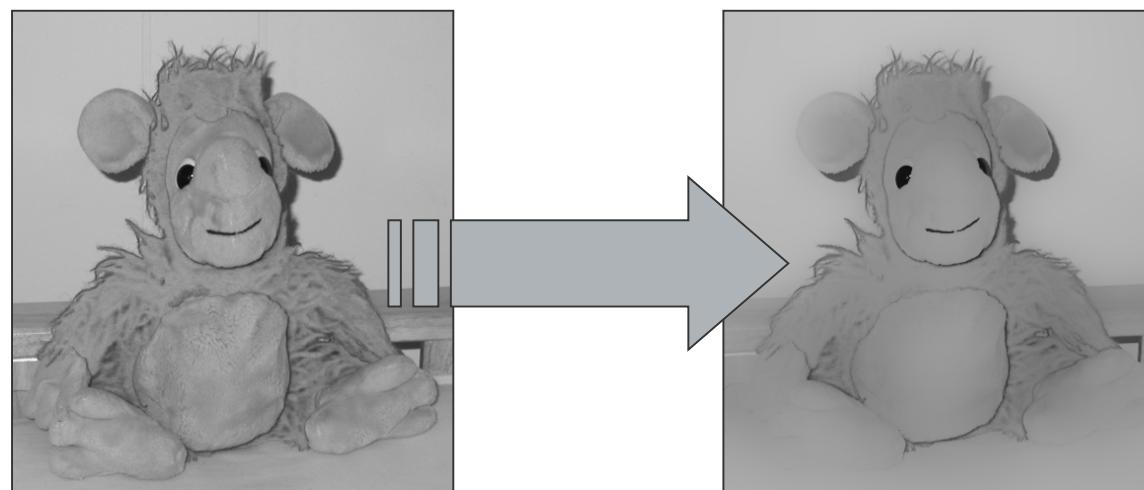
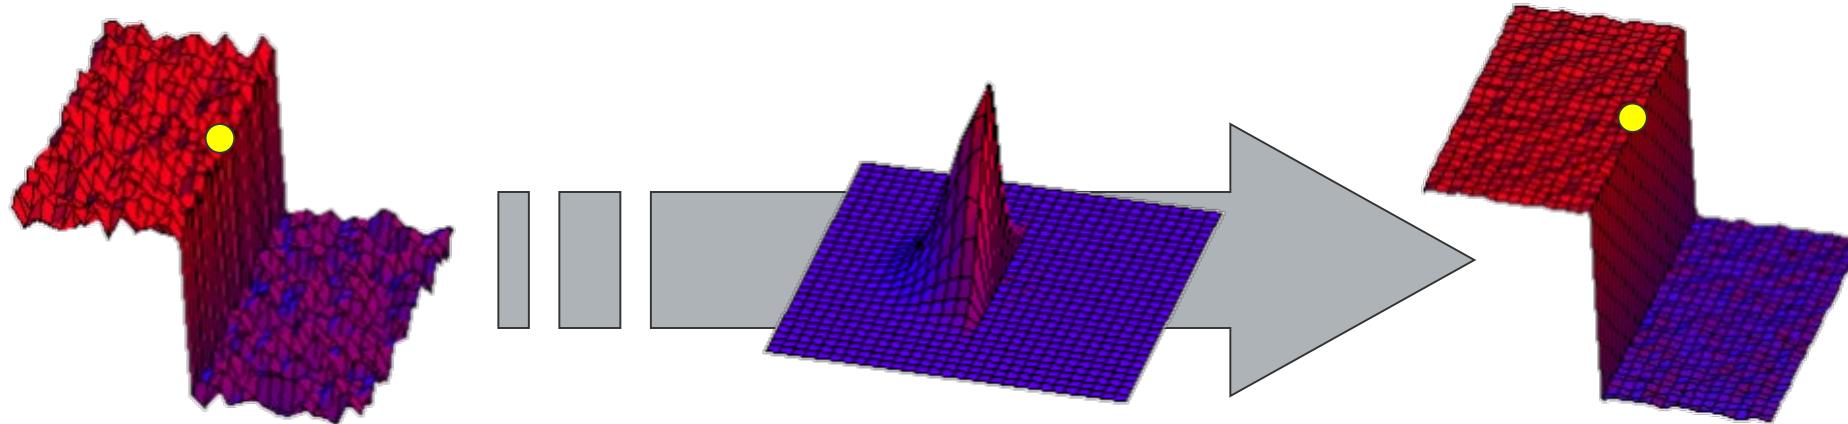
$$\bar{f}(x, y) = G_{(x,y)} \otimes f(x, y) = G_y \otimes (G_x \otimes f(x, y))$$





# Large-scale Layer

- **Bilateral filter – edge preserving filter**
- Smith and Brady 1997; Tomasi and Manducci 1998; Durand et al. 2002



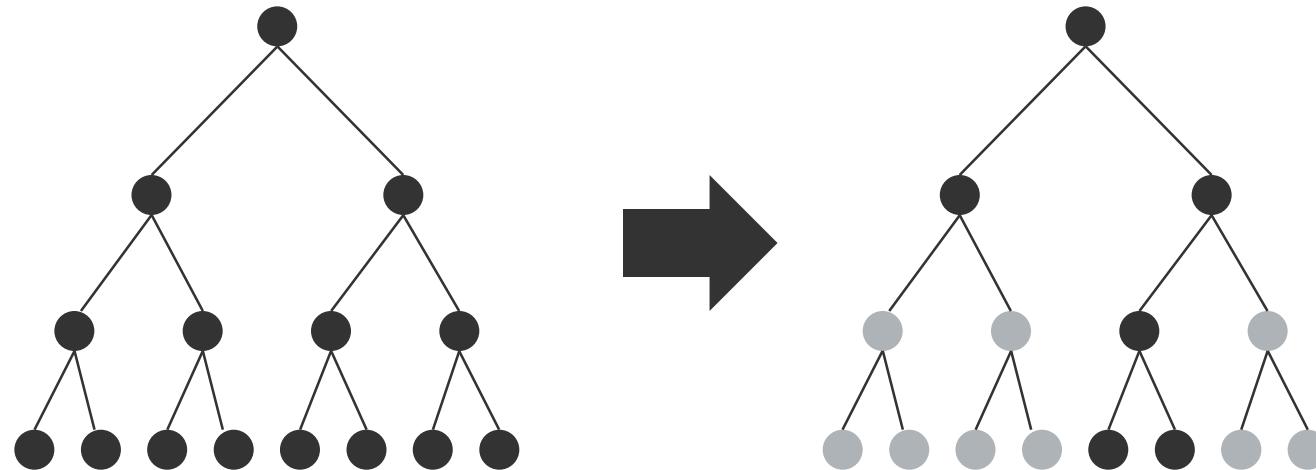
# Edge-Avoiding À-Trous Wavelets – Smoothing





# Sparse Representation

- First measure entire matrix
- Then compress by determining the sparse tree
  - since all information is given you are on the save side





# Summary

---

- Simple Filters
  - Gauss
  - Sobel
  - Median, ...
- Bilateral
  - Upsampling
  - Cross-modal Upsampling
- Wavelets
  - Haar
  - $\tilde{A}$ -Trous
  - Edge-avoiding  $\tilde{A}$ -Trous
  - Denoising