

Spatial Acceleration Structures

Why spatial acceleration structures?

- For naïve ray tracing you have to check for each ray all triangles in the scene
- this a lot to compute
 - sort primitives spatially in a hierarchy


Bounding Volumes

pack primitives into simple volumes → only check primitive intersection, if bounding volume is hit

- sphere: + easy intersection calculation
- inefficient, because too large
- aabb: + easy intersection calculation
- sometimes too large
- oriented bounding box (non-axis-aligned): + better fit
- complex intersection computation



Bounding Volume Hierarchy




organize bounding volumes hierarchically

- + very good adaptivity
- + efficient traversal $O(\log n)$
- how to arrange BVs?
 - ↳ avoid overlapping of bb on same level

Grid

partition scene with equal sized voxels

→ one object can be represented in multiple cells



+ trivial insertion of objects

+ easy construction

- high memory costs (a^4)

- expansive traversal

↳ a lot of empty voxels to traverse

↳ which voxel to traverse next? (Bresenham Algorithm)

- stop traversal after detecting intersection

↳ otherwise overlapping objects will be intersected twice

↳ alt. solution: mailboxing (store index/id of intersected primitive)

- scene dependent grid resolution ("Teapot in a stadium")


↳ solution: hierarchical grid

Quadtree


hierarchical subdivision into 4 cells each level

subdivide until cell is empty or has less than n primitives

→ tree traversal structure



1	2
3	4





Octree

- 3D subdivision into 8 Voxels
- complex traversal
- slow to refine complex regions

- + simple construction in $O(n)$
- + simple insertion of objects
- complex traversal $O(k \cdot \log n)$
- high memory consumption



Binary Space Partition (BSP)

- recursively split space into halves
- „arbitrary“ direction of planes



kdt-Tree

- same principle as BSP, but axis aligned
- planes are defined by
 - axis flag
 - split point
 - child pointer



BSP / kd-Tree Traversal

- „Front-to-Back“ traversal (start with root node)
- child nodes are traversed in order along the ray
- traversal stops with first intersection
- implementation with stack

function traverse(ray, root):

node = root

intersection == false

while intersection == false:

if node is leaf:

intersection = checkTriangleIntersection(node)

if intersection == false:

node = pullFromStack

if node is empty:

return "no intersection"

else: continue

else:

if node.far-child is hit by ray:

pushOnStack(node.far-child)



node = node.close-child

while end

return intersectionPoint(node)

traversal example



1.



Process	Stack
A	

depending on intersection point with plane A,
we know, we hit the outer bounding box
→ check both children, first the closer one



2.



Process	Stack
C	B

because the ray isn't intersecting plane C,
we know it only intersects the closer child



3.



Process	Stack
D	B

both objects in D could be intersected
→ check the closer one first, stack the others

4.





Process	Stack
1	B
	2,3

object intersection check is negative

→ pop next one from stack

5.





Process	Stack
2,3	B

object intersection check is negative

→ pop root child node B from stack



6.



Process	Stack
B	

intersection with plane B means both children are hit by the ray → first check closer objects

7.





Process	Stack
3,4	5,6

no intersection detected

→ pop other objects from stack

8.



Process	Stack
5,6	

intersection with object 5 detected

→ terminate traversal

Surface Area Heuristic (SAH)

build good kd-Trees by accounting probability and costs

- probability of a ray hits the cell is proportional to its size
- cost of the cell is given by triangle count → more triangles = higher costs

$$C(\text{cell}) = C_{\text{trav}} + p(\text{hitL}) \cdot C(L) + p(\text{hitR}) \cdot C(R)$$

⇒ produces large chunks of empty space

⇒ automatically and rapidly isolates complexity

Large Scenes

what if spatial data does not fit into memory?

- **Lazy build:** build subtree only when it's needed (potential intersection)
 - + no memory wasted
 - can be slow, same subtree has to be build over and over again
 - **Lazy build with caching:** each needed subtree will be stored when it's built for the first time (if memory is full, some subtrees will be deleted)
 - lots of tests
 - inefficient deletion of subtrees
 - **Multi-Level-Hierarchy:** first levels of tree are fixed, lazy build on lower levels
 - + often visited nodes (upper levels) don't have to be rebuild
- **Ray-Reordering:** reminds frequently hit elements and keeps them in memory

Dynamic Scenes

in moving scenes you have to update your spatial structure fast

- combine kdt-Trees with BHT in lower levels to bounding kdt-Trees